




KubeCon



CloudNativeCon

THE LINUX FOUNDATION



AI_dev
Open Source GenAI & ML Summit

China 2024



KubeCon



CloudNativeCon



China 2024

OS Migration Solution on Cloud

Jianlin Lv, eBay

Agenda



China 2024

- Why Migration and Challenge
- eBay Cloud Platform
- Migration Solution Introduction
- Identify and Solve Production Issues
- Upstream Contribution

Why Migration and Challenge




China 2024

- Licensing models, TCO
- Linux Distributions EOL
 - Lack of official support
 - Compatibility issues
 - Software security vulnerabilities
- Challenge
 - Data security
 - Transparent to customers
 - Quickly upgrade nodes in limited time window
 - Balance upgrade coverage and normal business operation

eBay Cloud Platform



China 2024



- Nodes 130K+
- Pods 2M+
- OS
 - Ubuntu 22.04
 - Ubuntu 20.04
- Monthly release/rollout

Immutable OS



KubeCon



CloudNativeCon

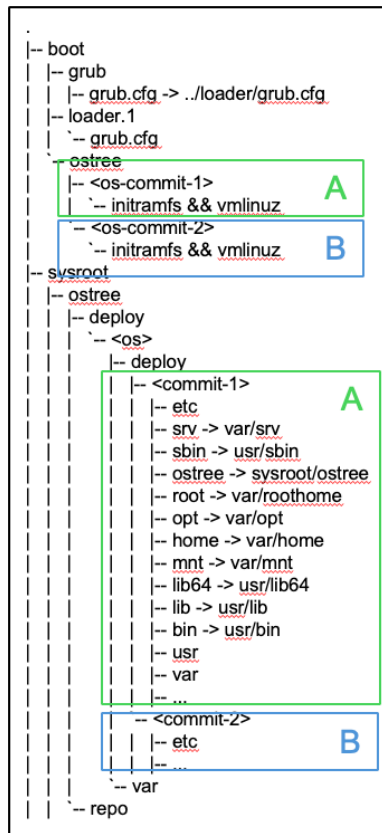


China 2024



- OSTree
- System update for Linux-based OS
- Deploy new OS
 - Atomic Upgrade/Rollbacks
 - OSTree commit create deployment at `/sysroot/ostree/deploy/<os>/deploy/<commit>`
 - /etc, 3-way merge, the changing will be propagated
 - Install kernel data at `/boot/ostree`
- System layout

Physical device



-
- ```

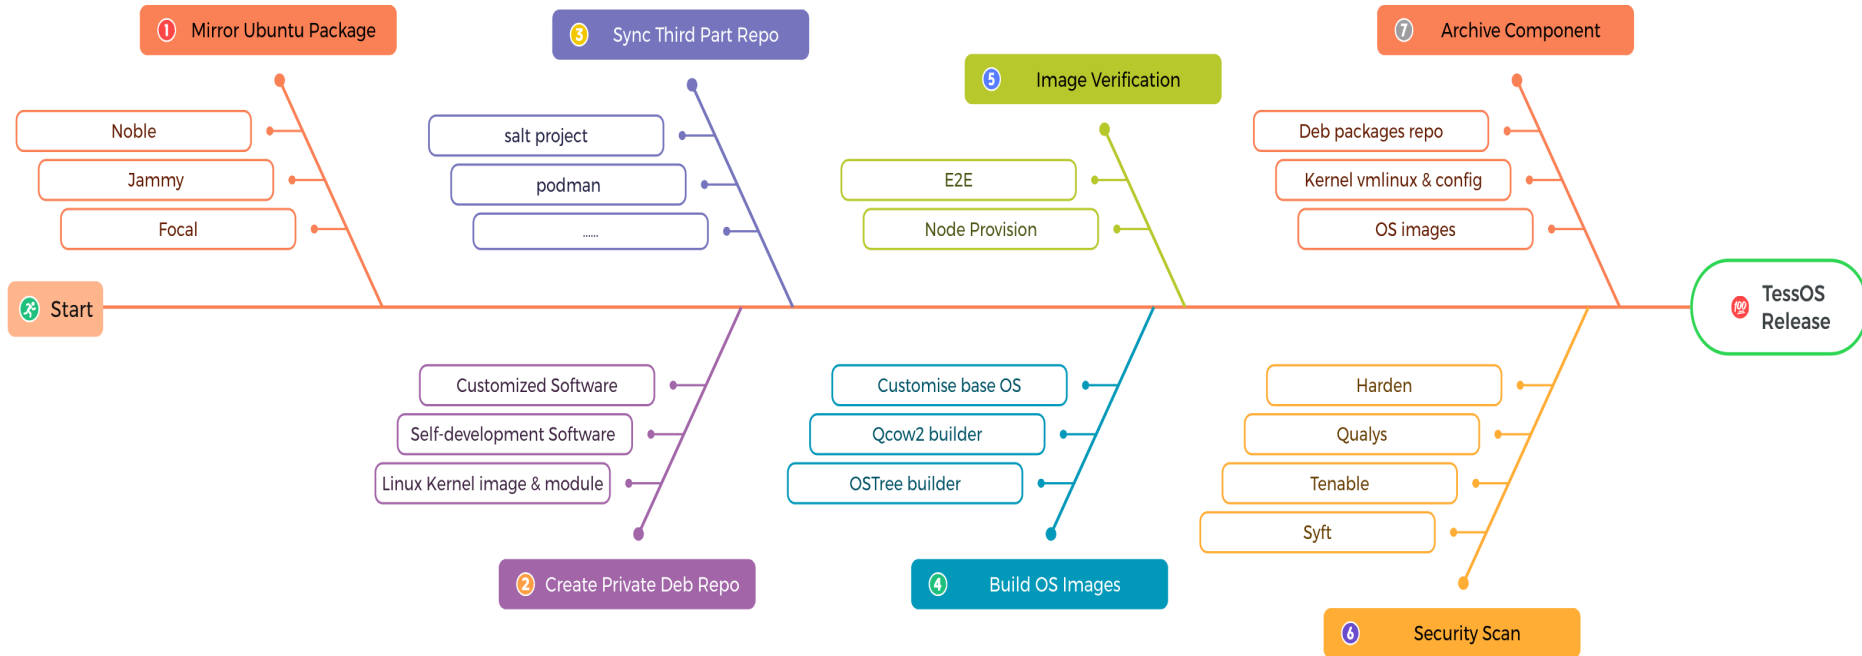
-- boot
-- grub
| |-- grub.cfg -> ../loader/grub.cfg
-- loader.1
| |-- grub.cfg
-- ostree
| |-- <os-commit-1>
| | |-- initramfs && vmlinuz
| |-- <os-commit-2>
| | |-- initramfs && vmlinuz
-- sysroot
| |-- ostree
| |-- deploy
| | |-- <commit-1>
| | | |-- etc
| | | |-- srv -> var/srv
| | | |-- sbin -> usr/sbin
| | | |-- ostree -> sysroot/ostree
| | | |-- root -> var/root/home
| | | |-- opt -> var/opt
| | | |-- home -> var/home
| | | |-- mnt -> var/mnt
| | | |-- lib64 -> usr/lib64
| | | |-- lib -> usr/lib
| | | |-- bin -> usr/bin
| | | |-- usr
| | | |-- var
| | | |-- ...
| | |-- <commit-2>
| | | |-- etc
| | |-- var
-- repo

```
- ```

bin -> usr/bin
boot
etc
home -> var/home
lib -> usr/lib
lib64 -> usr/lib64
mnt -> var/mnt
opt -> var/opt
ostree -> sysroot/ostree
...
root -> var/root/home
run
sbin -> usr/sbin
srv -> var/srv
sysroot
tmp
usr
  |-- bin
  |-- etc
  |-- include
  |-- lib
  |-- ...
  |-- local -> /var/usr/local
  |-- sbin
  |-- src
var
  |-- home -> ../sysroot/home
  |-- lib
  |-- local
  |-- lock -> ../run/lock
  |-- ...
  |-- root -> root
  |-- run -> ../run
  |-- usrlocal

```

• OS release workflow

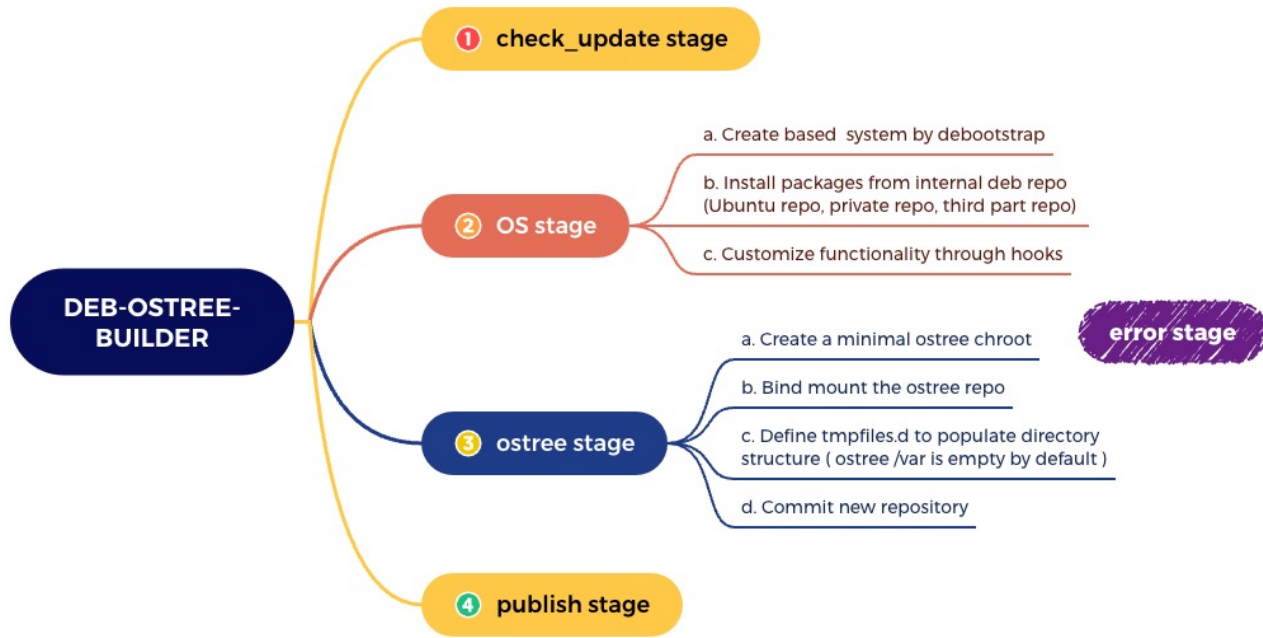


OSTree Builder



China 2024


- GitHub: [deb-ostree-builder](https://github.com/deb-ostree-builder/deb-ostree-builder)
- PR: [Support Ubuntu OSTree build](#)




TessOS Migration



China 2024



- Pre hook
 - Move user data (/var) to new deployment for Centos
 - Update grub file if need
- Migrate hook
 - Migrate config file (/etc) for application
- Post hook
 - Setup for new OS
- Monthly Patching
 - > 90% Nodes / Monthly



Identify and Solve Production Issues



China 2024

- Identify Problems
 - collection and standardization HW and OS metrics
 - data observability
 - real-time data analysis
- Dig into the Details
 - diagnostic tools(cpu, network, storage, bpf, etc.)
 - LnP testing environment
 - kernel debugging OS
- Develop Solutions

Upstream Contribution

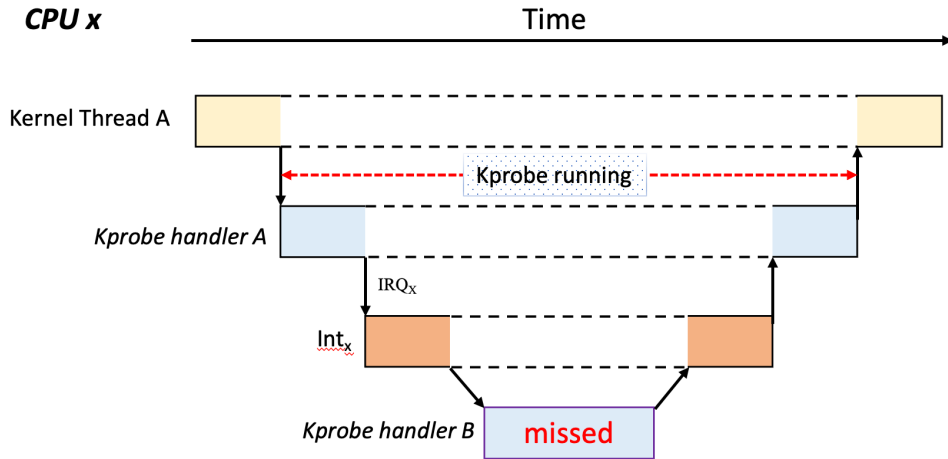


China 2024

- Support Ubuntu OSTree builds
- Fixed memory leak in audit module for Ubuntu 22.04
- SCSI Device Probing Non-deterministic
- Kprobe No Re-entrant

Kprobes Non-reentrant

- Kprobes event missing
- CONFIG_KPROBES_ON_FTRACE
- tracepoint: Allow livepatch module add trace event



```
kernel/tracepoint.c | 4 +--  
1 file changed, 2 insertions(+), 2 deletions(-)
```

```
diff --git a/kernel/tracepoint.c b/kernel/tracepoint.c  
index f23144af5743..8d1507dd0724 100644  
--- a/kernel/tracepoint.c  
+++ b/kernel/tracepoint.c  
@@ -571,8 +571,8 @@ static void for_each_tracepoint_range(  
    bool trace_module_has_bad_taint(struct module *mod)  
    {  
        return mod->taints & ~((1 << TAIN_T_OOT_MODULE) | (1 << TAIN_T_CRAP) |  
-           (1 << TAIN_T_UNSIGNED_MODULE) |  
-           (1 << TAIN_T_TEST));  
+           (1 << TAIN_T_UNSIGNED_MODULE) | (1 << TAIN_T_TEST) |  
+           (1 << TAIN_T_LIVEPATCH));  
    }  
}
```



KubeCon



CloudNativeCon



China 2024



Thank You