KubeCon CloudNativeCon

THE LINUX FOUNDATION

OPEN SOURCE SUMMIT

AI_dev

Open Source GenAI & ML Summit

China 2024

KubeCon | CloudNativeCon China 2024

THE LINUX FOUNDATION OPEN SOURCE SUMMIT

AI_dev Open Source GenAI & ML Summit

Towards an Unified Cloud-Native Infra for AI & HPC

TACC - A Five-Year Journey

Kaiqiang Xu: HKUST

Peter Pan : DaoCloud

TACC: Turing AI Computing Cloud

About Us

Kaiqiang Xu

- **PhD Researcher**
  in Computer Systems for Machine Learning
- **Hong Kong University of Science and Technology**

香港科技大學
THE HONG KONG
UNIVERSITY OF SCIENCE
AND TECHNOLOGY

Peter Pan

- **R&D Engineer Lead**
- **Open Source Advocate**
- **DaoCloud**

DaoCloud

TACC: Intro

High-performance and highly scalable AI computing infrastructure

TACC is an AI computing infrastructure designed for machine learning applications, supporting and accelerating the constantly evolving research in machine learning at both the software and hardware levels.

Due to system-level optimizations specifically targeted for ML/DL programs, TACC outperforms traditional HPC computing clusters in terms of both performance and stability.

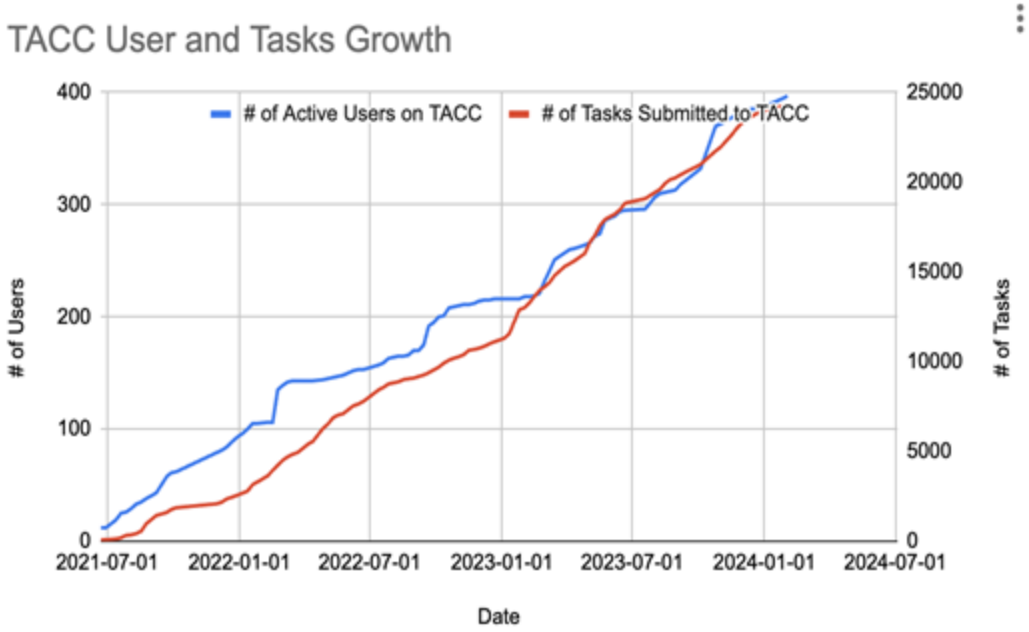

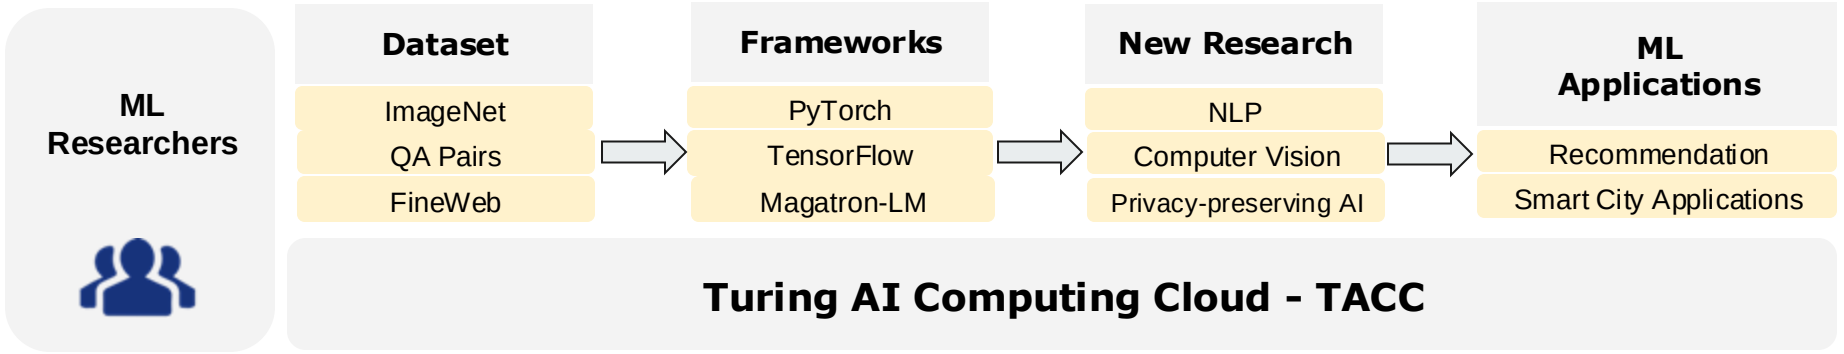

TACC: Intro

Underpinning Research

- **ML frameworks** enhance model development through advanced parallelization and distributed training, efficiently handling complex computations and large datasets.
- **Cluster resource scheduler** optimizes cluster-wide resource allocation across AI tasks, boosting overall job throughput and other efficiency factors in AI clusters.
- **AI-centric networking** improves data flow and reduces latency by efficiently managing large model transport and using FPGAs for compute offloading

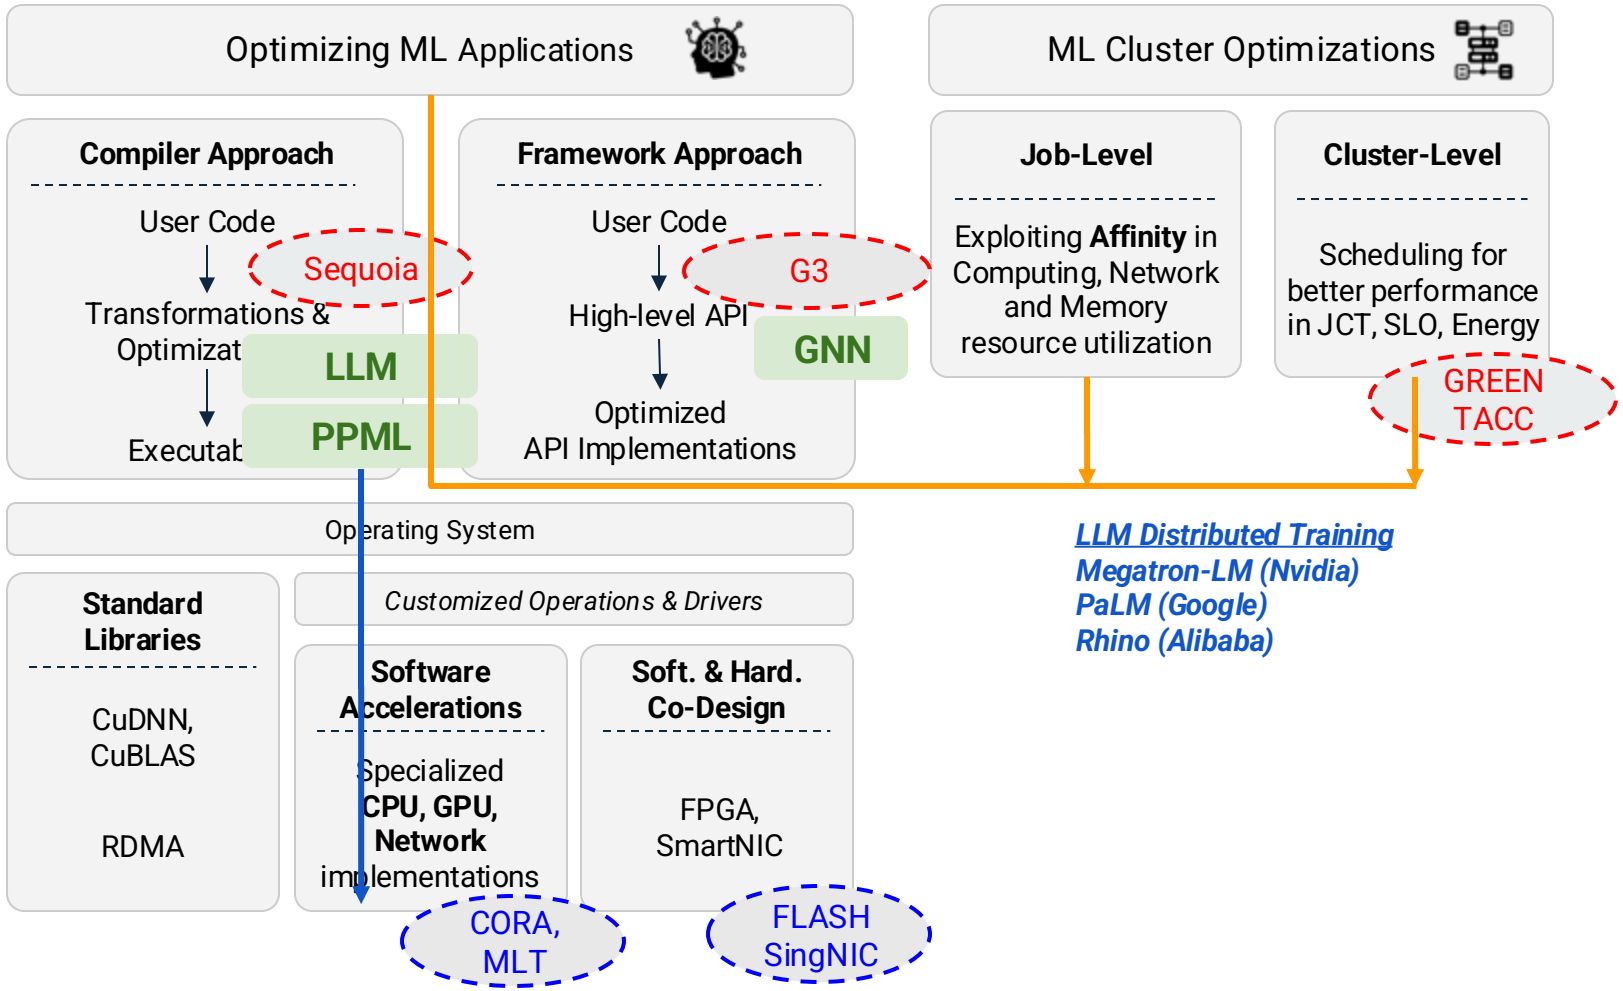

TACC: Intro

Tutorial

- https://github.com/turingaicloud/quickstart

Core Code

- https://github.com/turingaicloud/tcloud-sdk

TACC: Streamline Scientists‘ daily

Interface

- Scientist are familiar with slurm-like CLI
- Strightforward CLI & Config (tuxiv.config)

```
# tcloud submit          # submit job to remote cluster

# tcloud ps              # show job status

# tcloud upload          # upload to remote working dir
```

tuxiv.config

```
entrypoint:
  - python mnist.py ..
environment:
    channels:
      - nvidia
    dependencies:
      - pytorch=1.9.0
      - torchvision=0.10.0
      - cudatoolkit=11.1.74
datasets:
    -  $name
job:
    name: test
    general:
      - nodes=2
      - cpus-per-task=10
      - gres=gpu:2
```

entrypoint

Python packages

dataset source

Slurm Params

TACC: Streamline Scientists‘ daily

Remote Execution

- SCP to copy local code to user folder
- SSH to run script remotely
- Slurm agent in docker

Tenant & User

- Re-use Linux user system in servers
- Mount distributed storage to each /home/$user

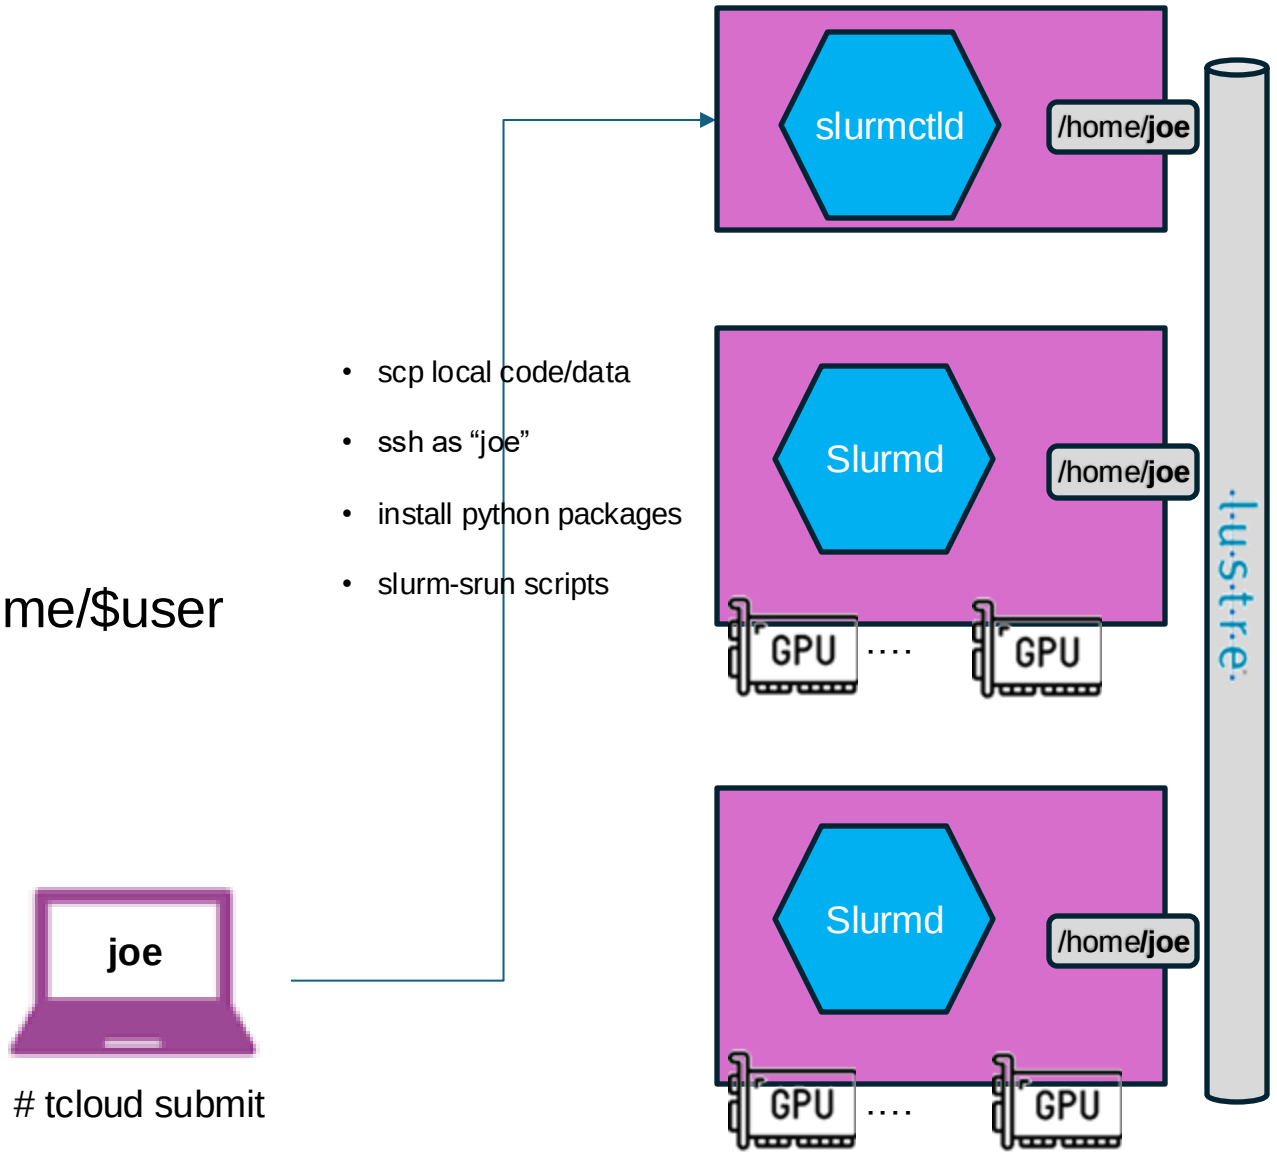

TACC: Streamline Scientists‘ daily

Python Env

⚠ “venv + pip install” every time is time-consuming !

- Individual Conda Env for each user
- Persistent & re-use env and packages in user folder

Dataset

- pre-fetch dataset to user folder
- Distributed shared storage (e.g. Lustre)

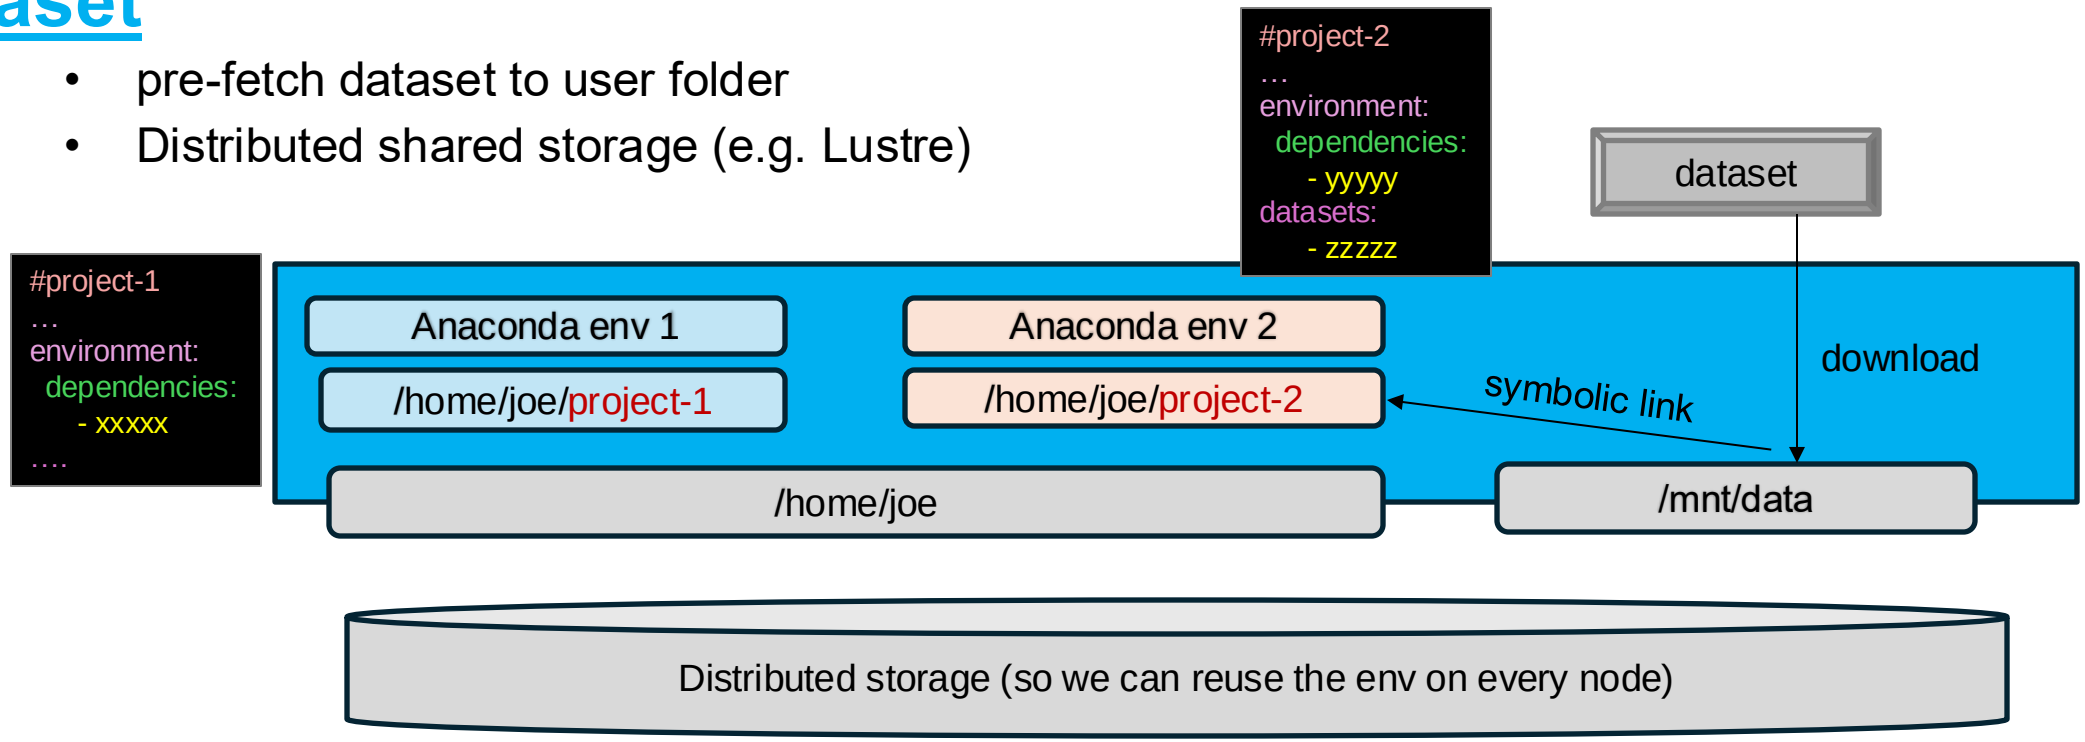

Kubernetes - Slurm

Survey:
Which resource-scheduling/job-management tool do you use as part of your AI/ML tech stack?

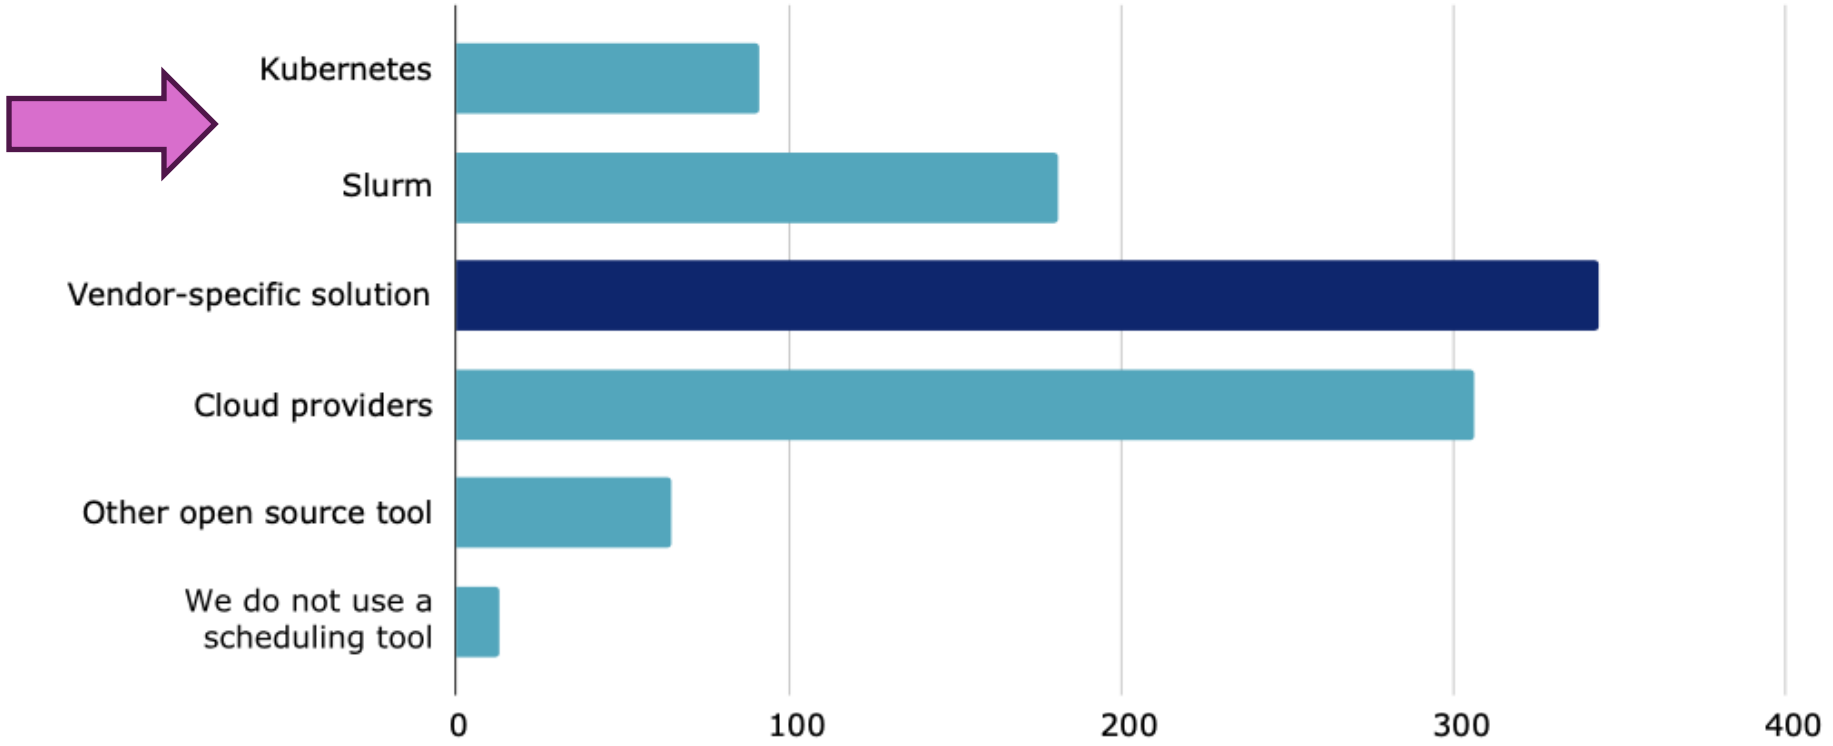

https://ai-infrastructure.org/wp-content/uploads/2024/03/The-State-of-AI-Infrastructure-at-Scale-2024.pdf

TACC powered by K8S

Kubernetes **v.s.** **Slurm**

- 10 years, still young
- not initially designed for ML workload, but scheduling capability keep evolving
- declarative (final-state-driven)
- docker build burden
- flexible and scalable
- both for training & serving

- 22 years old, mature for HPC
- Born for batch workload: Queue, Job, Scheduling
- Stability issue
- Just plain shell and files
- lack of auto-scaling
- OCI image limited supported

TACC: Streamline Scientists‘ daily

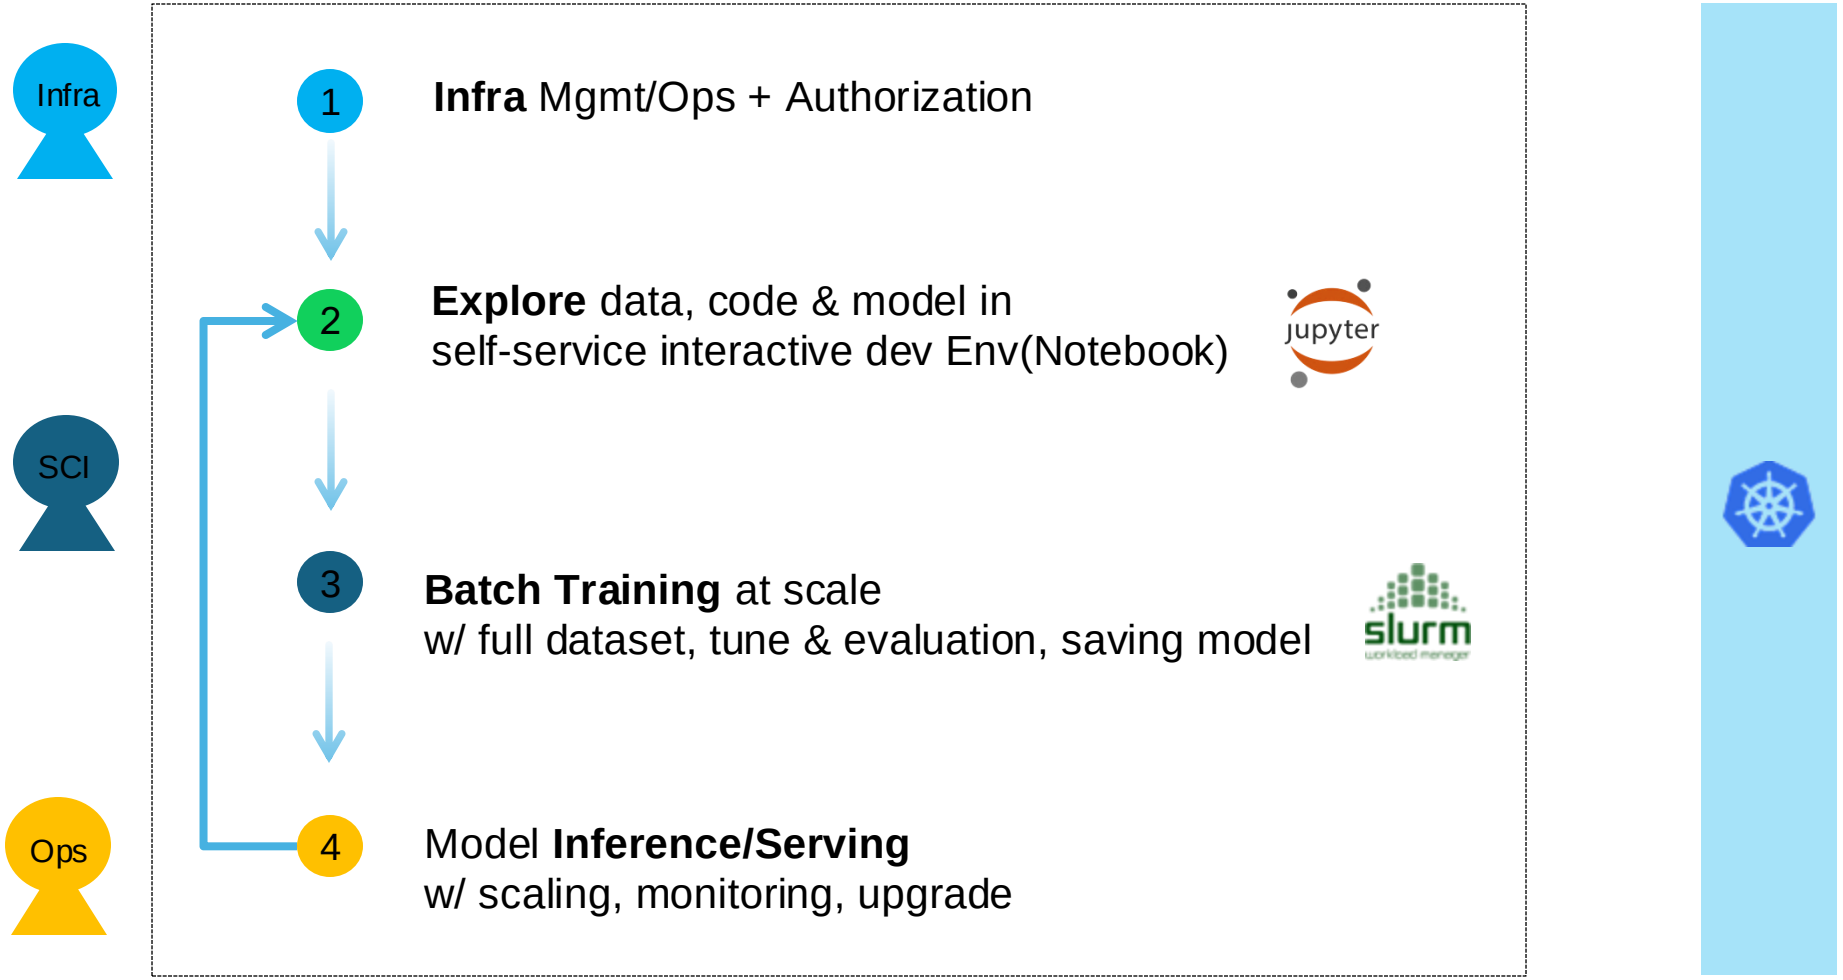

How K8S helps Infra. Engineer

Infra

Kubernetes	All machines in a Resource Pool
GPU-Operator	No worry about driver/plugin/runtime.. for each node
dcgm-exporter	GPU monitoring
MIG / HAMi	virtual GPU (partition)
Volcano / Kueue	Job & Queue Scheduling
Karmada	Multiple Cloud/Cluster

How K8S helps Infra. Engineer

When introducing other Heterogeneous xPUs

Infra

Kubernetes

GPU-Operator | xPU-Operator

dcgm-exporter | xPU-exporter

MIG / HAMi *

Volcano / Kueue

Karmada

How K8S helps Scientist

Multi-Tenant

Dev & Train

Scheduling

Monitor

Performance & Robustness

- k8s namespace for each $user
- Namespace Quota
- Queue Quota (Kueue concept)

```
# kubectl -n joe get resourcequotas  -o yaml

apiVersion: v1
kind: ResourceQuota
spec:
  hard:
    requests.nvidia.com/mig-1g.6gb: "2”
    requests.storage: 1000Gi
    limits.memory: 128Gi
```

```
# kubectl  get clusterqueue default -o yaml

apiVersion: kueue.x-k8s.io/v1beta1
kind: ClusterQueue
spec:
flavors:
    - resources:
      - name: cpu
        nominalQuota: "200"
      - name: memory
        nominalQuota: 1280Gi
      - name: nvidia.com/gpu
        nominalQuota: "10"
      - name: rdma/hca_shared_devices_a
        nominalQuota: "100"
```

How K8S helps Scientist

- Multi-Tenant
- Dev & Train
- Scheduling
- Monitor
- Performance & Robustness

code

Dataset / Modeling

Python Packages

Paints

- Should learn about “docker build”
- New image when any tiny change
- Huge image size

```
FROM python

RUN pip install torch==2.4.0  torchvision==0.19.0 transformers==4.4.1 datasets==2.21

RUN  git clone $code

RUN  wget  $dataset / $model
```

How K8S helps Scientist

- Multi-Tenant
- Dev & Train
- Scheduling
- Monitor
- Performance & Robustness

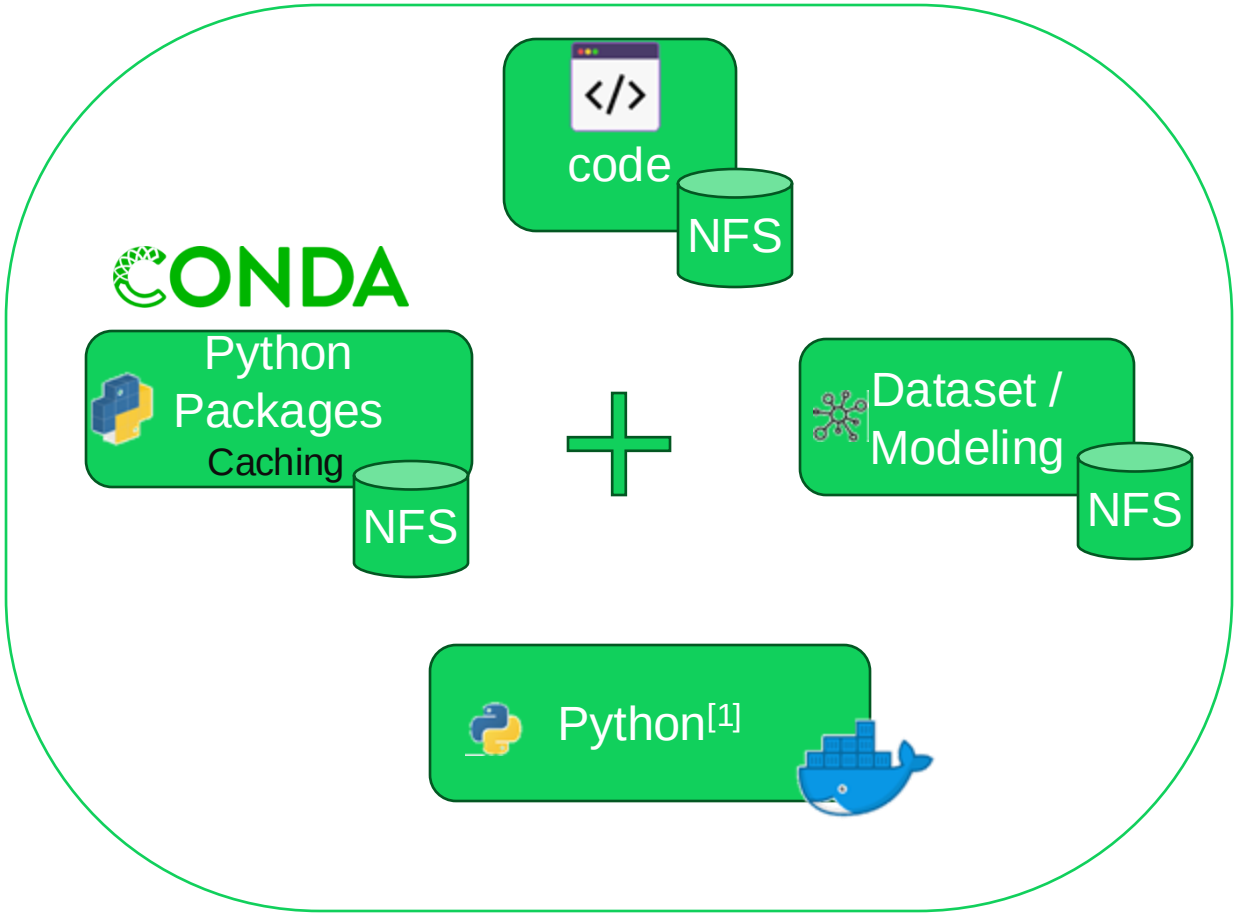

- Reusable conda/pip environment
- Data Isolated by namespaced PVC
- No worry about building ”docker image”

1. **Code**: can be either managed by Git then pre-fetched by init-container, or rsync to NFS shares

2. **Data/Model**: can either download from S3 and pre-fetch by init-container

3. **Python Packages**: with conda capability, pre-install into NFS shares and can be re-used

4. **Docker image**: can be very flexible. either include all required pip packages & cuda inside, or a just a bare image all relies on conda volume.

[1] Base Image: python itself can also be hosted in conda env.
But in practice in LLM age, Pytorch + CUDA

How K8S helps Scientist

- Multi-Tenant
- Dev & Train
- Scheduling
- Monitor
- Performance & Robustness

distributed training on K8s

Kubeflow's training-operator

Features

- Distributed Training (e.g. PyTorchJob)
- All-Reduce Style Training with MPI
- High Performance Computing (HPC) with MPI
- Job Scheduling with Volcano, Kueue
- Elastic Training

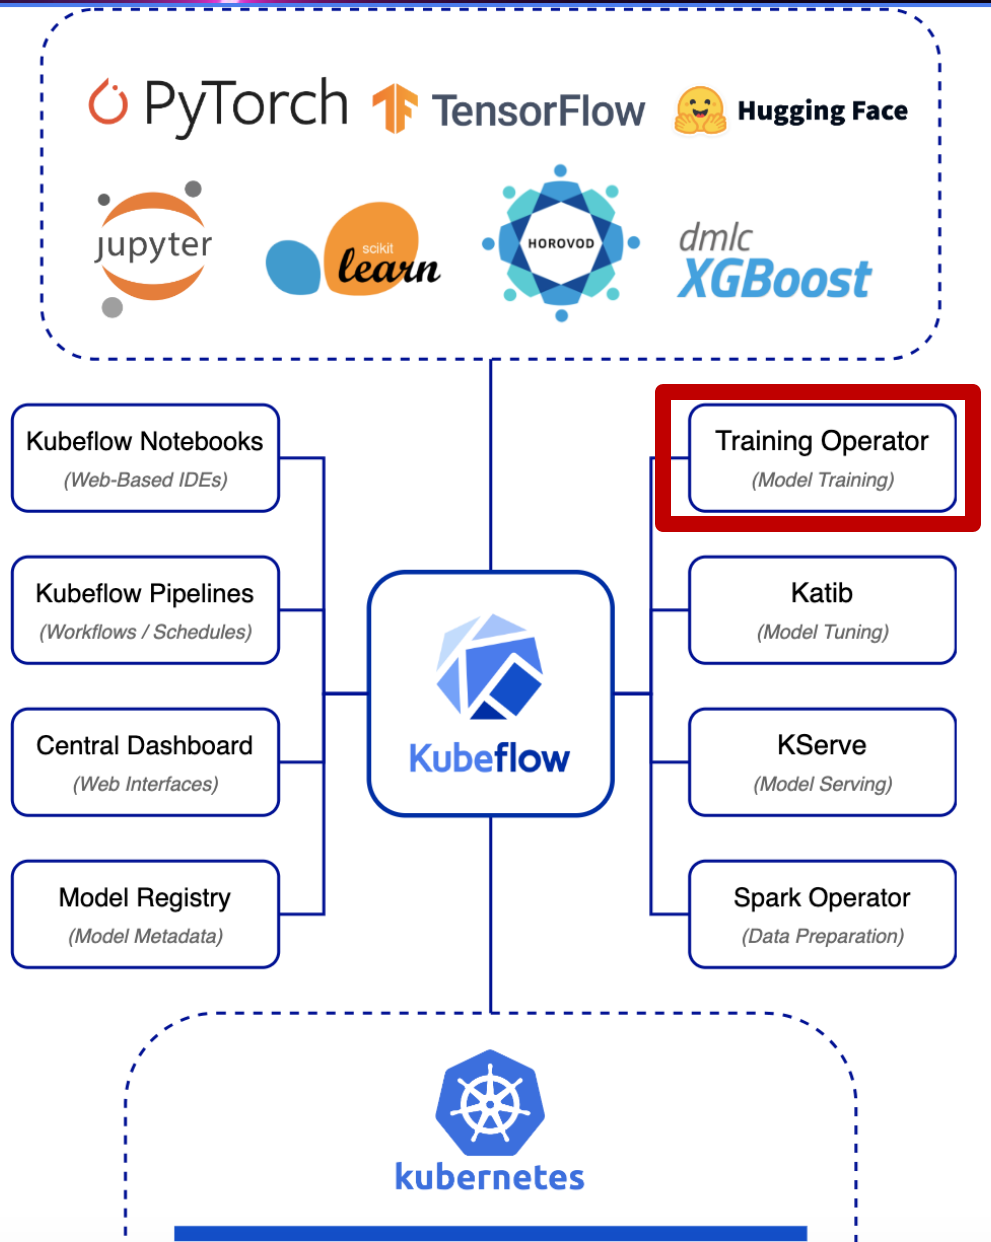

Web UI

- Multi-Tenant
- Dev & Train
- Scheduling
- Monitor
- Performance & Robustness

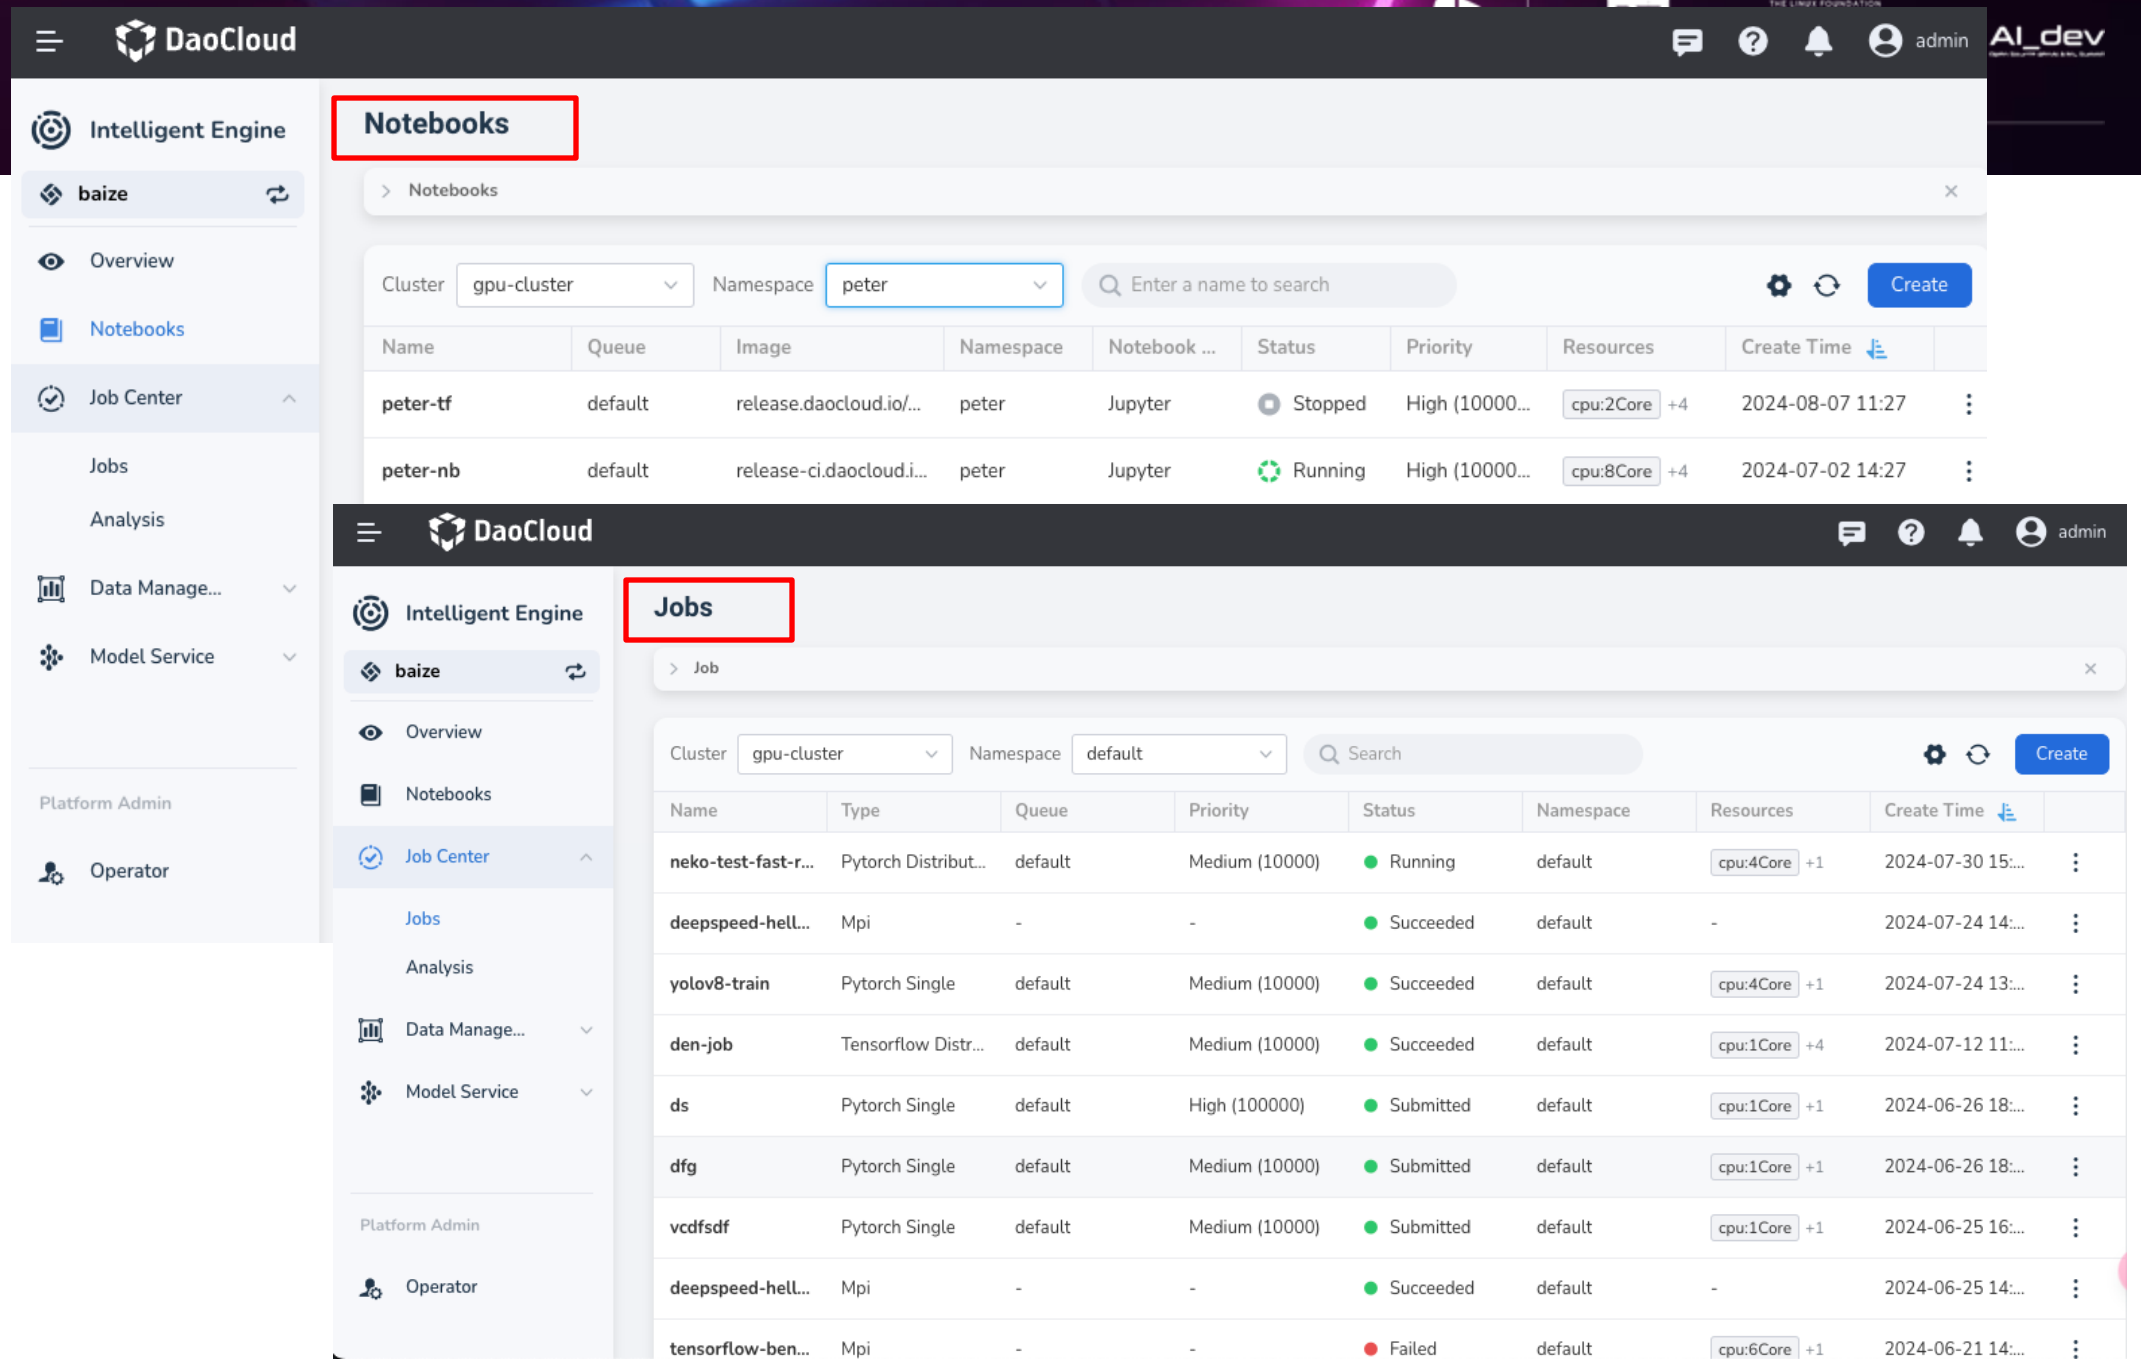

Scheduling is key demand

- Multi-Tenant
- Code & Data
- Scheduling
- Monitor
- Performance & Robustness

Q: What types of solutions does your organization currently lack in your AI/ML tech stack?

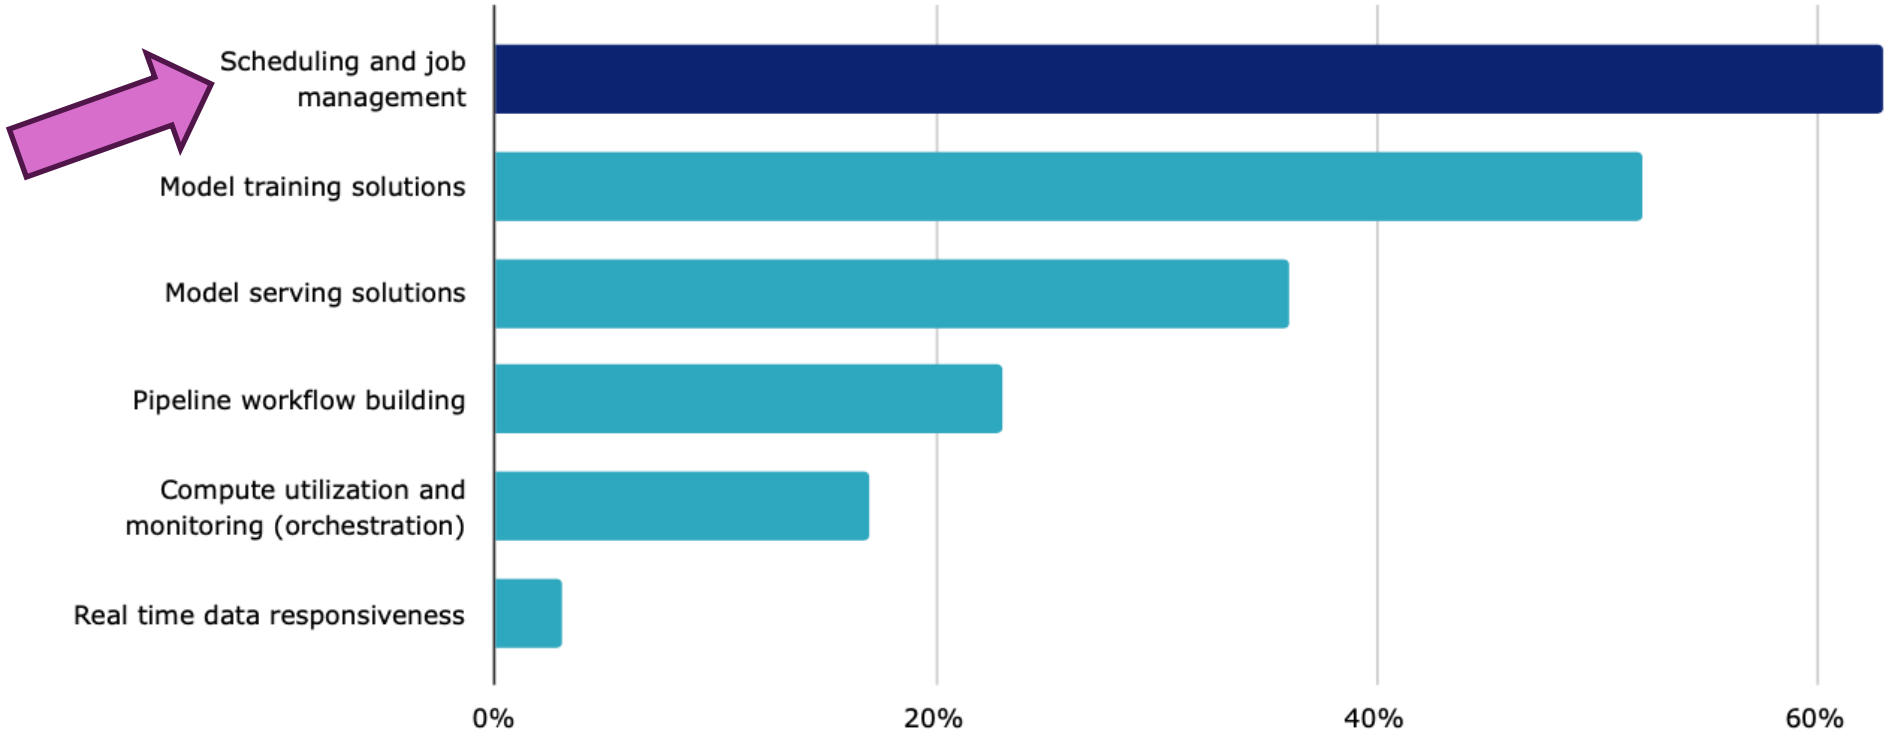

https://ai-infrastructure.org/wp-content/uploads/2024/03/The-State-of-AI-Infrastructure-at-Scale-2024.pdf

How K8S helps Scientist

SCI

- Multi-Tenant
- Code & Data
- Scheduling
- Monitor
- Performance & Robustness

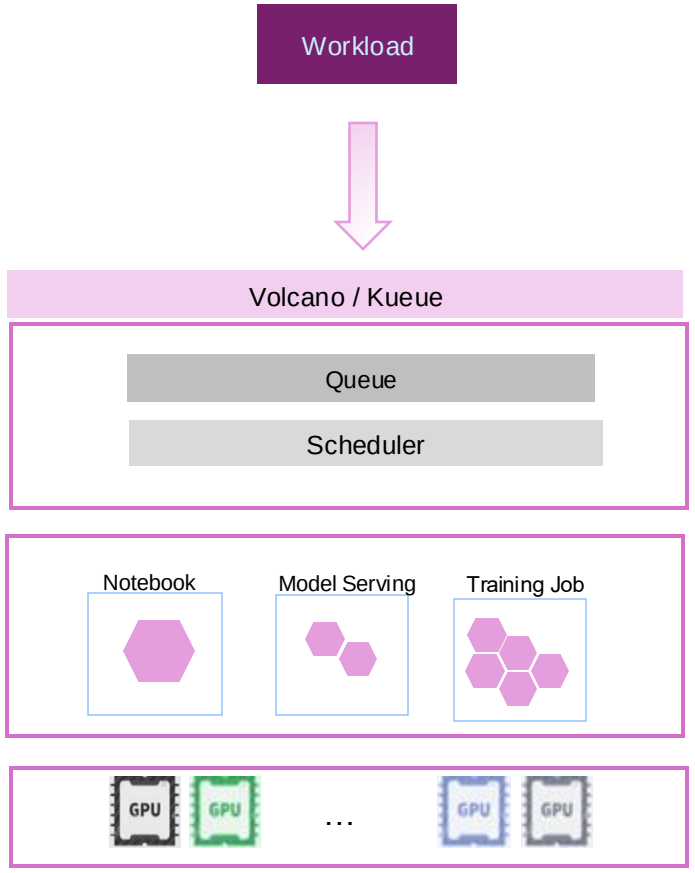

Various Scheduling Policy supported

Gang

- avoid waste/ deadlock
- Big jobs starve smalls

Preempt

- Diverse SLO

Binpack

- Avoid fragments
- Small jobs starve biggers

Affinity

- inner-node performance first

How K8S helps Scientist

SCI

- Multi-Tenant
- Code & Data
- Scheduling
- Monitor
- Performance & Robustness

- GPU Inventory
- GPU Healthy
- GPU Utilization *
- Inference Queue/Latency
- Billing for tenants
- Training Failure
- Alerting

- TensorBoard

arena submit xxx --tensorboard

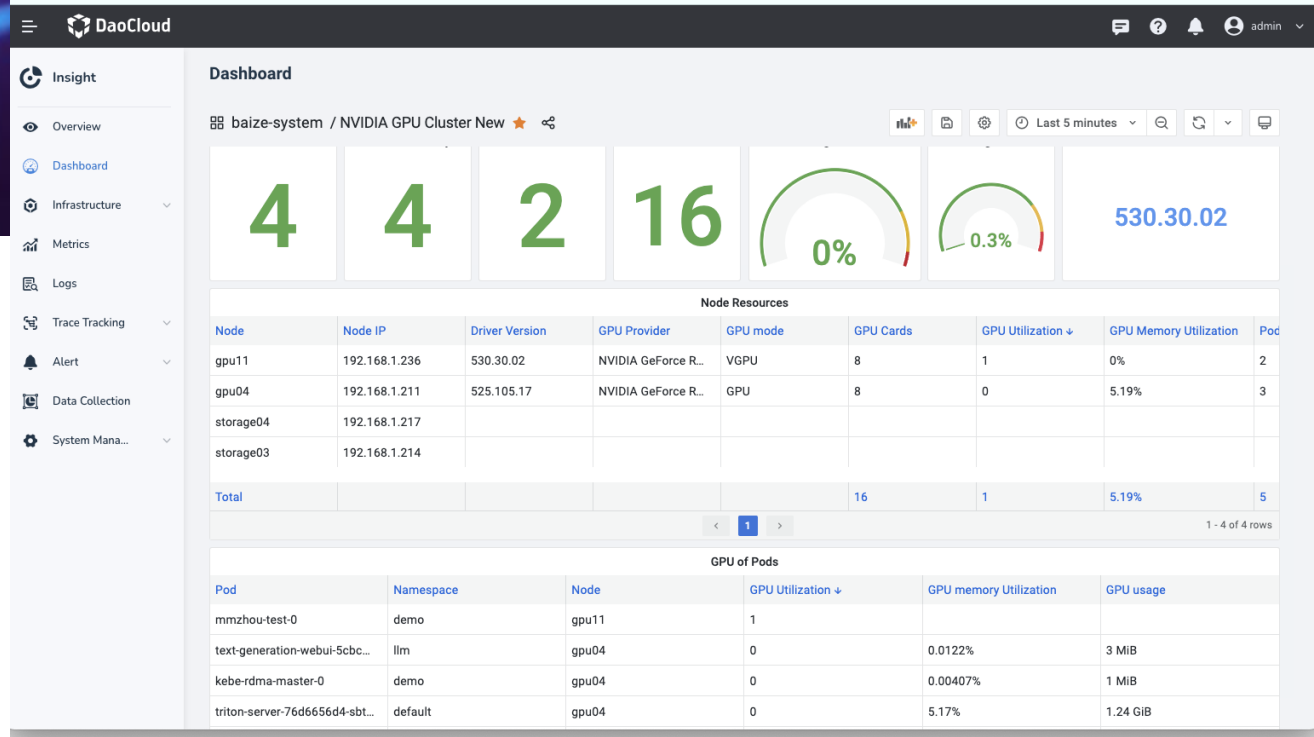

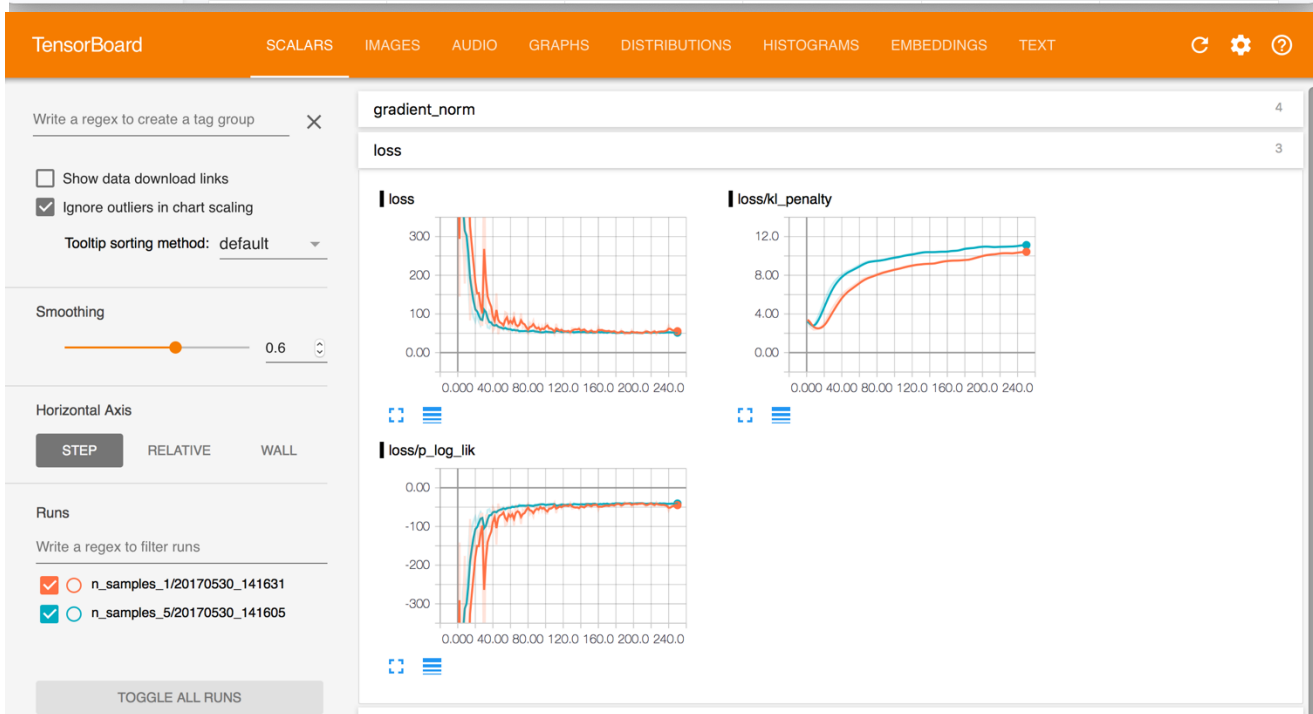

How K8S helps Scientist

Multi-Tenant

Code & Data

Scheduling

Monitor

Performance & Robustness

Bottom necks

- Dataset loading --- caching
- Model checkpoint saving --- disk I/O bandwidth
- Network interconnect --- RDMA full utilized + best topology
- Intermedia failure --- waste of GPU time

How K8S helps Scientist

Multi-Tenant

Code & Data

Scheduling

Monitor

Performance & Robustness

- **Distributed Storage**
  - Lusture
  - BeeGFS
  - MinIO
- **Caching**
  - JuiceFS
  - Alluxio
- **Local Storage**
  - HwameiStor

- **RoCE**
- **InfiniBand**

SpiderPool

- SRIOV-CNI
- GPU-Direct-RDMA
- GPU-NIC Topology Scheduling

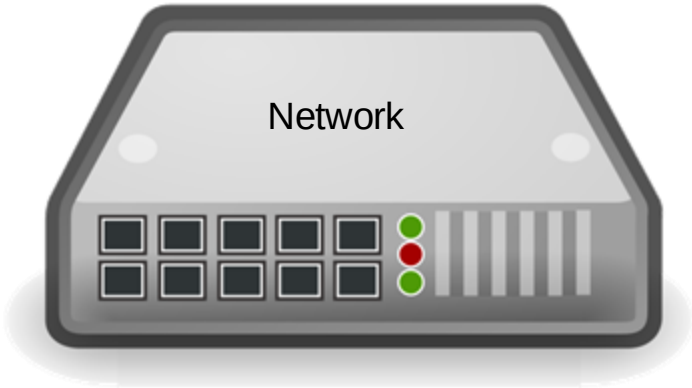

How K8S helps Scientist

SCI

- Multi-Tenant
- Code & Data
- Scheduling
- Monitor
- Performance & Robustness

Meta:

during a 54-day Llama 3 405B training,
Average interruptions **every 3 hours**,
>50% of the 419 unexpected
were caused by issues with **GPUs** or their onboard HBM3 memory.

aware, locate, diagnostic

↓

mitigation , resume

Sit Back and Relax with Fault Awareness and Robust Instant Recovery for Large Scale AI Workloads | 坐和放宽，了解大规模 AI 负载场景下的故障感知和健壮的快速故障恢复 - Fanshi Zhang & Kebe Liu, DaoCloud

https://sched.co/1eYY2

Wednesday August 21, 2024 15:35 - 16:10 HKT

Level 1 | Hung Hom Room 3

How to Unify K8S & Slurm

Challenges of co-hosting K8S & Slurm on the same cluster

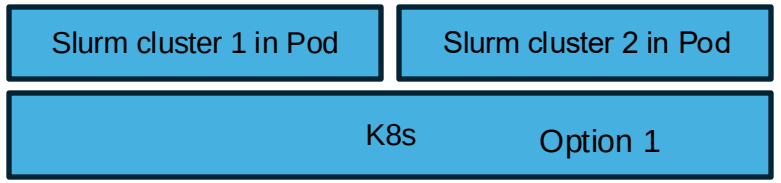

1. How to deploy slurmd(slurm agent) and container-runtime?

- Option 1: slurm cluster in pod →
- Option 2: slurm/K8s live in the same place 👍

  slurmd does not conflict with containerd or kubelet

2. How to make them aware of each other's resource usage ?

- **problem**: if a GPU been occupied by K8S Pod, Slurm isn't aware of it

PS: "SUNK" is not available yeah.. https://slurm.schedmd.com/SC23/Slurm-and-or-vs-Kubernetes.pdf

So TACC's way as below:

Co-host Deployment

every GPU machine is ready to handle either Slurm or K8S workloads

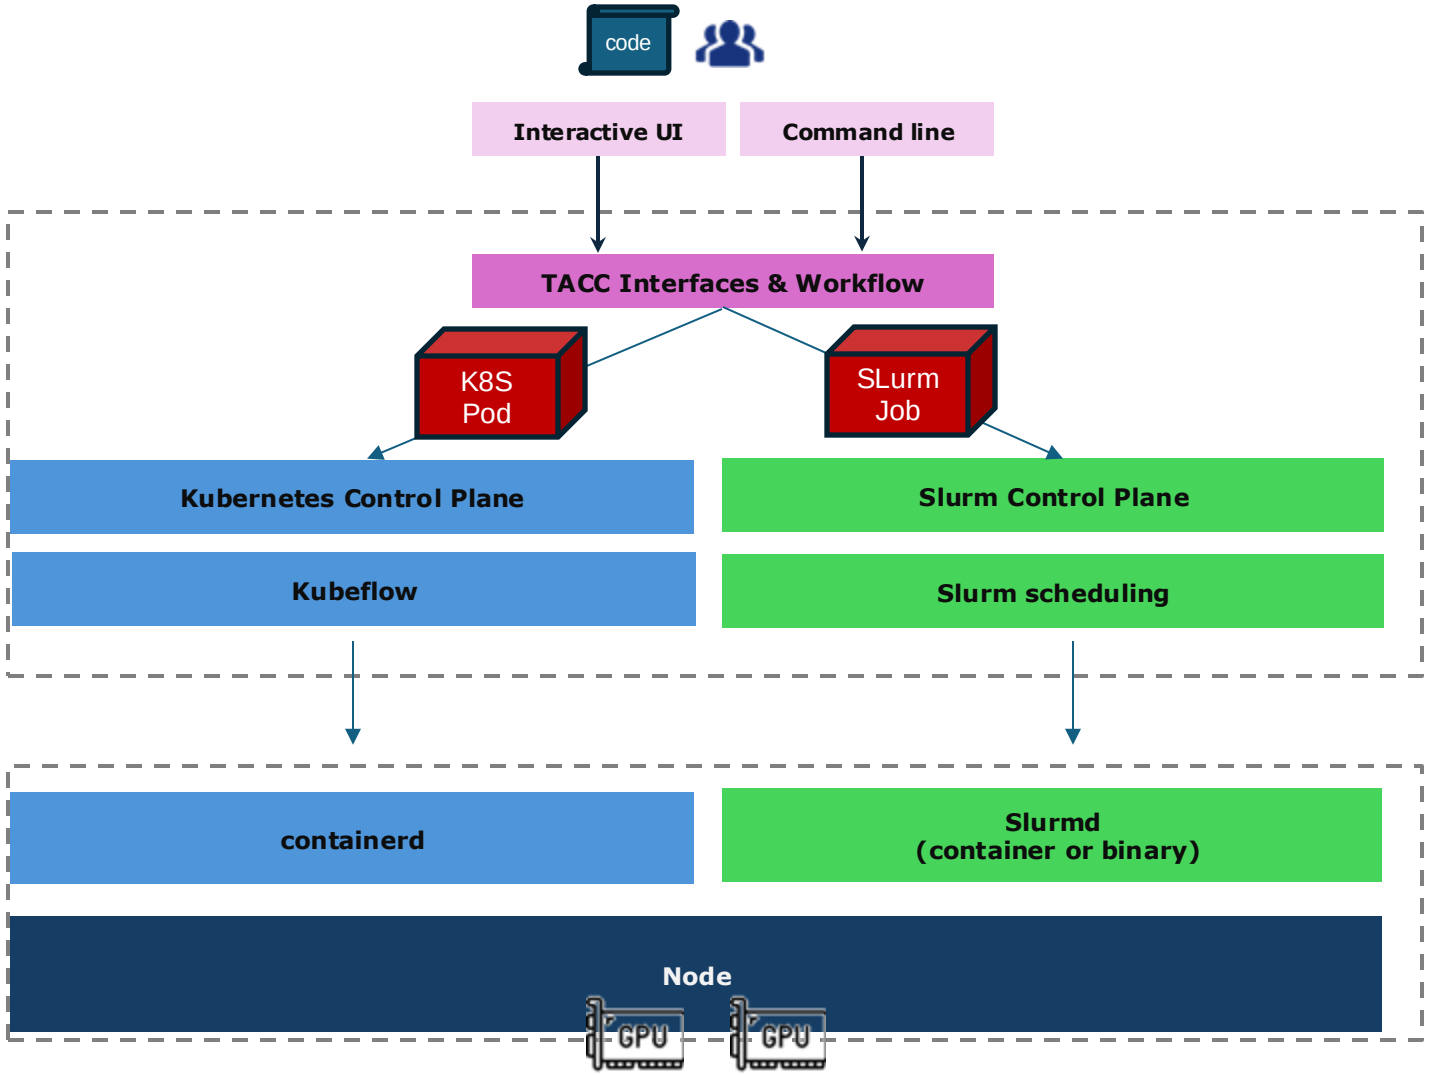

k8s will not use those resource Slurm already consumes, and vice versa.

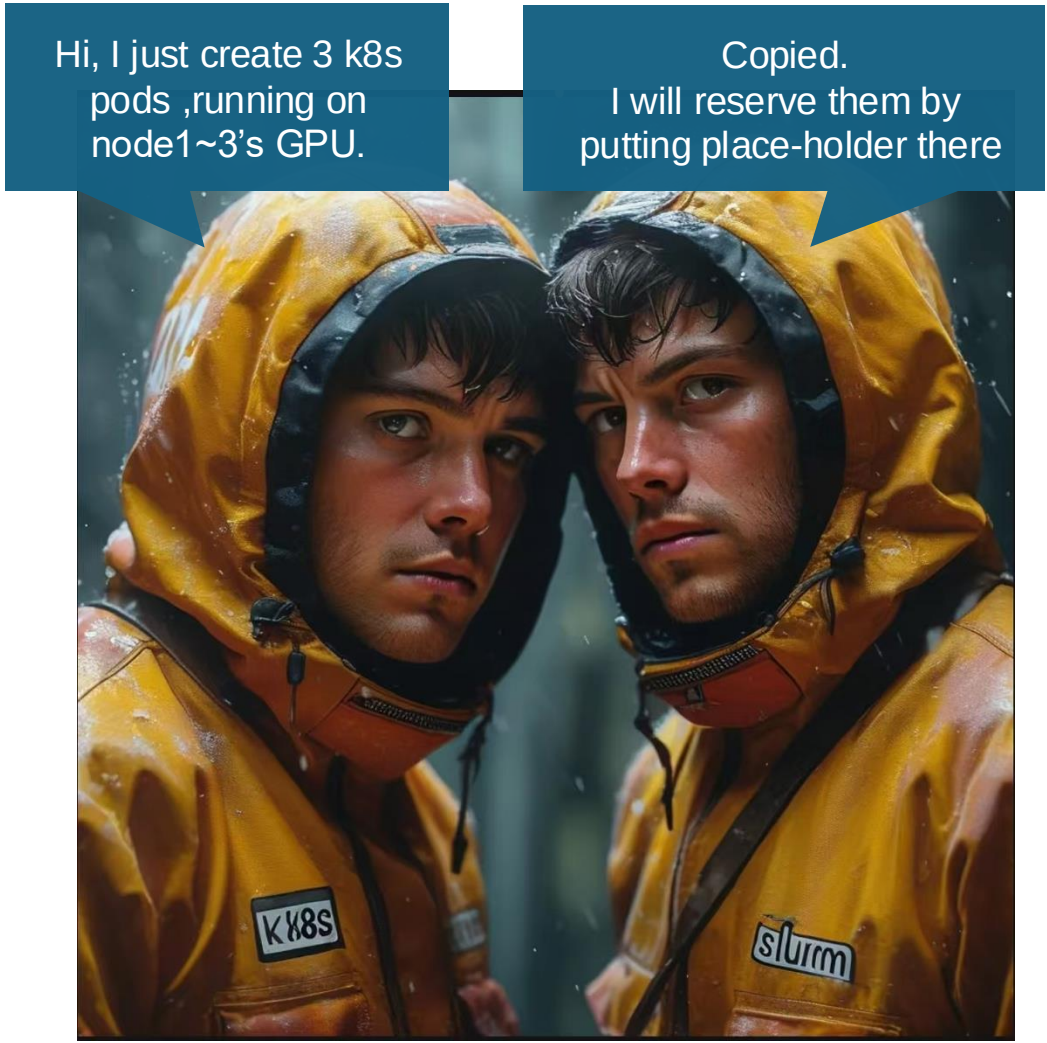

Shadow Resource Place Holder

Slurm job as place holder but should not really use the GPU

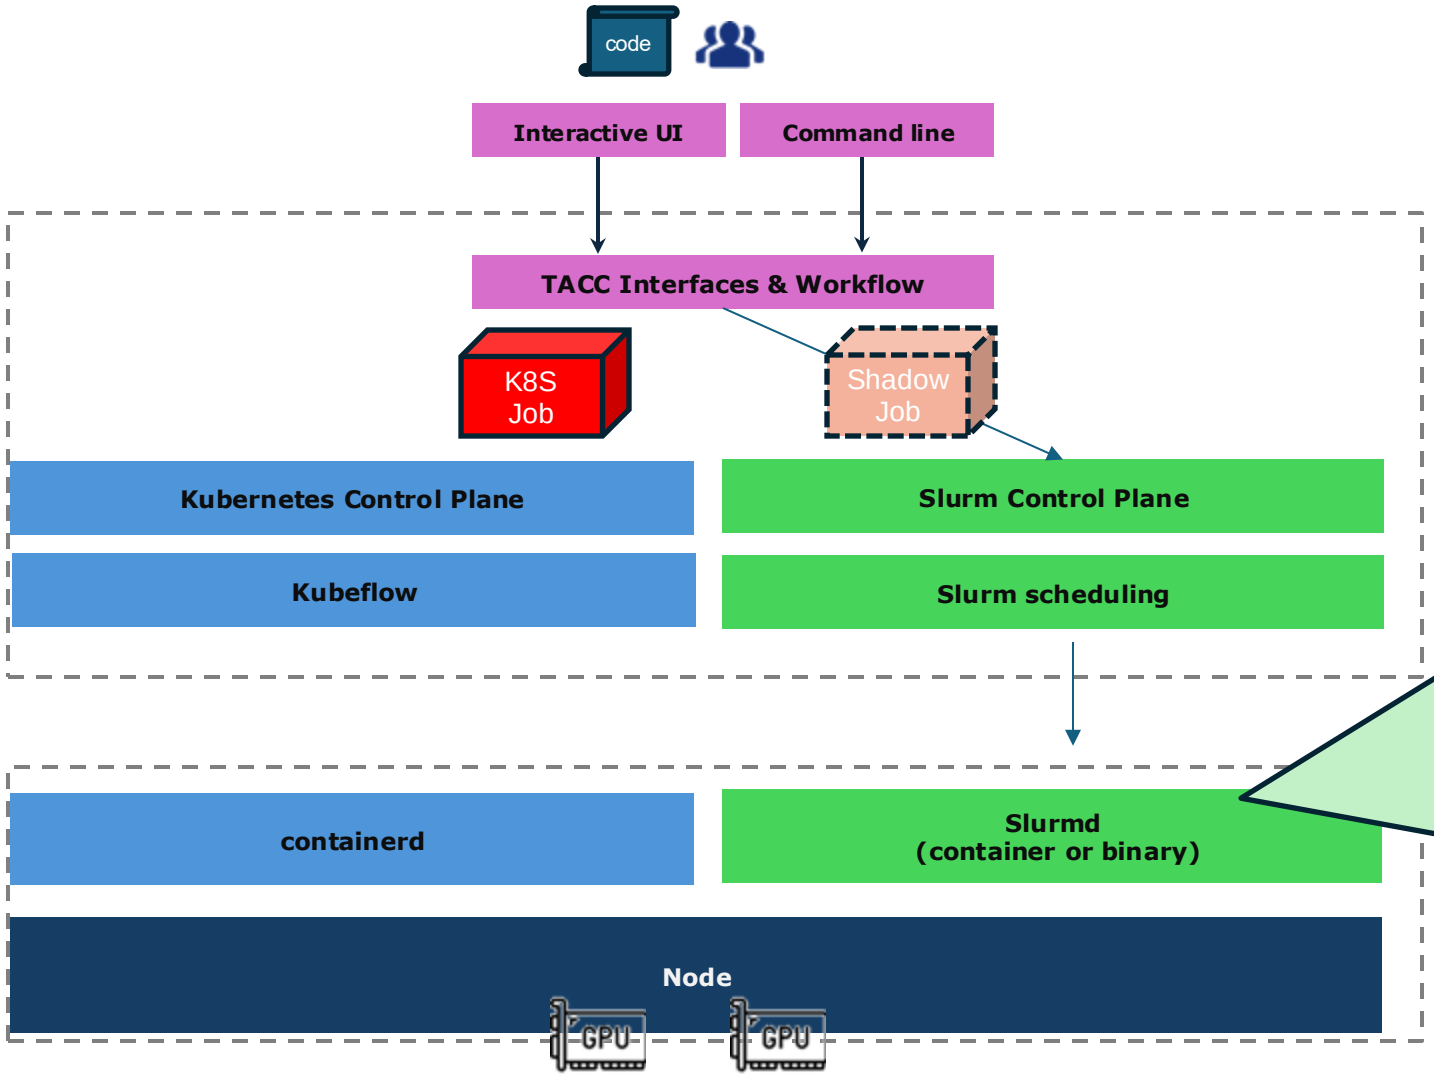

Example:

K8s Pod (actual workload)

- 2C/4G/1GPU
- 4 replicas scheduled to Node 1/3/4/5

Shadow Slurm job :

- Command:
- sleep and wait job $aboveJob (`kubectl wait `)
- -N 4 –n 4
- --cpus-per-task=2
- --mem=4G
- -gres=gpu:1
- --nodelist node[1,3,4,5]

Vice versa

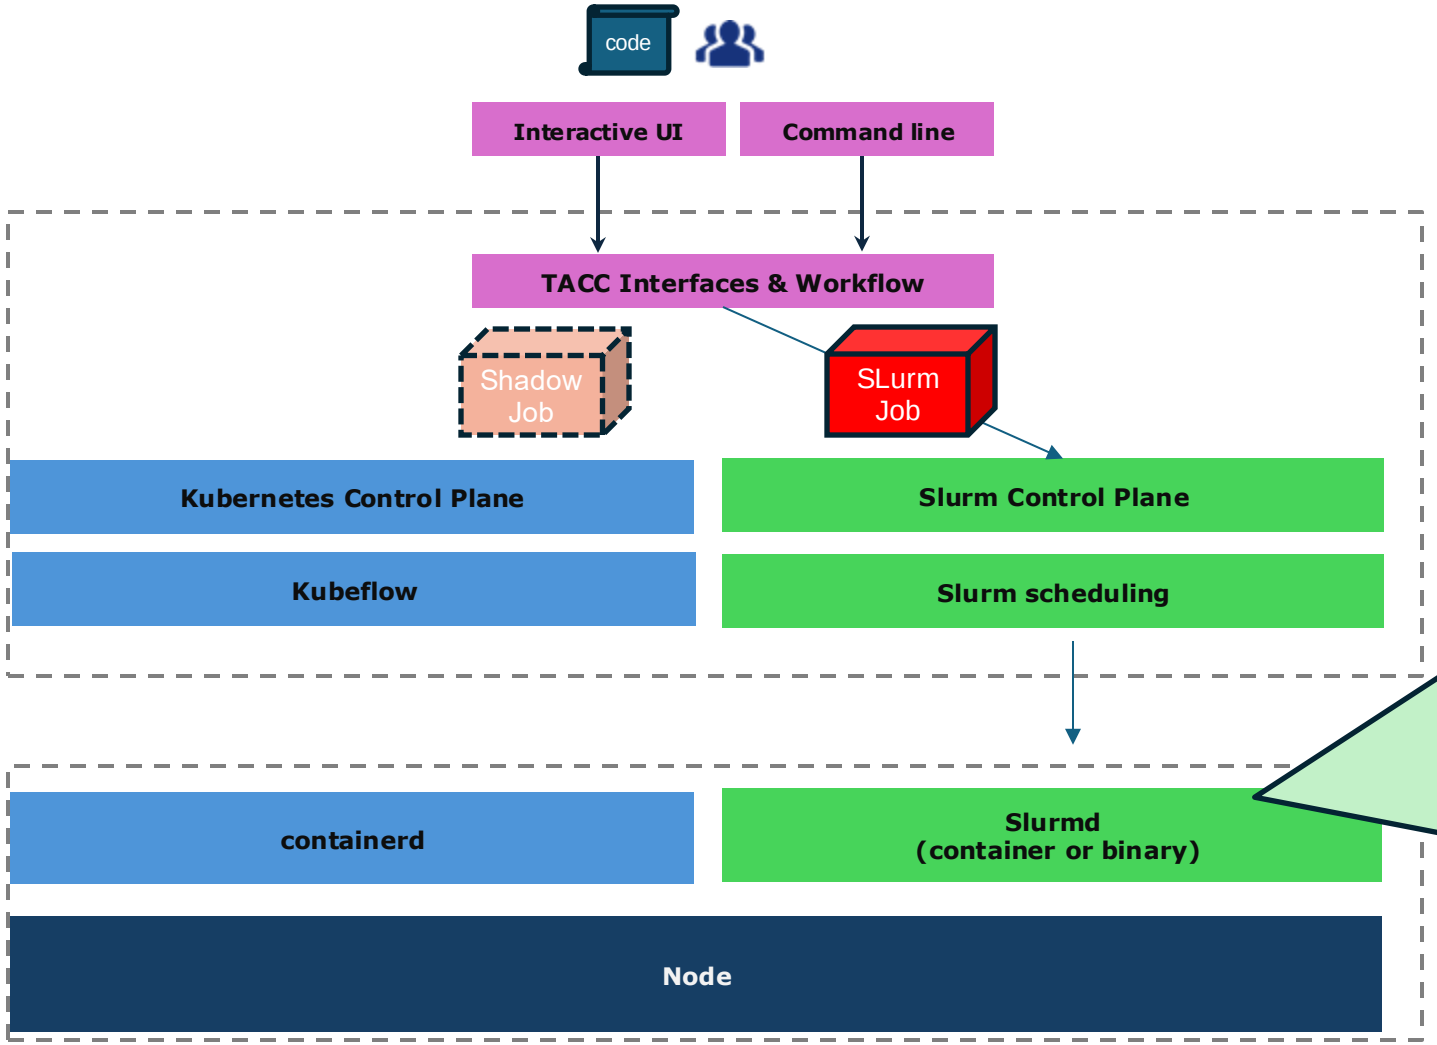

Example:

Slurm Job (actual workload)

- 4 process scheduled to Node 1/3/4/5

Shadow K8S Pods :

- Command: `squeue` to wait job completion
- Request same resources as slurm job for each

Demo Code

- https://github.com/turingaicloud/tcloud-sdk/tree/k8s-slurm

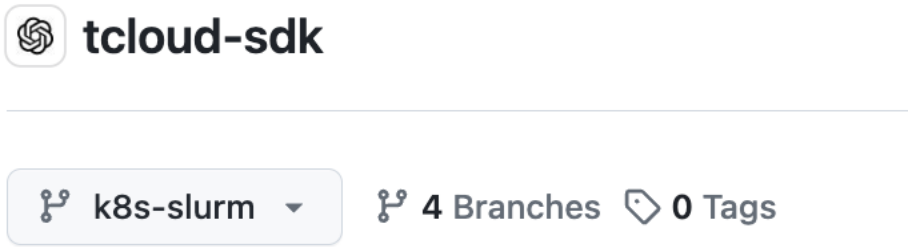

Demo

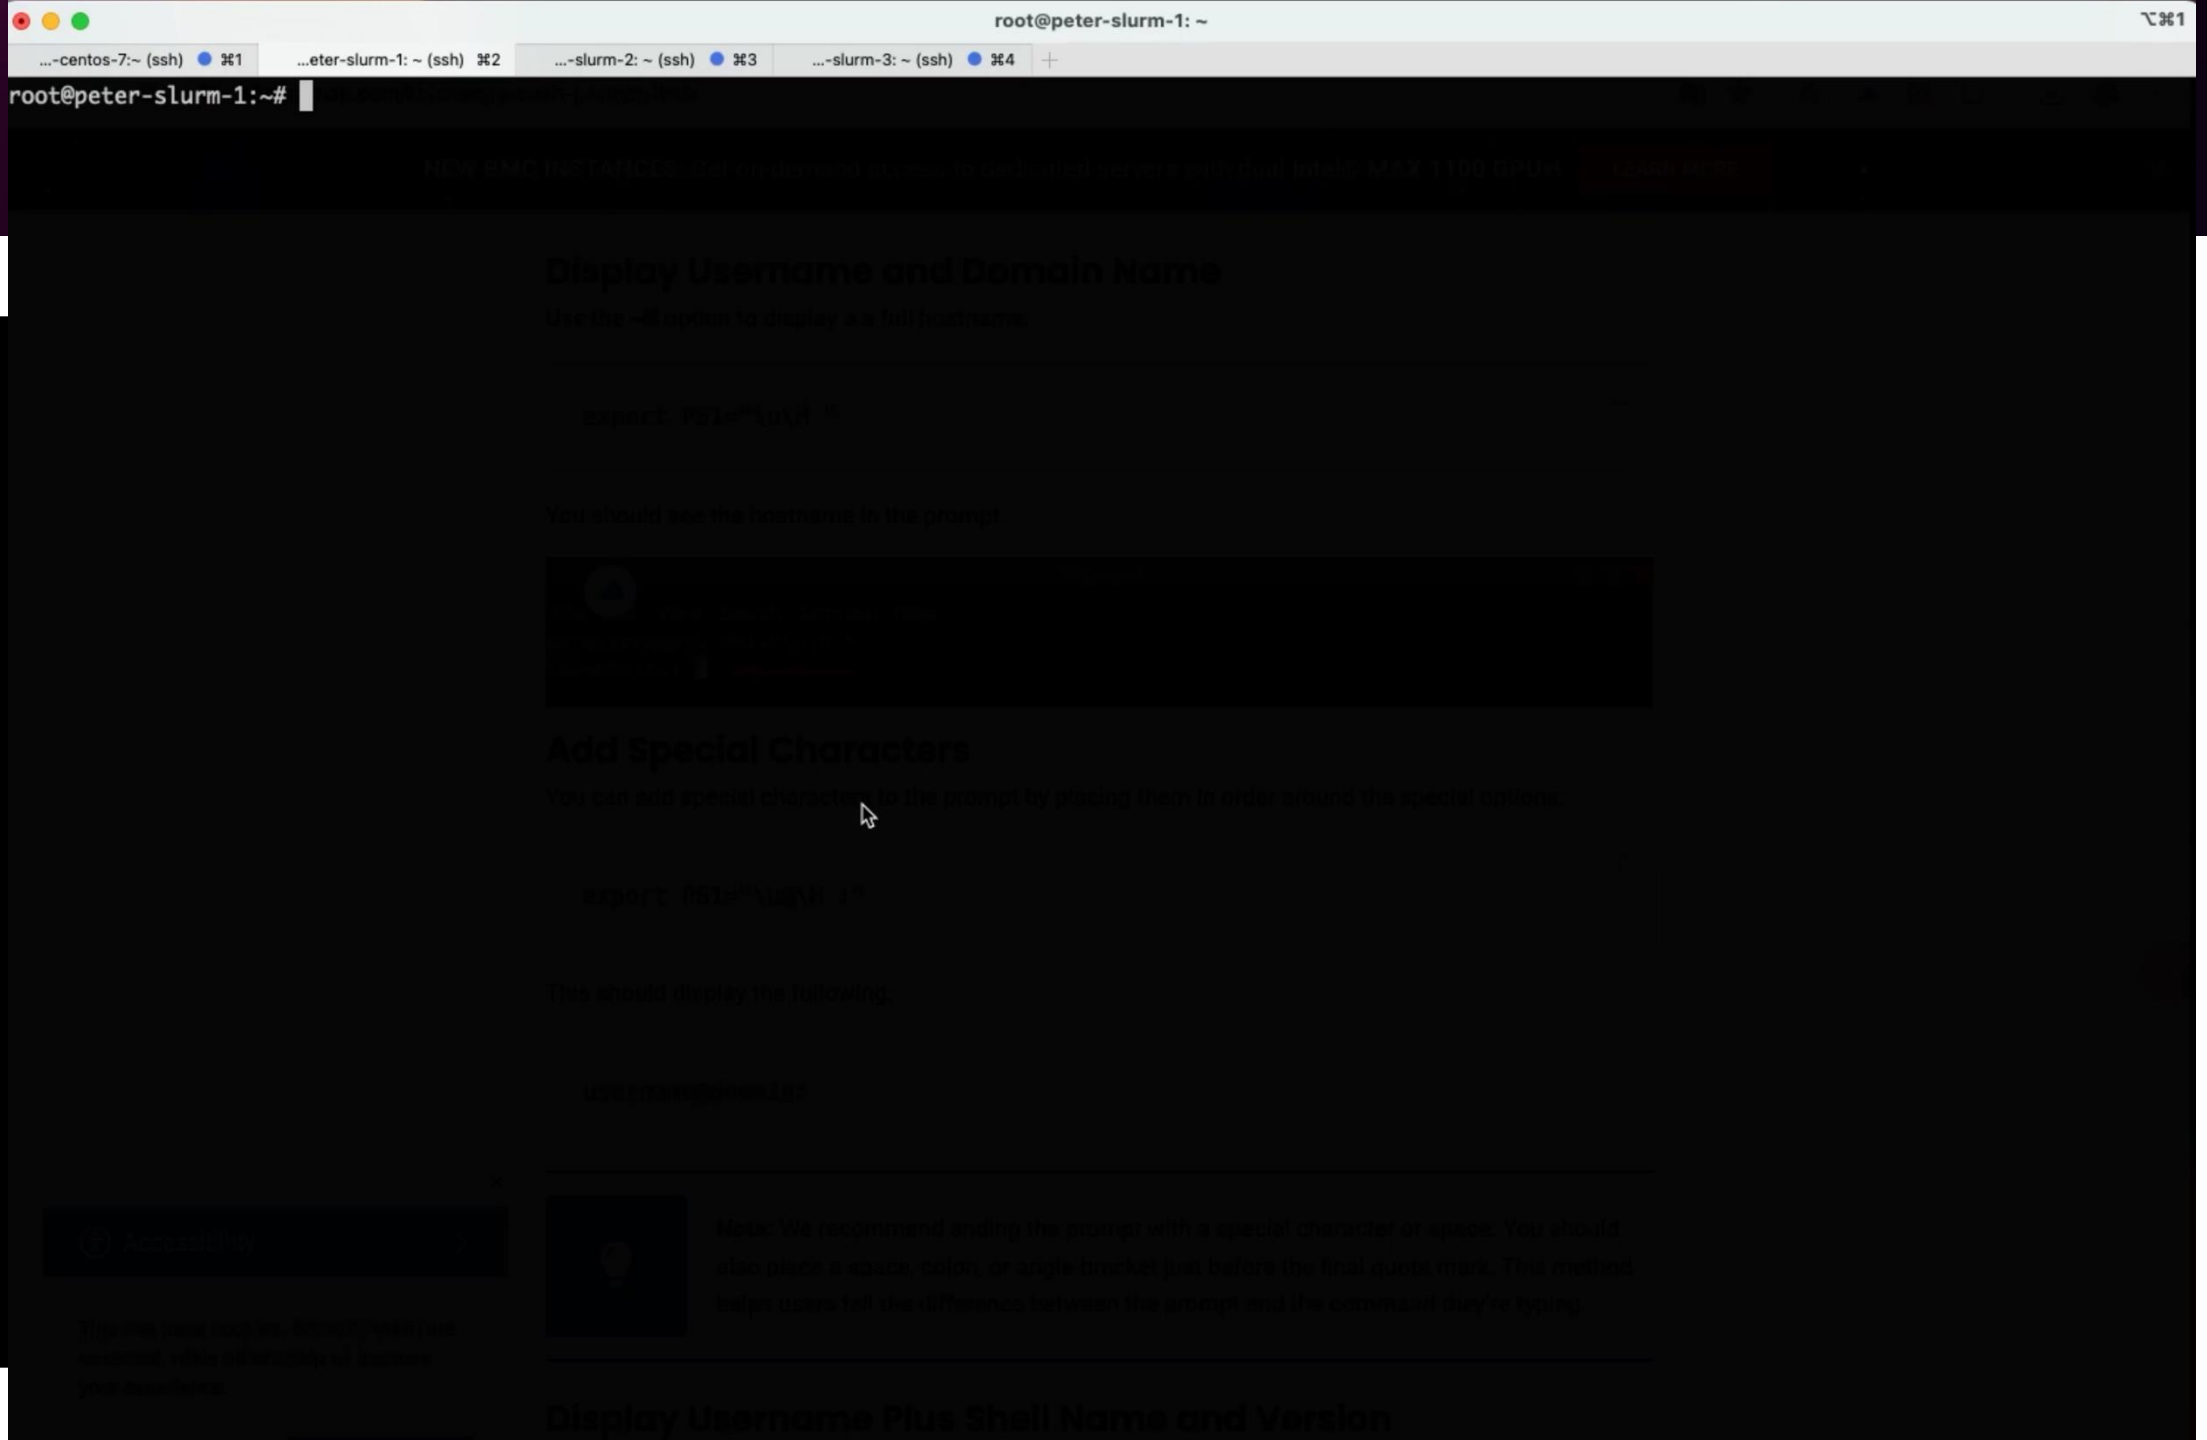
root@peter-slurm-1: ~
⌥⌘1
...-centos-7:~ (ssh) ⌘1
...eter-slurm-1: ~ (ssh) ⌘2
...-slurm-2: ~ (ssh) ⌘3
...-slurm-3: ~ (ssh) ⌘4
root@peter-slurm-1:~#

FAQ

Q1: Race Condition?

A1: good catch. So the best way out is either to do this in each's scheduler level, or control the entrance in a central manner . This is in future consideration.

Q2: place-holder shows node-level grain size , how about GPU-level(in one node)?

A2: Both slurm GRES-plugin and K8s device-plugin schedule those GPU on the same node in order, so it still works in GPU-in-node grant size.

Thank you

We hope **TACC**'s practice can inspire the community and industry, about how to streamline HPC and K8S workload together, and build a scientist friendly toolset and platform, with CNCF stacks.

香港科技大學
THE HONG KONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

DaoCloud