



**KubeCon**



**CloudNativeCon**

THE LINUX FOUNDATION



**AI\_dev**  
Open Source GenAI & ML Summit

---

**China 2024**

---



KubeCon



CloudNativeCon

THE LINUX FOUNDATION



China 2024

# Build Container Runtime Based on Sandbox API of Containerd

Cai Wei, DaoCloud  
Shaobao Feng, HuaweiCloud



KubeCon



CloudNativeCon



China 2024

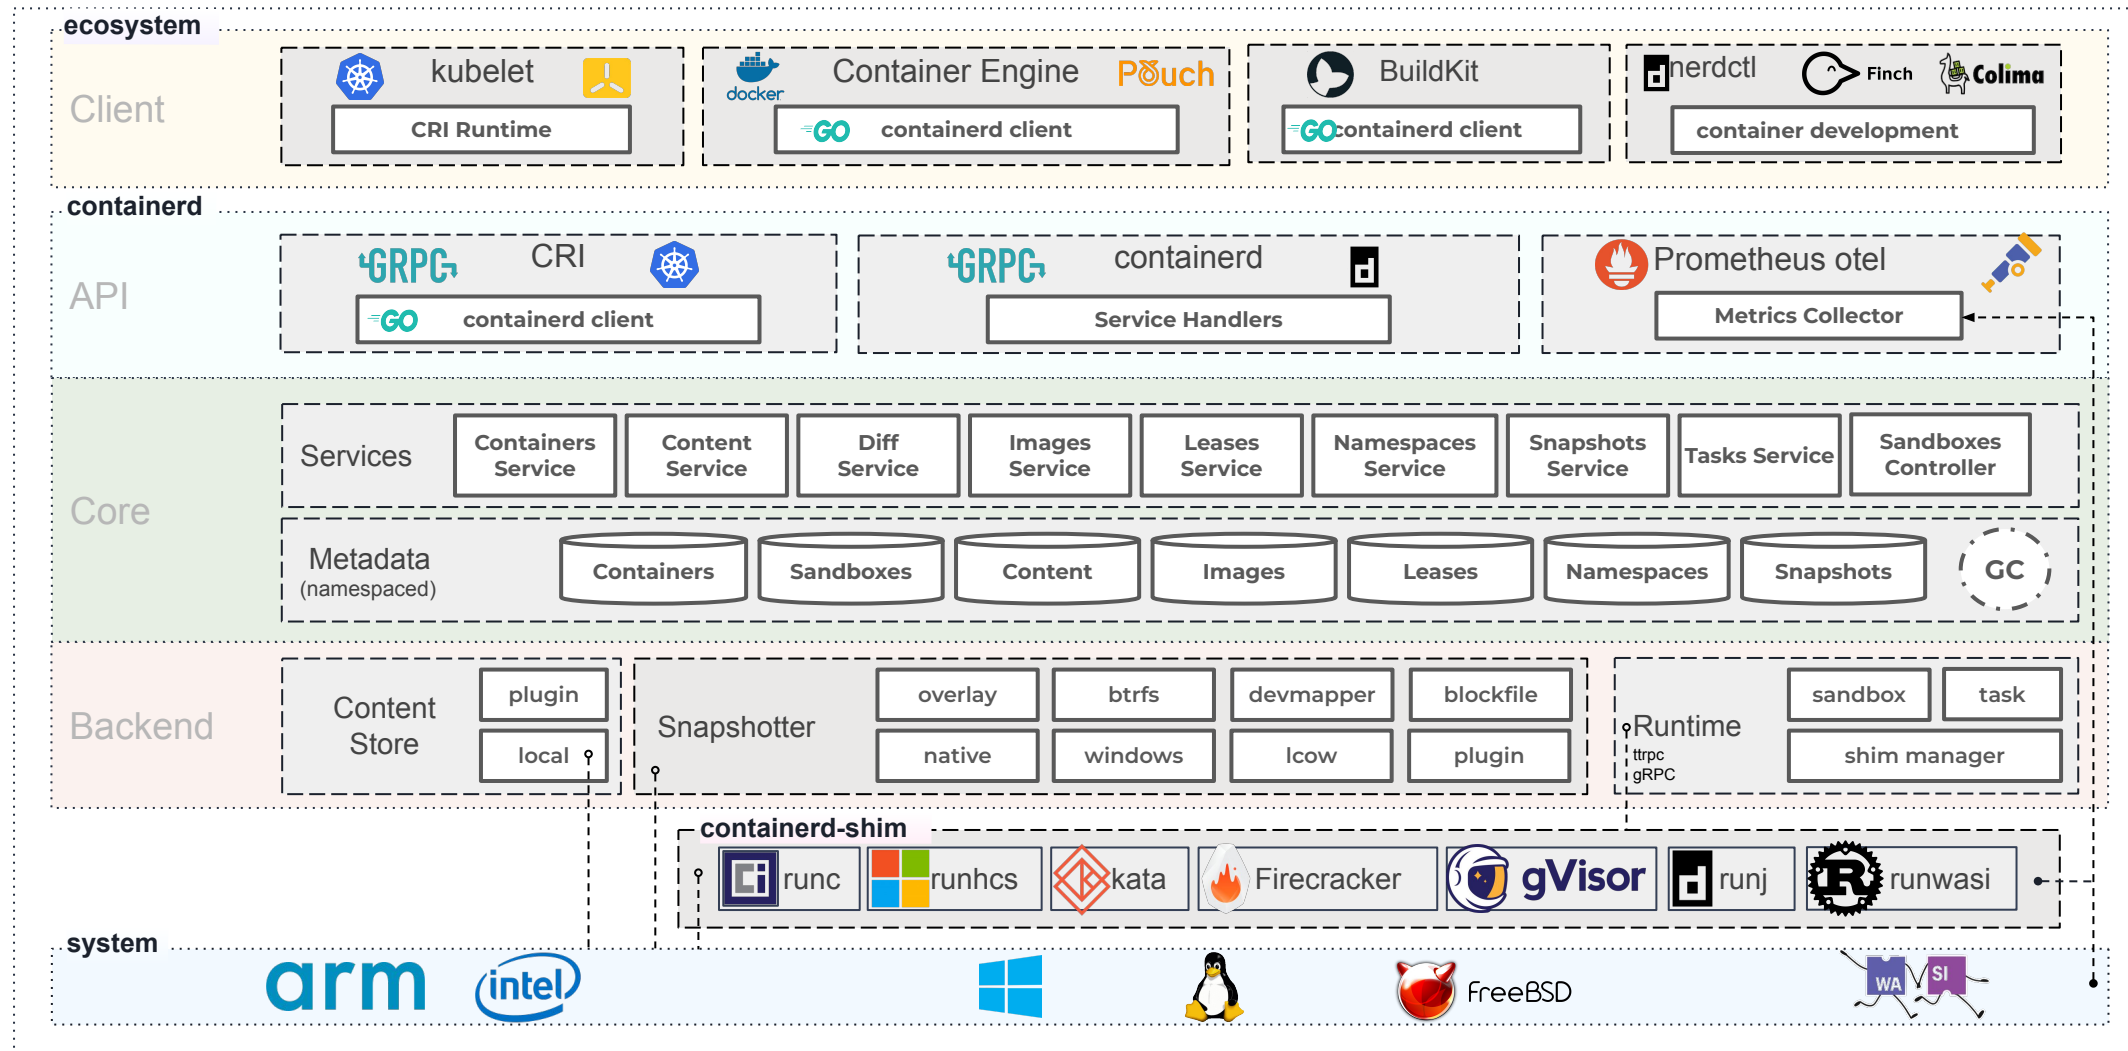


# Build **Container** Runtime based on **Sandbox** API of Containerd

# Contained 2.0



China 2024

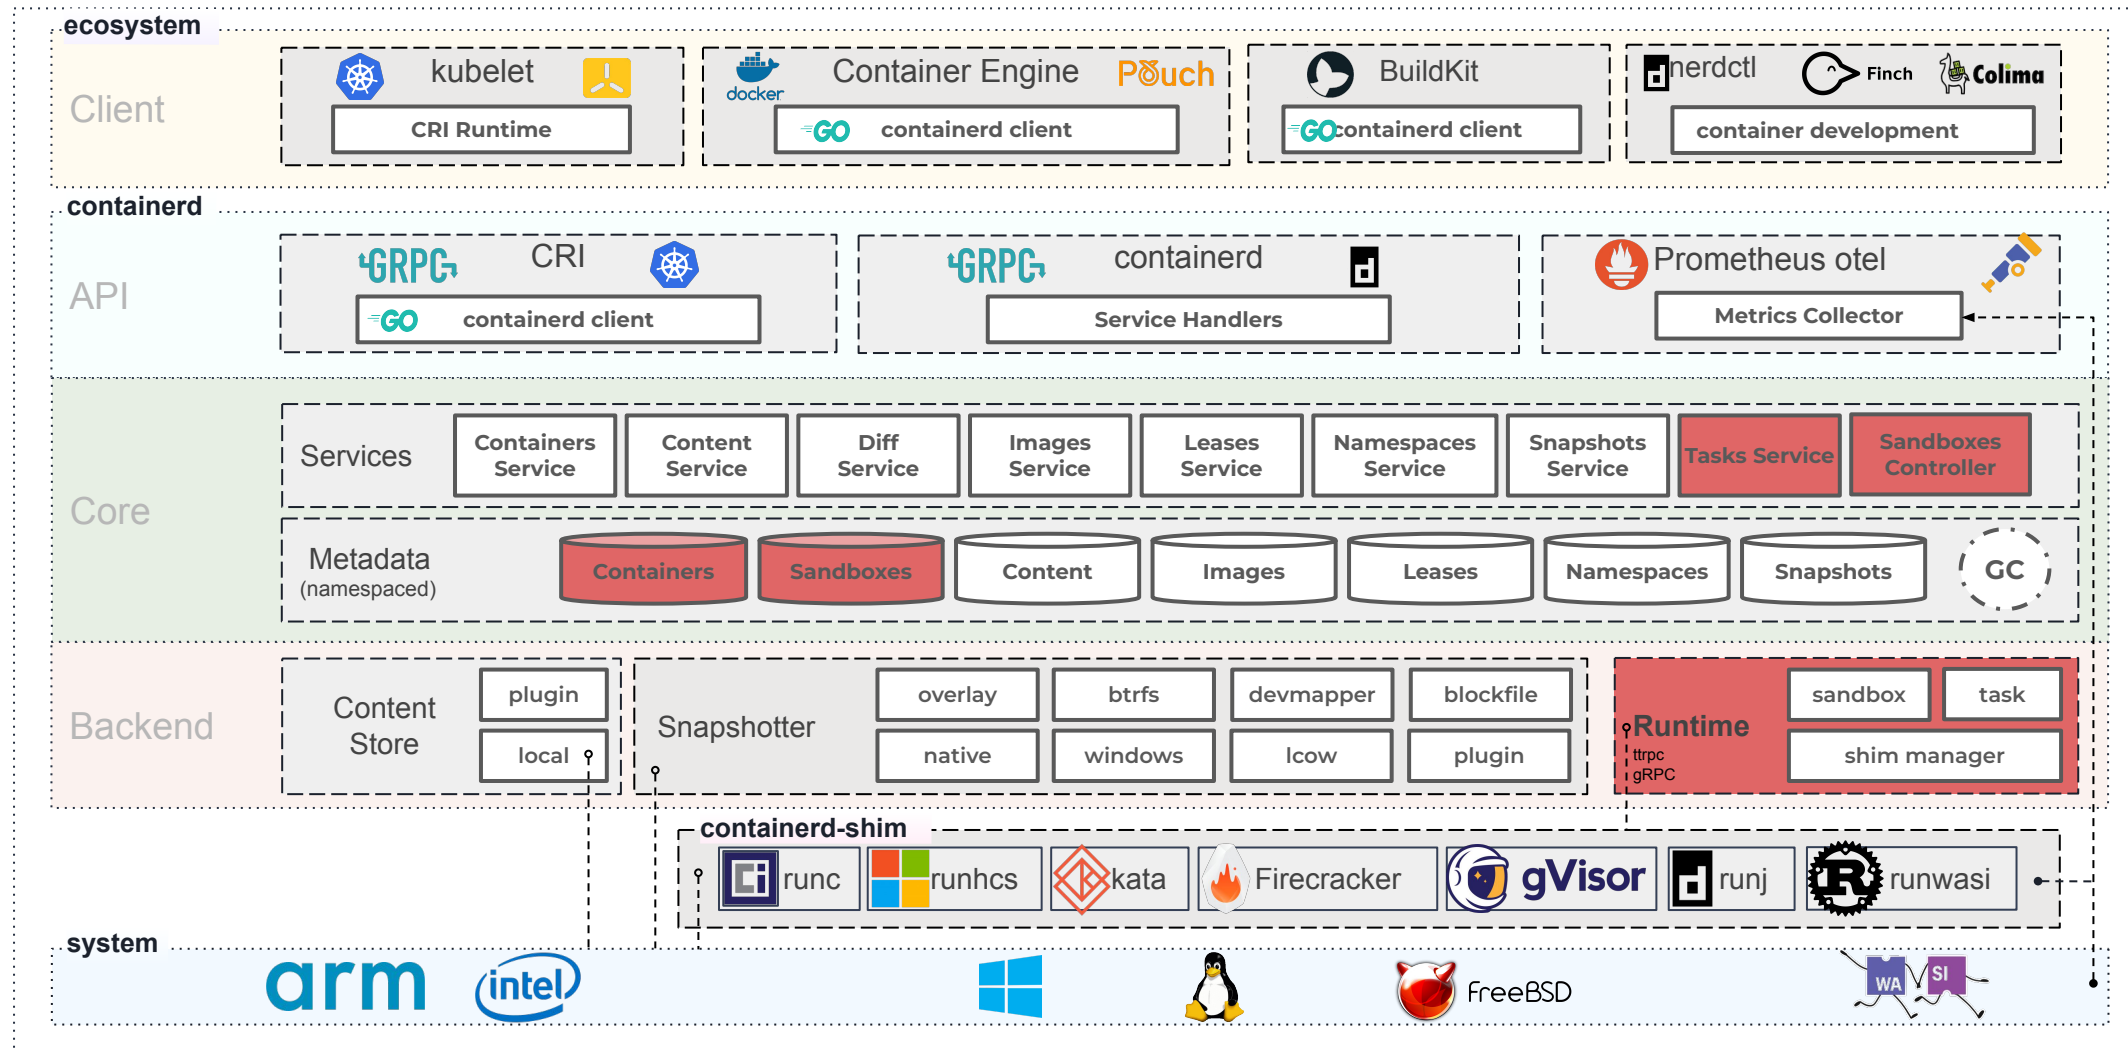




# Sandbox & Container



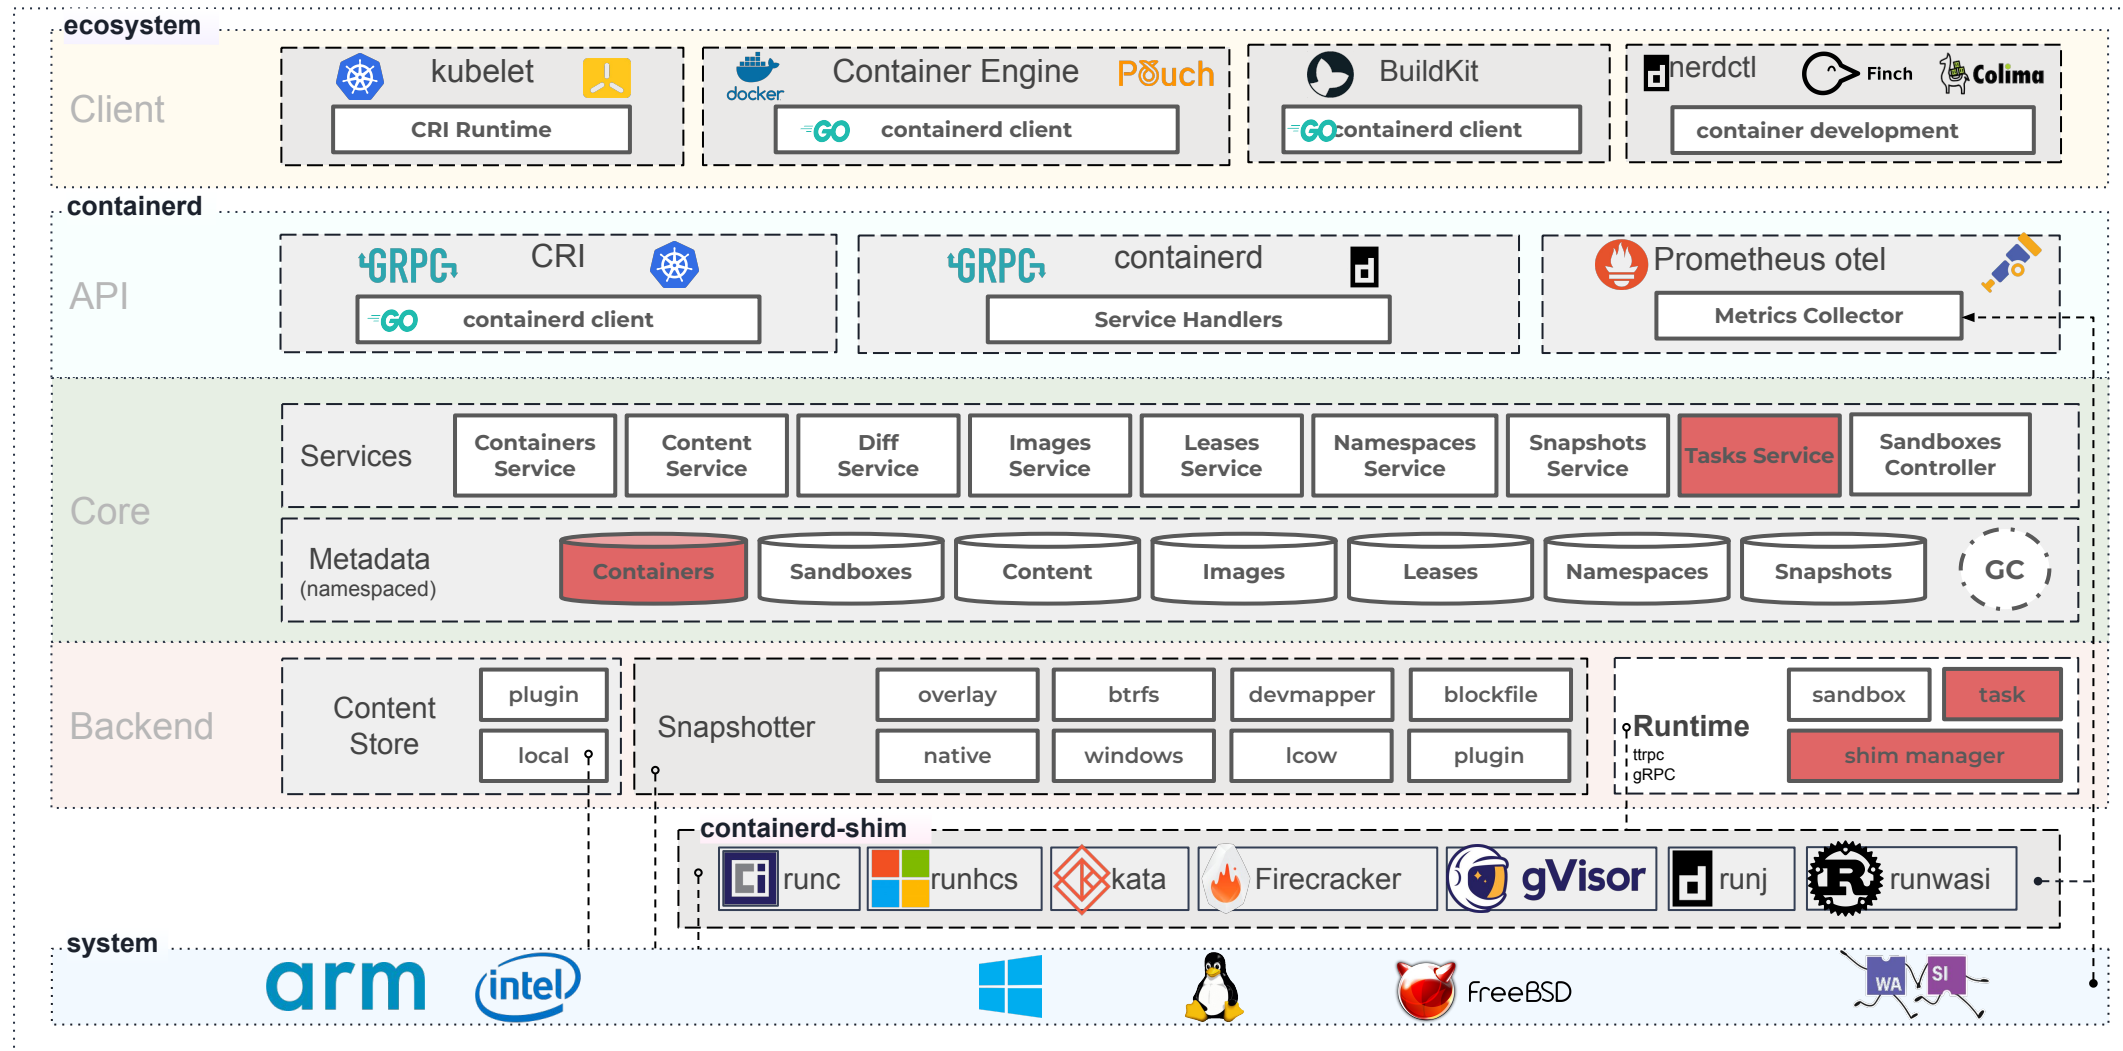
China 2024



# Container in Contained



China 2024

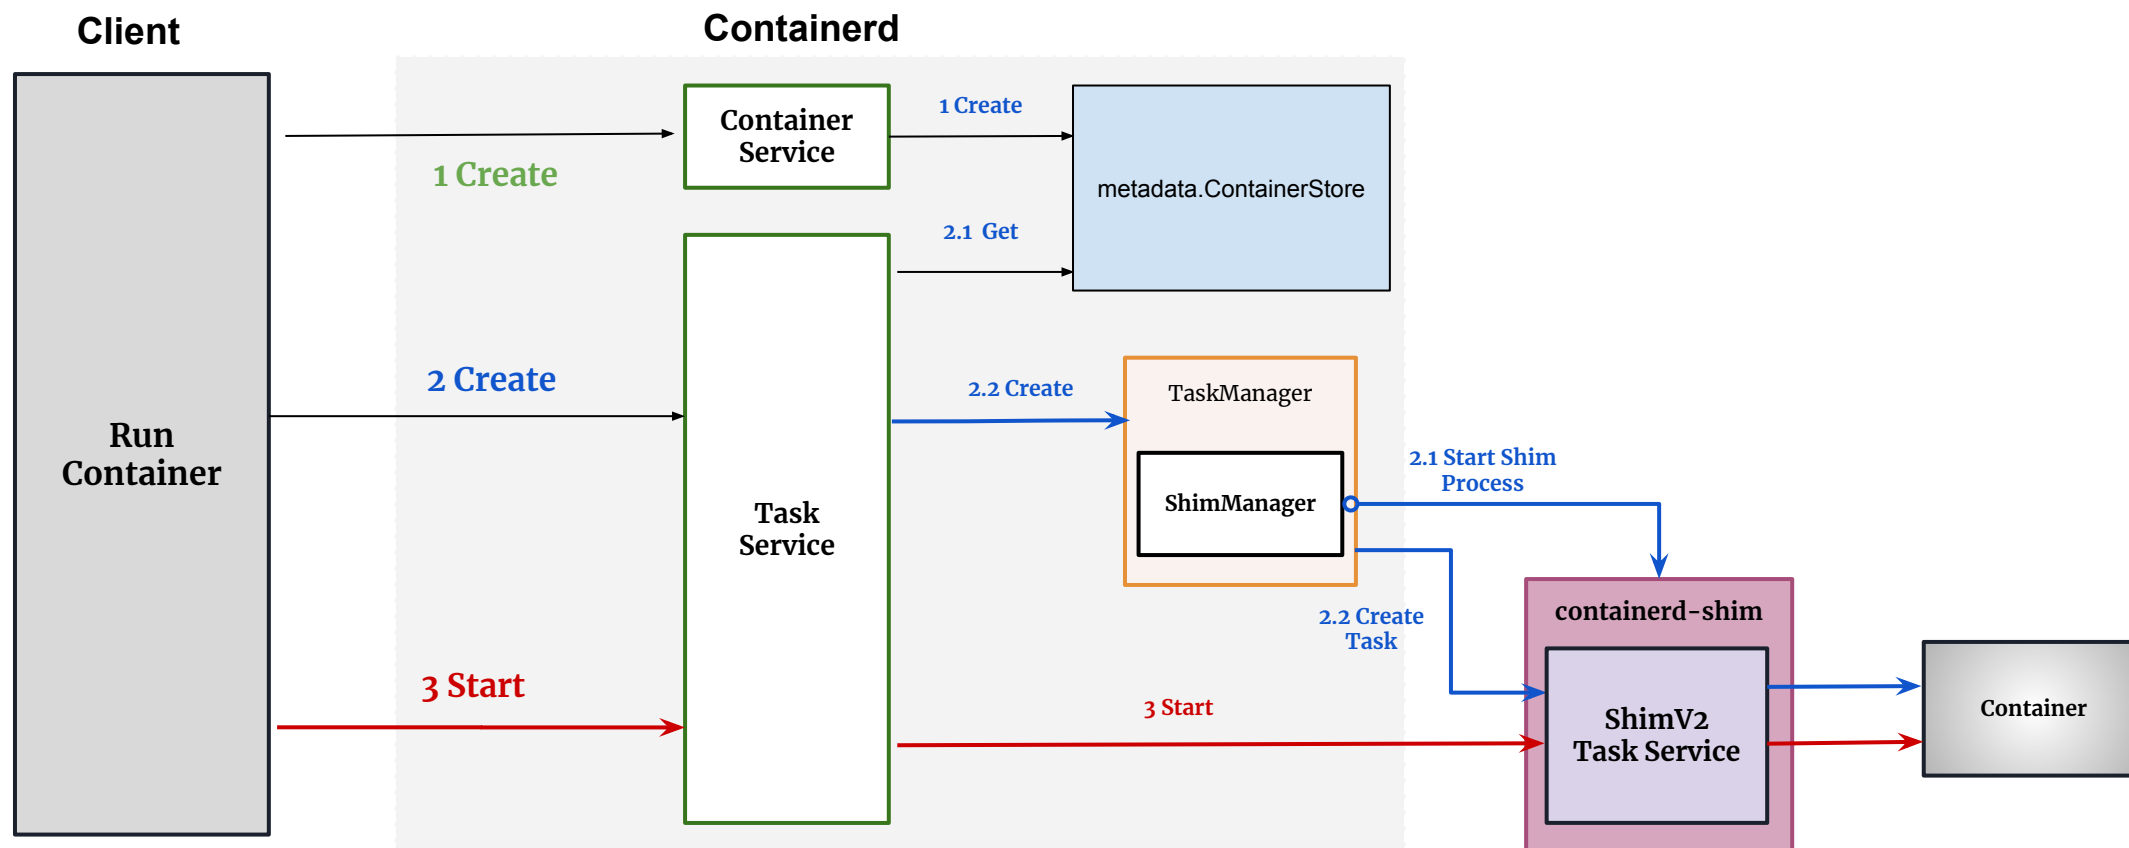


# Container & Task



China 2024

**Run a Container = Create Container + Create Task + Start Task**



# Task & Shim



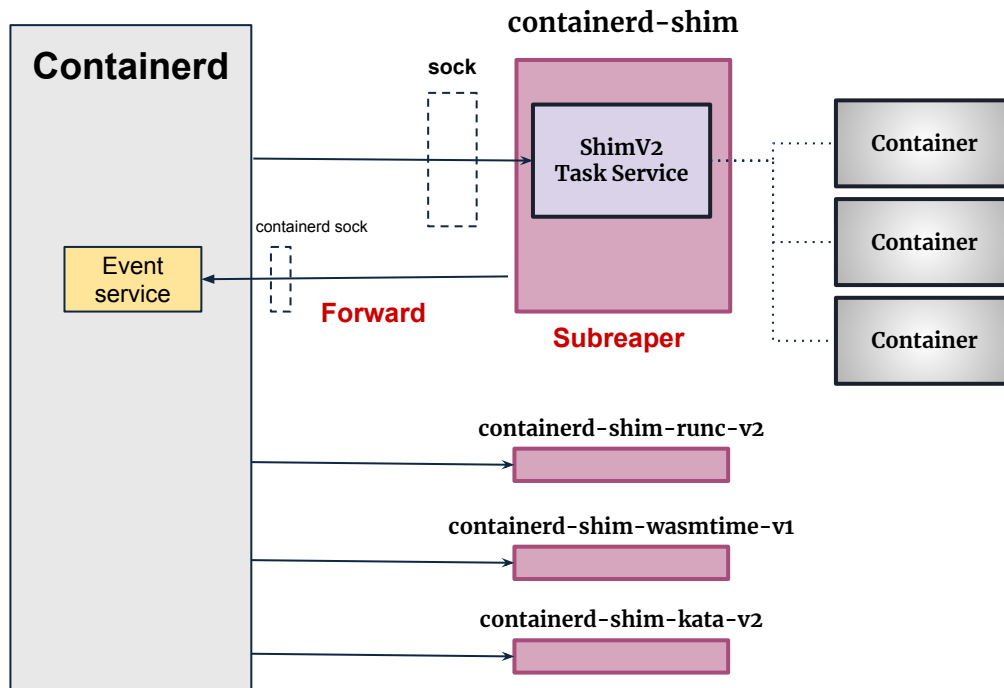
China 2024

```

→ pstree
systemd--agetty
        |--containerd--11*[{containerd}]
        |--containerd-shim--dumb-init--node--node--2*[esbuild--7*[{esbuild}]]
        |                                     |
        |                                     |--11*[{node}]
        |                                     |--10*[{node}]
        |
        |--11*[{containerd-shim}]
        |--containerd-shim--frpc--6*[{frpc}]
        |                                     |
        |                                     |--11*[{containerd-shim}]
        |--containerd-shim--node--10*[{node}]
        |                                     |
        |                                     |--12*[{containerd-shim}]
        |--containerd-shim--vaultwarden--12*[{vaultwarden}]
        |                                     |
        |                                     |--11*[{containerd-shim}]
    
```

```

service Task {
  rpc State(StateRequest) returns (StateResponse);
  rpc Create(CreateTaskRequest) returns (CreateTaskResponse);
  rpc Start(StartRequest) returns (StartResponse);
  rpc Delete(DeleteRequest) returns (DeleteResponse);
  rpc Pids(PidsRequest) returns (PidsResponse);
  rpc Pause(PauseRequest) returns (google.protobuf.Empty);
  rpc Resume(ResumeRequest) returns (google.protobuf.Empty);
  rpc Checkpoint(CheckpointTaskRequest) returns
    (google.protobuf.Empty);
  rpc Kill(KillRequest) returns (google.protobuf.Empty);
  rpc Exec(ExecProcessRequest) returns (google.protobuf.Empty);
  rpc ResizePty(ResizePtyRequest) returns (google.protobuf.Empty);
  rpc CloseIO(CloseIORequest) returns (google.protobuf.Empty);
  rpc Update(UpdateTaskRequest) returns (google.protobuf.Empty);
  rpc Wait(WaitRequest) returns (WaitResponse);
  rpc Stats(StatsRequest) returns (StatsResponse);
  rpc Connect(ConnectRequest) returns (ConnectResponse);
  rpc Shutdown(ShutdownRequest) returns (google.protobuf.Empty);
}
    
```



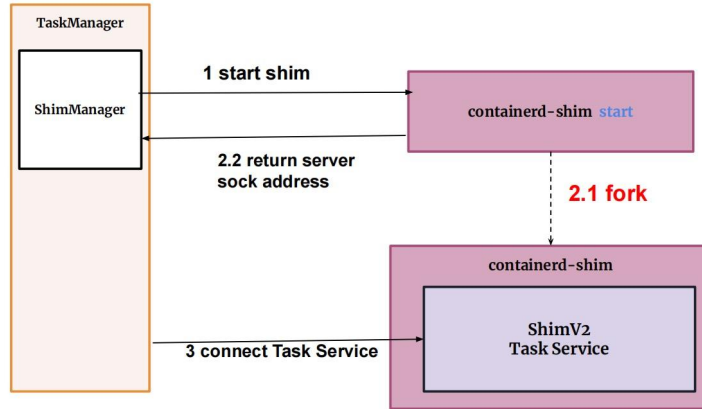


# Shim & Container Group

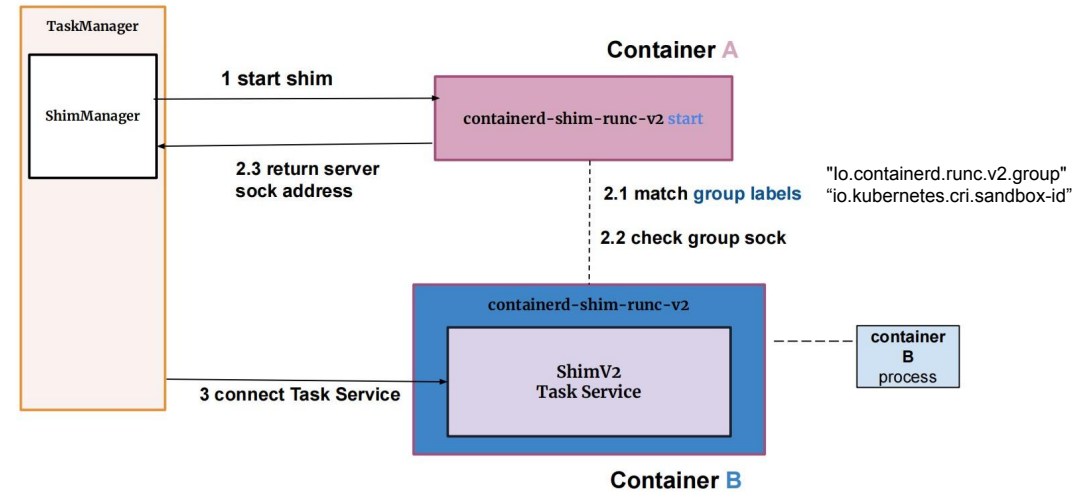


China 2024

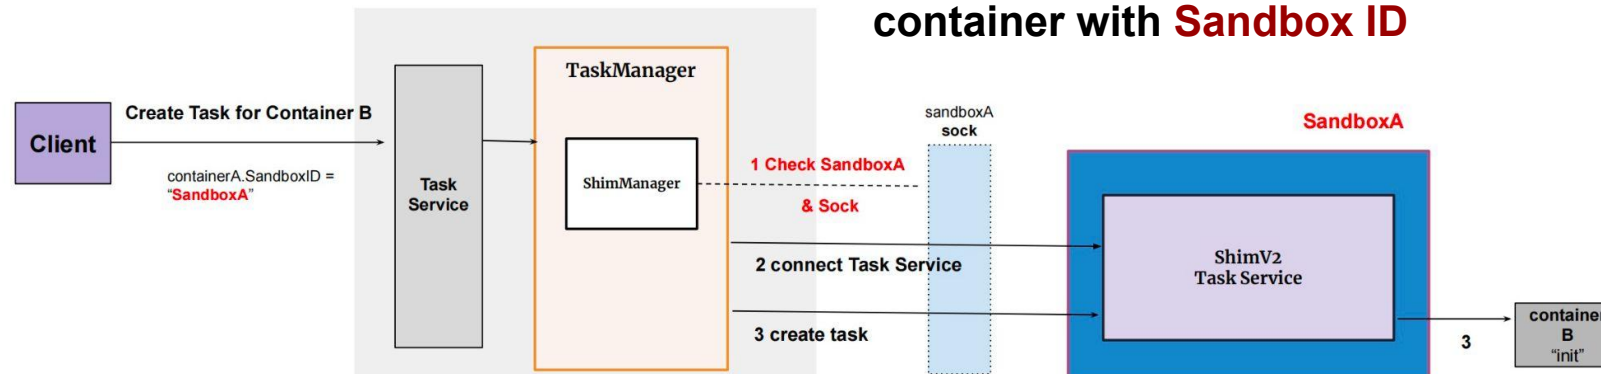
## Shim Start



## container with **group labels**



## Containerd

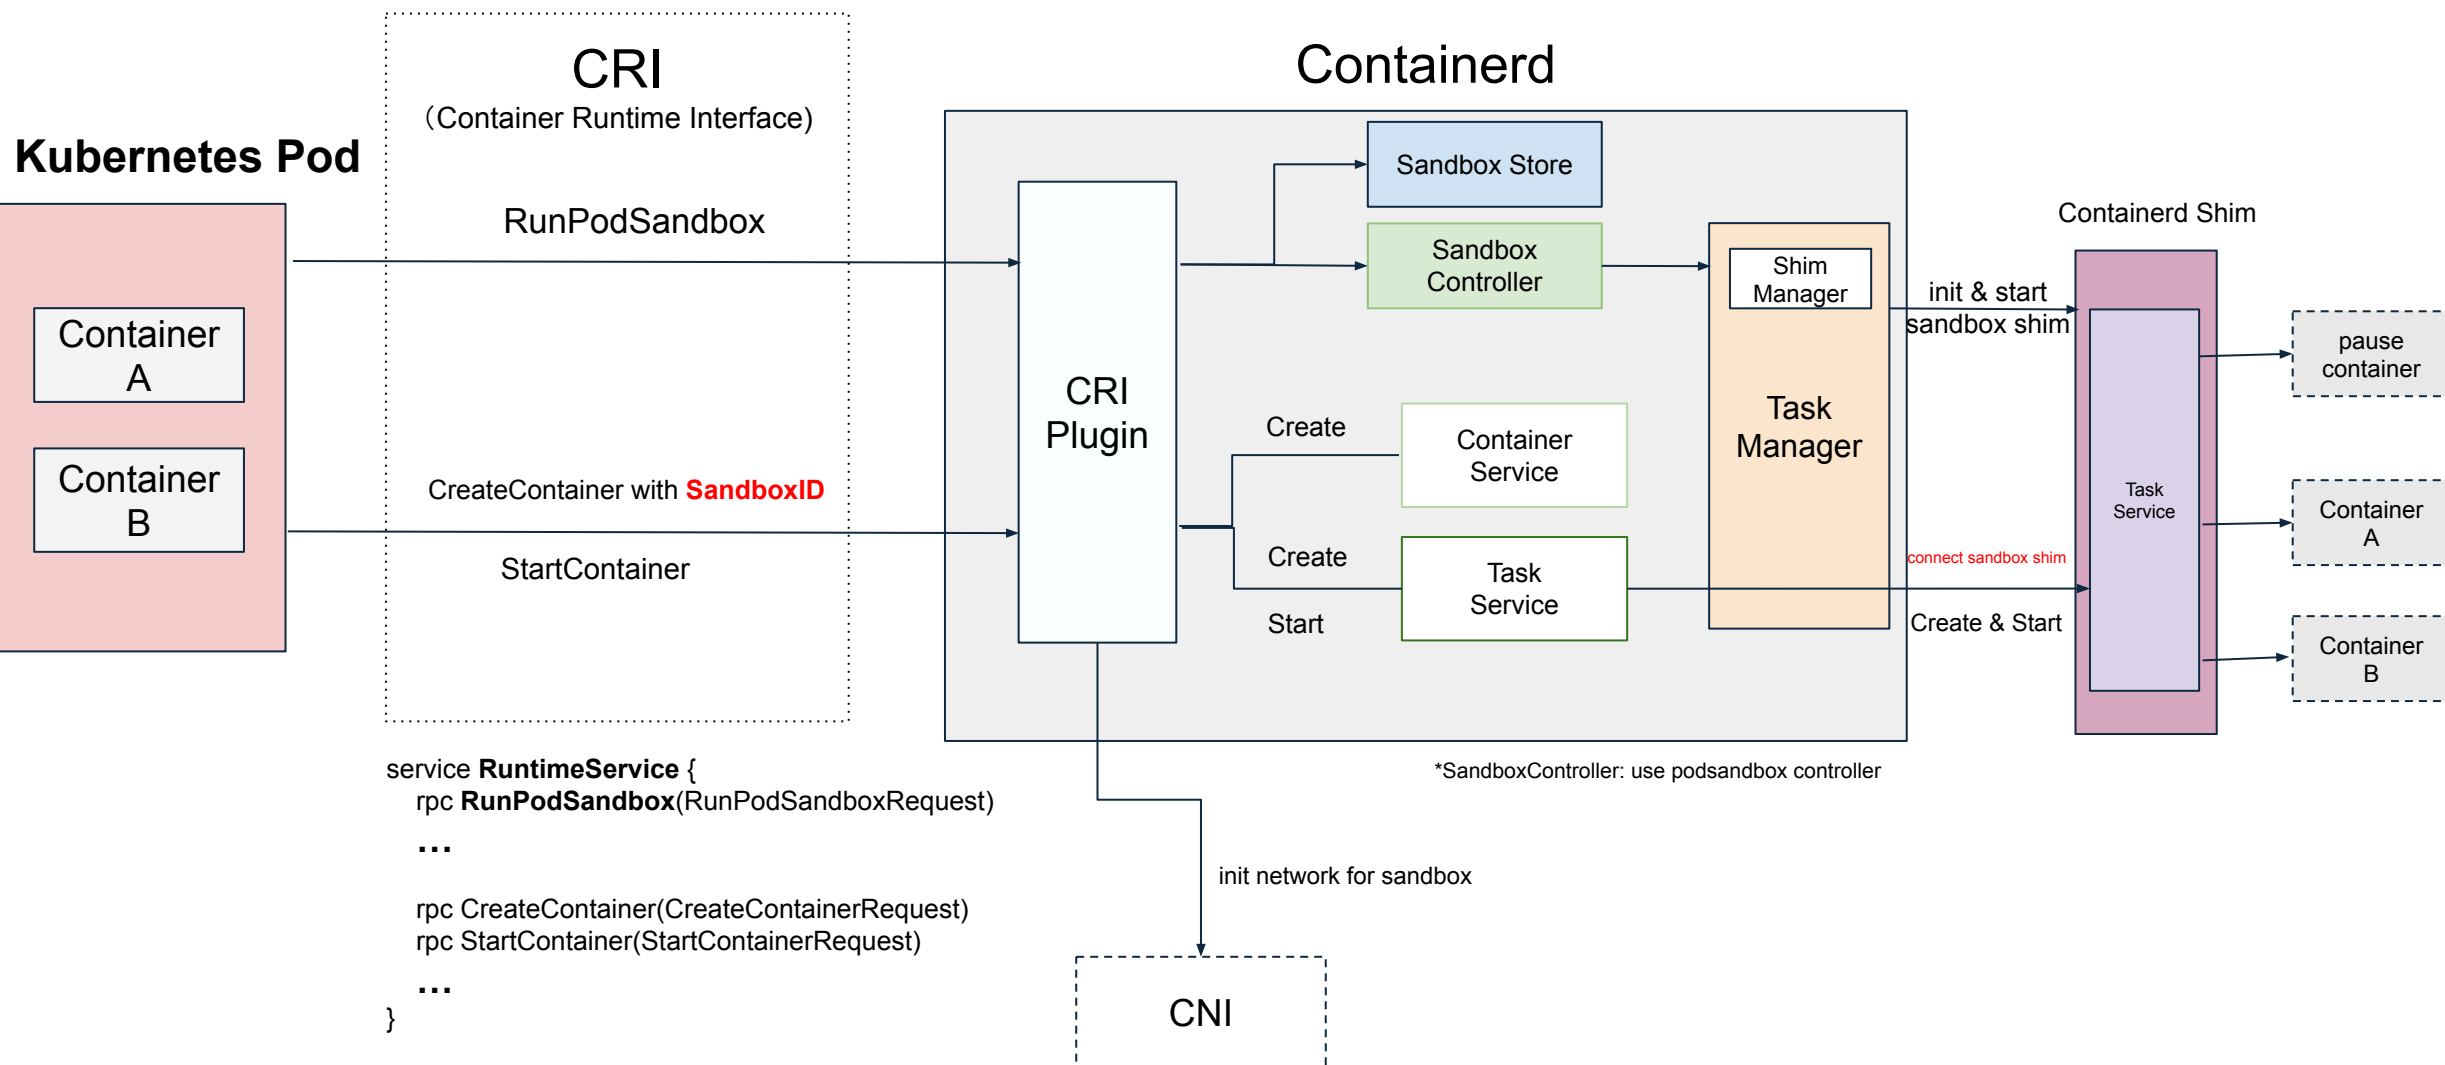


## container with **Sandbox ID**

# Container Group & Sandbox



China 2024



# Sandbox & Sandbox API



China 2024

## CRI Service Plugin



CreateContainer

Create

SandboxStore Service

metadata.SandboxStore

Create & Start Sandbox

SandboxController Plugin

"podsandbox"

"shim"

proxy plugin

Container Service

1 Create

2 Create

3 Start

Task Service

TaskManager

ShimManager

2.1 Start Shim Process

2.2 Create Task

containerd-shim

ShimV2 Task Service

Pause Container

Container Process

2 Create

2 Start

2 Create

3 Start

2 Create

3 Start

2.1 connect task service

2.2 Create Task

2.2 init

3 start

# Sandbox API



China 2024

## Sandbox Controller

<https://github.com/containerd/containerd/blob/main/core/sandbox/controller.go>

```
// Controller is an interface to manage sandboxes at runtime.
// When running in sandbox mode, shim expected to implement 'SandboxService'.
// Shim lifetimes are now managed manually via sandbox API by the containerd's client.
type Controller interface {
    // Create is used to initialize sandbox environment. (mounts, any)
    Create(ctx context.Context, sandboxInfo Sandbox, opts ...CreateOpt) error
    // Start will start previously created sandbox.
    Start(ctx context.Context, sandboxID string) (ControllerInstance, error)
    // Platform returns target sandbox OS that will be used by Controller.
    // containerd will rely on this to generate proper OCI spec.
    Platform(ctx context.Context, _sandboxID string) (imagespec.Platform, error)
    // Stop will stop sandbox instance
    Stop(ctx context.Context, sandboxID string, opts ...StopOpt) error
    // Wait blocks until sandbox process exits.
    Wait(ctx context.Context, sandboxID string) (ExitStatus, error)
    // Status will query sandbox process status. It is heavier than Ping call and must be used whenever you need to
    // gather metadata about current sandbox state (status, uptime, resource use, etc).
    Status(ctx context.Context, sandboxID string, verbose bool) (ControllerStatus, error)
    // Shutdown deletes and cleans all tasks and sandbox instance.
    Shutdown(ctx context.Context, sandboxID string) error
    // Metrics queries the sandbox for metrics.
    Metrics(ctx context.Context, sandboxID string) (*types.Metric, error)
    // Update changes a part of sandbox, such as extensions/annotations/labels/spec of
    // Sandbox object, controllers may have to update the running sandbox according to the changes.
    Update(ctx context.Context, sandboxID string, sandbox Sandbox, fields ...string) error
}
```

## Runtime Sandbox Service

```
service Sandbox {
    // CreateSandbox will be called right after sandbox shim instance launched.
    // It is a good place to initialize sandbox environment.
    rpc CreateSandbox(CreateSandboxRequest) returns (CreateSandboxResponse);

    // StartSandbox will start a previously created sandbox.
    rpc StartSandbox(StartSandboxRequest) returns (StartSandboxResponse);

    // Platform queries the platform the sandbox is going to run containers on.
    // containerd will use this to generate a proper OCI spec.
    rpc Platform(PlatformRequest) returns (PlatformResponse);

    // StopSandbox will stop existing sandbox instance
    rpc StopSandbox(StopSandboxRequest) returns (StopSandboxResponse);

    // WaitSandbox blocks until sandbox exits.
    rpc WaitSandbox(WaitSandboxRequest) returns (WaitSandboxResponse);

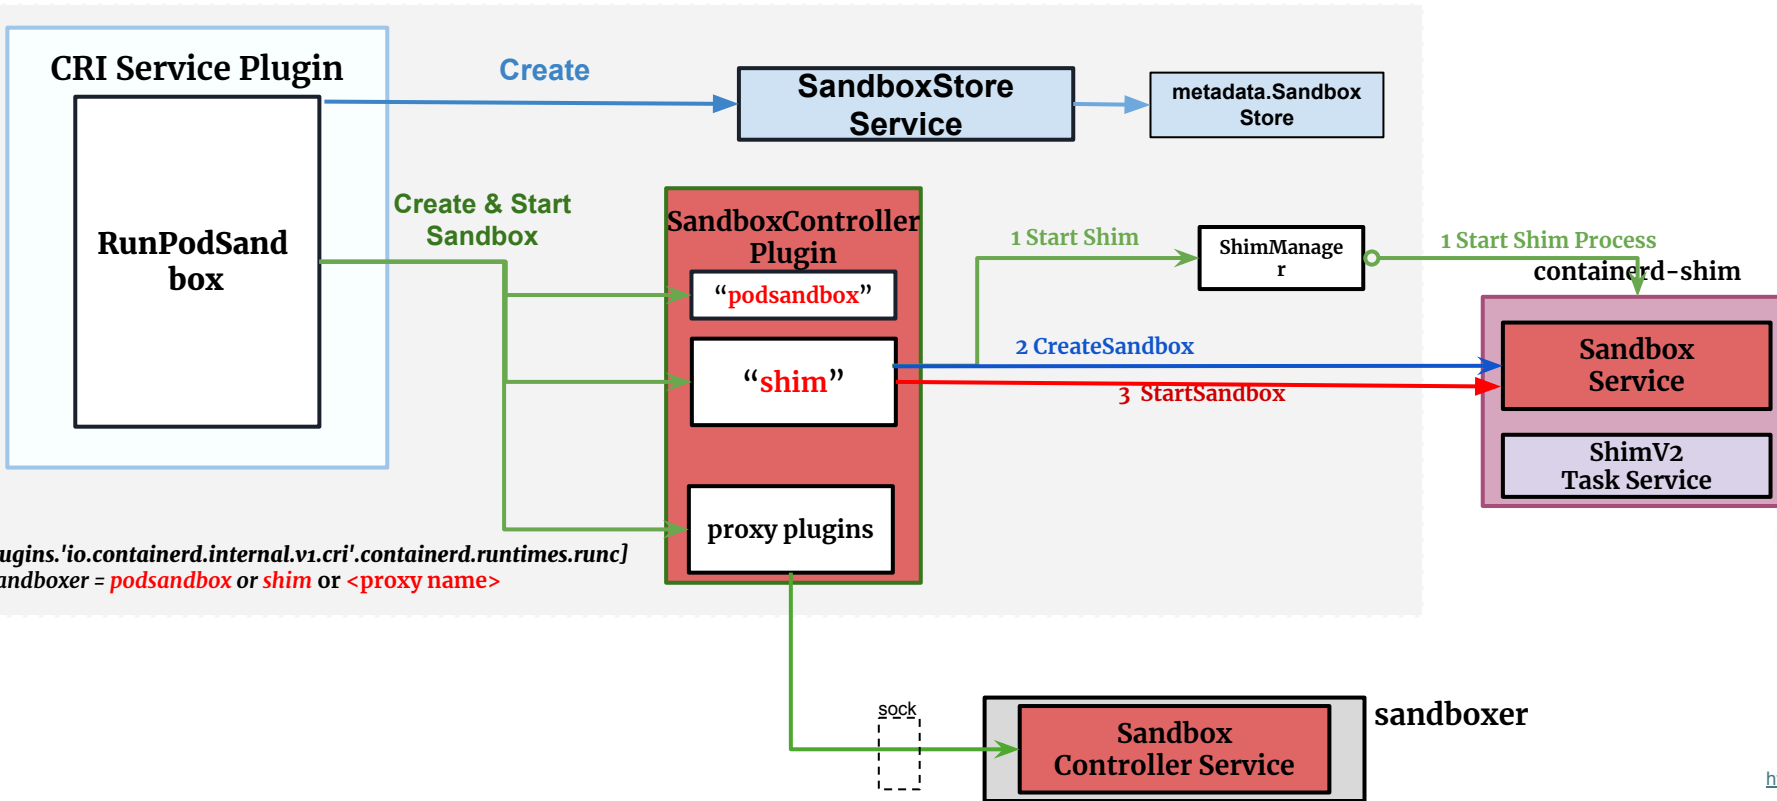
    // SandboxStatus will return current status of the running sandbox instance
    rpc SandboxStatus(SandboxStatusRequest) returns (SandboxStatusResponse);

    // PingSandbox is a lightweight API call to check whether sandbox alive.
    rpc PingSandbox(PingRequest) returns (PingResponse);

    // ShutdownSandbox must shutdown shim instance.
    rpc ShutdownSandbox(ShutdownSandboxRequest) returns (ShutdownSandboxResponse);

    // SandboxMetrics retrieves metrics about a sandbox instance.
    rpc SandboxMetrics(SandboxMetricsRequest) returns (SandboxMetricsResponse);
}
```

<https://github.com/containerd/containerd/blob/main/api/runtime/sandbox/v1/sandbox.proto>

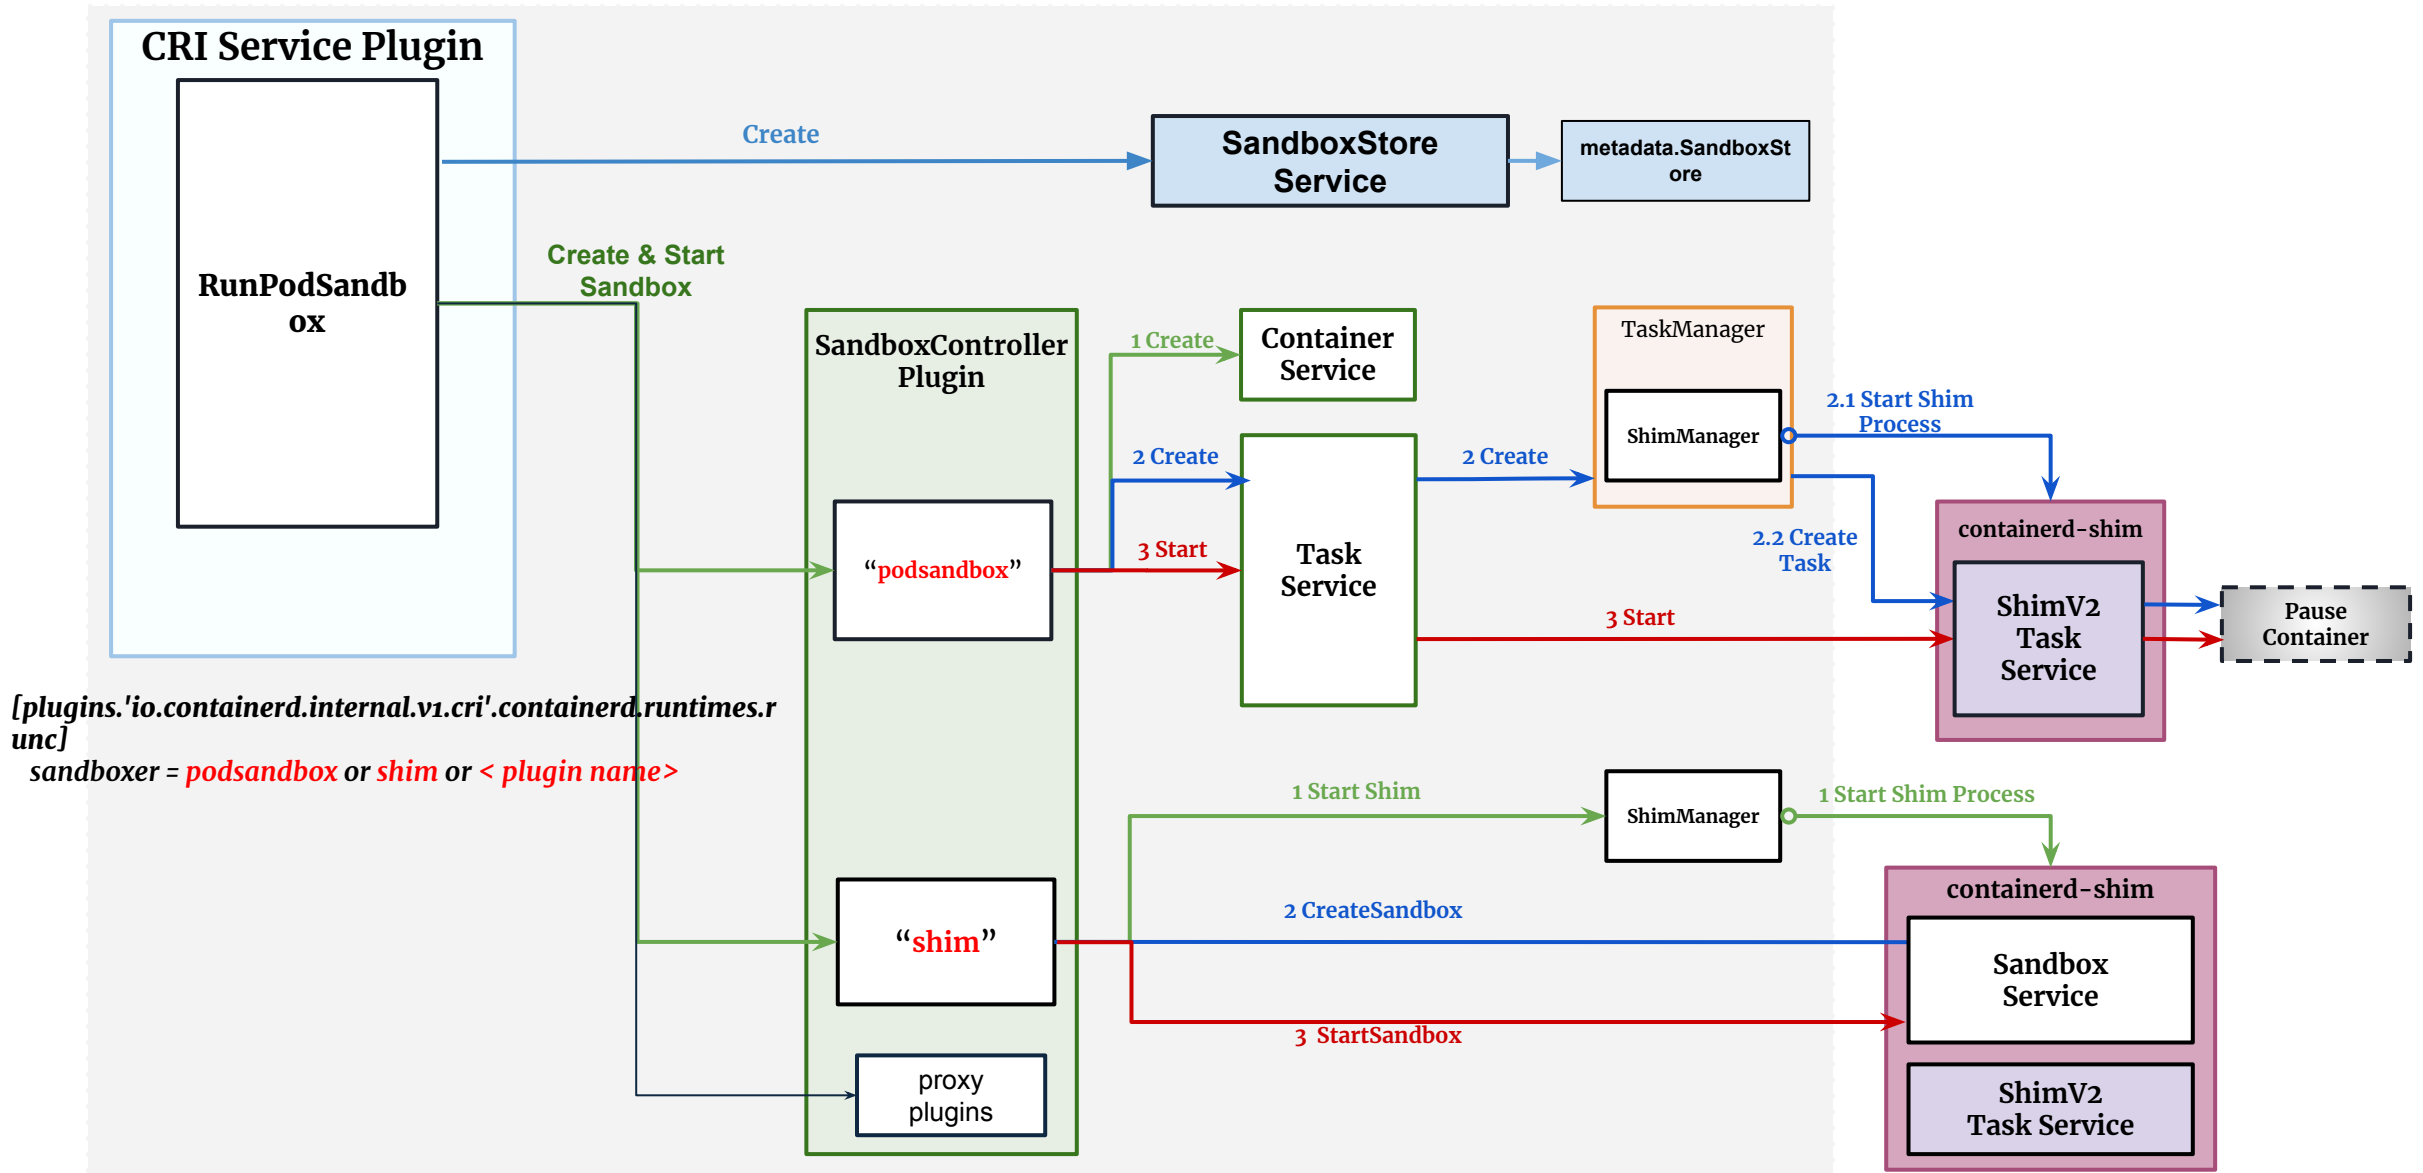


## Sandbox Controller Service

```
service Controller {
    rpc Create(ControllerCreateRequest) returns (ControllerCreateResponse);
    rpc Start(ControllerStartRequest) returns (ControllerStartResponse);
    rpc Platform(ControllerPlatformRequest) returns (ControllerPlatformResponse);
    rpc Stop(ControllerStopRequest) returns (ControllerStopResponse);
    rpc Wait(ControllerWaitRequest) returns (ControllerWaitResponse);
    rpc Status(ControllerStatusRequest) returns (ControllerStatusResponse);
    rpc Shutdown(ControllerShutdownRequest) returns (ControllerShutdownResponse);
    rpc Metrics(ControllerMetricsRequest) returns (ControllerMetricsResponse);
    rpc Update(ControllerUpdateRequest) returns (ControllerUpdateResponse);
}
```

<https://github.com/containerd/containerd/blob/main/api/services/sandbox/v1/sandbox.proto>

# Sandbox API & Shim



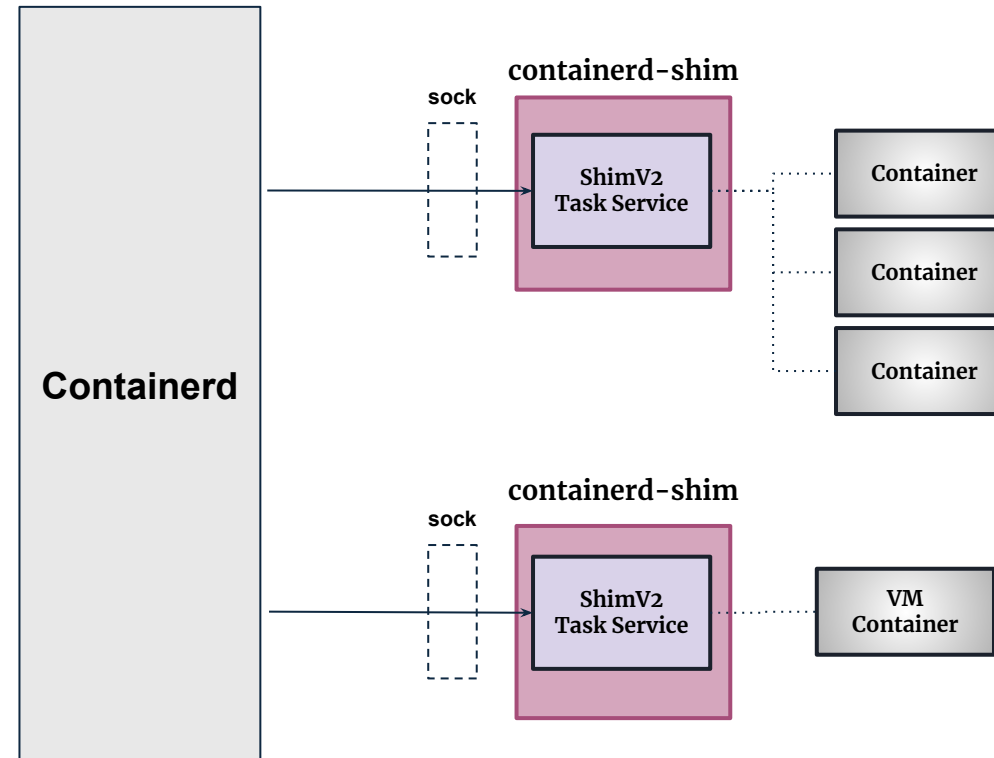


# Do we **have to** have Shim?



China 2024

```
↳ pstree
systemd--agetty
        |--containerd--11*[{containerd}]
        |--containerd-shim--dumb-init--node--node--2*[{esbuild--7*[{esbuild}]]
        |                                     |--11*[{node}]
        |                                     |--10*[{node}]
        |                                     |--11*[{containerd-shim}]
        |--containerd-shim--frpc--6*[{frpc}]
        |                                     |--11*[{containerd-shim}]
        |--containerd-shim--node--10*[{node}]
        |                                     |--12*[{containerd-shim}]
        |--containerd-shim--vaultwarden--12*[{vaultwarden}]
        |                                     |--11*[{containerd-shim}]
```

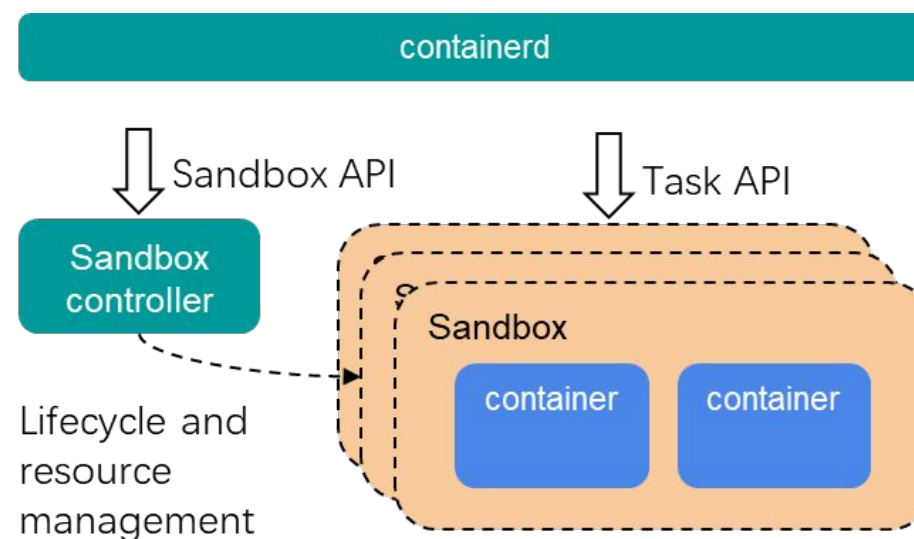


# Sandbox In Containerd



China 2024

- An isolated environment to run Tasks(containers) in it.
  - It can be a black box
  - It exposes **Task API** (through uds/vsock/tcp...).
  - It can be highly integrated(one process/thread per sandbox).
  - It can be managed by **Sandbox Controller API**.



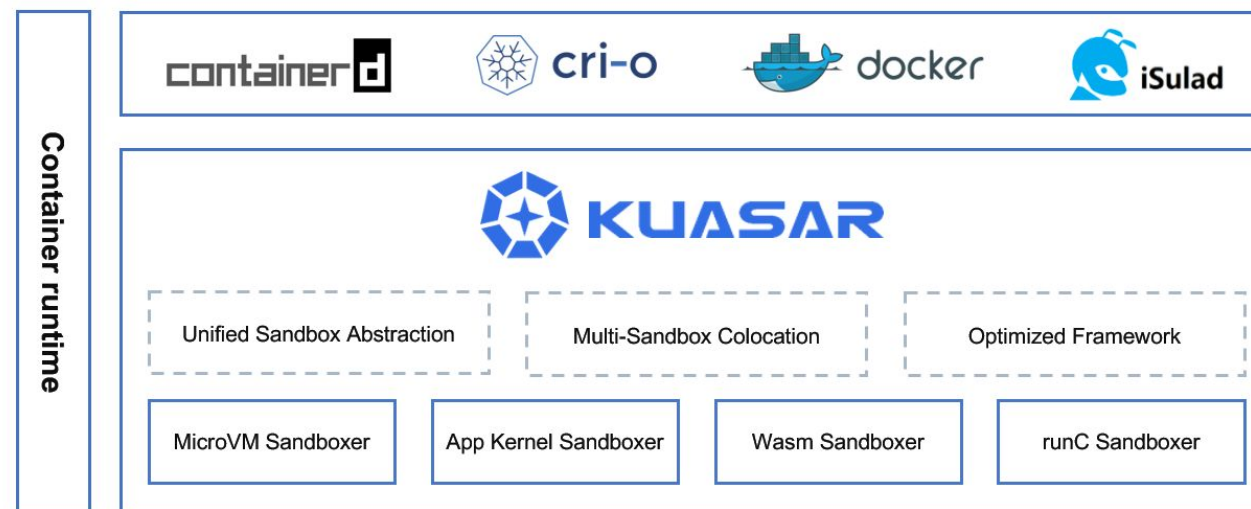
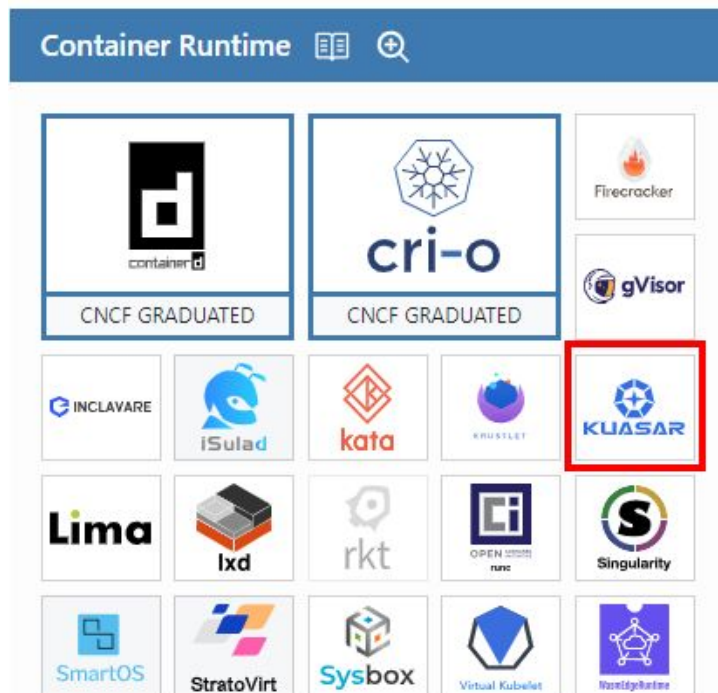
# Kuasar: A CNCF Sandbox Project



China 2024

A low-level container runtime that provides different kind of sandbox controllers(or sandboxers).

- A pure rust framework to provide sandboxer implementations to containerd
- A set of sandboxers based on the framework: MicroVM/App Kernel/WebAssembly/runC...

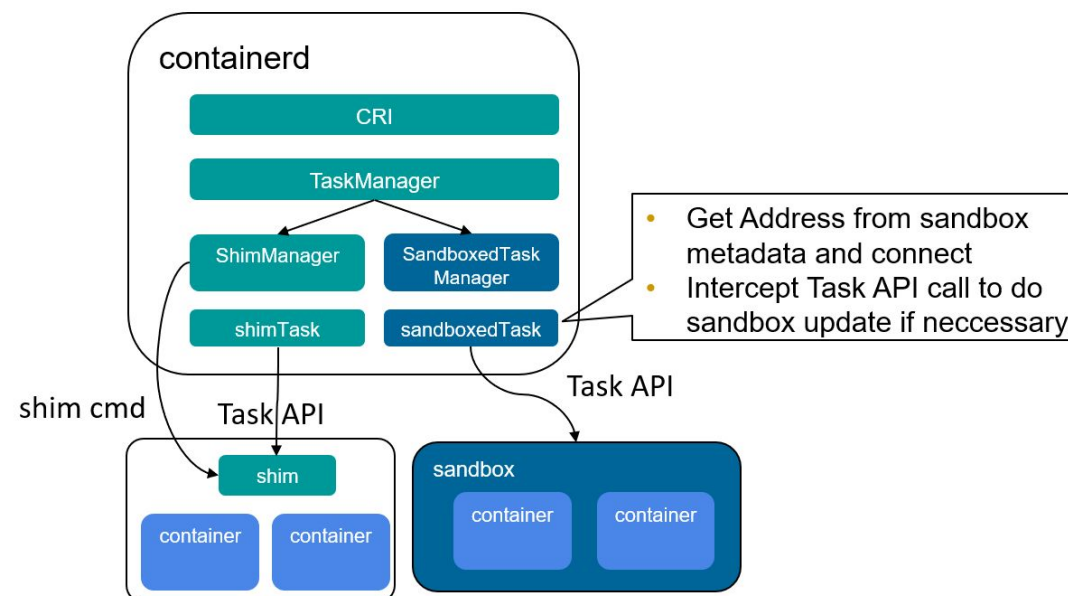


# Kuasar and Containerd

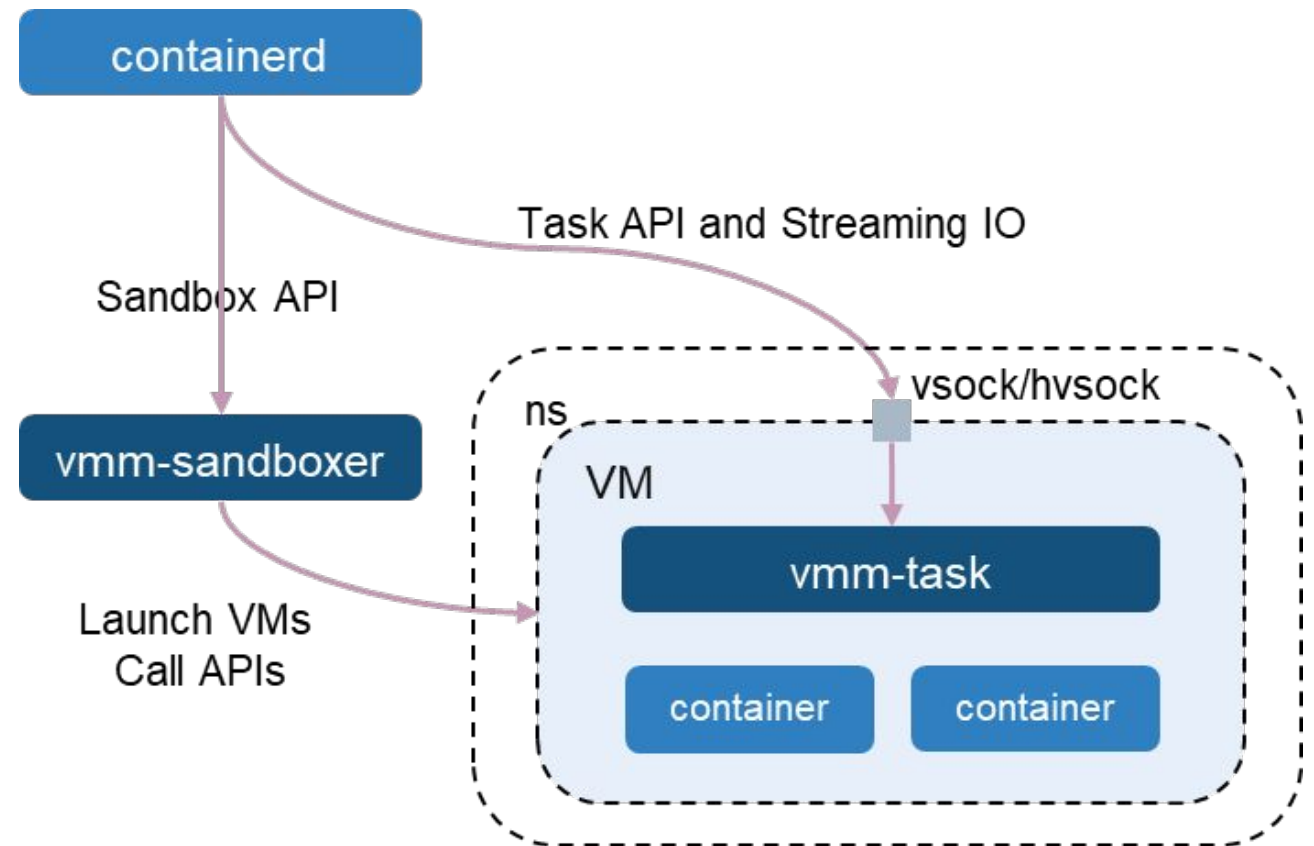


China 2024

- Kuasar is the first runtime implemented with proxy plugin of sandbox controller.
- 17 PRs submitted to containerd, to improve sandbox API and framework
  - Decouple PodSandbox(pause container) controller and cri plugin
  - Decouple shim plugin and task plugin
  - Add Endpoint address of sandbox
  - Support Vsock connection to task api
  - Support io by Streaming API
  - Retry for wait to remote sandbox controller
- 4 PRs is expecting to be merged
  - Sandboxed task
  - Call Update API in container lifecycle



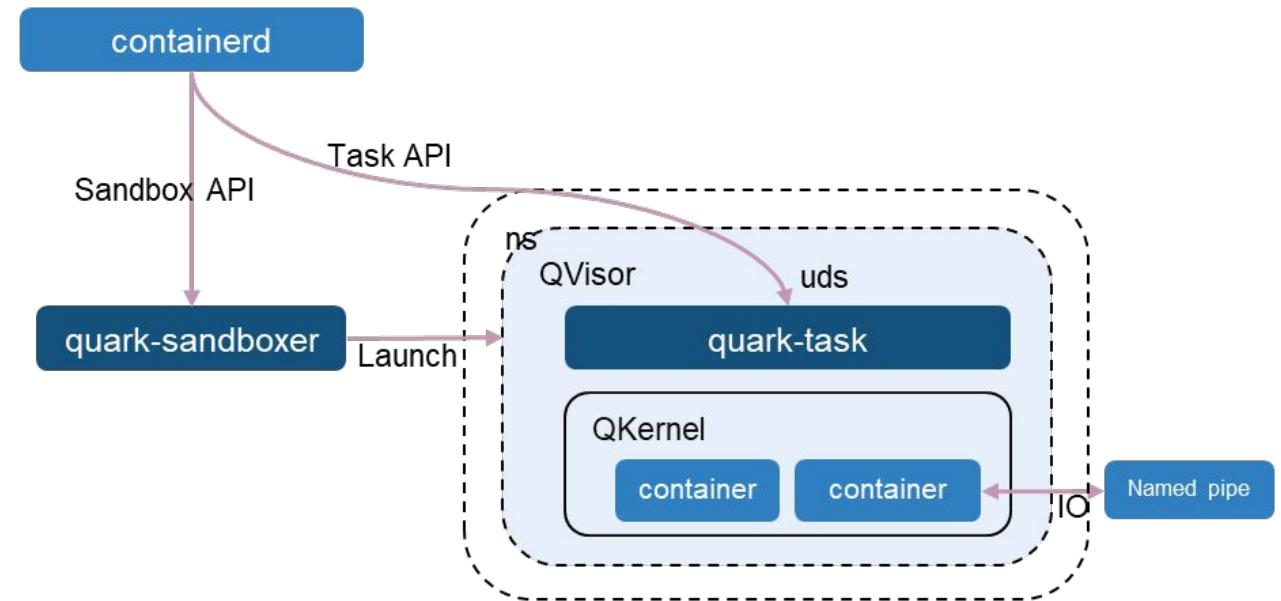
- **Vmm-sandboxer** focus on the vm lifecycle and resource management.
  - Support launch vm by cloud-hypervisor/qemu/stratovirt.
- **Task Service** is running inside the VM as the PID 1 process.
  - Containerd connect it with vsock/hvsock
- **Streaming IO**: IO stream transferred by the Streaming API.
- **No shim process**
  - Arch and maintenance is simplified.
  - Performance is improved.



\* Containerd is a forked version as Update API is not merged



- **Quark:** Application kernel sandbox.
  - Hypervisor(Qvisor) and kernel(Qkernel) rewritten in rust.
  - Secure container with kernel isolation.
  - Significant performance.
- **Task Service** is directly running inside the Qvisor.
  - No extra shim process.
- **No Shim process**



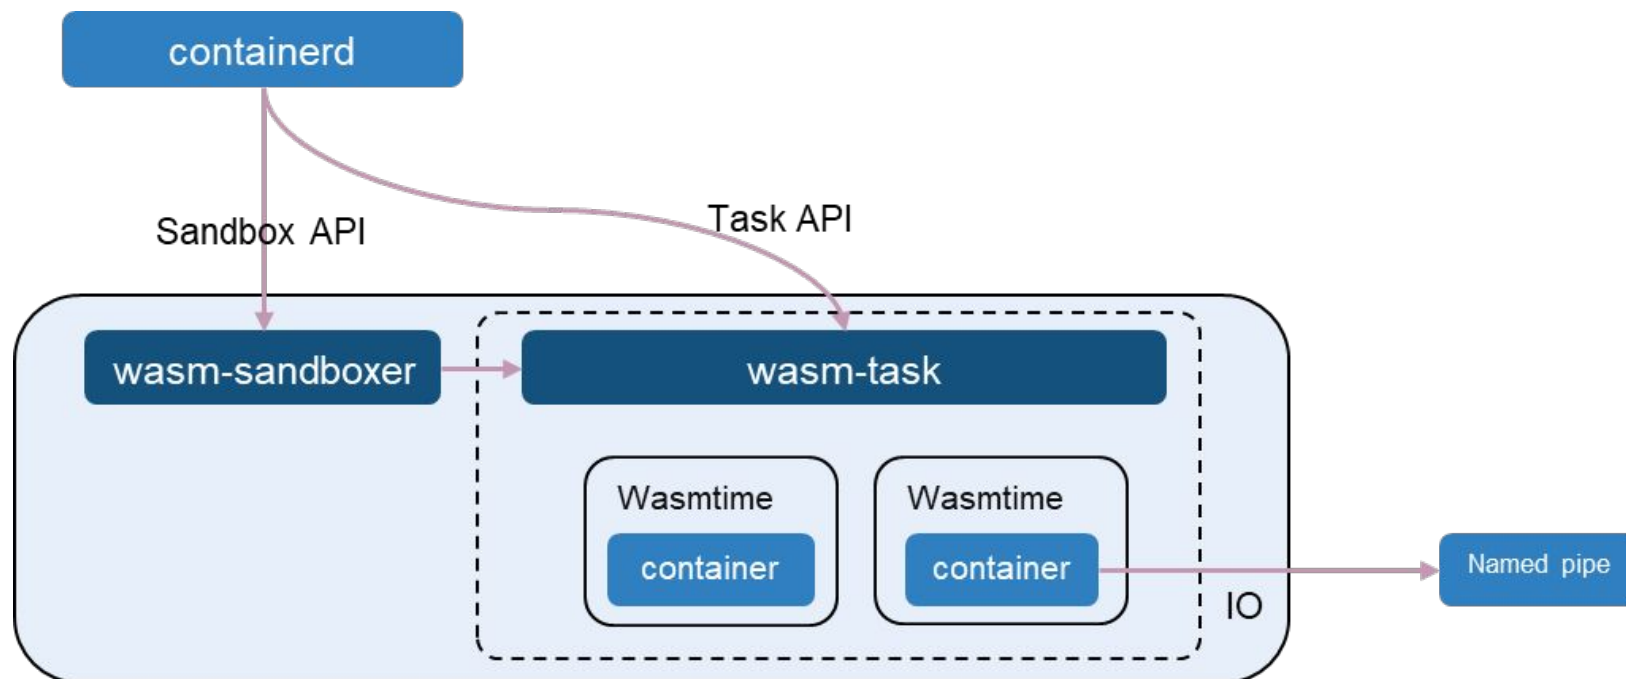
\* Containerd is a forked version as Update API is not merged

# Kuasar-wasm



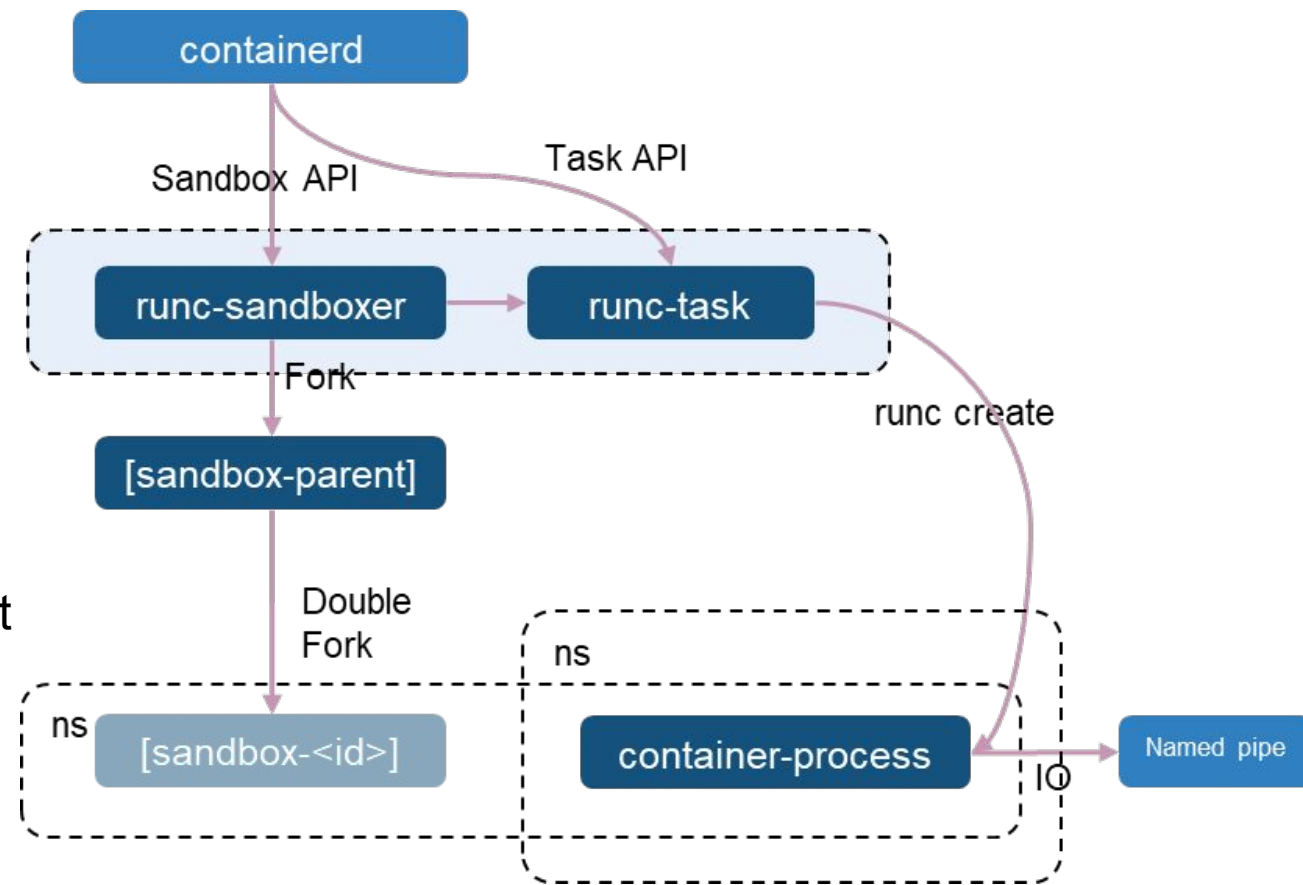
China 2024

- All in a single process.
  - Low start latency for wasm app.
  - Low overhead for each container.
- Wasmtime/WasmEdge supported
- **No shim process**



\* Containerd is a forked version as Update API is not merged

- **Sandbox process:** forked by an empty sandbox parent.
  - Exist only when shareProcessNamespace.
  - Memory overhead smaller than 100k.
- **Task Service forked by Sandboxer**
  - Fork a new task service everytime sandboxer restart.
  - The task service only exit when all containers it manages is removed.
- **No shim process.**
  - Only runc created containers



\* Containerd is a forked version as Update API is not merged

# How to use



China 2024

- 1. Start sandboxer
  - /usr/local/bin/vmm-sandboxer --listen /run/vmm-sandboxer.sock --dir /run/kuasar-vmm
- 2. Configure the /etc/containerd/config.toml
- 3. Start the sandbox with --runtime or create RuntimeClass in kubernetes

```
[plugins."io.containerd.grpc.v1.cri".containerd.runtimes.kuasar-vmm]
  runtime_type = "io.containerd.kuasar-vmm.v1"
  sandboxer = "vmm"
  io_type = "streaming"
  privileged_without_host_devices = true
  base_runtime_spec = "/etc/containerd/default_runtime_spec.json"

[proxy_plugins.vmm]
  type = "sandbox"
  address = "/run/vmm-sandboxer.sock"

[plugins."io.containerd.grpc.v1.cri".containerd.runtimes.kuasar-wasm]
  runtime_type = "io.containerd.kuasar-wasm.v1"
  sandboxer = "wasm"

[proxy_plugins.wasm]
  type = "sandbox"
  address = "/run/wasm-sandboxer.sock"
```

# Use Case 1: Lightweight Secure Container

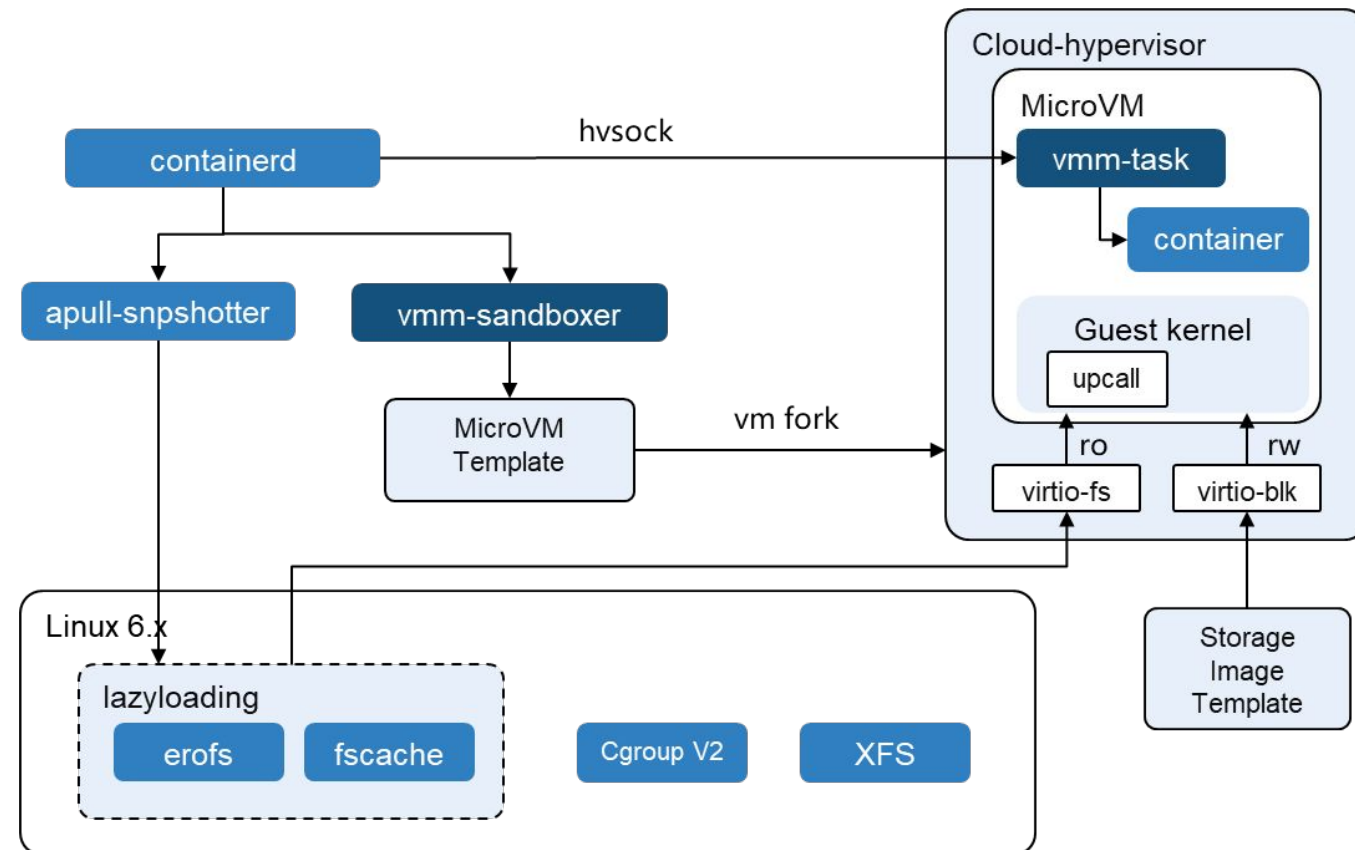


China 2024

- 1. VM Template based on mmap and copy-on-write
- 2. Condensed and pre-patched guest
- 3. Upcall in guest to remove acpi
- 4. Parallel cgroup creation

Memory overhead: **100MB** → **17MB**

Start latency: **850ms** → **12ms**



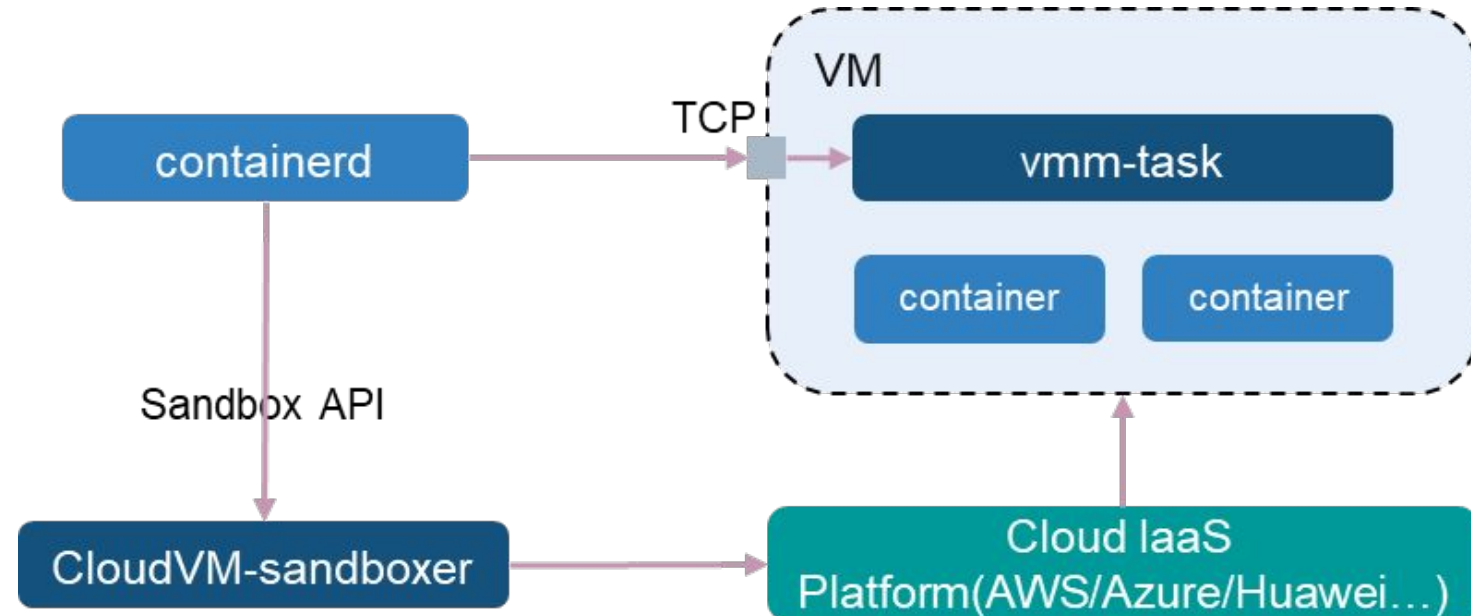


# Use Case 2: Remote sandbox



China 2024

- 1. Remote VM as a sandbox
- 2. Call IaaS API to create/delete VM
- 3. Task API and Streaming IO by TCP
- 4. Make a unified platform for VM and Serverless Container

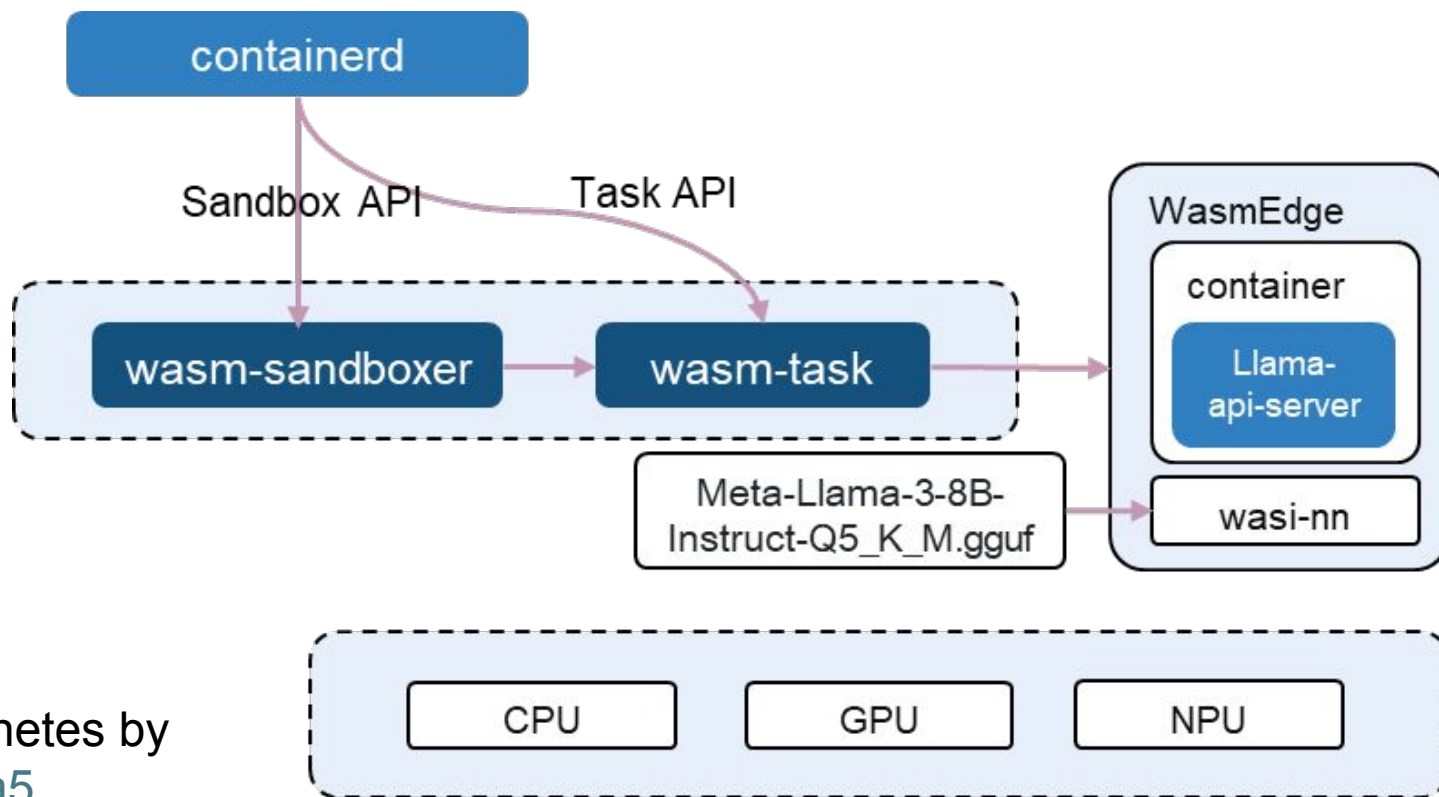


# Use Case 3: LLM running in wasm



China 2024

- 1. Fast and lightweight deploy
- 2. Cross cpu/gpu and OSes
- 3. Zero python dependency



Keynote: Deploying LLM Workloads on Kubernetes by WasmEdge and Kuasar: <https://sched.co/1eYa5>

<https://github.com/kuasar-io/kuasar/blob/main/docs/wasm/How-to-run-Llama-3-8B-with-Kubernetes.md>

# Welcome to join us!



China 2024

Containerd: <https://github.com/containerd/containerd>

Containerd community meeting:  
<https://docs.google.com/document/d/1Q8KyVJd26oAQ3MafbnkVBgJCopPI2Bw45H9L9br9Vus/edit#heading=h.bhv4ajd2ibx0>

Kuasar: <https://github.com/kuasar-io/kuasar>

Kuasar homepage: <https://kuasar.io/>

Kuasar slack:  
<https://cloud-native.slack.com/archives/C052JRURD8V>

# containerd

