



KubeCon



CloudNativeCon

THE LINUX FOUNDATION



AI_dev
Open Source GenAI & ML Summit

China 2024



KubeCon



CloudNativeCon



China 2024

Constructing 10^x Efficiency of Cloud-Native AI Infrastructure

Qiuping Dai

Peter Pan



KubeCon



CloudNativeCon



China 2024



Constructing 10^x Efficiency of Cloud-Native AI Infrastructure

Qiuping Dai

Peter Pan

About us



Qiuping Dai

Product Manager of Cloud Native

was major designer **of d.run**

skilled at network, storage, GPU area

Peter Pan

Cloud Native **Dev Lead**

employed by **DaoCloud**

open source advocate



Background



Opportunity

Government

MIIT + State Council 《 Action Plan for the High-quality Development of Computing Power Infrastructure 》

By 2025, up to **300EFLOps**, AI computing reaches **35%** ,

工业和信息化部、国务院国资委等6部门联合印发《算力基础设施高质量发展行动计划》，提出建设规划：2025年算力规模超过300EFLOPS，智能算力占比达到35%，

Industry

Analytical report 《AIDC industrial trend in China》
AIDC > 50% growth rate


Challenge

Low Utilization and efficiency

GPU Supply-Demand Imbalance

Technical Obstacle


Immature Operating Experience



Challenges - GPU Supply & Utilization

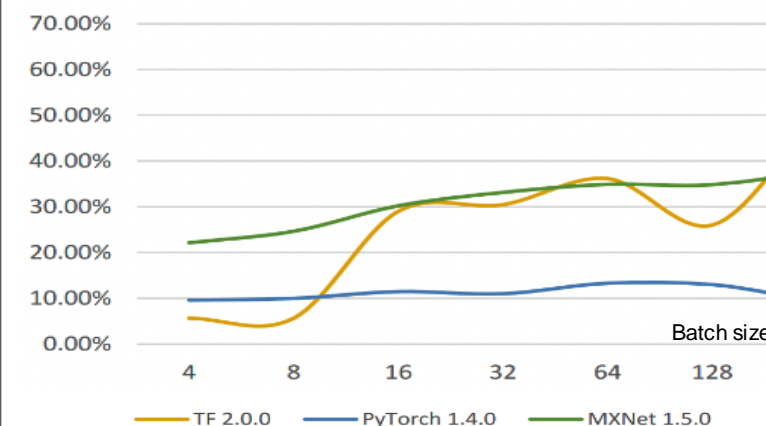


GPU Supply Shortage

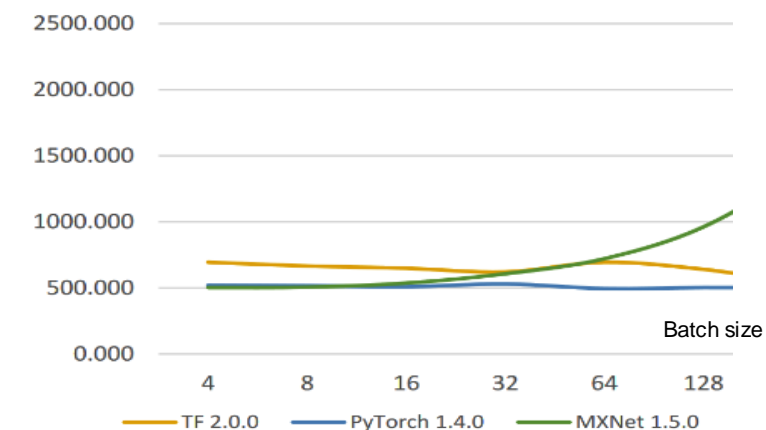


Low Utilization

GPU Power Usage



GPU Mem Usage (MiB)



Low Efficiency - Reasons



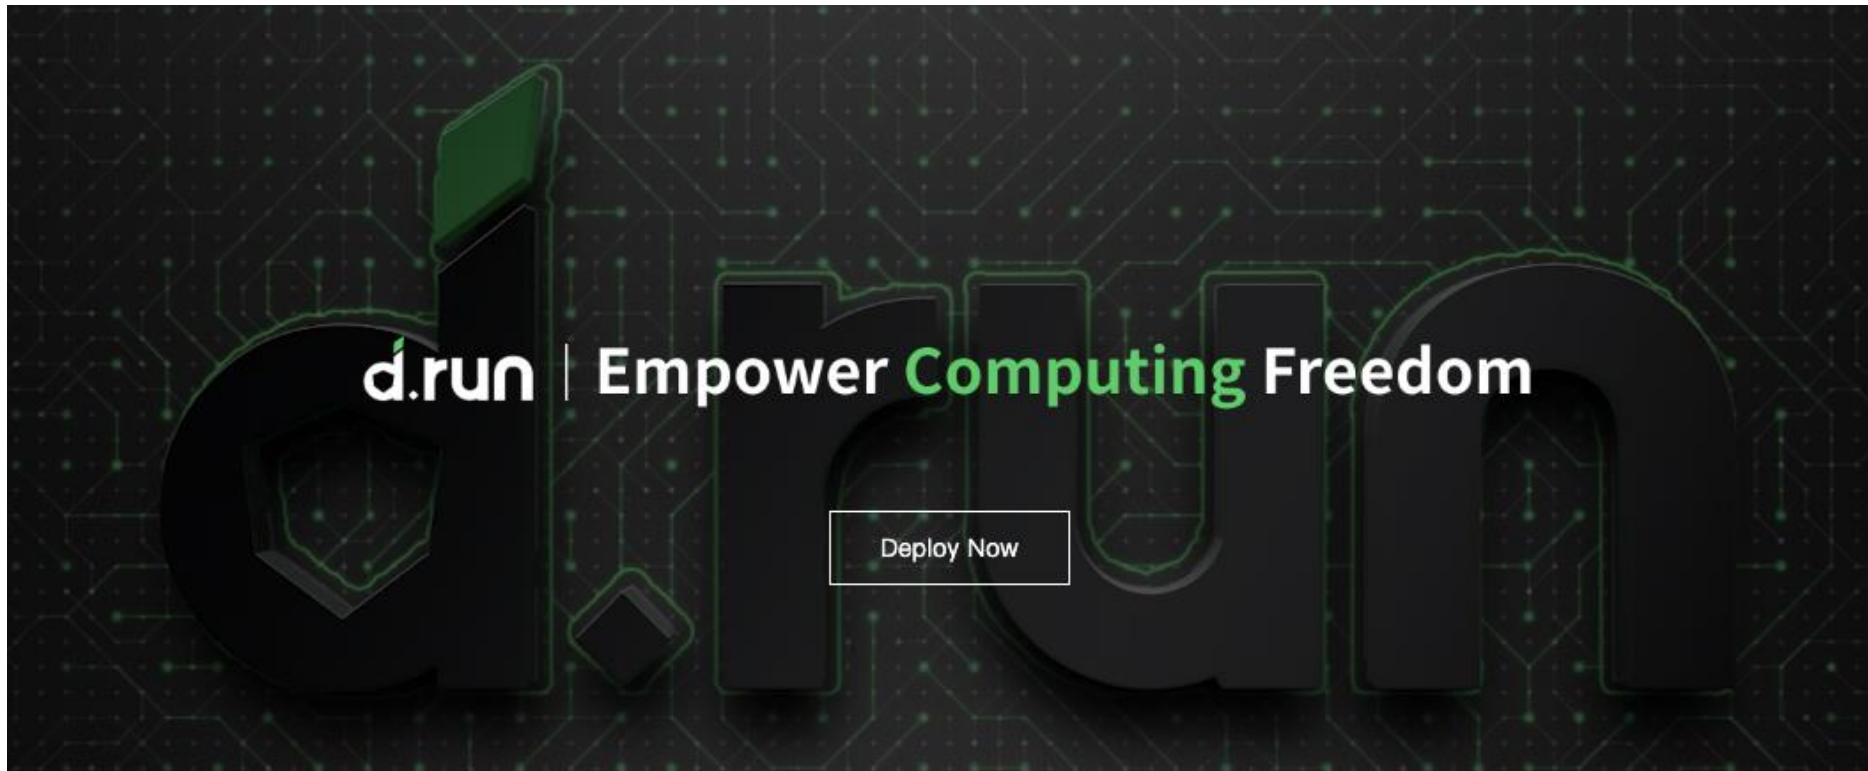
1. Network Bottom Neck
 - Data communication in low efficiency
2. Storage Read & Write Efficiency
 - Dataset load is slow
 - checkpoint saving is slow
3. GPU Allocation is not Optimal
 - GPU Scheduling is not matching the AI scenario
 - Allocated GPU is more than required in Inference scenario
4. GPU Fault waste Training time

Success Story



d.run as a computing-hub for AIDC across China. empowe

Both SaaS and On-Premises

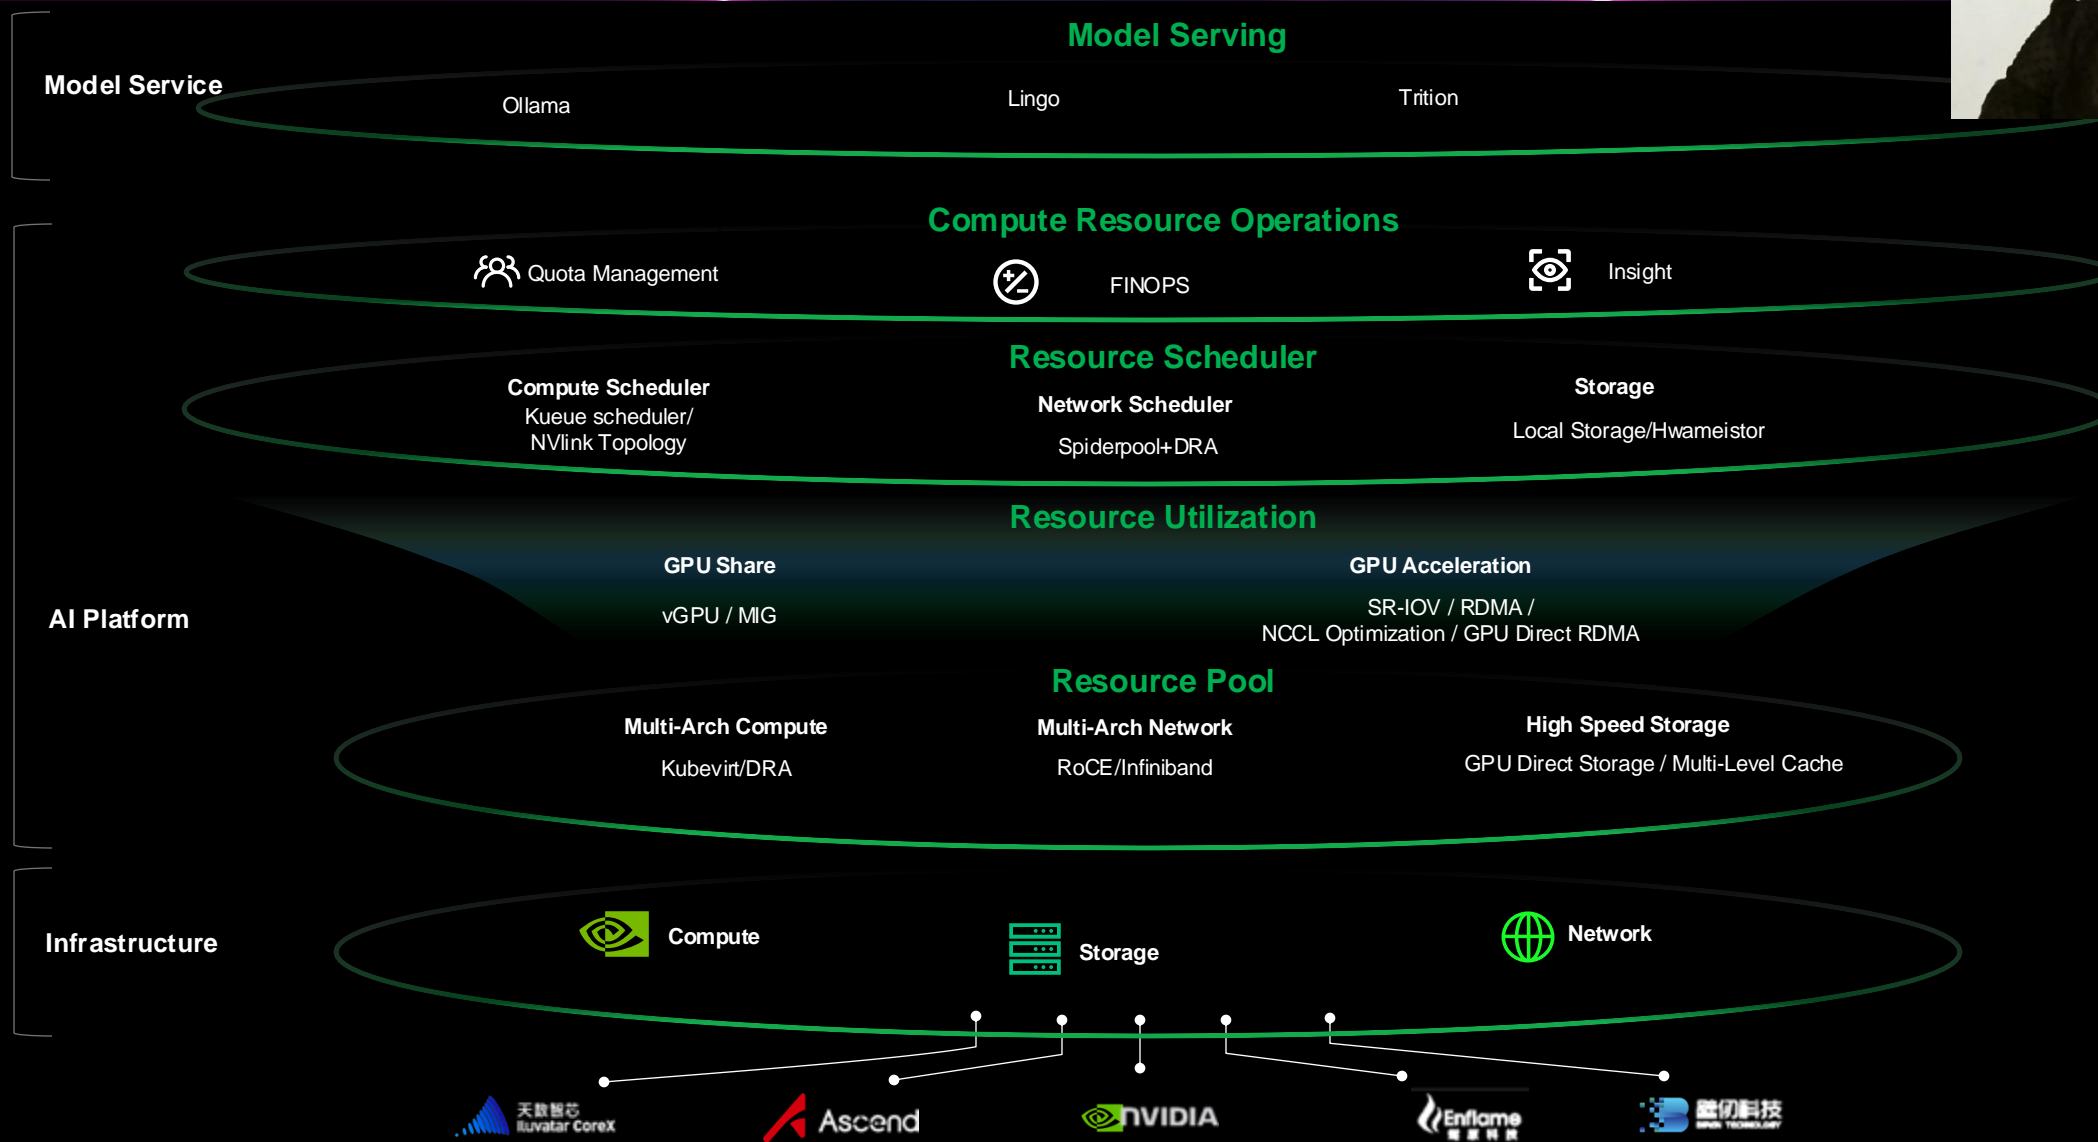


GPU cost: 48% saving

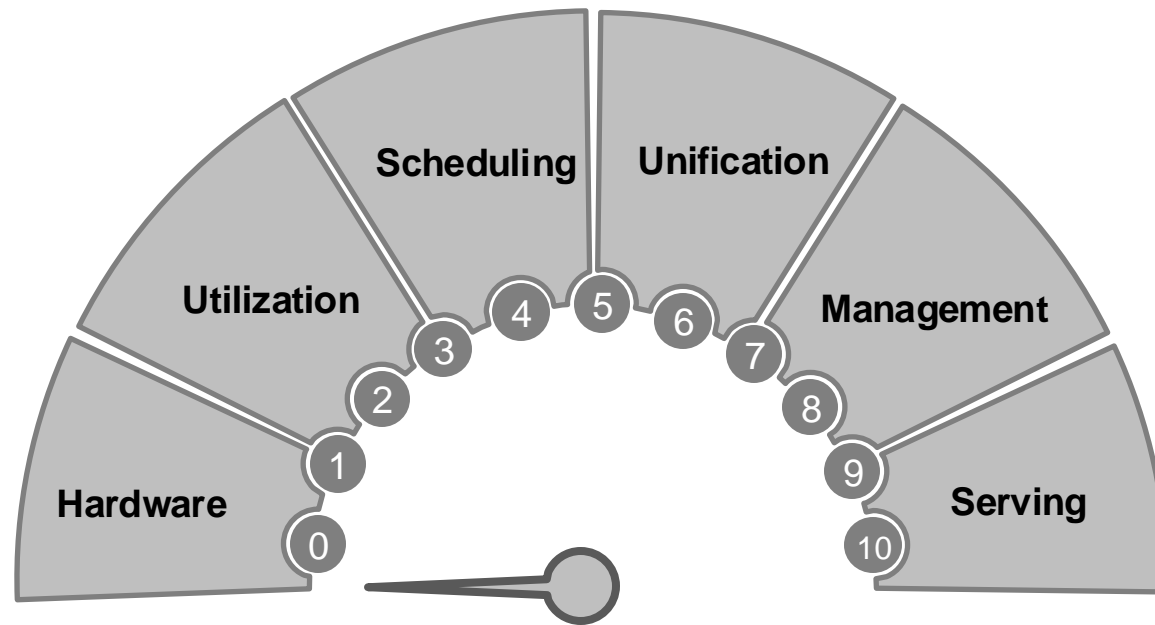
GPU avg utilization : 25% --> 54% [\[1\]](#)

[\[1\]](#) based on statistic of two A100 clusters managed by d.run

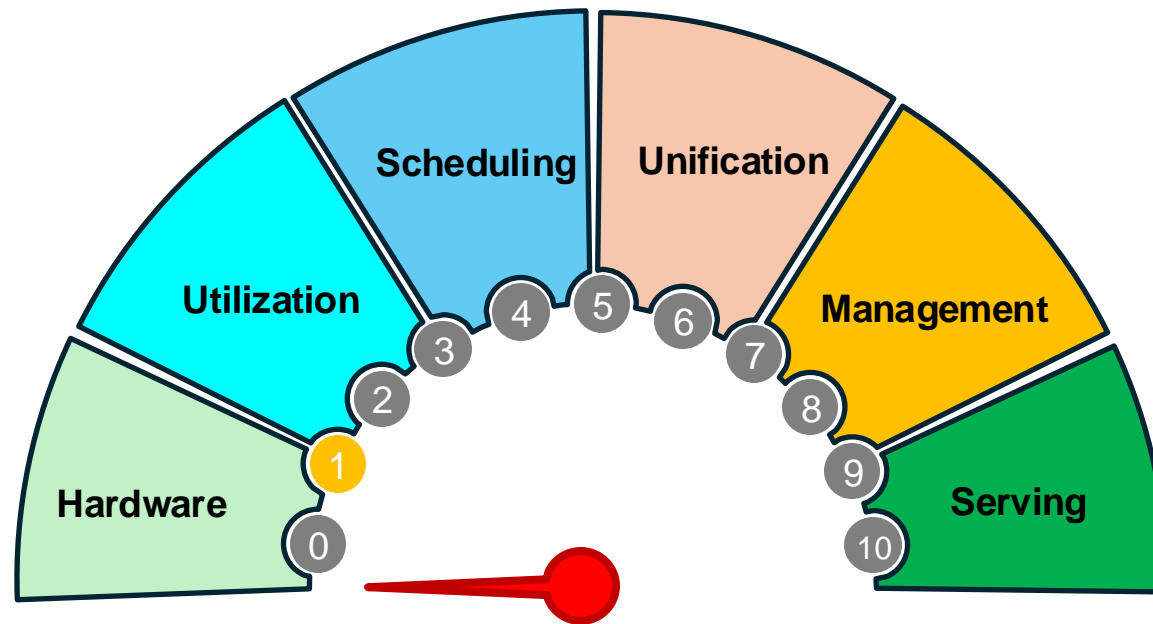
Optimization in different Layers



Journey starts !



1. Hardware Acceleration



Network : GPU Acceleration


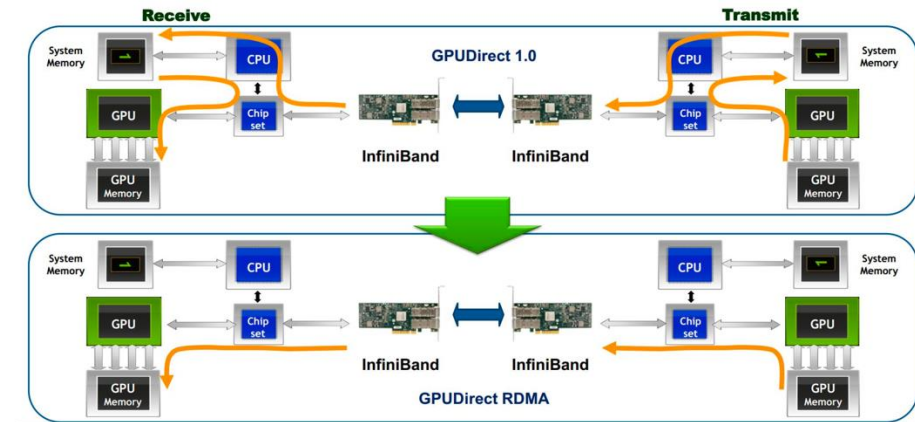


1. GPU Director RDMA to accelerate the GPU and RDMA Interface
2. Spider Pool Project Support different CNIs to work with RoCE/Infiniband

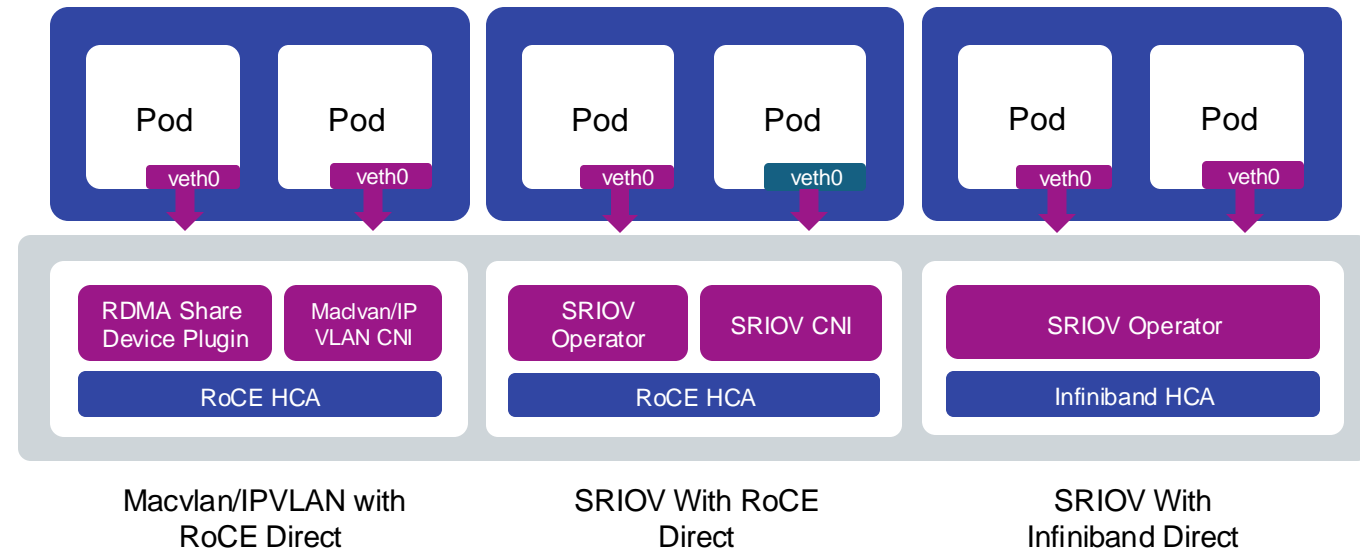


<https://github.com/spidemet-io/spiderpool>

- Based on Macvlan/IPVLAN CNI to Communicate by RoCE
- Based on SRIOV CNI to Communicate by RoCE
- Based on SRIOV CNI to Communicate by Infiniband



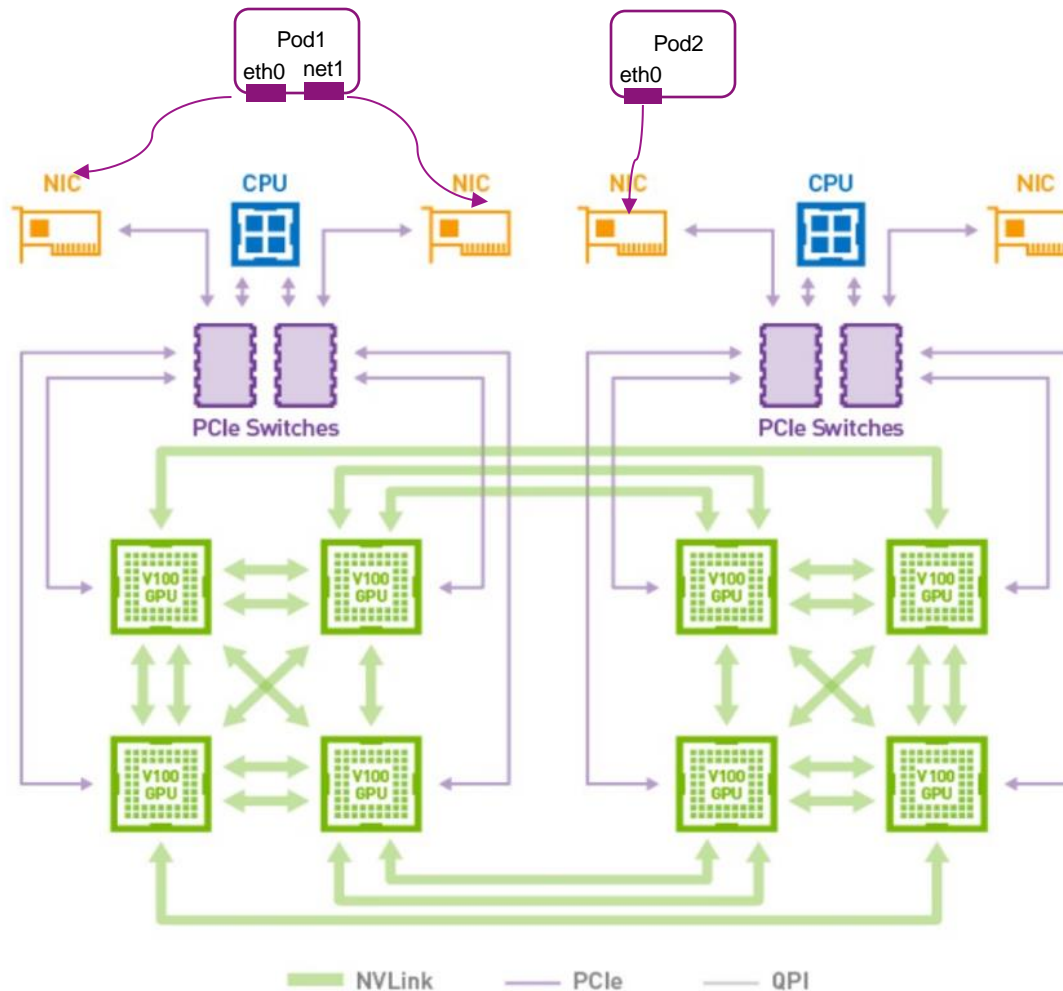
End-to-end communication latency of different technologies



Network : GPU topology-aware scheduling



1. Pod Allocated with RDMA HCA affinity with same PCIe Switch
2. Based on DRA to dynamically select the RDMA device with best topology



Spiderpool

+ DRA

<https://github.com/spidernet-io/spiderpool>

Storage : fully leveraged



Use Legacy Inventory 利旧

- CSI enabling : External Storage, HPC Storage(Lustre, BeeGFS, HDFS)

Scenario Fit

- Dataset: mostly R/O -> S3 + JuiceFS caching, Labeling -> write to S3
- Checkpoint Saving: IO Throughput(reduce GPU idle time), distributed storage(for resume training)
- Intermediate data: leverage local storage (HwameiStor , drbd HA is optional)
- Model: OCI image/raw file in S3
- Inference history: S3

Storage : Accelerate



Fluid - K8s Data Acceleration (backed by Alluxio/JuiceFS..etc)

1. Define Dataset

```
apiVersion: data.fluid.io/v1alpha1
kind: Dataset
metadata:
  name: llama3
spec:
  mounts:
  - mountPoint: minio://llama-3.1-8b
    # minio://<bucket name>
    name: minio
    options:
      minio-url: http://$minio-sever:9000
```

2. Mount PVC, trigger caching

```
#Before Caching
#kubectl get dataset minio-demo
```

| NAME | UFS TOTAL SIZE | CACHED | CACHE CAPACITY | CACHED PERCENTAGE | PHASE | AGE |
|--------|----------------|--------|----------------|-------------------|-------|-----|
| llama3 | 16.2GiB | 0.00B | 30.00GiB | 0.0% | Bound | 2h |

```
#After Caching
#kubectl get dataset minio-demo
```

| NAME | UFS TOTAL SIZE | CACHED | CACHE CAPACITY | CACHED PERCENTAGE | PHASE | AGE |
|--------|----------------|---------|----------------|-------------------|-------|-----|
| llama3 | 16.2GiB | 16.2GiB | 30.00GiB | 100.0% | Bound | 2h |


3. x10 faster w/ cluster-local caching

```
#Before Caching
root@nginx-648b57989d-4rfvs:/tmp# time cp /mnt/llama-3/* .
real    0m54.592s
user    0m0.001s
sys     0m0.232s

#After Caching
root@nginx-648b57989d-4rfvs:/tmp# time cp /mnt/llama-3/* .
real    0m5.609s
user    0m0.001s
sys     0m0.146s
```

10 times faster !

2. Resource Utilization



2. Maximize Utilization



Challenges

- Workload mis-match (e.g. run a Phi-3-mini (3.8B) on a H100-80G)
- Occupied GPU is idle (e.g. occupied by notebook but just coding)
- GPU fragments (e.g. small task)
- Training recovery from failure (e.g.: waste in re-train from scratch)

2. Max Util - Training Recover



Recover from GPU failure

1. Detect : from log & nvidia diag
2. Recovery: re-schedule + re-run from checkpoint

```
=== 1. Running diagnostics on Node gpu04 with DCGM Pod ID nvidia-dcgm-exporter-kgvtr ===

+----- Deployment -----+
| Denylist                  | Pass
| NVML Library              | Pass
| CUDA Main Library         | Pass
| Permissions and OS Blocks | Pass
| Persistence Mode          | Pass
| Environment Variables     | Pass
| Page Retirement/Row Remap | Pass
| Graphics Processes        | Pass
| Inforom                   | Pass
+-----+
+----- Integration -----+
| PCIe                      | Fail - All
| Warning                   | GPU 0 Error using CUDA API cudaDeviceGetByPCI
|                            | BusId Check DCGM and system logs for errors.
|                            | Reset GPU. Restart DCGM. Rerun diagnostics. '
|                            | forward compatibility was attempted on non su
|                            | pported HW' for GPU 0, bus ID = 00000000:1B:0
|                            | 0.0
| Warning                   | GPU 0 Error using CUDA API cudaDeviceGetByPCI
|                            | BusId Check DCGM and system logs for errors.
|                            | Reset GPU. Restart DCGM. Rerun diagnostics. '
|                            | forward compatibility was attempted on non su
|                            | pported HW' for GPU 0, bus ID = 00000000:1B:0
|                            | 0.0
+-----+
+----- Hardware -----+
| GPU Memory                | Skip - All
+-----+
+----- Stress -----+
+-----+

command terminated with exit code 226\n
=== Diagnostics finished ===

=== 1.0. Failure detected, killing master 0 torchrun ===
pod "nemo-distributed-demo-master-0" deleted
```

```
NVRM: GPU at PCI:0000:5d:00: GPU-f1906b9b-557a-e961-045c-9fe4be3ce012
NVRM: GPU Board Serial Number: 1653923026510
NVRM: Xid (PCI:0000:5d:00): 79, pid='<unknown>', name='<unknown>' GPU has fallen off the bus.
NVRM: GPU 0000:5d:00.0: GPU has fallen off the bus.
NVRM: GPU 0000:5d:00.0: GPU serial number is 1653923026510.
NVRM: A GPU crash dump has been created. If possible, please run
NVRM: nvidia-bug-report.sh as root to collect this data before
NVRM: the NVIDIA kernel module is unloaded.
```

2. Max Util - Optimal GPU Fit



- Example: There's a tiered GPU cluster: mixing H/A100, A6000, 4090
 - LLM **training** : H/A100 with best NVLink connect
 - LLM **Inference**: A6000/4090 with better cost performance

"nvidia.com/gpu: 1" will not tell ->



```
apiVersion: apps/v1
kind: Deployment
...
spec:
  resourceClaims:
  - name: a100
    source:
      resourceClaimTemplateName: a100
```

Workload specific A100

```
apiVersion: gpu.resource.nvidia.com/v1alpha1
kind: GpuClaimParameters
metadata:
  name: a100
spec:
  count: 1
  ...
  - productName: "*a100*"
```

DRA defines A100 scale offering

```
apiVersion: gpu.resource.nvidia.com/v1alpha1
kind: GpuClaimParameters
metadata:
  name: a6000
spec:
  count: 1
  ...
  - productName: "*a6000*"
  - orExpression:
    - index: 0
    - index: 1
    - index: 3
    - index: 5
```

DRA defines A6000 offering

More example and detail in
<https://github.com/NVIDIA/K8s-dra-driver>


Moreover, the gpu0/1/3/5 on each node

2. Max Utilization : vGPU



GPU Partition

Case1 :



Application



GPU Server

GPU
Virtualization



Useful in Inference /developing/study Scenario




training




Inference




Develop



GPU 1



vGPU 2



vGPU 3

GPU Server

1 GPU Used by 1 App
Utilization is lower than 20%

1 GPU Used by N App
Utilization improved to 70%

~~GPU~~
上班摸鱼



2. Max Utilization : vGPU



Case 2: Jupyter notebook 80%~90% time for coding or non-gpu works

apiVersion: [kubeflow.org/v1](https://kubeflow.org/docs/v1)
kind: **Notebook**
spec:
 template:
 ...
 ...
 resources:
 limits:
 nvidia.com/gpu: 1

```
4 from scipy import integrate
5
6 def solve_lorenz(sigma=10.0, beta=8./3, rho=28.0):
7     """Plot a solution to the Lorenz differential
8     equations."""
9
10    max_time = 4.0
11    N = 30
12
13    fig = plt.figure()
14    ax = fig.add_axes([0, 0, 1, 1], projection='3d')
15    ax.axis('off')
```

Some ways out below, but not user-friendly

- auto-shutdown of idle notebook
- using arena CLI to submit async job on demand



- Split GPU at will
- Over-commit of GPU Memory

2. Max Utilization : vGPU



Solution 1

M.I.G - GPU partition (physical)

- Good performance/isolation, but limited on fixed scale/max 7 slices
- **Dynamic MIG** w/ **DRA** supported (create MIG “device” on demand)

<https://github.com/NVIDIA/K8s-dra-driver/blob/main/demo/specs/quickstart/gpu-test4.yaml>

Pod

```
resourceClaims:
  - name: mig-enabled-gpu
    source:
      resourceClaimTemplateName: mig-enabled-gpu
  - name: mig-2g
    source:
      resourceClaimTemplateName: mig-2g.10gb
```


2. Max Utilization : vGPU & Over-Commit



Solution 2

GPU Middleware – **HAMi**   (logical partition)

1. GPU slicing size - “at will”



resources:

limits:

`nvidia.com/gpucores` : “50” # 50% GPU time

`nvidia.com/gpumem`: 3000 # 3GB HBM

2. HBM sharing & over-commit



On single A100,

But 2 pods can request both 80G at the same time, and use/share HBM at diff time.

When two pods use both 80G, “swap” happens: system memory will help

more details in



Is Your GPU Really Working Efficiently in the Data Center? N Ways to Improve GPU Usage | 您的GPU在数据中心真的高效工作吗? 提高GPU使用率的N种方法 - Xiao Zhang, DaoCloud & Wu Ying Jun, China Mobile

Wednesday August 21, 2024 13:50 - 14:25 HKT

Level 1 | Hung Hom Room 3

Binpack



Resource fragments



node1



node2



But no way to schedule it
unless App4 went to node1

Solution #1:

K8s scheduling-plugin “NodeResource”

1. **MostAllocated** policy : favoring nodes with higher allocation.
2. GPU with more **weight** than others

```
apiVersion: kubescheduler.config.K8s.io/v1
kind: KubeSchedulerConfiguration
profiles:
- pluginConfig:
  - name: NodeResourcesFit
    args:
      scoringStrategy:
        type: MostAllocated
      resources:
        - name: nvidia.com/gpu
          weight: 3          # more favor to GPU
        - name: cpu
          weight: 1
        - name: memory
          weight: 1
```

Solution #2:  VOLCANO

<https://volcano.sh/en/docs/plugins/#binpack>

Global config of volcano to enable binpack

Solution #3  HAMi

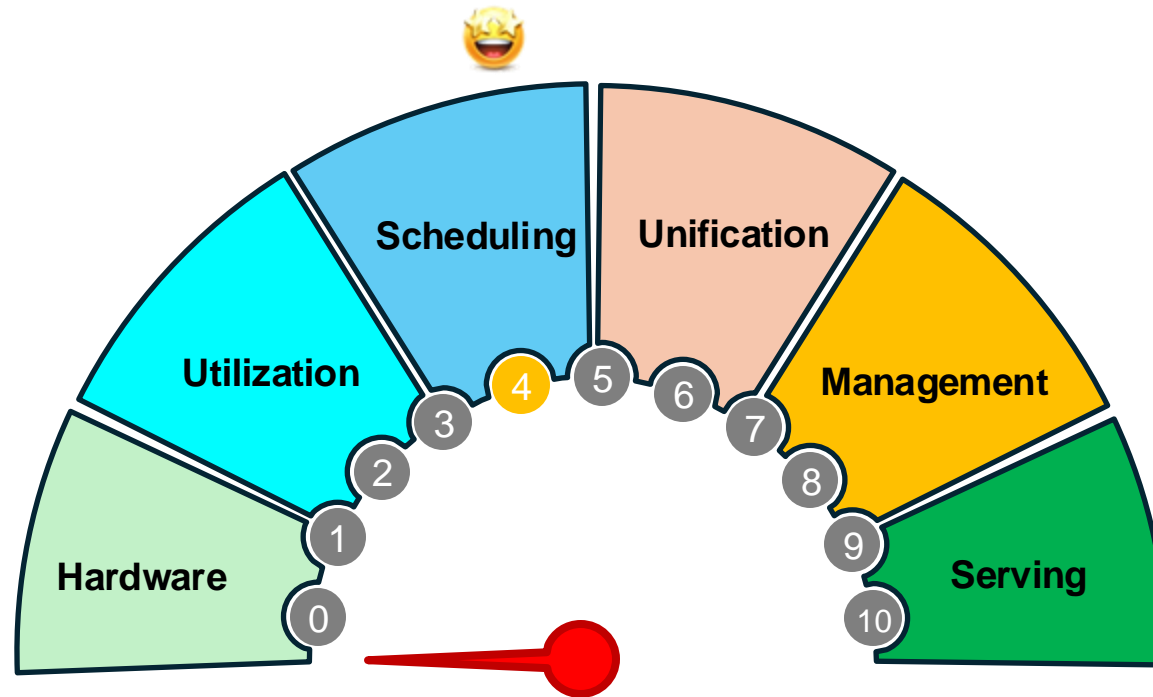
<https://github.com/Project-HAMi/HAMi>

vGPU-level binpack

(put 0.5 vGPU + 0.3 vGPU pods on the same physical GPU)

<https://kubernetes.io/zh-cn/docs/concepts/scheduling-eviction/resource-bin-packing/>
<https://github.com/kubernetes-sigs/scheduler-plugins/tree/master/pkg/noderesources>

3. Task Queue. & Scheduling




3. Task Mgmt : co-location



China 2024

1. **Put co-locating jobs deployed** in one queue
2. **Setting Pod Priority Class**, inference Job seizes the CPU with high priority, and training task is evicted first in case of OOM
3. **Setting Job Priority in Queue**, inference Job use the GPU with high priority, Training Job will be suppressed
4. **Setting CPU Burst** (kernel 5.14+), It can help the latency sensitive workload to "break" CPU limit sometimes.
5. **Namespace Elastic Quota**, different NS can share GPU quota, to improve the GPU usage
6. **Fair Schedule**, avoid some workload never get its requirement fulfilled , and "starve to death"

3. Task Mgmt : co-location



When App4(H) comes in, it “jumps the queue”
App3(L) be preempted or suspended


- Jobs/deployments mixed deployed in one queue
- Critical job sets high priority

```
apiVersion: scheduling.k8s.io/v1
kind: PriorityClass
name: my-high-priority
preemptionPolicy: PreemptLowerPriority
value: 100000
```

(1) Pod's `PriorityClass`

```
apiVersion: kueue.x-k8s.io/v1beta1
kind: WorkloadPriorityClass
metadata:
  name: medium-priority
  value: 10000
```

(2) Kueue's `WorkloadPriorityClass`



17-9c8d-c19b2cefd8ff", "ctrname": "container-1", "deviceuuid": "GPU-26a583dd-542e-09bb-5dd1-9cc5bd6eb552", "endpoint": "monitorport", "exported_name": "GPU-26a583dd-542e-09bb-5dd1-9cc5bd6eb552"

(3) `CUDA_TASK_PRIORITY` in HAMi

3. Task Mgmt : Elastic Quota & HPA

Training & Inference workload in different Namespace

Total: **10** GPU

NS Quota A (min:4, **max:6**) → App-A

NS Quota B (min:6, **max:8**) → App-B

AppA use 4GPU (min), AppB use 3GPU, everything is OK




AppA use 6GPU (max), AppB use 3GPU, everything is OK



AppB need another 3 GPUs, resource eviction happened , AppA retain the GPU back , AppB get 6 GPUs (min)



HPA rule for Inference Application



Metrics:

GPU utilization may not be good practice

Try “Triton inference Queue-Time”
[nv_inference_queue_duration_us](#)

Elastic Quota by Capacity Scheduler:

<https://github.com/kubernetes-sigs/scheduler-plugins/tree/master/pkg/capacityscheduling>

https://docs.nvidia.com/deeplearning/triton-inference-server/user-guide/docs/user_guide/metrics.html

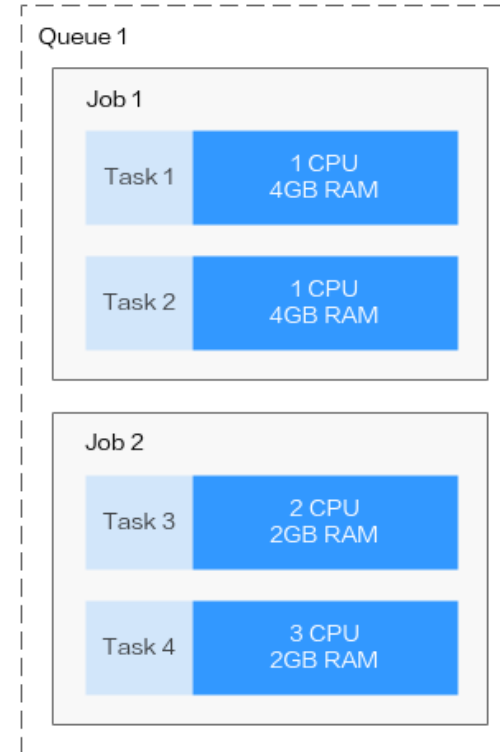
3. Task Mgmt: Fair Schedule

[Volcano DRF]



when in resource shortage, small tasks are fulfilled as possible, Then left part allocated fairly to other tasks

- Resource is allocated on an increasing demand basis
- Each user receives no more resources than needs
- Unsatisfied users share the remaining resources equally



Job 1:

Total CPU: 2 CPU

CPU Share = $2/10 = 0.2$

Total Memory: 8GB

Memory Share = $8/20 = 0.4$

Dominant resource is Memory

Share = 0.4

Job 2:

Total CPU: 5 CPU

CPU Share = $5/10 = 0.5$

Total Memory: 4GB

Memory Share = $4/20 = 0.2$

Dominant resource is CPU

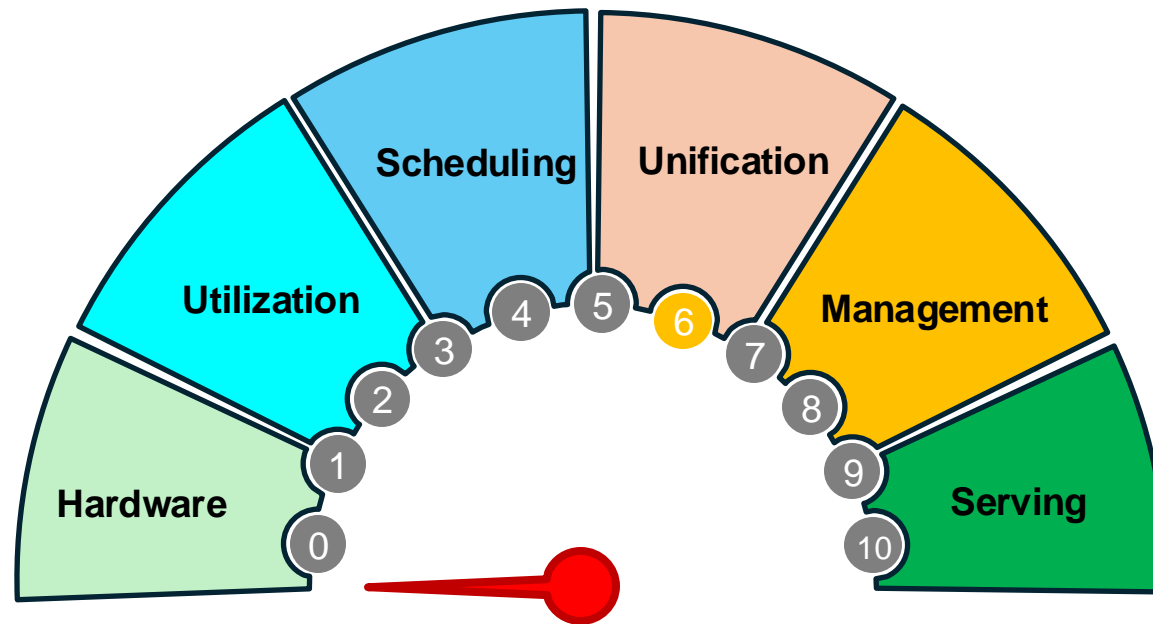
Share = 0.5

Cluster Resources: 10 CPU, 20GB RAM, 0 GPU

Share = Total Request / Cluster Resources

Ways to calculate the jobs "share" value

4. Unified Architecture



4. Unified Architecture - VM



China 2024

- Background : VM still lives, due to legacy apps / mgmt-policy
- Cure : **KubeVirt** can help out
- Benefits
 - Reuse the same underlying hardware
 - Reduce the IaaS software cost

just 3 Steps:

1. **Nvidia gpu-operator**: support container and VM

<https://docs.nvidia.com/datacenter/cloud-native/gpu-operator/latest/gpu-operator-kubevirt.html#install-the-gpu-operator>

- Passthrough GPU
- Nvidia vGPU (License required)

2. **Config KubeVirt CR**

add specific gpu device

3. **VM instance**:

device yaml specific GPU

```
spec:
  configuration:
    developerConfiguration:
      featureGates:
        - GPU
        - DisableMDEVConfiguration
    permittedHostDevices:
      pciHostDevices:
        - externalResourceProvider: true
          pciVendorSelector: 10DE:1BB3
          resourceName: nvidia.com/GP104GL_TESLA_P4
      mediatedDevices:
        - mdevNameSelector: NVIDIA A10-24Q
          resourceName: nvidia.com/NVIDIA_A10-24Q
```

Physical GPU

Nvidia vGPU

```
apiVersion: kubevirt.io/v1
kind: VirtualMachine
...
spec:
  domain:
    devices:
      gpus:
        - deviceName: nvidia.com/GP104GL_TESLA_P4
          name: gpu0
...
```

4. Unified Architecture - Heterogeneous

- diff GPU model/brands together
- Training:
 - Mixed training is a great challenges - Buckets Effect

Inference

- Individual cluster for each xPU type - ☒
- Mixed cluster, diff xGPU in diff node - ☒
- Mixed node is challenging (DRA or change device-plugin)

Create Training Job

Resource Spec

i Pytorch Distributed job, in which one Master node is automatically elected and the rest are Worker nodes.

Specs * 8 Core 16 G

GPU ☒ On

GPU Type *

Select GPU Type

Nvidia vGPU

Nvidia MIG

Iluvatar VGPU BI V100

Ascend 310

Ascend 910

Manage GPU Types

Python Code

CUDA compatible or import
diff framework

Libraries in Image

Usually using specific base
image even CUDA compatible.

Workload Yaml

Provides interface for Scientist,
to fit their image to xPU type,
auto adding `resources.limit`

Device plugin

Each node carries their own
device-plugin/xPU-exporter,
with taint for xpu-operator.

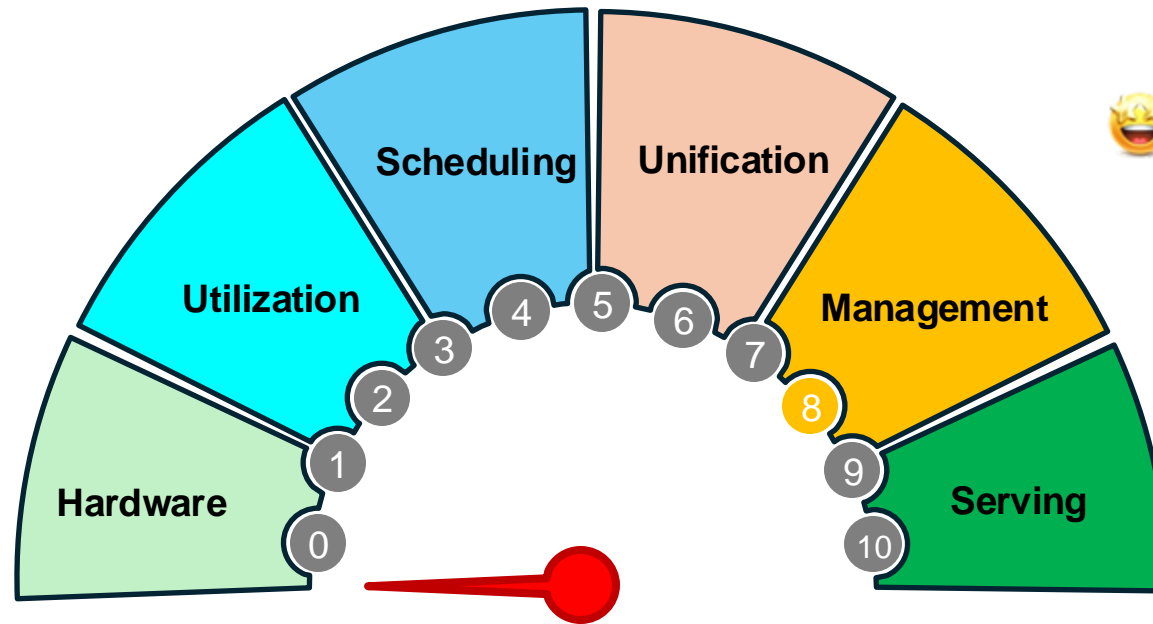
xPU hardware

K8s

Maybe in the future:
Auto detect the image to fit the xPU
Auto detect the code to use correct
image

- nvidia.com/gpu : 1
- Huawei.com/ascend910: 1
- metax-tech.com/gpu: 1
- iluvatar.ai/gpu: 1

5. Management Philosophy



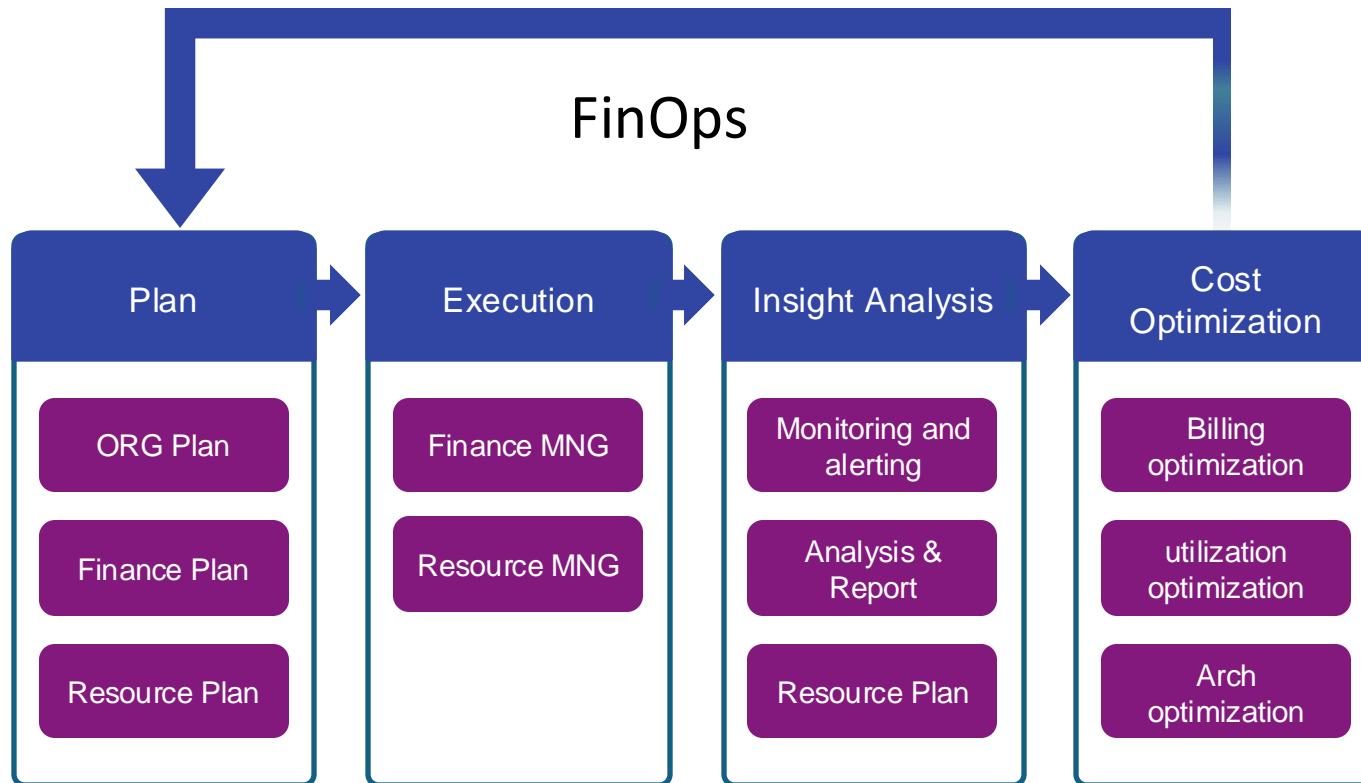
5. Management & Operation



China 2024

- **Organization Management** 组织管理
 - Organization mapping with namespace and privilege
 - Tenant isolation & resource quota
- **Operation Management** 运营管理
 - Org Billing
 - Cost Management – FinOps

5. Mgmt - FinOps Philosophy




- Measure the return on investment (ROI)
- Choose appropriate resources and billing methods
- Introduce elasticity mechanisms for application workloads
- Continuously monitor and optimize

5. Mgmt - OpenCost

FinOps measurement

<https://github.com/opencost>



Cost Allocation

Cost Allocation

Cloud Costs

Yesterday by namespace

31 October 2023 by Namespace

Date Range

Yesterday

Breakdown

Namespace


Resolution

Entire window

Currency

USD

Namespace Cost statistic



| Name | CPU | RAM | PV | Efficiency | Total cost |
|-------------|--------|--------|--------|------------|------------|
| Totals | \$0.58 | \$0.06 | \$0.01 | 24.5% | \$0.65 |
| kube-system | \$0.53 | \$0.03 | \$0.00 | 5.8% | \$0.56 |
| | \$0.00 | \$0.02 | \$0.01 | Inf% | \$0.03 |
| | \$0.02 | \$0.01 | \$0.00 | Inf% | \$0.03 |
| | \$0.02 | \$0.00 | \$0.00 | 12.3% | \$0.02 |
| | \$0.01 | \$0.00 | \$0.00 | Inf% | \$0.01 |

nippo-system namespace

12H

24H

7D

30D

FILTERS (0)

Total Cost (12h)

RAM

CPU

Network

LB

PV

Shared @

GPU

US\$0.17

US\$0.13

US\$0.04

US\$0.00

US\$0.00


US\$0.00

US\$0.00

US\$0.00

Cost Efficiency Factor

Historical Cost



Cost Efficiency

60.30%



Bottomline rule for our team, e.g.:
Cost Efficiency $\geq 40\%$
Mem Efficiency $\leq 130\%$

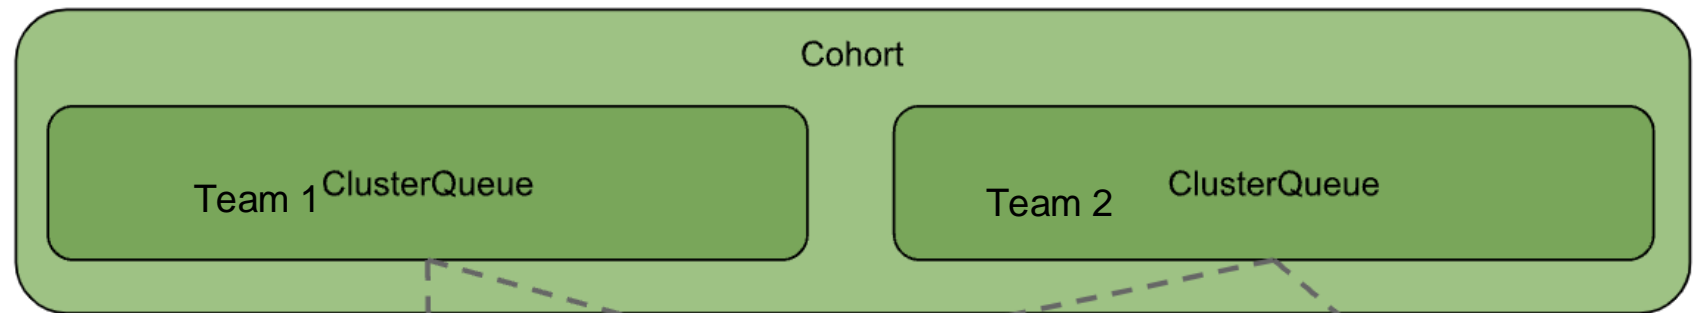
5. Mgmt - Tenant



China 2024

1. Namespace – `ResourceQuota`
2. Queue Quota – Kueue's `ClusterQueue` `nominalQuota`
3. “Cohort” – flexible quota between teams

- Borrow / Lend quota to others
- Reclaim/Preempt from others



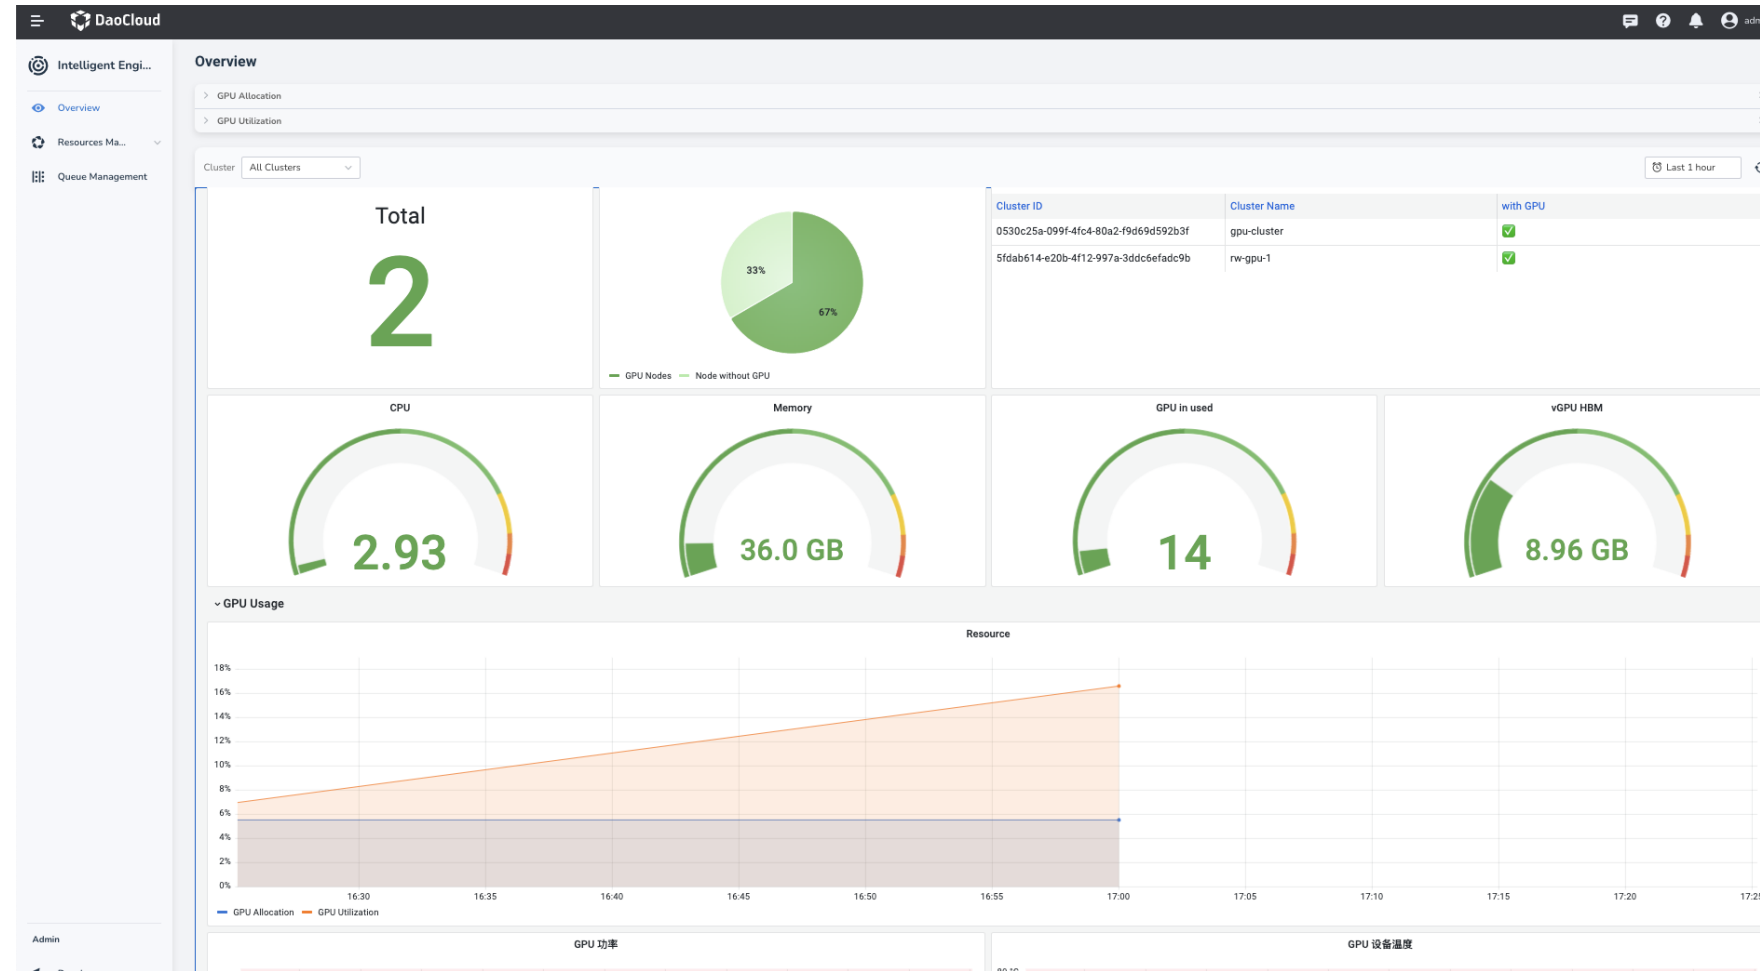
5. Mgmt – Role base view



China 2024

Separate **Operator** & Scientist

- **Operators** focus on :
 - Inventory
 - Resource Utilization
 - Tenant quota usage
 - Alerts / Monitor



5. Mgmt – Role base view



China 2024

Separate Operator & Scientist

Pros: Privilege & user-friendliness

- **Scientist** focus on :
 - Notebook
 - Training jobs
 - Dataset
 - Tenant quota total/usage

DaoCloud

Notebooks

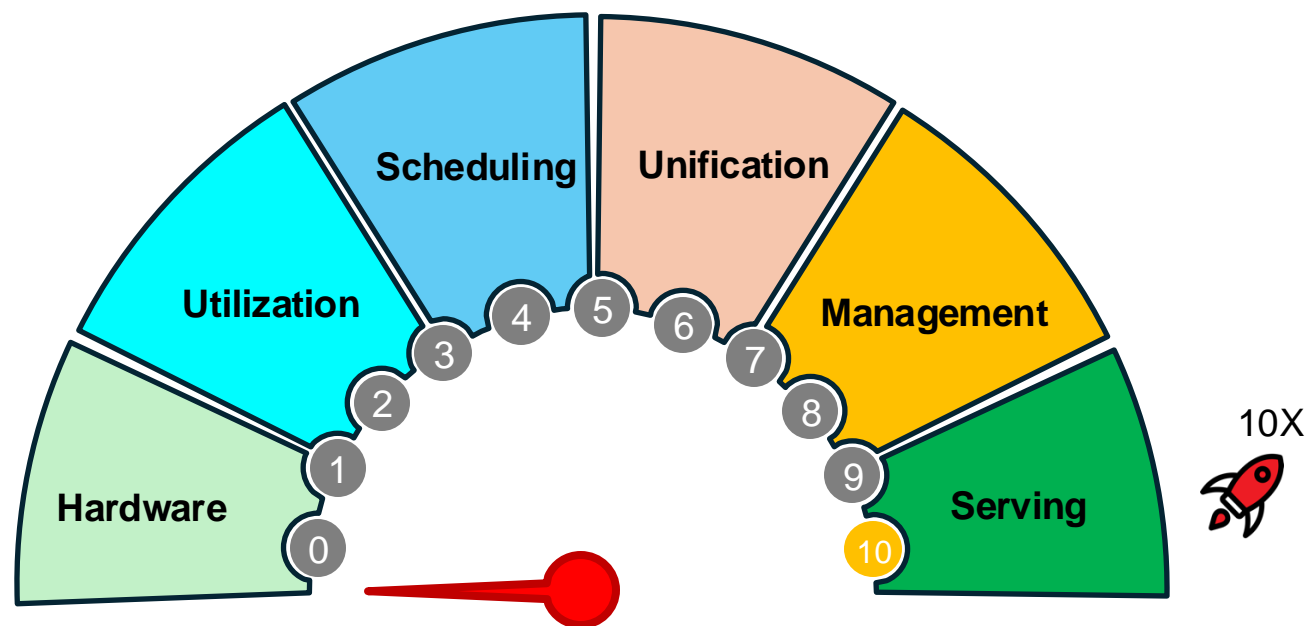
> Notebooks

Cluster: Namespace:

| Name | Queue | Image | Namespace | Notebook Type | Status |
|-------------------|-----------|------------------------------------|-----------|---------------|---------|
| test-01 | default | 10.6.135.168:30002/arm/kotf:l... | demo | Jupyter | Pending |
| baize-demo | queue-jxj | release-ci.daocloud.io/baize/ba... | jxj | Jupyter | Stopped |
| demo | default | release-ci.daocloud.io/baize/ba... | demo | Jupyter | Stopped |
| zxw-test-notebook | default | release-ci.daocloud.io/baize/ba... | kebe | Jupyter | Stopped |

Total 4 records

6. Boost Model Serving !



6. Model Serving



China 2024

Paints of Efficiency

- **Complex** to deploy LLM model:
 - Need to Download model -> write transformer/llama code -> python env/run
- Slow model **download** from time to time
 - How to pack and reuse/cache the model in cluster
- Models **occupy** HBM but no request:
 - Knative can scale-to/from-zero , but somehow heavy (istio required)

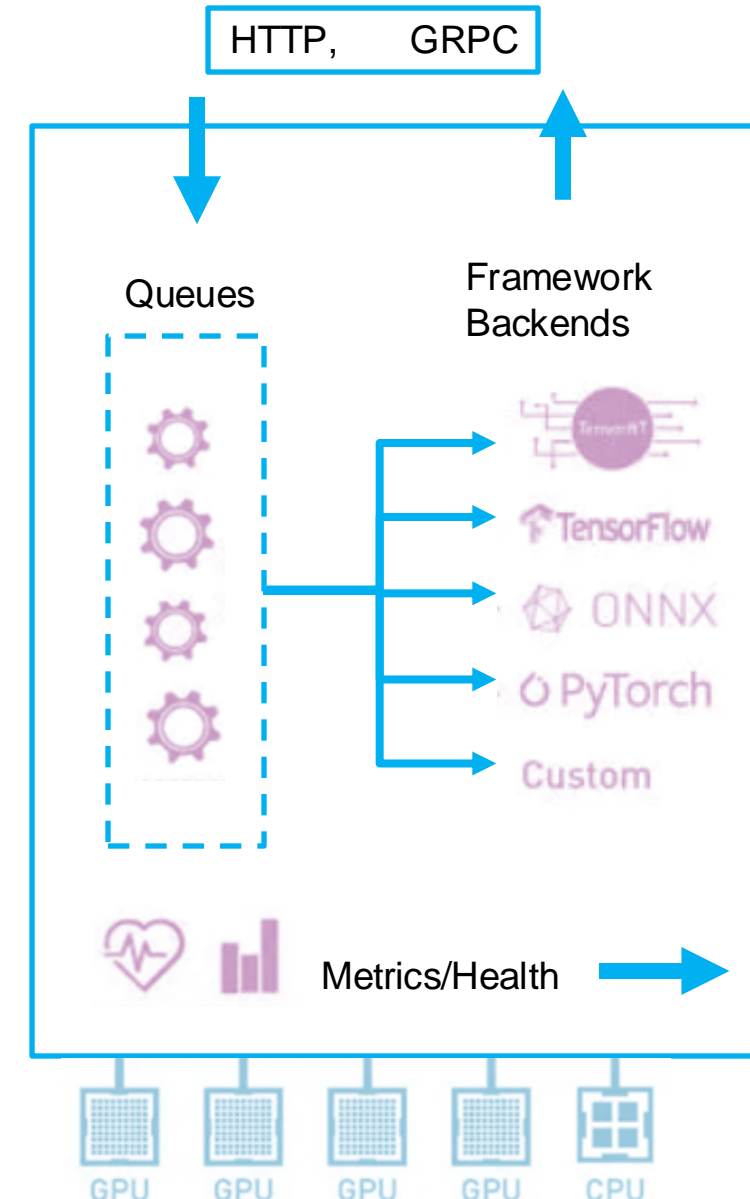
6. Model Serving – Triton

• Triton

- Handy: `docker run -v ${model-path} nvcr.io/nvidia/tritonserver`
- All-in-one : `serving / api-gateway / metrics / scheduling / batching`
- Wide range of backend supported :
 - `vLLM`、`TensorFlow`、`PyTorch`、
 - `ONNX`、`OpenVINO`、`TensorRT`..
- Cons : lack of OpenAI format API

Our Practice in K8S : Triton-operator

```
apiVersion: serving.baize.io/v1alpha1
kind: Inference
spec:
  framework:
    triton:
      backend: vllm
  models:
    - modelPath: model/mnist_cnn/1/model.pt
      config: ...
```



6. Model Serving – OCI Volume



China 2024

```
apiVersion: serving.baize.io/v1alpha1
kind: Inference
spec:
  framework:
    triton:
      backend: vllm
  models:
  - modelPath: model/mnist_cnn/1/model.pt
    config: ...
```

OCI Artifacts Volume ^{alpha} in Kubernetes 1.31

```
spec:
  containers:
  - name: test
    image: nvcr.io/nvidia/tritonserver
    volumeMounts:
    - name: volume
      mountPath: /model
  volumes:
  - name: volume
    image:
      reference: quay.io/crio/artifact:v1
```

After 10 years long last, finally it comes

Image volumes and container volumes #831



thockin opened this issue on Aug 8, 2014 - 147 comments



thockin commented on Aug 8, 2014

This would map closely to Docker's native volumes support, and allow people to build and version p containers. Maybe read-only? Haven't thought that far...



98

6. Model Serving - ORAS



China 2024

```
\[root]ls llama-3/*
llama-3/model-00001-of-00004.safetensors  llama-3/model-00003-of-00004.safetensors  llama-3/model.safetensors.index.json
llama-3/model-00002-of-00004.safetensors  llama-3/model-00004-of-00004.safetensors  llama-3/tokenizer.json
\[root]oras push --plain-http release.daocloud.io/peter/llama-3-weights:v3.1-8b  llama-3:application/octet-stream
✓ Uploaded llama-3
  ↳ sha256:0b4d055901b51b24a62e3fcca15a60f100239803dd5c07fe49489b168f468261
✓ Uploaded application/vnd.oci.empty.v1+json
  ↳ sha256:44136fa355b3678a1146ad16f7e8649e94fb4fc21fe77e8310c060f61caaff8a
✓ Uploaded application/vnd.oci.image.manifest.v1+json
  ↳ sha256:e5145fa1529423daa154cbda91176b4a081bde51878087b783caa084731fd95
Pushed [registry] release.daocloud.io/peter/llama-3-weights:v3.1-8b
ArtifactType: application/vnd.unknown.artifact.v1
Digest: sha256:e5145fa1529423daa154cbda91176b4a081bde51878087b783caa084731fd95
```



Distribute Artifacts Across OCI Registries With Ease

Previous

1. `oras pull` to pull/extract artifacts into a local disk path
2. Mount the local path to the pod

```
k8s>mkdir downloads
k8s>cd downloads/
k8s>oras pull release.daocloud.io/peter/llama-3-weights:v3.1-8b
✓ Pulled llama-3
  ↳ sha256:0b4d055901b51b24a62e3fcca15a60f100239803dd5c07fe49489b168f468261
✓ Pulled application/vnd.oci.image.manifest.v1+json
  ↳ sha256:e5145fa1529423daa154cbda91176b4a081bde51878087b783caa084731fd95
Pulled [registry] release.daocloud.io/peter/llama-3-weights:v3.1-8b
Digest: sha256:e5145fa1529423daa154cbda91176b4a081bde51878087b783caa084731fd95
k8s>ls
llama-3
```

Now - in Pod's OCI volume

```
spec:
  containers:
  ....
  volumes:
  - name: volume
    image:
      reference: release.daocloud.io/peter/llama-3-weight:v3.1-8b
```


6. Model Serving – Ollama



China 2024

Ollama - most simple LLM tool (backed by llama.cpp)

1. Basic Usage

```
# ollama run llama3.1
pulling manifest
pulling 8934d96d3f08... 100% 3.8 GB
pulling 170370233dd5... 100% 4.1 GB
....
>>> How is the weather today?
```

#ollama serve → serve as open-AI format API server

2. Modelfile ☺

```
FROM llama3

PARAMETER temperature 1
PARAMETER num_ctx 4096
assistant SYSTEM $system-prompt
```

#ollama create my-image -f Modelfile

#ollama run my-image

3. Operator in K8s

```
apiVersion: ollama.ayaka.io/v1
kind: Model
metadata:
  name: phi
spec:
  image: phi
  persistentVolume:
    accessMode: ReadWriteOnce
```

<https://github.com/nekomeowwww/ollama-operator>

● No More Runtime Setup! Let's Bundle, Distribute, Deploy, Scale LLMs Seamlessly with Ollama Operator | 无需运行时设置! 让我们使用Ollama Operator轻松捆绑、分发、部署、扩展LLMs - Fanshi Zhang, DaoCloud

Click here to add to My Schedule.

Friday August 23, 2024 15:00 - 15:35 HKT

Level 1 | Hung Hom Room 2

6. Model Serving – Lingo



China 2024

Lingo - LLM gateway



<https://github.com/substratusai/lingo>



Starred 103



- scale-to-zero & scale-from-zero
- Re-use same GPU for N models
- Lightweight , no dependency
- Queue, but no batch
- alternative to knative

6. Model Serving – Lingo



China 2024

1. Serve

- MiniLM-L6
- Mistral-7b

```
# kubectl get deployment
```

| NAME | READY | UP-TO-DATE | AVAILABLE | AGE |
|--------------------------|-------|------------|-----------|-----|
| mistral-7b-instruct-vllm | 0/0 | 0 | 0 | 21d |
| minilm-l6-v2 | 0/0 | 0 | 0 | 85d |

```
curl http://$IP:8080/v1/completions -d '{"model": "text-embedding-ada-002", "prompt": "..."}'
```

2. miniLM scale from 0

```
# kubectl get deployment
```

| NAME | READY | UP-TO-DATE | AVAILABLE | AGE |
|--------------------------|-------|------------|-----------|-----|
| mistral-7b-instruct-vllm | 0/0 | 0 | 0 | 21d |
| minilm-l6-v2 | 1/1 | 1 | 1 | 85d |

3. miniLM scale to 0

```
# kubectl get event
```

109s ... Scaled up replica set minilm-l6-v2-5dd7bb4fc8 to 1 from 0

a while later

15s ... Scaled down replica set minilm-l6-v2-5dd7bb4fc8 to 0 from 1

```
curl http://$IP:8080/v1/completions -d '{"model": "mistral-7b-instruct-v0.1", "prompt": "..."}'
```

4. Mistral-7b wake up

```
# kubectl get deploy
```

| NAME | READY | UP-TO-DATE | AVAILABLE | AGE |
|--------------------------|-------|------------|-----------|-----|
| mistral-7b-instruct-vllm | 1/1 | 1 | 1 | 21d |
| minilm-l6-v2 | 0/0 | 0 | 0 | 85d |

Thanks

- d.run
- HAMi
- JuiceFS
- HwameiStor
- kubeFlow
- dcgmi-diag
- Kueue
- Volcano
- OpenCost
- ORAS
- Triton
- Ollama
- Lingo

