



KubeCon



CloudNativeCon

THE LINUX FOUNDATION



AI_dev
Open Source GenAI & ML Summit

China 2024



KubeCon



CloudNativeCon



China 2024

Dragonfly: Intro, Updates and Ant Group's Practice of Accelerating Model Distribution in Ray Serving

Qixiang Chen @ AntGroup & Wenbo Qi(Gaius) @ AntGroup


Introduction




China 2024



Container Registry



Container Runtime



What is Dragonfly?

Provide efficient, stable and secure file distribution and image acceleration based on p2p technology to be the best practice and standard solution in cloud native architectures.

It is hosted by the Cloud Native Computing Foundation(CNCF)as an *Incubating* Level Project.

There are more than **100** contributors, and maintainers come from *Ant Group, Alibaba Group, ByteDance, Intel, Baidu Group, Zhipu AI* and *Dalian University of Technology*.

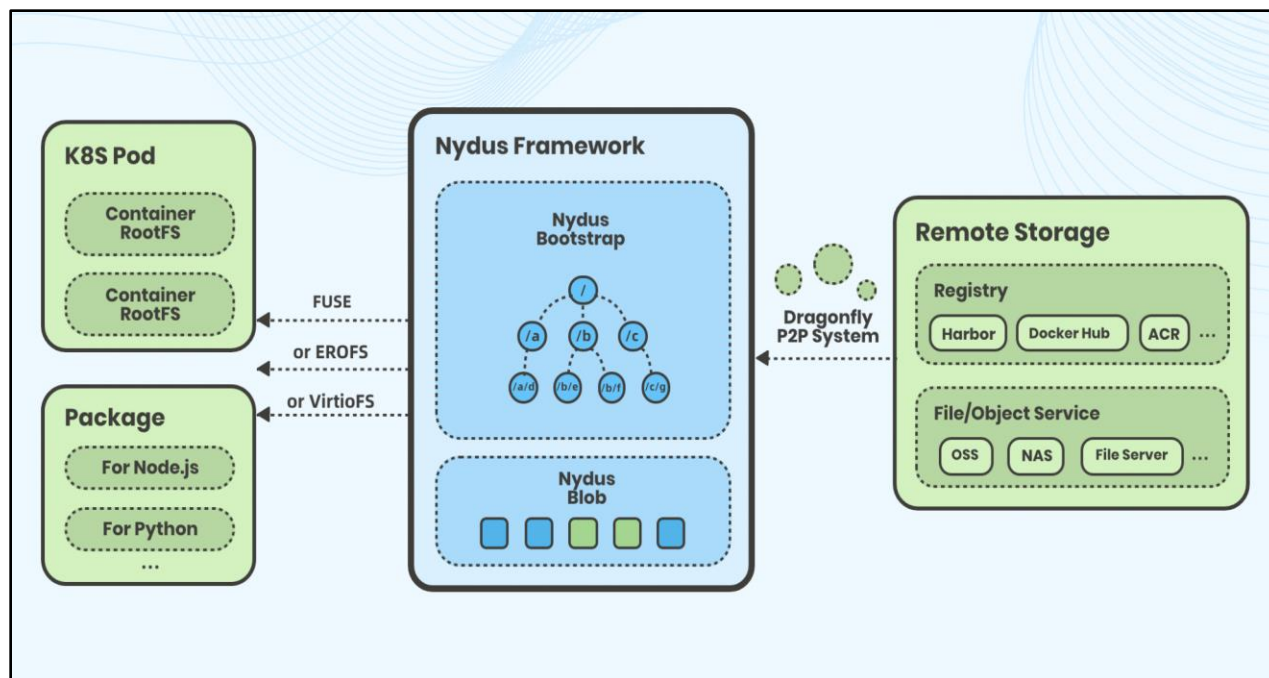
Public cloud users include *Alibaba Cloud(Aliyun), Google Cloud Platform (GCP), Volcano Engine, Baidu AI Cloud*, etc.



Introduction



China 2024



What is Nydus?

Provide a content-addressable file system on the RAFS format, which enhances the current OCI image specification by improving container launch speed, image space and network bandwidth efficiency, and data integrity.

It is a *sub-project of Dragonfly*. It can reduce the end-to-end cold start time of containers from *minutes to seconds*, and supports the creation of *millions of containers* every day in production.



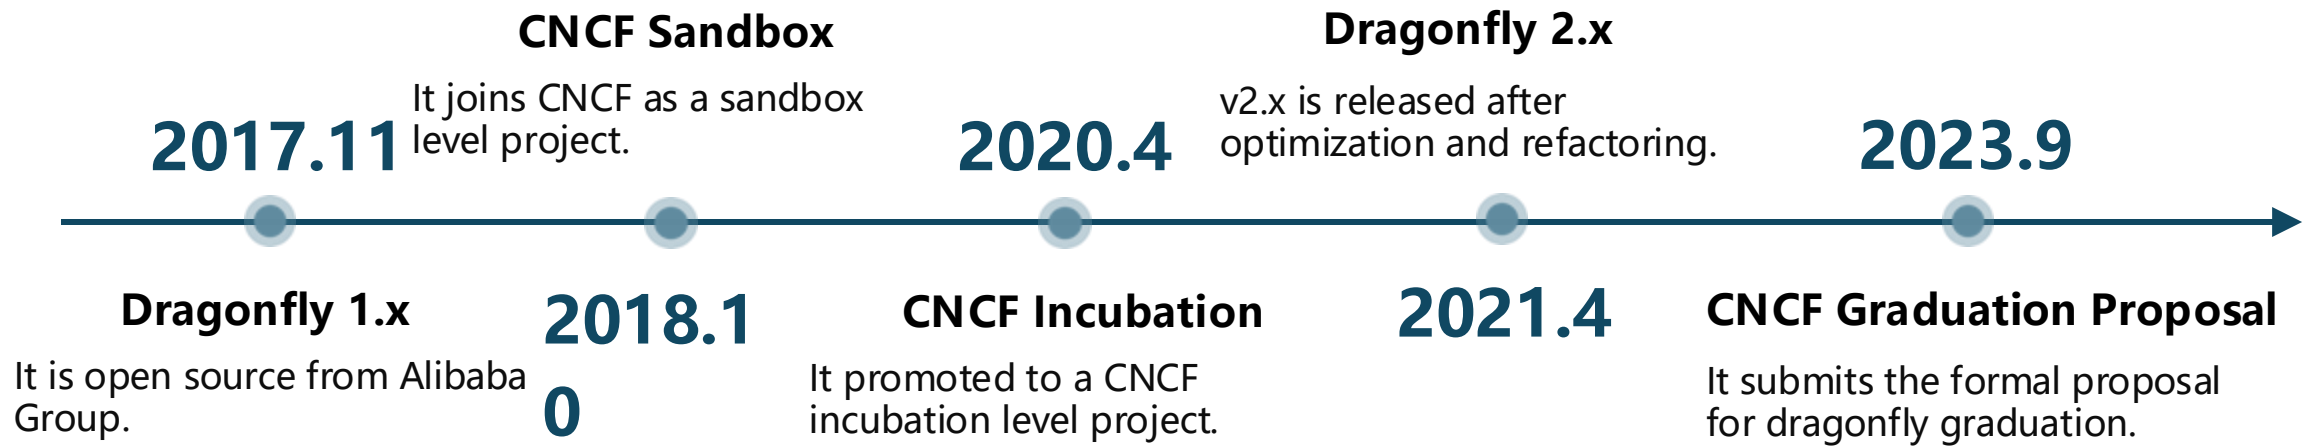
CLOUD NATIVE
COMPUTING FOUNDATION



Introduction

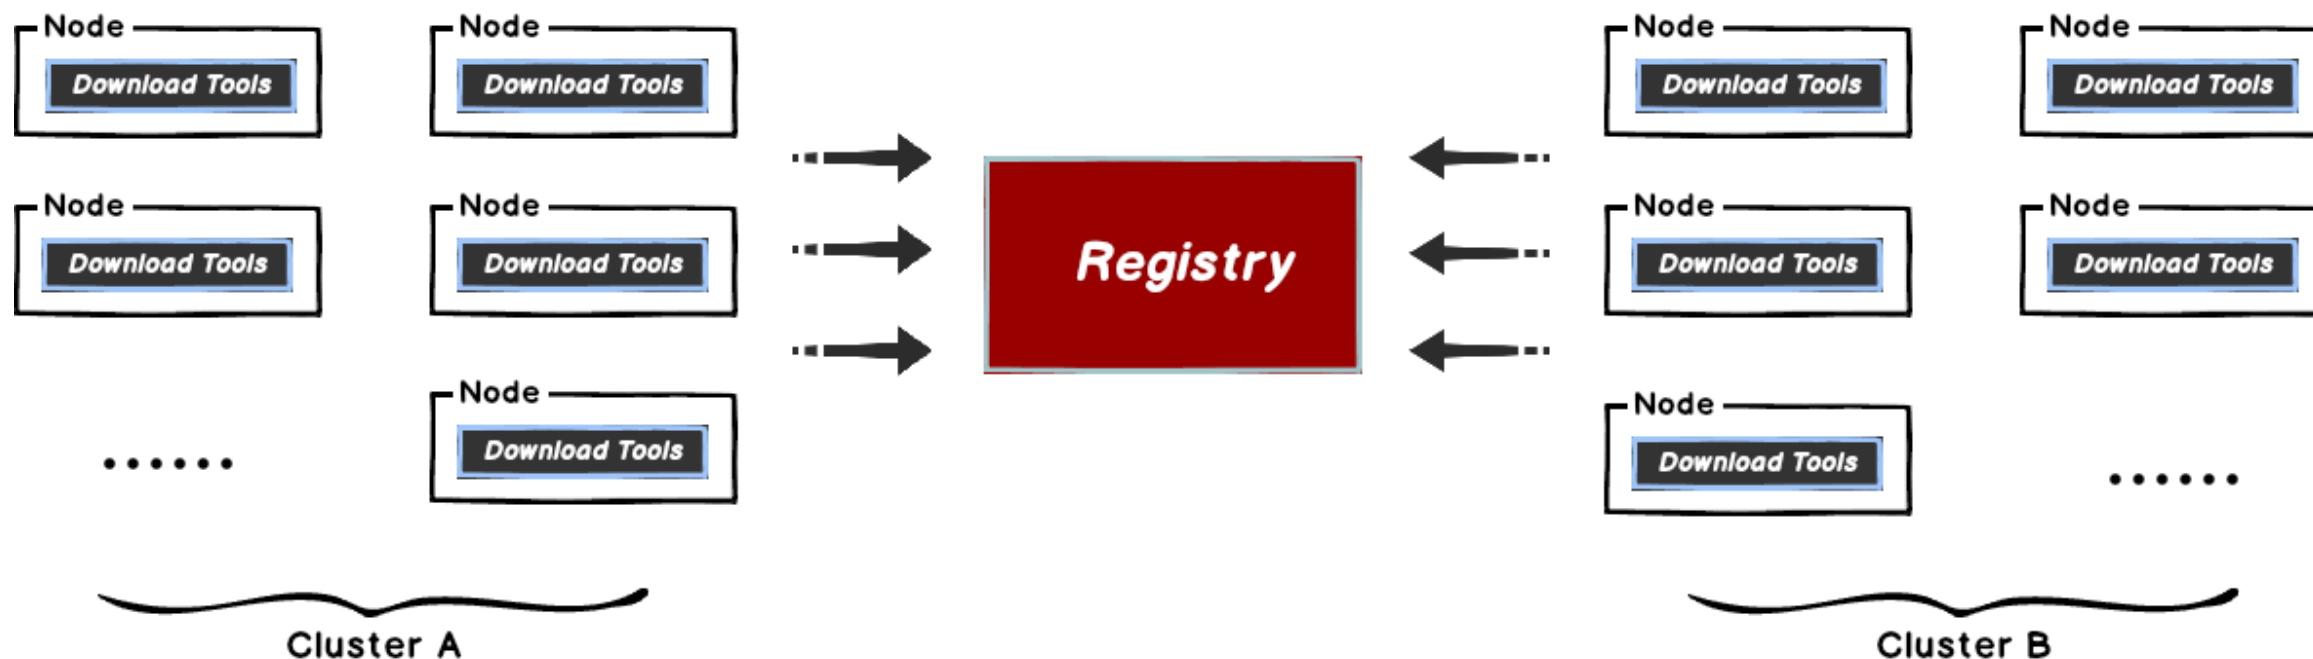


China 2024



Introduction

When downloading files in large batches, the storage bandwidth can easily reach the limit.



How to solve the problem?

**Increase storage
bandwidth**



**P2P eliminates the
impact of bandwidth
limitations**



**Reduce download
size**

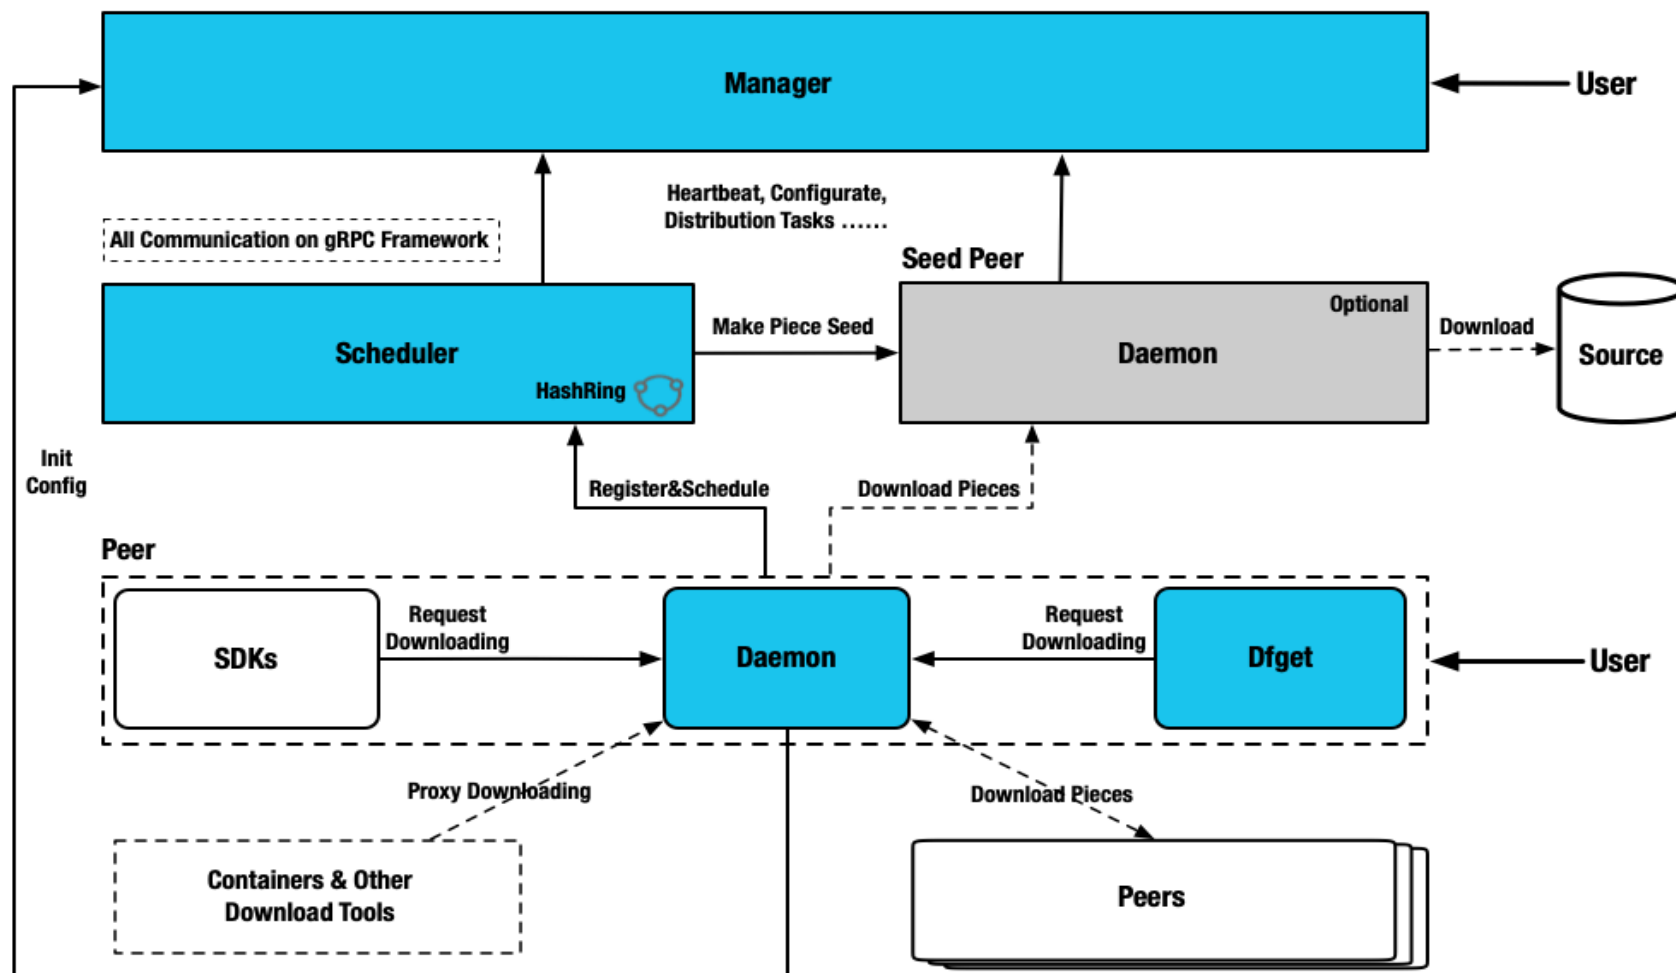


Introduction

Manager manages cluster relationships, dynamic configuration and provides a console.

Scheduler schedules a set of parent nodes for the download node.

Peer provides upload and download capabilities.



Exciting Updates

Document

Dragonfly

Documentation

Blog

Video

Community

Nydux

Introduction

Getting Started

Quick Start

Kubernetes

Multi-cluster Kubernetes

Installation

Helm Charts

Binaries

Operations

Best Practices

Deployment Best Practices

Observability

Security

Deployment

Architecture

Applications

Integrations

Reference

Configuration

Commands

Advanced Guides

Web Console

Preheat

Personal Access Tokens

Next

English

Search

This is unreleased documentation for Dragonfly **Next** version.

For up-to-date documentation, see the **latest version** (v2.1.x).

Getting Started

Quick Start

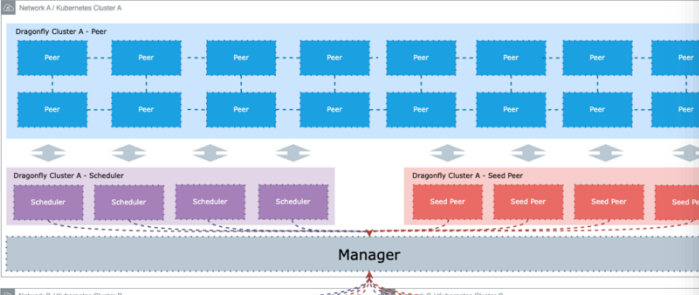
Multi-cluster Kubernetes

Version: Next

Multi-cluster Kubernetes

Documentation for deploying Dragonfly on multi-cluster kubernetes using helm. A Dragonfly cluster manages cluster within : network. If you have two clusters with disconnected networks, you can use two Dragonfly clusters to manage their own cluster.

The recommended deployment in a multi-cluster kubernetes is to use a Dragonfly cluster to manage a kubernetes cluster, a centralized manager service to manage multiple Dragonfly clusters. Because peer can only transmit data in its own Dragonfly cluster, if a kubernetes cluster deploys a Dragonfly cluster, then a kubernetes cluster forms a p2p network, and internal peer only schedule and transmit data in a kubernetes cluster.



Runtime

Setup kubernetes cluster

Kind loads Dragonfly image

Create Dragonfly cluster A

Create Dragonfly cluster A based on helm charts

Create NodePort service of the manager REST service

Visit manager console

Dragonfly

Documentation

Blog

Video

Community

Nydux

Introduction

Getting Started

Quick Start

Kubernetes

Multi-cluster Kubernetes

Installation

Helm Charts

Binaries

Operations

Best Practices

Deployment Best Practices

Observability

Security

Deployment

Architecture

Applications

Integrations

Reference

Configuration

Commands

Advanced Guides

Web Console

Preheat

Personal Access Tokens

Development Guides

Next

English

Search

This is unreleased documentation for Dragonfly **Next** version.

For up-to-date documentation, see the **latest version** (v2.1.x).

Operations

Best Practices

Deployment Best Practices

Version: Next

Deployment Best Practices

Documentation for setting capacity planning and performance tuning for Dragonfly.

Capacity Planning

A big factor in planning capacity is: highest expected storage capacity. And know the memory size, CPU core count, and disk capacity of each machine.

For predicting your capacity, you can use the estimates from below if you don't have your capacity plan.

Manager

The resources required to deploy the Manager depends on the total number of peers.

Run a minimum of 3 replicas.

Total Number of Peers	CPU	Memory	Disk
1K	8C	16G	200Gi
5K	16C	32G	200Gi
10K	16C	64G	200Gi

Capacity Planning

Manager

Scheduler

Client

Cluster

Performance tuning

Rate limits

Concurrency control

GC

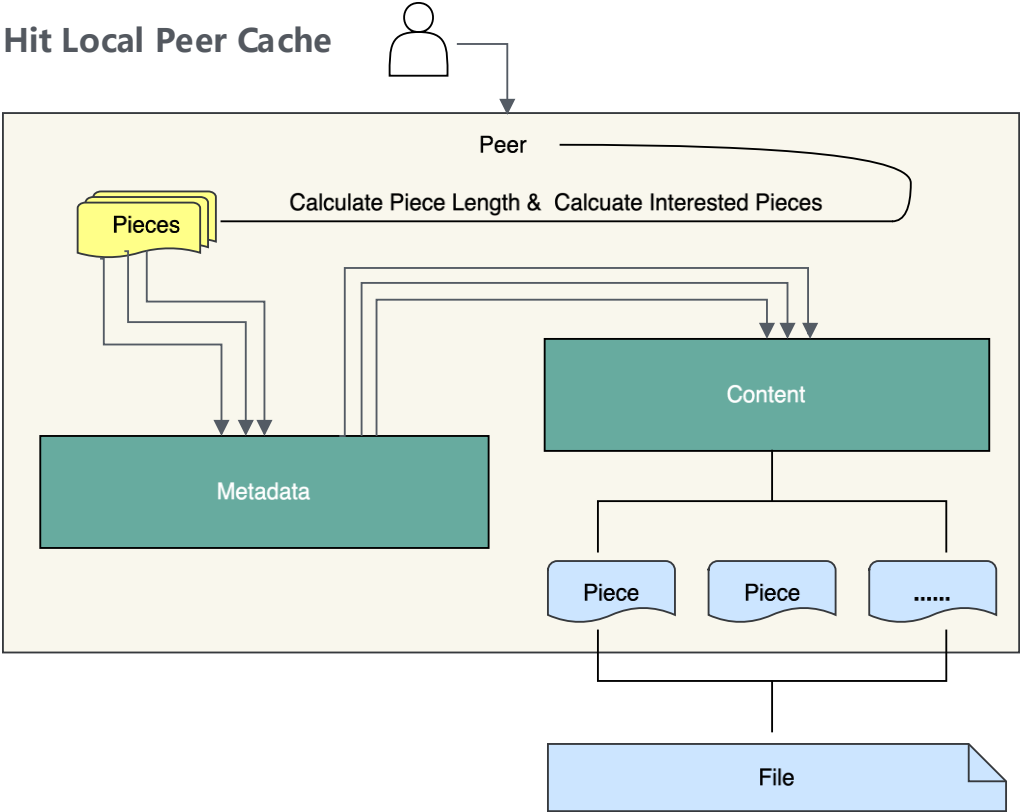
Nydux


Exciting Updates



China 2024

Rust Client



 client

Public

main

4 Branches


98 Tags

Go to file

t

Add file


















<> Code

 Liam-Zhao

Feat: add tar package build (#659)

175be38 · 6 hours ago

591 Commits

 .github	Feat: add tar package build (#659)	6 hours ago
 ci	Feat: dfdaemon is started using systemd (#648)	4 days ago
 dragonfly-client-backend	chore(deps): Bump libloading from 0.8.4 to 0.8.5 (#634)	last week
 dragonfly-client-config	feat: when send to scheduler failed, join_set detach all ta...	19 hours ago
 dragonfly-client-core	chore(deps): Bump libloading from 0.8.4 to 0.8.5 (#634)	last week
 dragonfly-client-init	feat: remove header in mirror config	2 days ago
 dragonfly-client-storage	feat: add log for deleting piece (#656)	3 days ago
 dragonfly-client-util	feat: support crc32 for calculating piece content (#649)	4 days ago
 dragonfly-client	feat: when send to scheduler failed, join_set detach all ta...	19 hours ago
 .gitignore	chore: change gitignore for vendor (#415)	4 months ago
 CONTRIBUTING.md	docs: fix typo in README.md (#381)	4 months ago
 Cargo.lock	feat: when send to scheduler failed, join_set detach all ta...	19 hours ago
 Cargo.toml	feat: when send to scheduler failed, join_set detach all ta...	19 hours ago
 LICENSE	Initial commit	last year
 README.md	Add license scan report and status (#611)	3 weeks ago
 codecov.yml	chore: add codecov for testing (#3)	last year
 rust-toolchain.toml	feat: add request timeout for connection (#627)	2 weeks ago

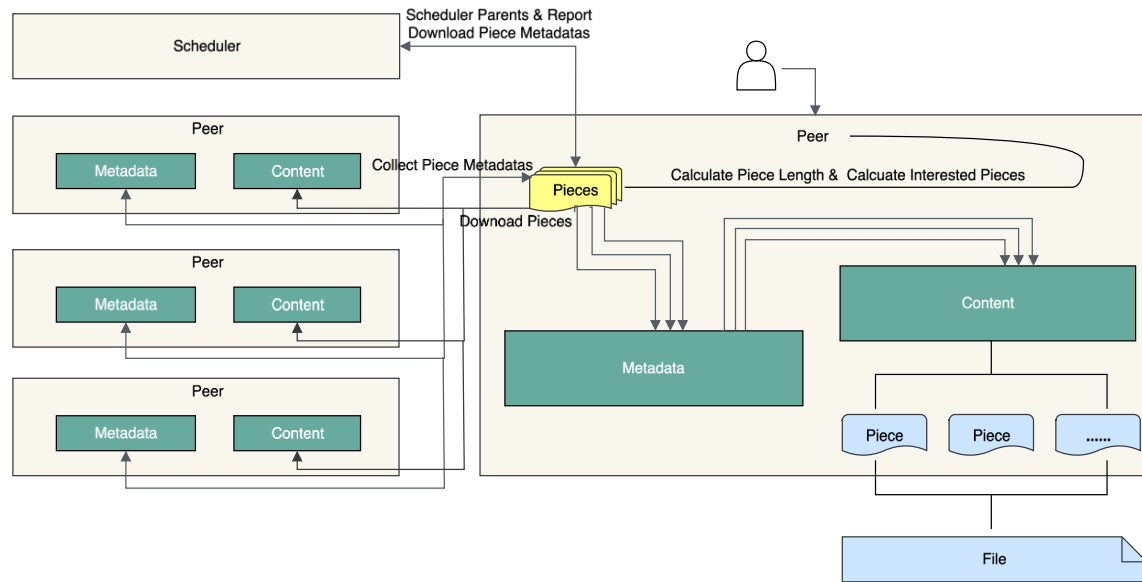
Exciting Updates



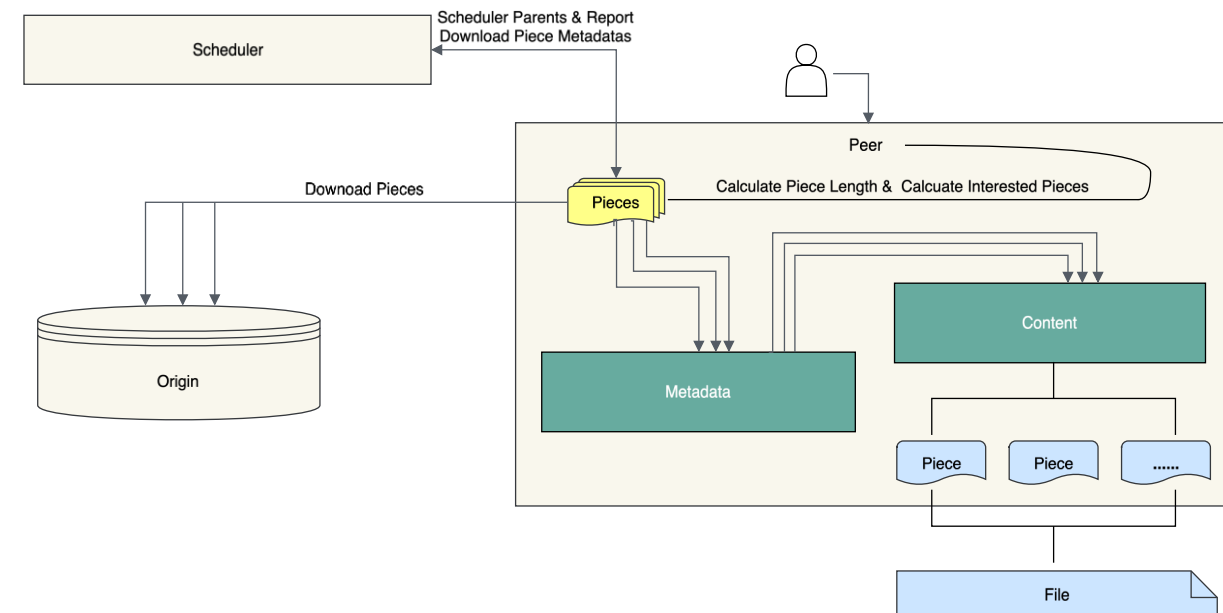
China 2024

Rust Client


Hit Remote Peer Cache



Download From Origin



Exciting Updates



China 2024

Kubernetes 1.31: Read Only Volumes Based On OCI Artifacts (alpha)

By **Sascha Grunert** | Friday, August 16, 2024

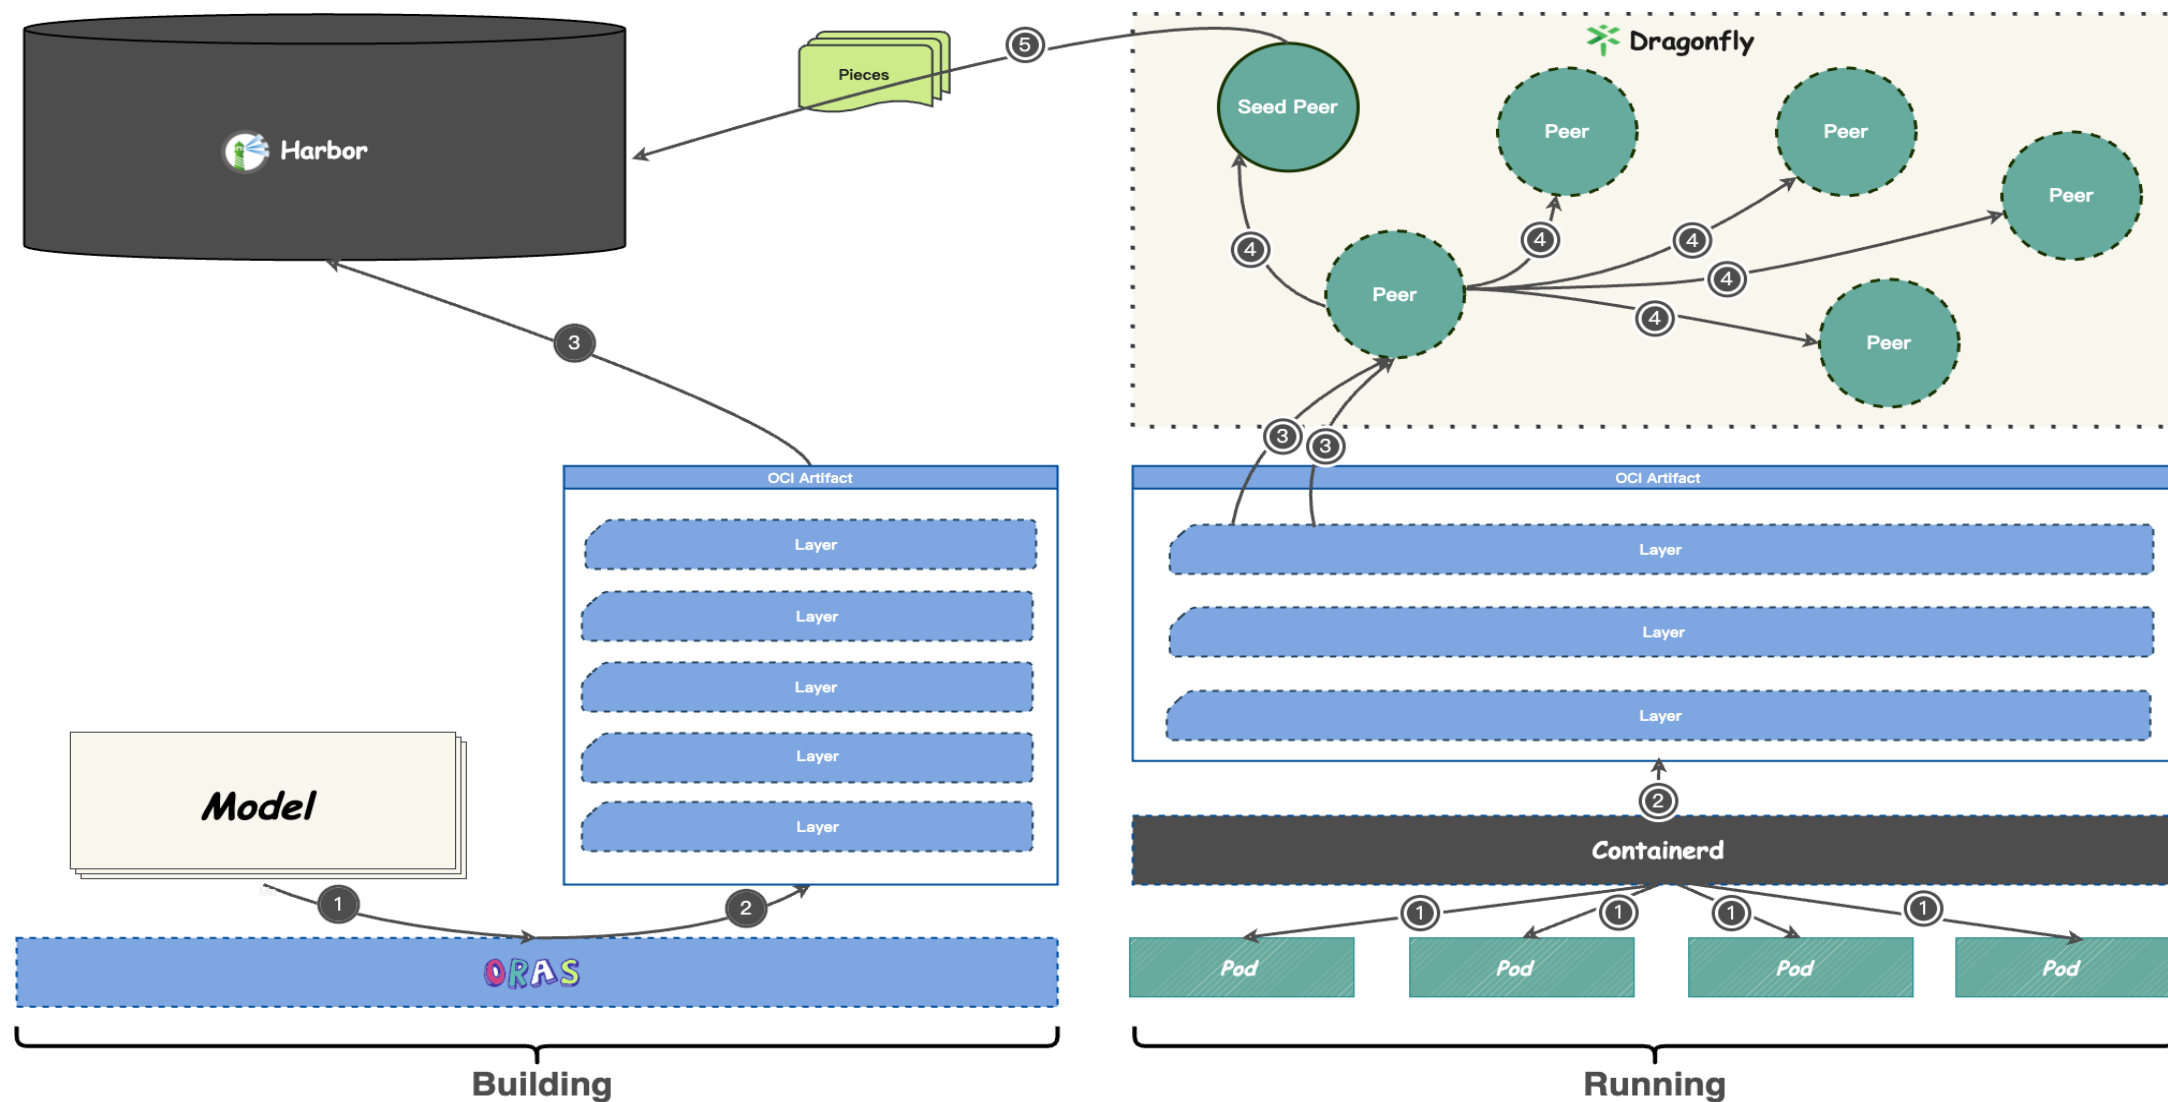
The Kubernetes community is moving towards fulfilling more Artificial Intelligence (AI) and Machine Learning (ML) use cases in the future. While the project has been designed to fulfill microservice architectures in the past, it's now time to listen to the end users and introduce features which have a stronger focus on AI/ML.

One of these requirements is to support [Open Container Initiative \(OCI\)](#) compatible images and artifacts (referred as OCI objects) directly as a native volume source. This allows users to focus on OCI standards as well as enables them to store and distribute any content using OCI registries. A feature like this gives the Kubernetes project a chance to grow into use cases which go beyond running particular images.

Given that, the Kubernetes community is proud to present a new alpha feature introduced in v1.31: The Image Volume Source ([KEP-4639](#)). This feature allows users to specify an image reference as volume in a pod while reusing it as volume mount within containers:

```
apiVersion: v1
kind: Pod
metadata:
  name: pod
spec:
  containers:
    - name: test
      image: registry.k8s.io/e2e-test-images/echoserver:2.3
      volumeMounts:
        - name: model
          mountPath: /model
  volumes:
    - name: model
      image:
        reference: harbor.example.com/example/llama:v1.0.0
        pullPolicy: IfNotPresent
```

Exciting Updates



Introduction



China 2024

What is Ray?

An open source framework to build and scale your ML and python APP easily

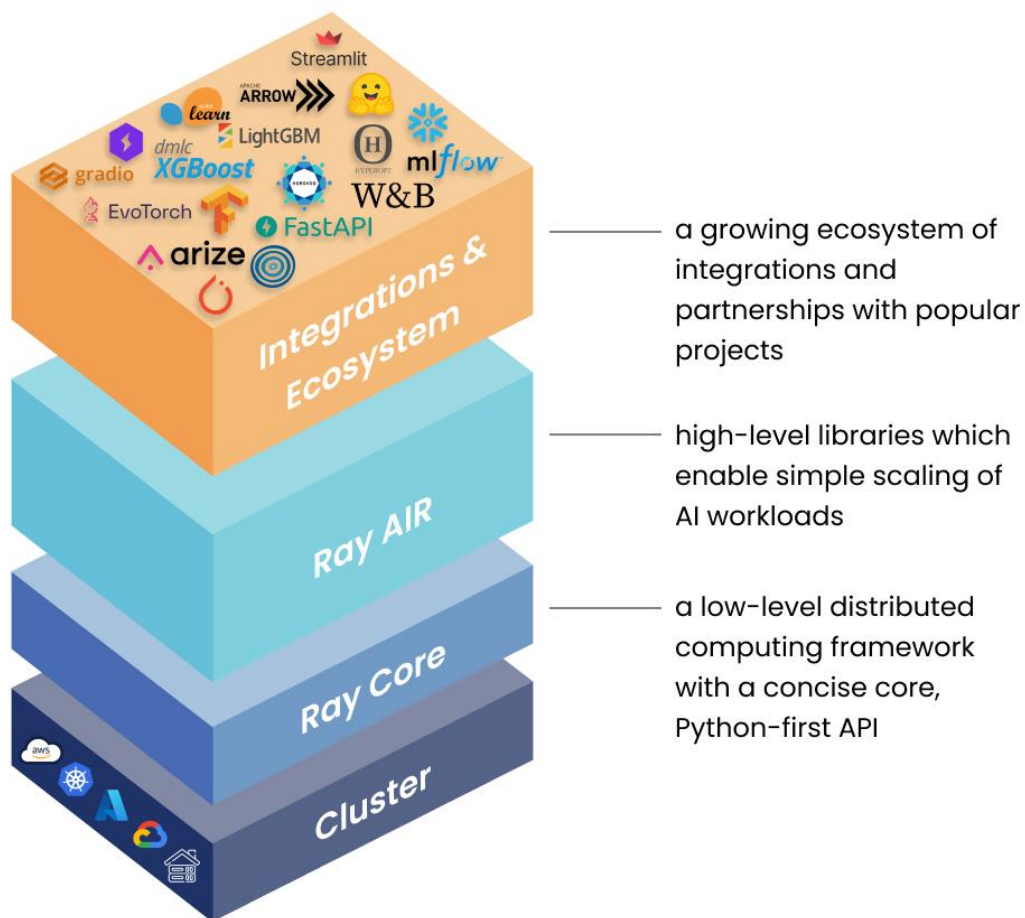
Originated from RISELab, 32.6k github star, supporting **100w+** CPU core in AntGroup

RayCore provides distributing primitives. **RayAIR** provides high-level library in AI workloads, e.g. training, parameter tuning, model serving, batch processing. A rich and growing **ecosystem** of integrations.

OpenAI

"At OpenAI, we are tackling some of the world's most complex and demanding computational problems. Ray powers our solutions to the thorniest of these problems and allows us to iterate at scale much faster than we could before."

Greg Brockman, CTO and cofounder, OpenAI



Introduction



China 2024

What capabilities a typical distributed system have?



Component
RPC



Data
Storage



Task
Scheduling



Deployment



Failover



Monitor
&
Maintenance

RPC How nodes are communicated. grpc or brpc ? FlatBuffer or ProtoBuffer?

Data Storage What flows between nodes. Memory, Disk, Cloud ? Row-wise, column wise? Schema?

Task scheduling Making the distribution worthy. Yarn, K8S?


Deployment Let the system physically run

Failover The necessary part for a distributed system

Monitor Important for complex system

These are what Ray ***precisely*** provides

Architecture overview



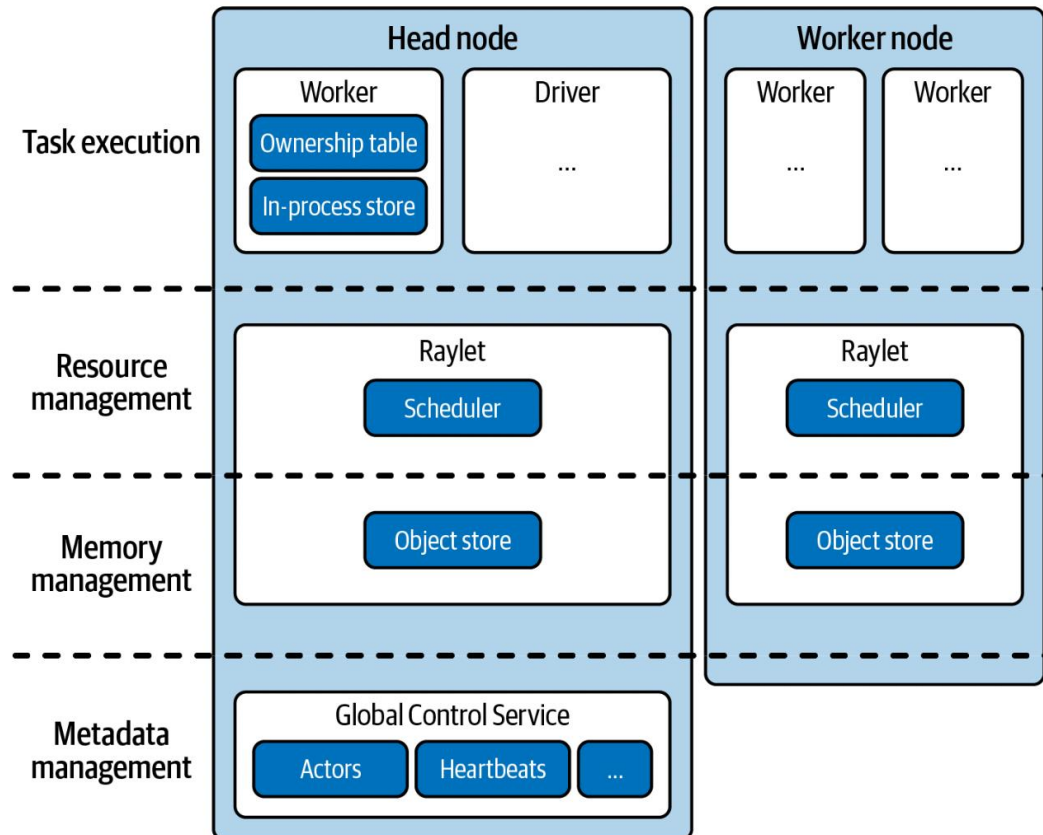
Head node: global state, top-level scheduling, job submission, resource management

Worker node: second-level scheduling, task execution.

- *Scheduler*: task scheduling & failover
- *Object Store*: data storage and transfer
- *Worker*: task execution

Introduction

Ray Architecture




Head node: global state, top-level scheduling, job submission, resource management

Worker node: second-level scheduling, task execution.

- *Scheduler*: second-level scheduling
- *Object Store*: data storage and transfer
- *Worker*: task execution

Introduction



China 2024

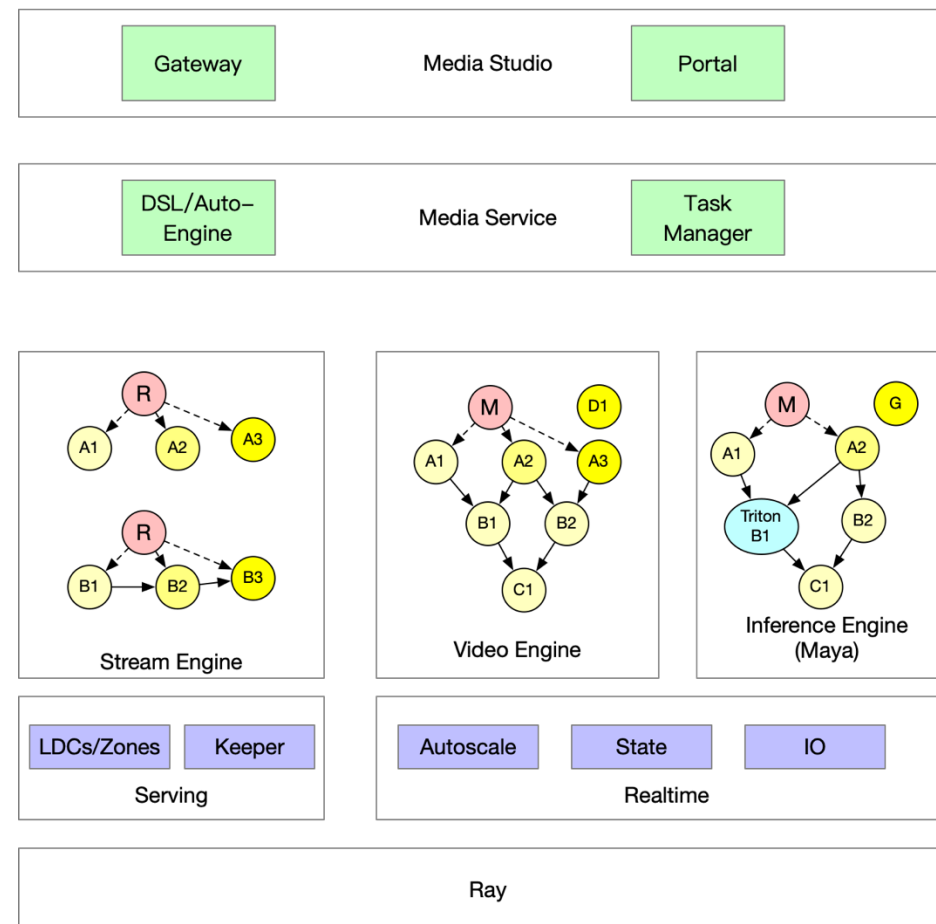
Media file processing

Stream Engine for live stream,
using RayServe. *Online*

Video Engine for media file processing,
e.g. IO, render, decode, APTS. *Near
online*

Inference Engine for multi-model DL,
feature extraction, high-performance
inference. *Offline*

The data are all disk files




Motivation

Remote large file download Media file 1MB ~ 500MB,
Model 1GB~20GB, cluster-level distribution

Directory download Model is actually a directory with
model states, structure, config, etc.

File life-cycle management File is either large or frequently
distributed to cluster, piling and leaking are not allowed

Natural failover The workflow may time consuming,
better not start over when Pod failed



- The whole copy is transferred even though chunked
- 1-to-1 even though other nodes may have the same copy

RayFile - Ray & Dragonfly

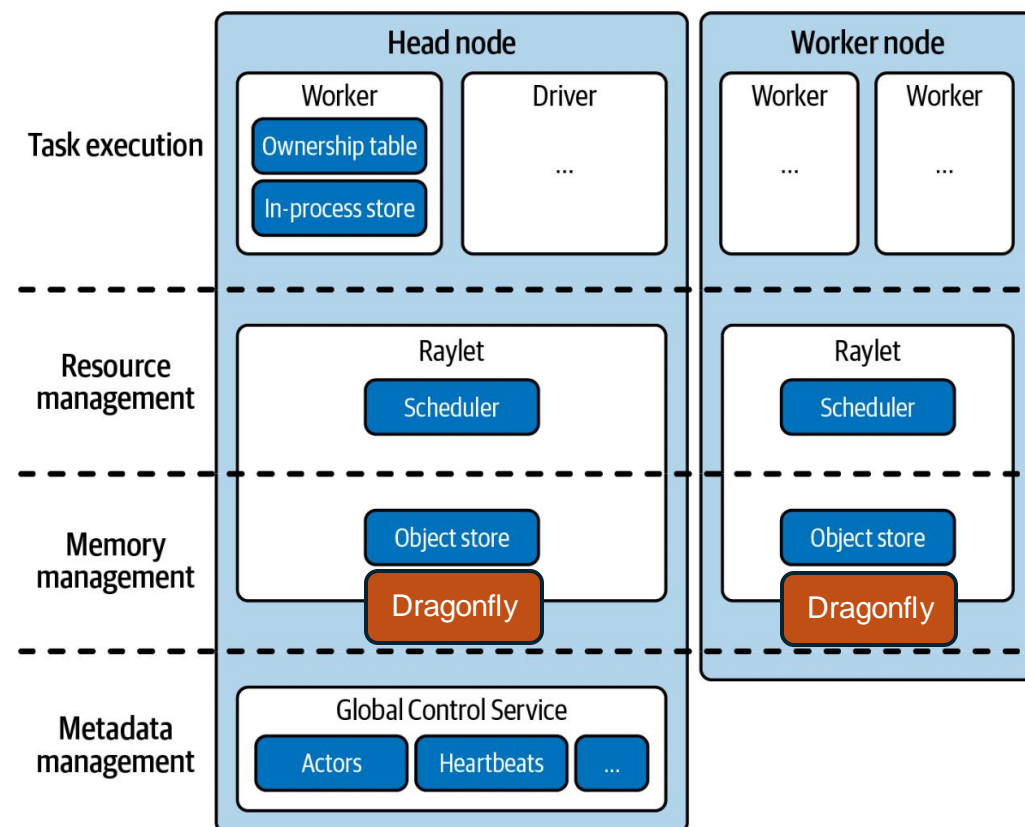
P2P transfer Data(File) chunk is distributed in multiple nodes

Native API Adapt dragonfly SDK to Ray native API

Hierarchical Cache Worker -> Job -> Ray node -> D7Y peer


GC protocol File's lifecycle follows the same GC protocol with Ray Object

File upload local disk file, e.g. ffmpeg output file, can be uploaded to the Ray(d7y) cluster



Dragonfly v.s. ObjectStore

The large file size or more number of copy is, the more efficient D7Y is, max to *10x* bandwidth boost



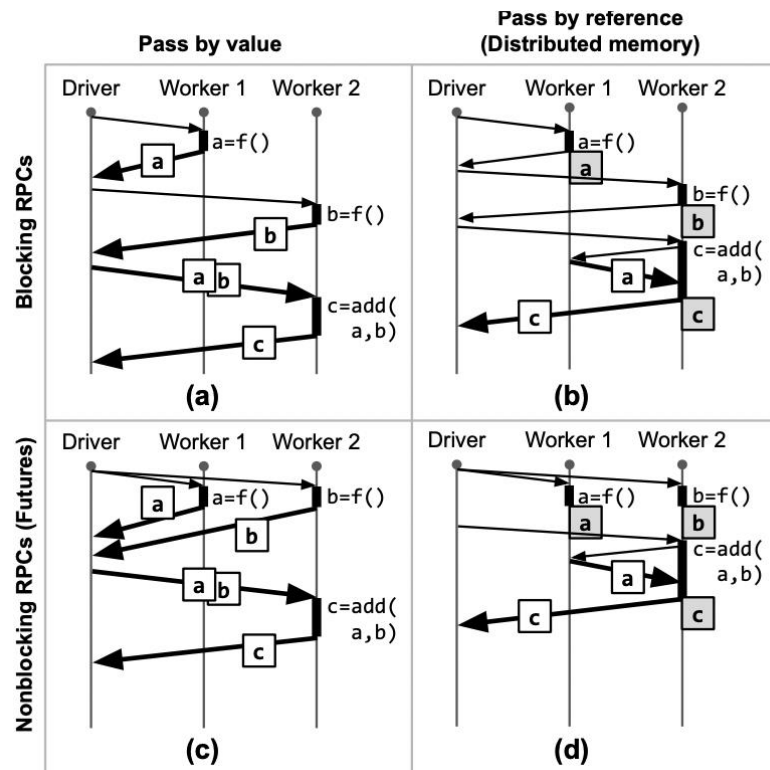
Stress test – Audio extraction scenario

File size(MB)	File amount	avg. bandwidth(MB/s)	P50 latency(s)	P90 latency(s)	P99 latency(s)
0~1	4814	35.75	0.14	1.04	30.82
16~32	2208	402.33	16.31	31.96	48.4
256-512	159	251.71	29.97	55.92	57.02
512-1024	27	114.16	49.26	50.27	50.39
1024-4096	7	170.26	49.27	49.57	49.57

Download latency is *stable* even file size grows from 512MB to 4GB ! But *unstable* when <64MB

RayFile/Dragonfly is incline to *large file* from both bandwidth and lantency perspective.

Remote Distributed Future



In Ray, the data that flows between tasks are just the *future reference* whose value are resolved in daemon. A task will only be executed by the time its input data's future is resolved

RayFile can potentially implement the mechanism since d7y is also transferring data in daemon and provides future reference, e.g. *cid*



KubeCon



CloudNativeCon



China 2024

Thanks!



Dragonfly