



**KubeCon**



**CloudNativeCon**

THE LINUX FOUNDATION



**AI\_dev**  
Open Source GenAI & ML Summit

---

**China 2024**

---



KubeCon



CloudNativeCon



China 2024

# Inplace-update: the past, present and future

Zhen Zhang & Yuxin Yuan Alibaba Cloud & OpenKruise maintainer



# Inplace-update: the origin



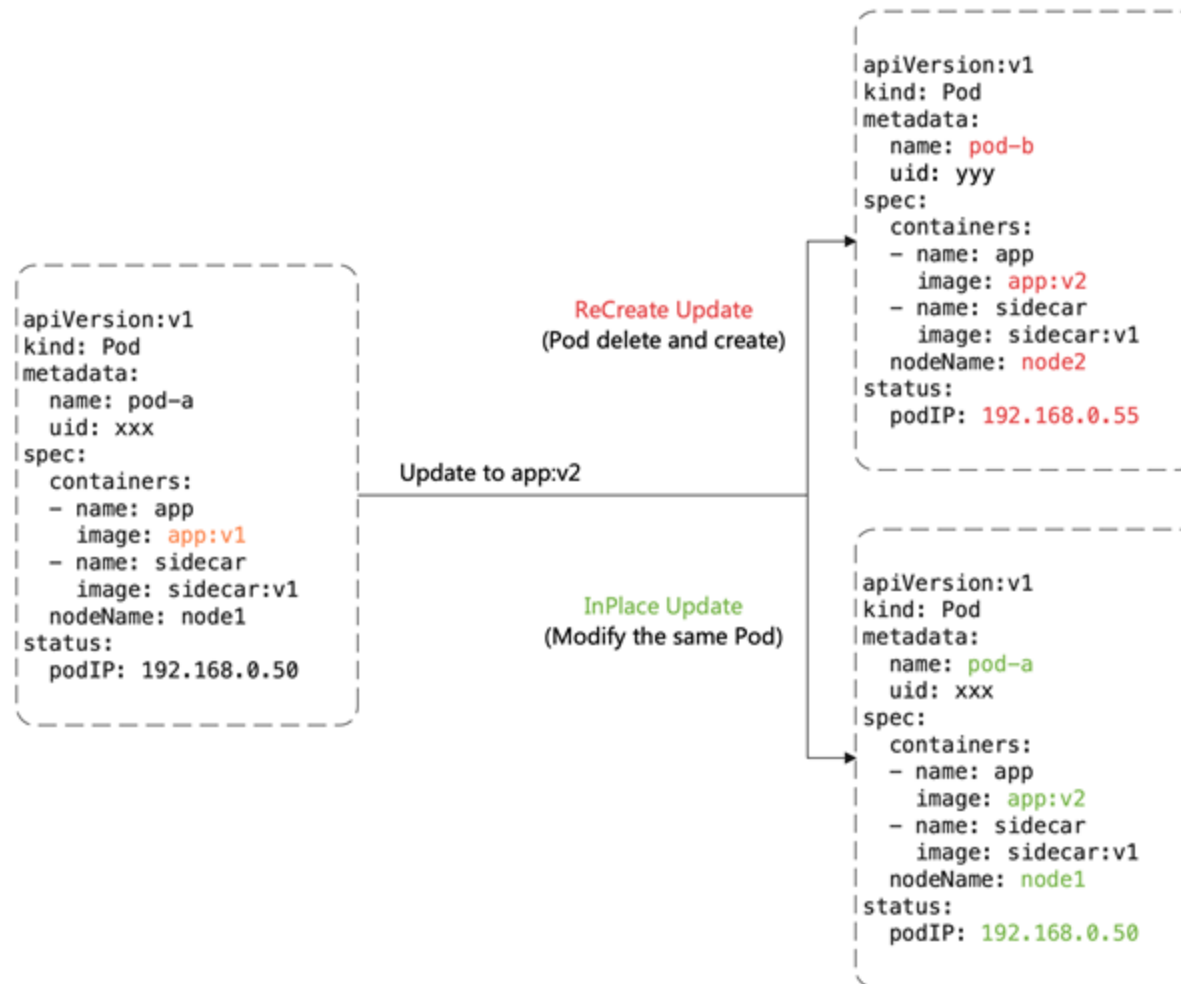
China 2024



1. Cluster is locked after the load test
2. Resource is highly consolidated
3. Image of rich container is huge
4. Service discovery of large scale pod recreate is challenging

Recreate rolling update is not an option !  
And we are not alone

# In-place Update Concept



- **Avoid additional cost** of scheduling, allocating IP, allocating and mounting volumes, load volumes with data
- **Faster image pulling**, because of we can re-use most of image layers pulled by the old image and only to pull several new layers
- When a container is in-place updating, **the other containers in Pod will not be affected and remain running.**

Upstream [KEP: In-place rolling update](#)

# Inplace-update & Immutable infra.



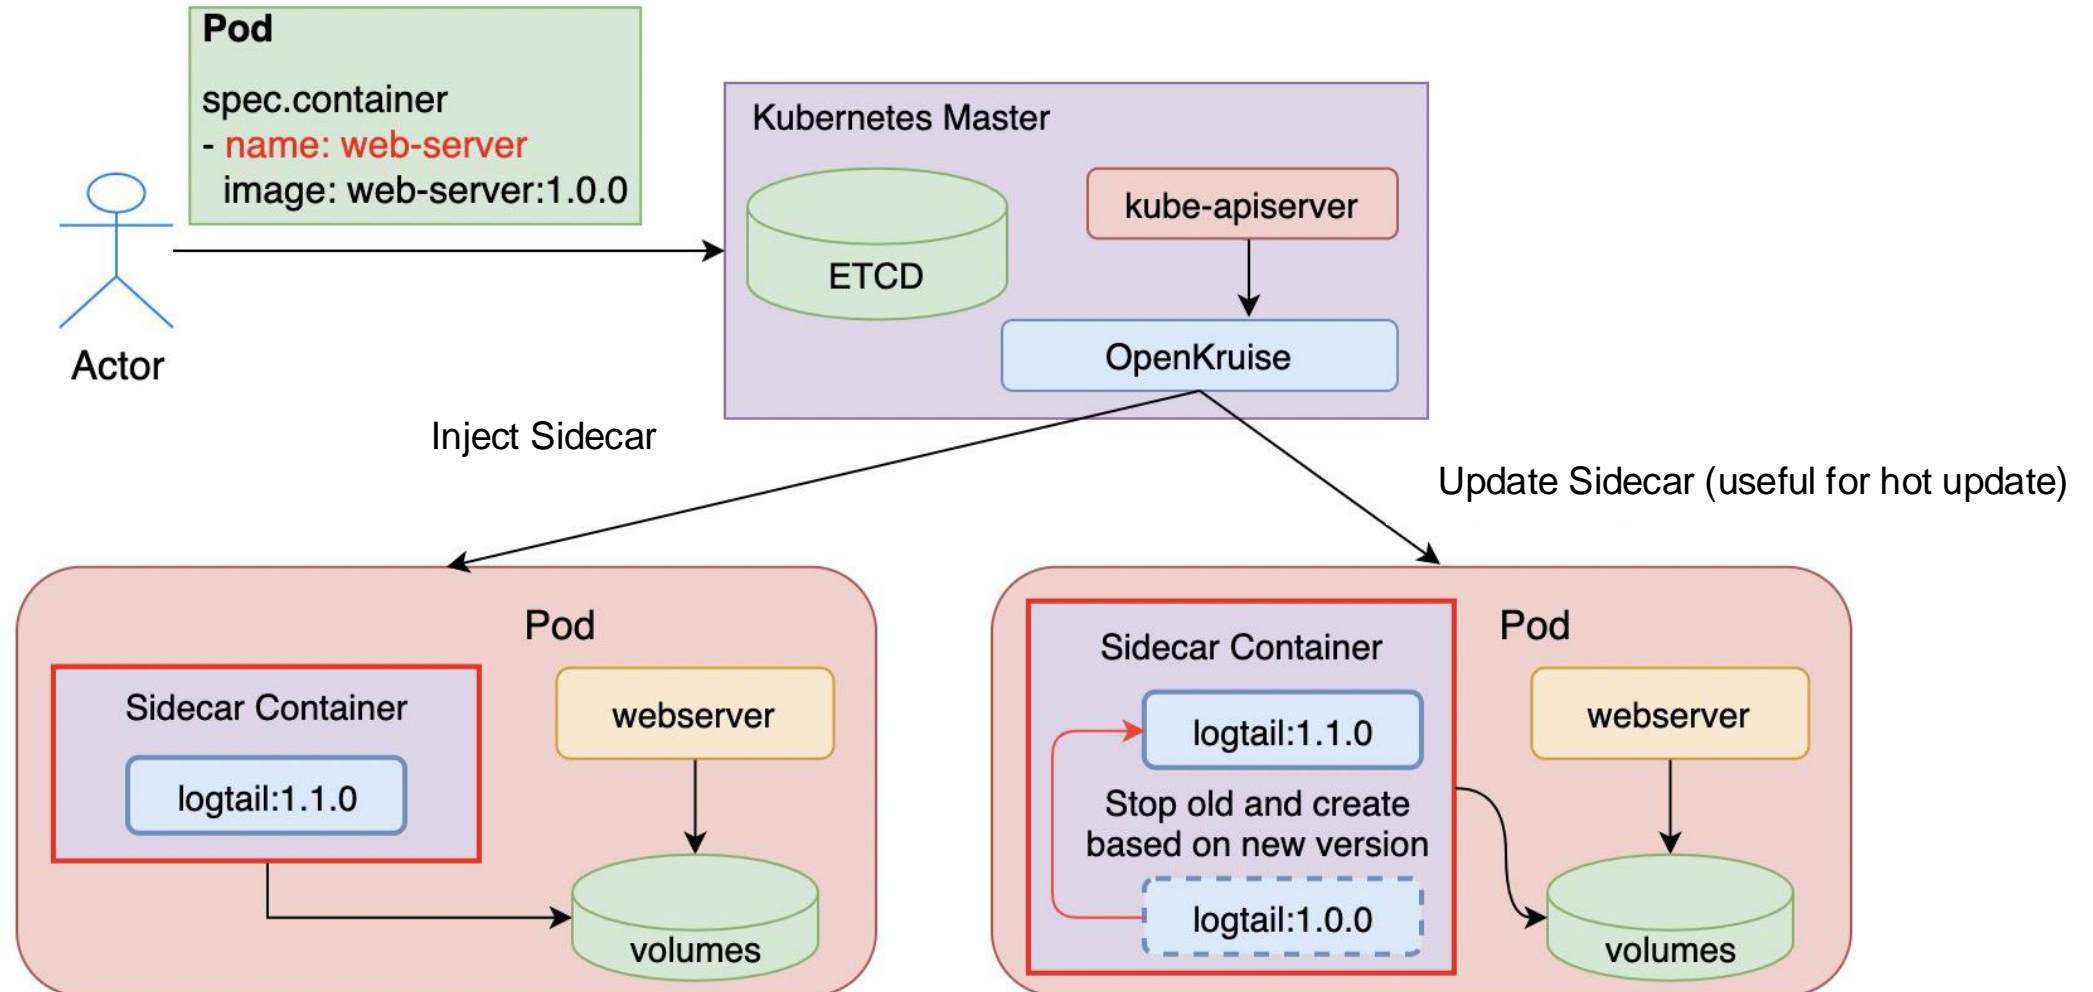
China 2024

- Immutable Infra with replaceable part
  - ~~Pod~~ Container as the element of Immutable infra
  - Only inplace-update if possible
  - Limit the replaceable part to:
    - Image
    - Resource

```
apiVersion: apps.kruise.io/v1alpha1
kind: CloneSet
spec:
  # ...
  updateStrategy:
    type: InPlaceIfPossible
    inplaceUpdateStrategy:
      gracePeriodSeconds: 10
```



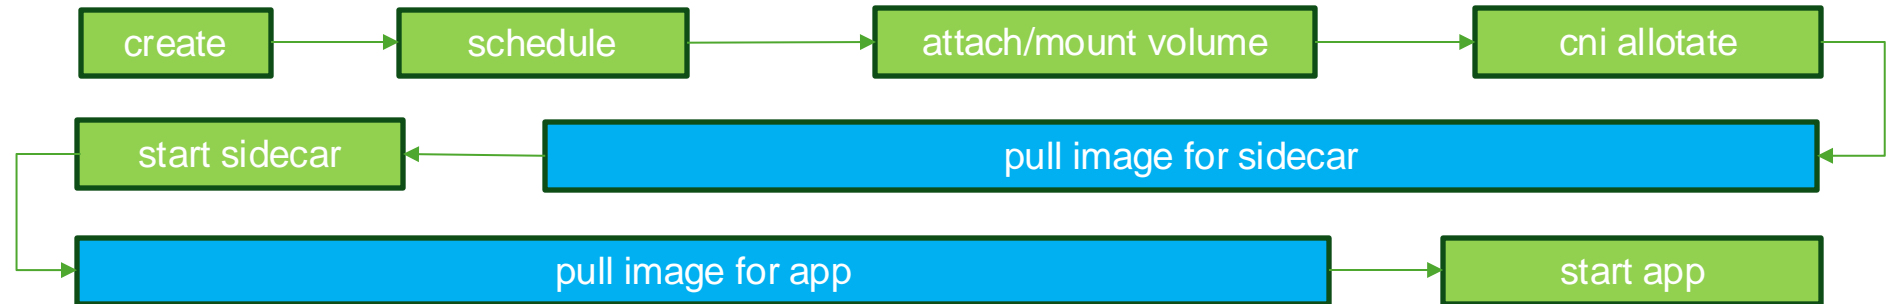
# Use case – Sidecar management






# Use case: Image pre-download

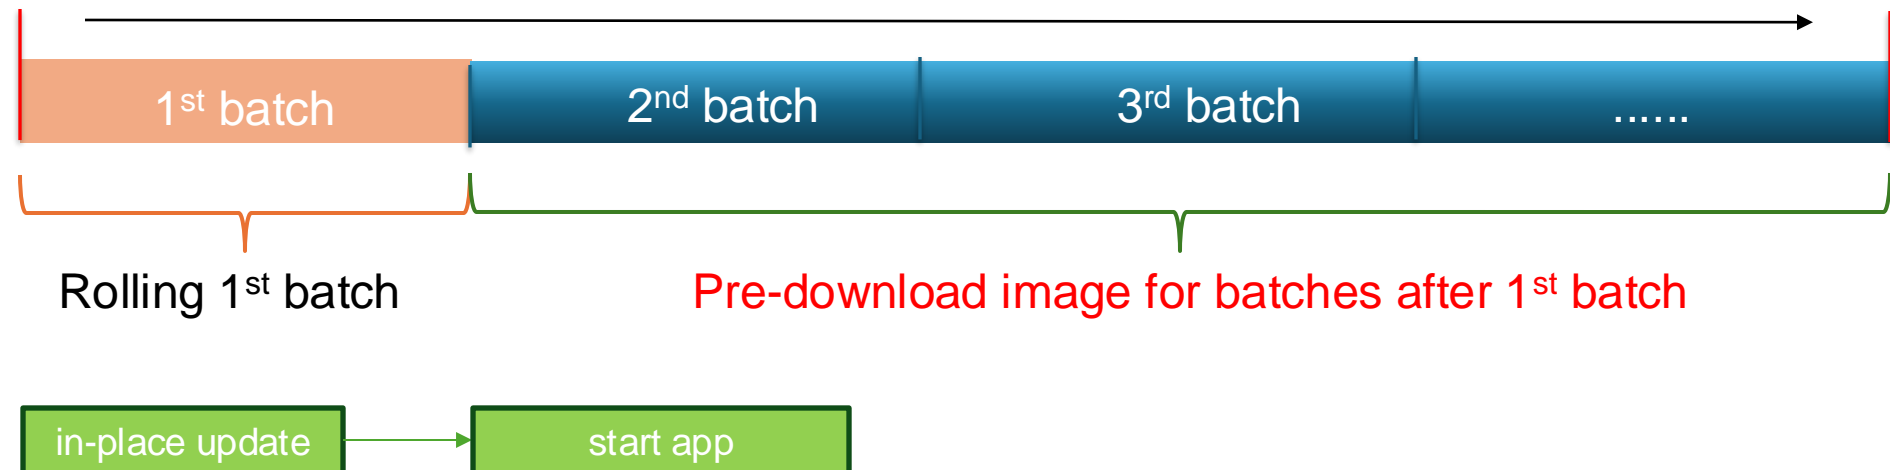
(1) Vanilla Pod creation :



(2) Pod creation with image pre-download



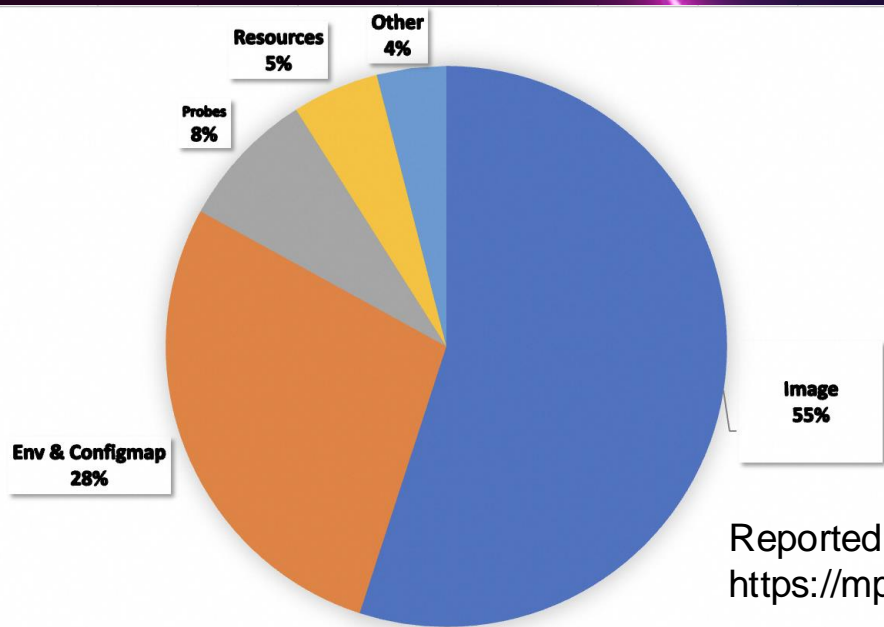
(3) inplace-update with Pre-download



# challenges and solution: updatable fields



China 2024



Reported by one user:  
[https://mp.weixin.qq.com/s/hRvZz\\_bZfchmP0tkF6M2OA](https://mp.weixin.qq.com/s/hRvZz_bZfchmP0tkF6M2OA)

fields	Mutable	Reported	Trigger update
Image	Y	Y	Y
Annotation	Y	N	N
Environment Command Argument	N	N	N
Resource	Y(since 1.27)	Y	Y

Immutable:

- 1. Change Apiserver 🐱
- 2. Wire with downward API 😊

Not trigger update

- 1. Update along with image
- 2. Kruise-daemon trigger update

Limit the updatable fields is **by design**

```
apiVersion: apps.kruise.io/v1alpha1
kind: CloneSet
spec:
  template:
    metadata:
      annotations:
        app-config: "... the real env value ..."
    spec:
      containers:
        - name: app
          image: app-image:v1
          env:
            - name: APP_CONFIG
              valueFrom:
                fieldRef:
                  fieldPath: metadata.annotations['app-config']
```

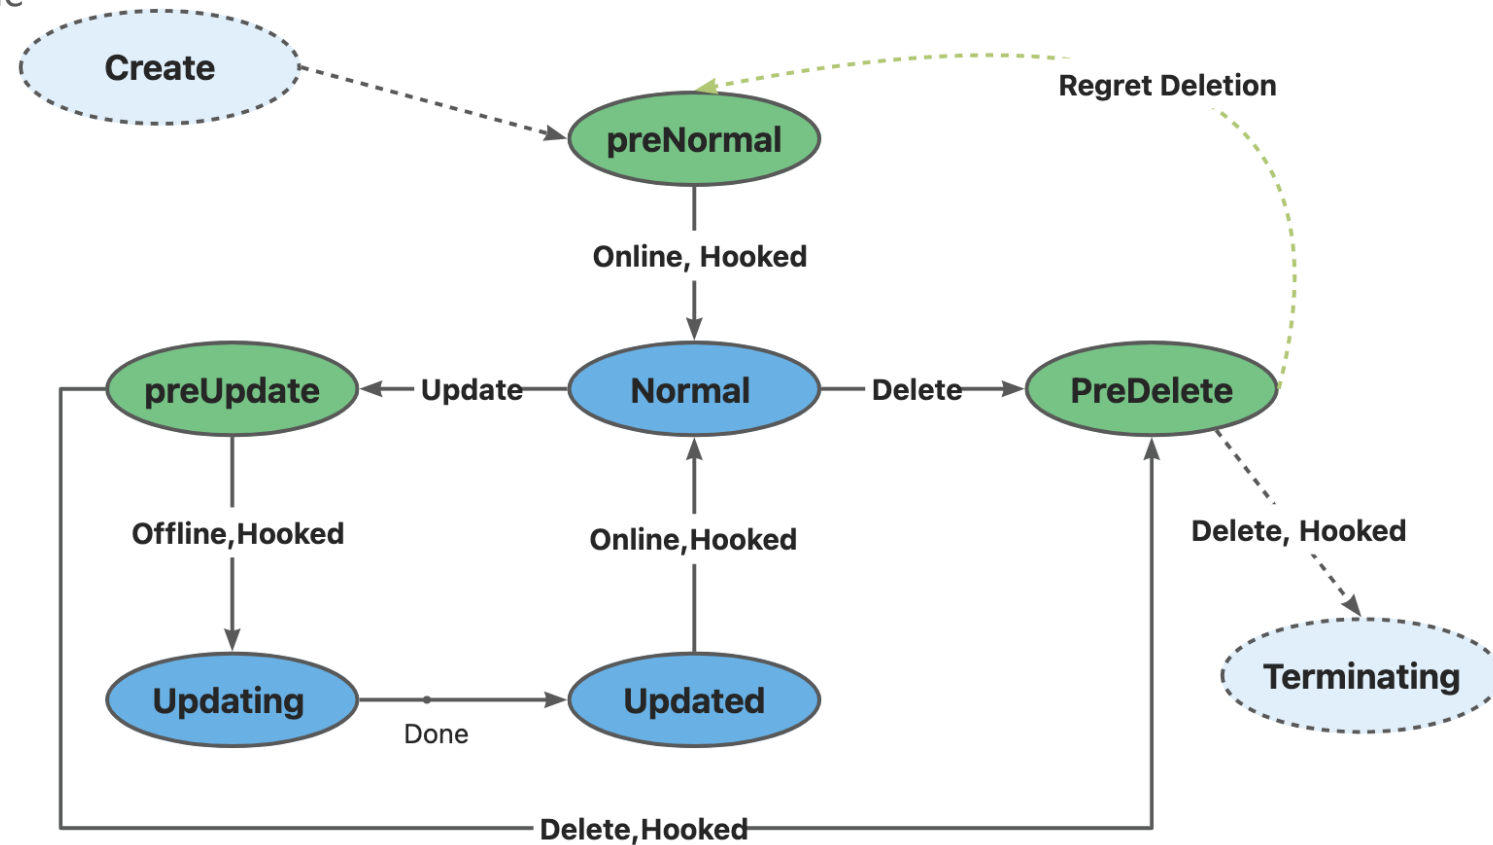


# challenges and solution: pod lifecycle

## Usecases:

- Remove traffic by mark pods not ready before inplace-update
- Ensure graceful shutdown before pod deletion
- Ensure graceful startup before pod becomes available

```
apiVersion: apps.kruise.io/v1alpha1
kind: CloneSet
spec:
  lifecycle:
    inplaceUpdate:
      markPodNotReady: true
    preDelete:
      finalizersHandler:
        - example.io/unready-blocker
    preNormal:
      finalizersHandler:
        - example.io/unready-blocker
```



# challenges and solution: restart counts



China 2024

- Limitation: Restart counts includes inplace-update count

Restart count since last update = restart count – restart count before last update

reported in pod status by kubelet

Recorded in pod annotations by kruise-manager before inplace-update

apiVersion: v1

kind: pods

metadata:

annotations:

apps.kruise.io/restart-count-before-last-update: '{"nginx":{"revision":"sample-d97f89dcf","restartCount":1}}'

apps.kruise.io/pod-restart-count-before-last-update: "2"

# Any chance to upstream ?



China 2024

- Top concerns
  - Security (Kritis & Notary)
  - Complexity
  - Immuable Conviction
- Opputunity
  - Good news: Inplace vertical pod scaling merged.
  - Bad news: **5 years** to merge, still alpha and not integrated by workload

# New opportunity : resource resizing



China 2024

## Use Cases

- load handled by the Pod has increased significantly, and current resources are not sufficient,
- load has decreased significantly, and allocated resources are unused,
- resources have simply been set improperly.



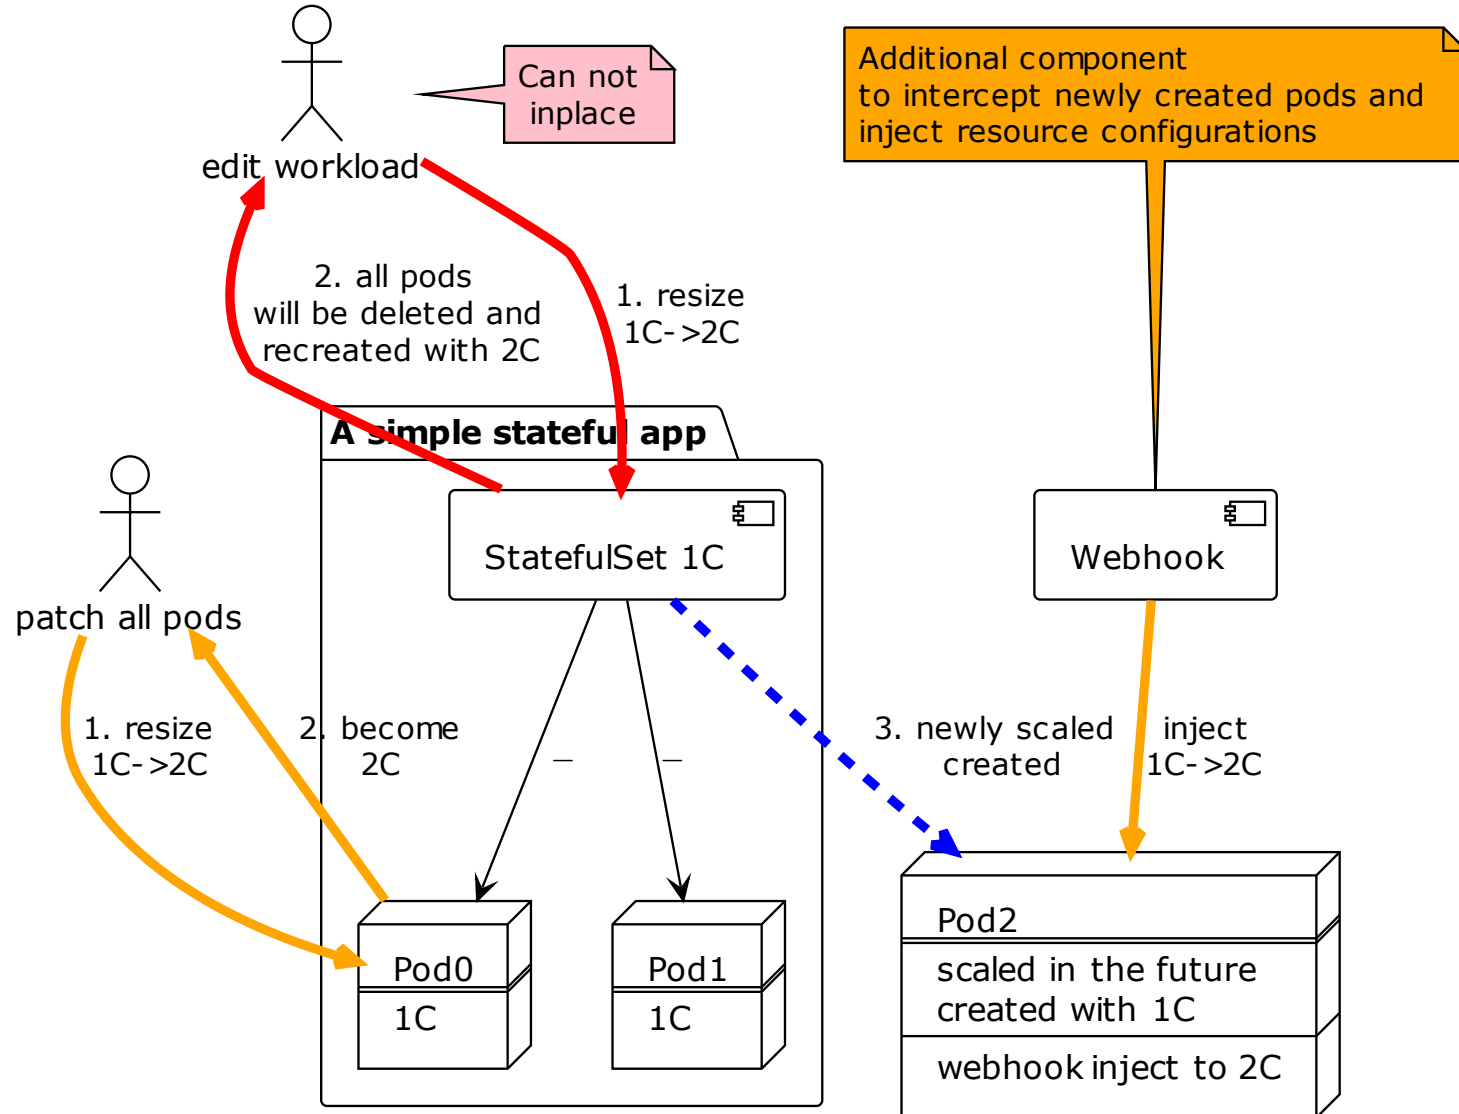


# New opportunity : resource resizing

## Background

### Kubernetes 1.27: In-place Resource Resize for Kubernetes Pods (alpha)

However, users cannot directly utilize this capability through workloads; directly modifying the resources of workload will still result in the reconstruction of pods.



# New opportunity : resource resizing

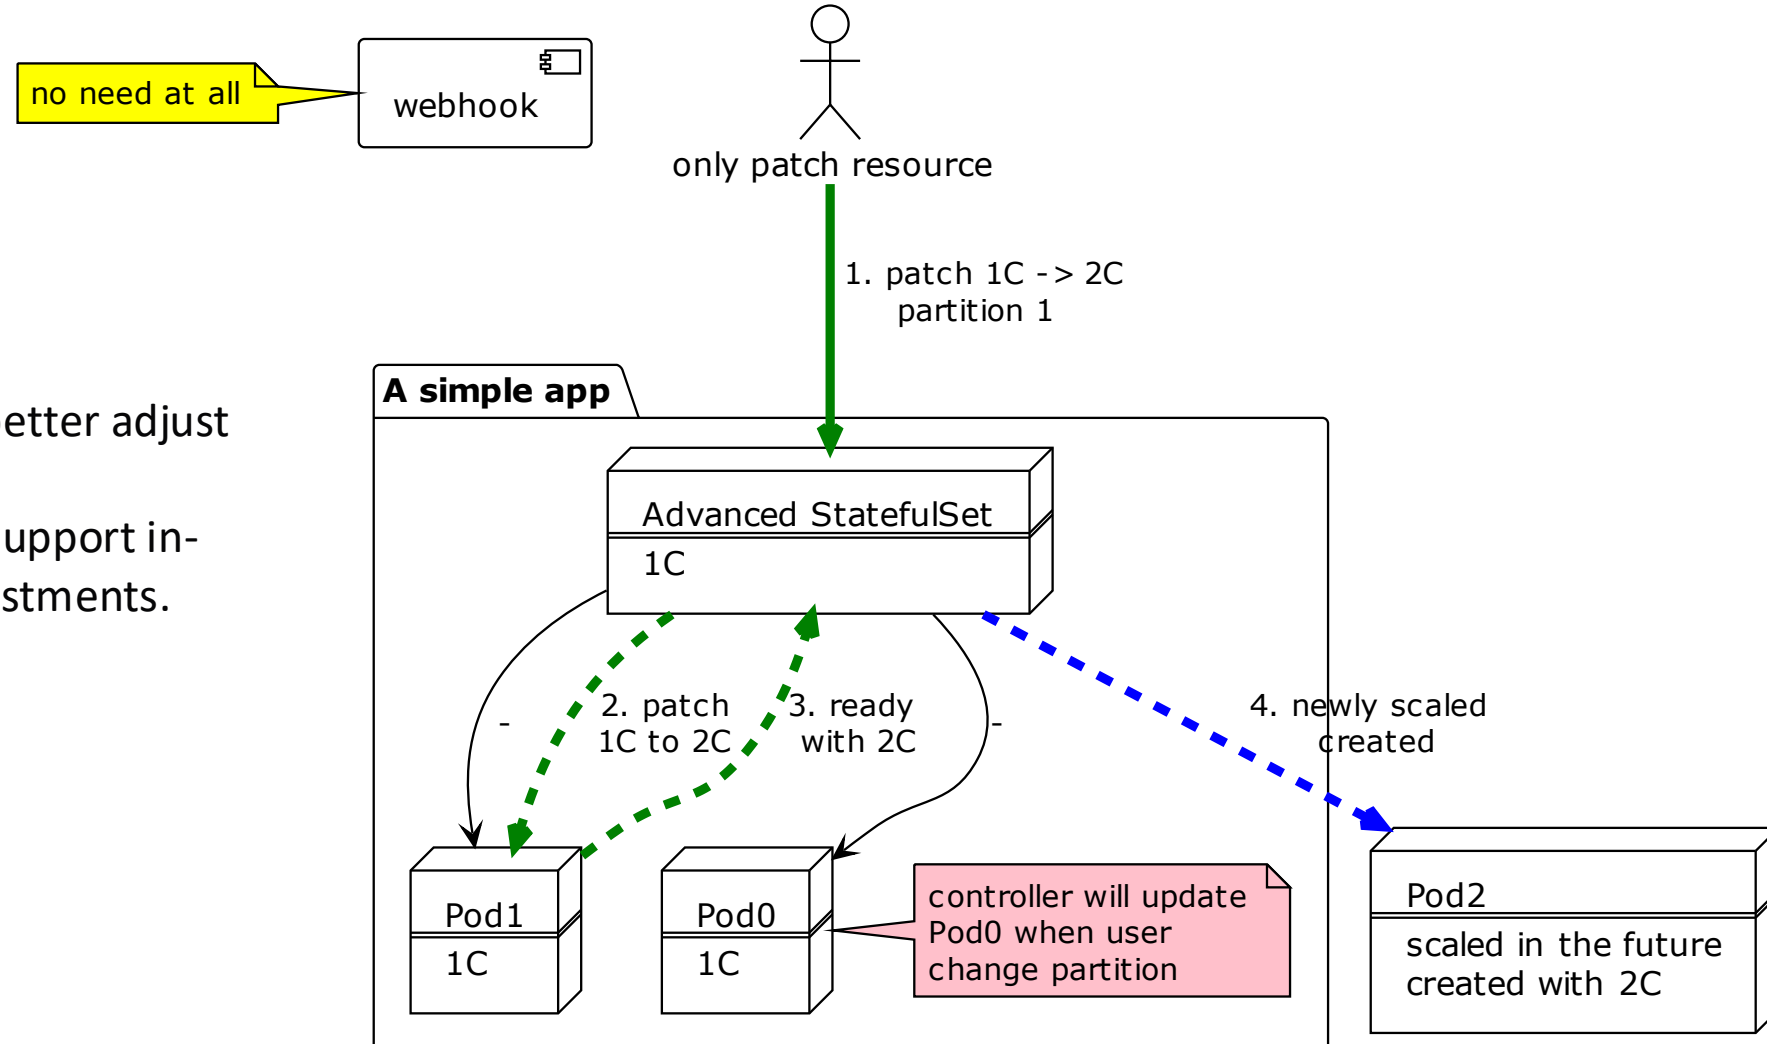


China 2024

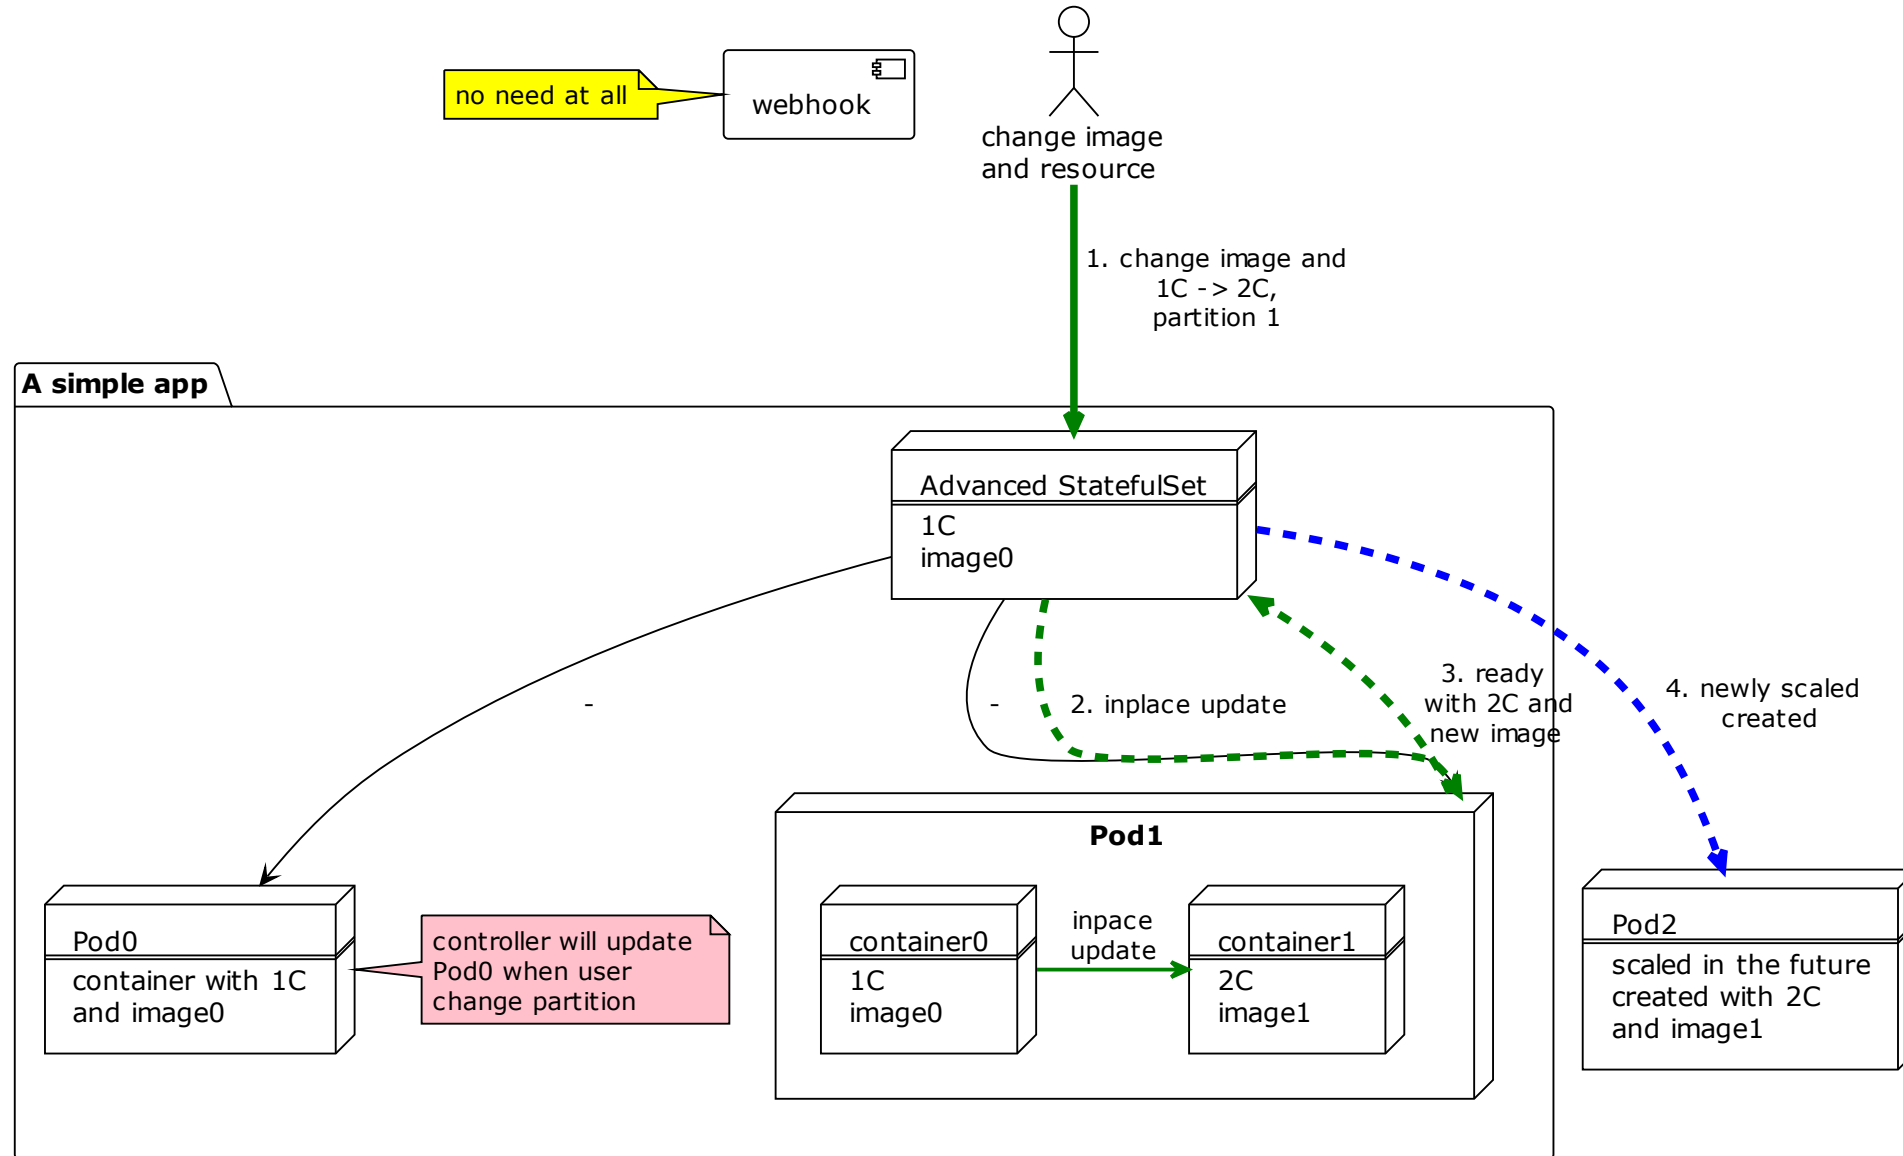
## Solution

How can we enable users to better adjust resource configurations?

-- Let our workloads natively support in-place resource allocation adjustments.



# New opportunity : resource resizing



# New opportunity : volume resizing

## Use Cases:

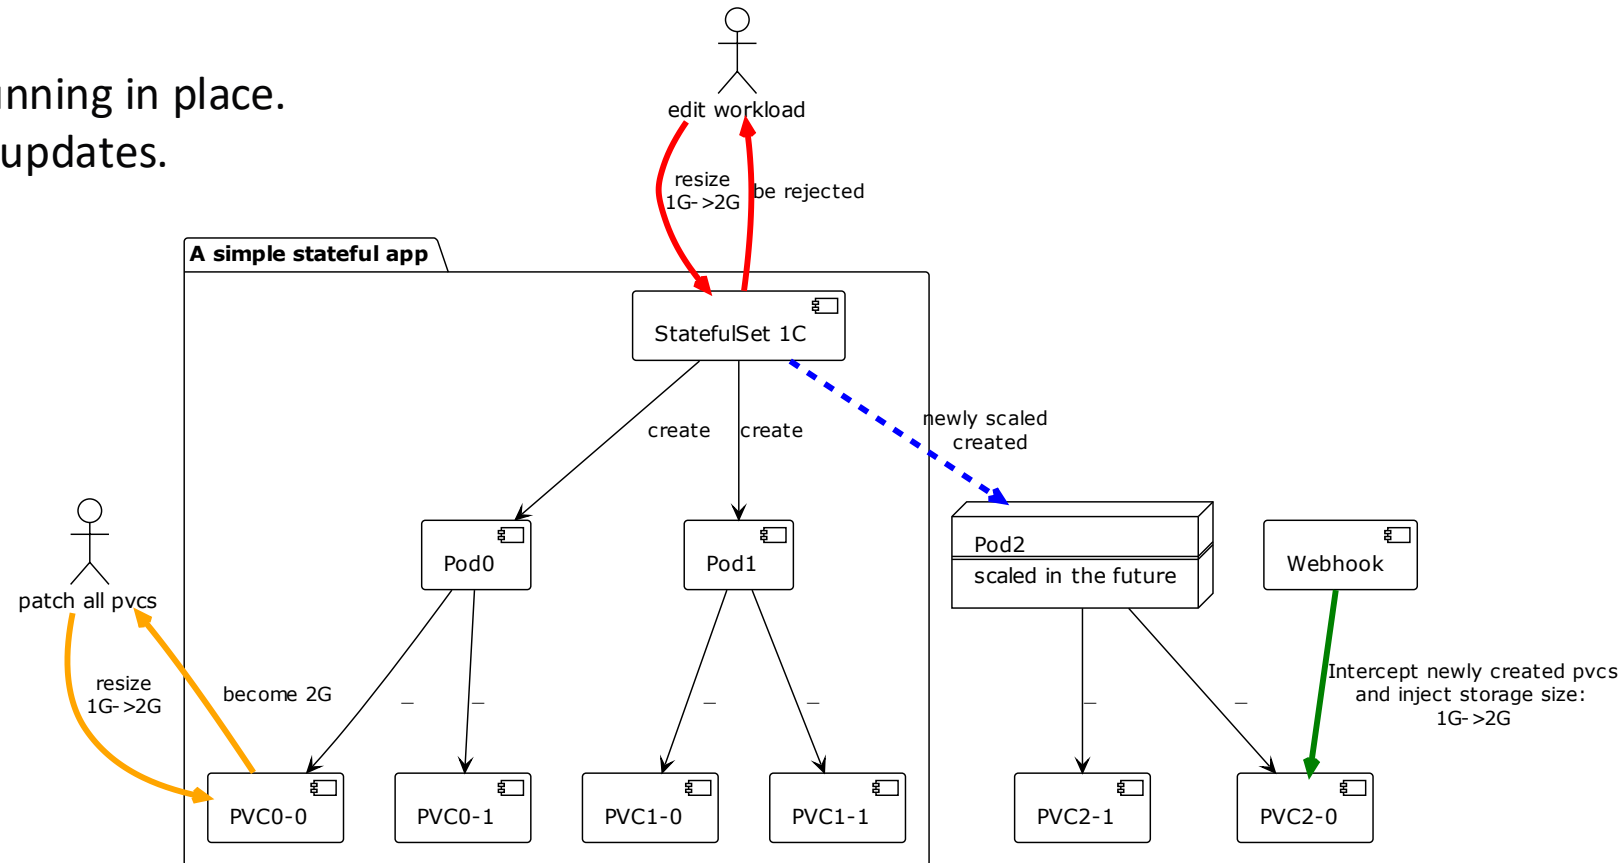
- Expand PVC size while the pod is running in place.
- Expand PVC size during pod rolling updates.

## Background :

In Kubernetes v1.11+, Volume expansion is enabled by default.

However, StatefulSet does not permit modification of this field.

Advanced StatefulSet allows modification but only takes effect for newly created PVCs.



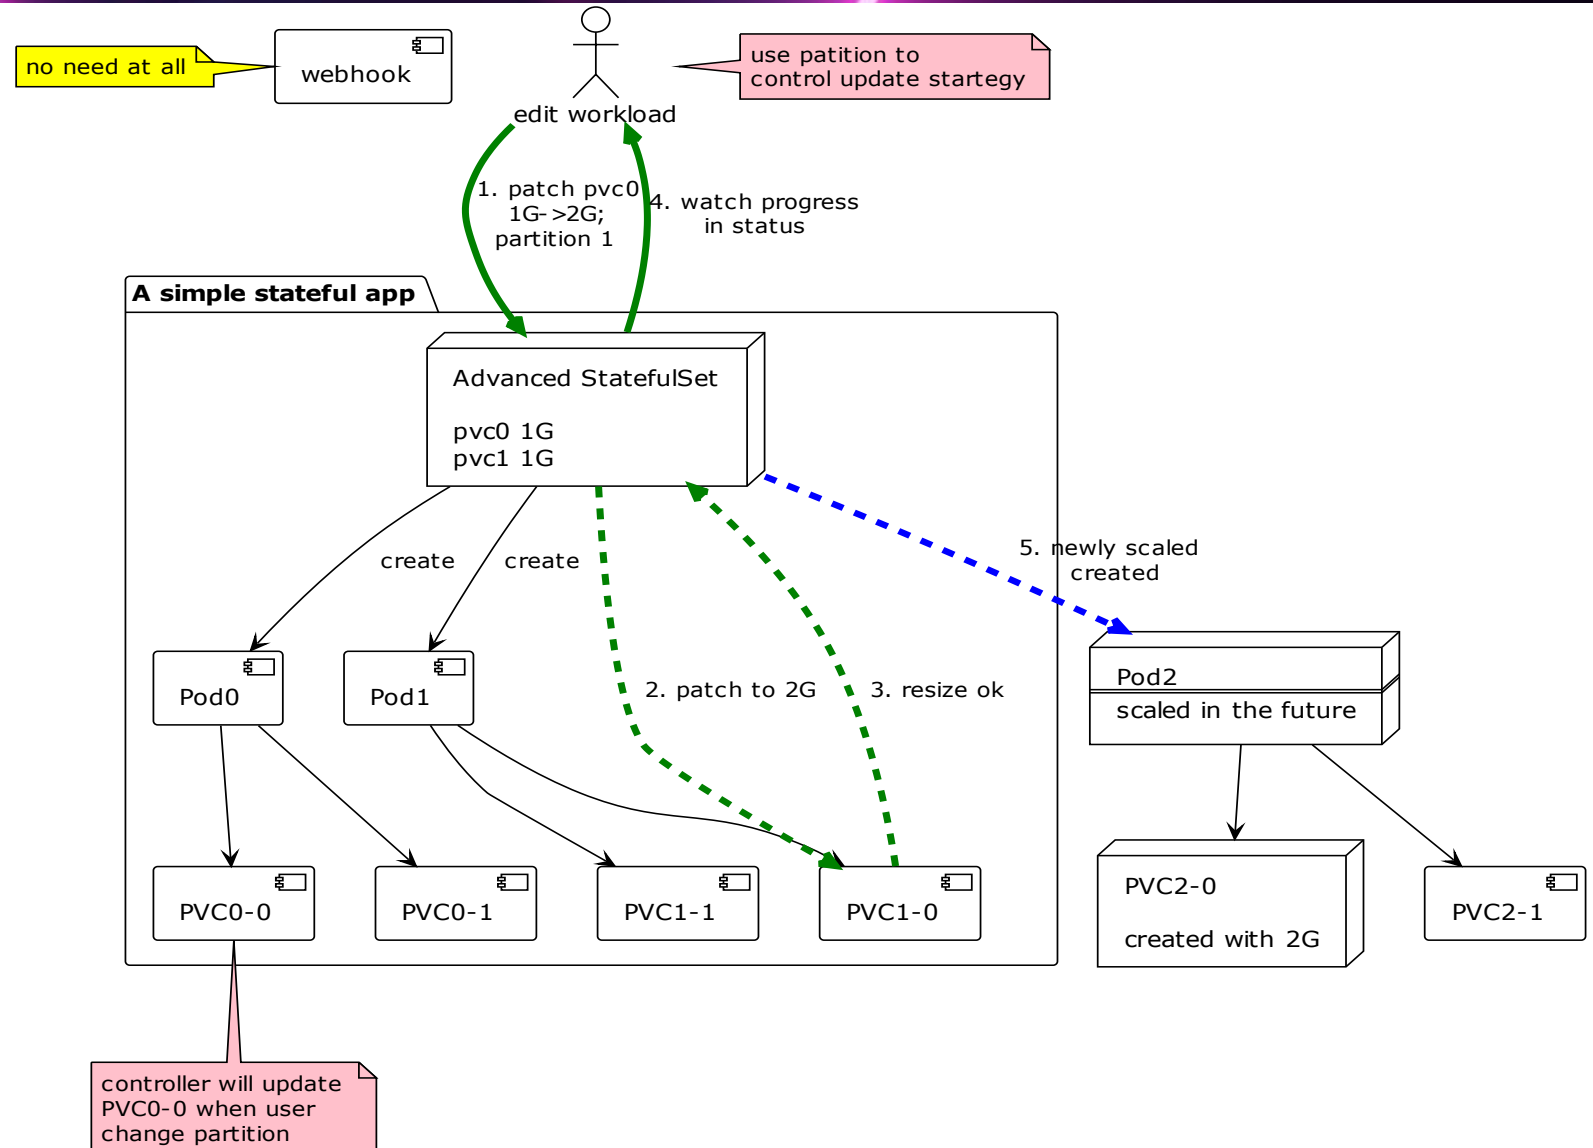


# New opportunity : volume resizing

## Solution :

How can we enable users to better resize pvc?

-- Let our workloads natively support in-place volume resizing.



# New opportunity : volume changes




China 2024

## Use Cases

- Modify other PVC configurations, such as reducing PVC size.
  - Users are willing to ensure data synchronization on their own
    - Use snapshot backup and restore methods to reconstruct PVCs.
- Use volume snapshots to restore PVCs managed by the workload.

## Not only Resize?

# New opportunity : volume changes



China 2024

## Solutions

- OnDelete allows editing templates of pvc and recreate newest pvc when user delete it.
- New PVC will be newest pvc

```
apiVersion: apps.kruise.io/v1beta1
kind: StatefulSet
metadata:
  name: vol-resize-test
spec:
  volumeClaimTemplates:
  - metadata:
      name: data0
    spec:
      accessModes:
      - ReadWriteOnce
      resources:
        requests:
          storage: 3Gi
      storageClassName: standard
    status: {}
  volumeClaimUpdateStrategy:
    # 可选值有 OnPodRollingUpdate 和 OnDelete
    # OnPodRollingUpdate : controller 会在 pod 升级过程中自动扩容 pvc
    # OnDelete : controller 只会在 pvc 删除后使用新的 volume claim
    # 重建 pvc
    # 即 kruise 1.7 版本之前默认的 pvc 更新策略。
    type: OnDelete
  status:
    replicas: 3
    readyReplicas: 2
    updatedReplicas: 3
    updatedReadyReplicas: 2
    volumeClaims:
    - compatibleReadyReplicas: 2
      compatibleReplicas: 3
      volumeClaimName: data
```

# New opportunity : volume changes



China 2024

## Background

In Kubernetes v1.17, VolumeSnapshot was introduced.

It can be used with the PVC's dataSource to create new PVC from snapshots.

However, users cannot directly use snapshots to restore PVCs through the workload, and manual handling is also inconvenient.

## Solutions

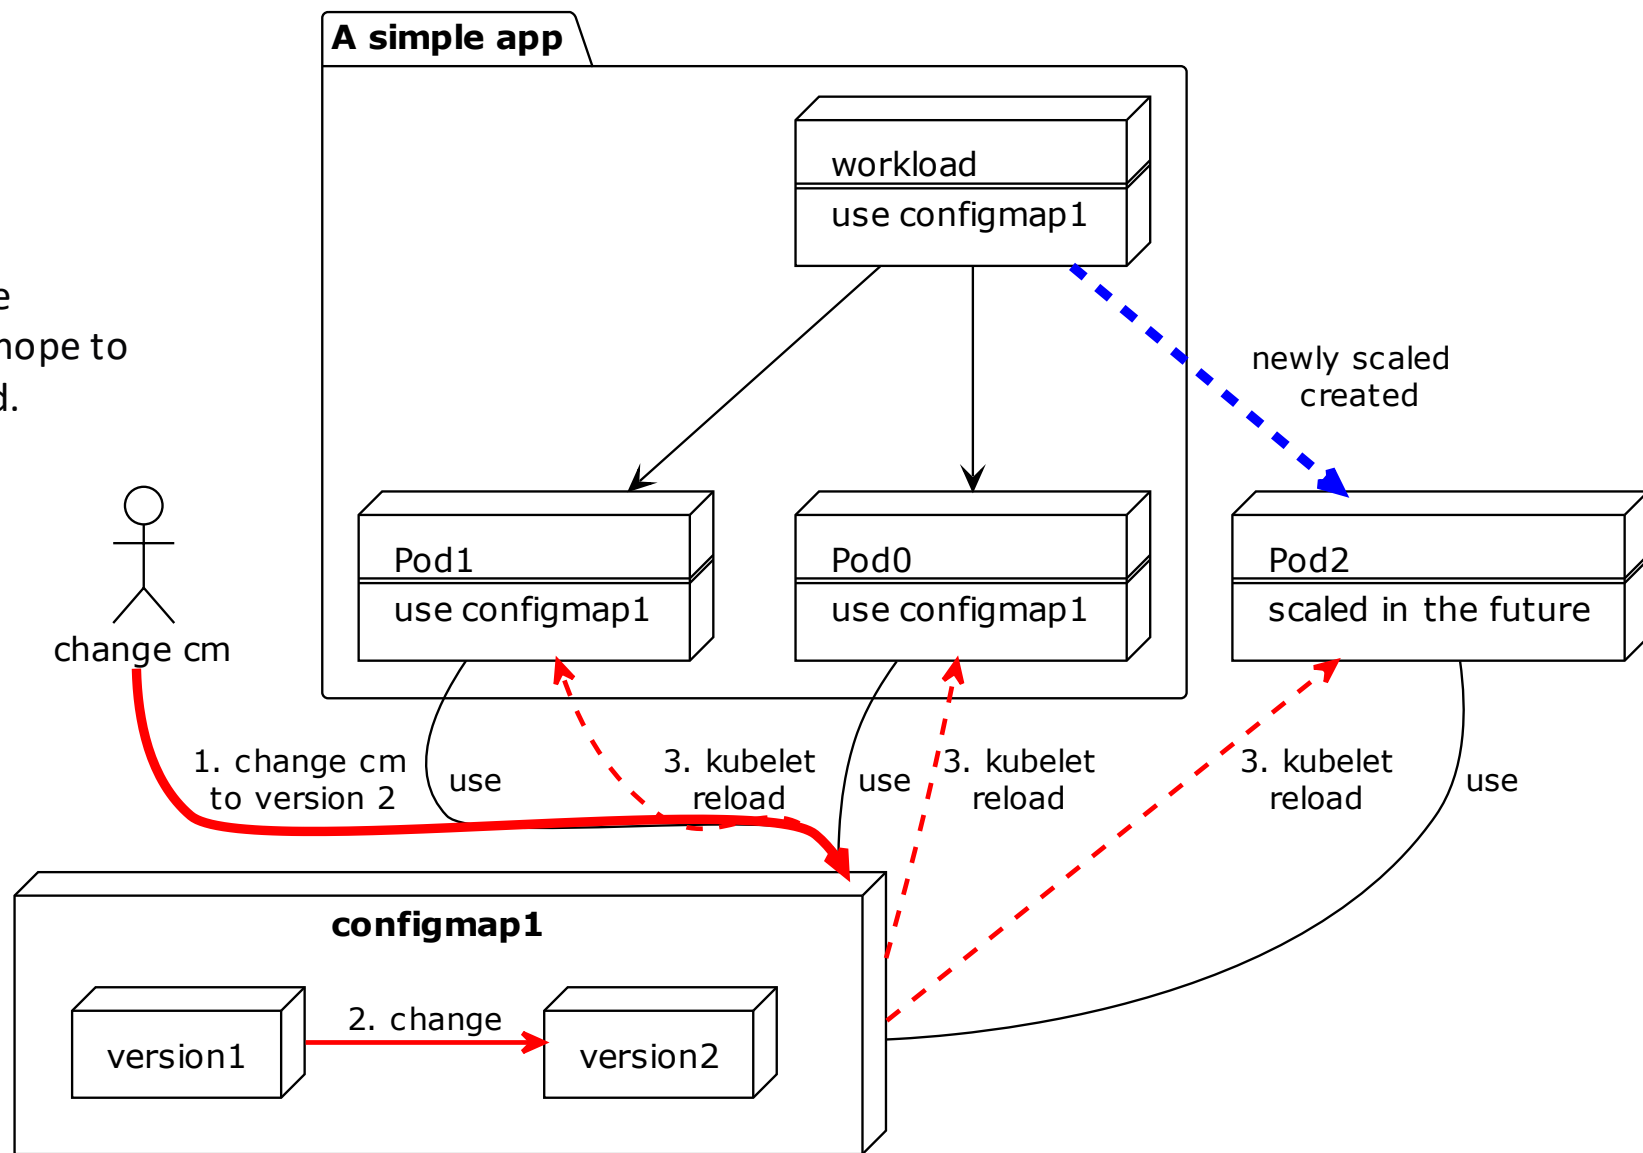
- [Plan] Define VolumeSnapshot information in VolumeClaimTemplates, maybe in annotations.
- [Plan] Implement a volume claim update strategy type about auto snapshot and restore from it.



# New opportunity : progressive configmap reload

## Use Cases

- Users use ConfigMaps to store configuration information and hope to implement a progressive reload.



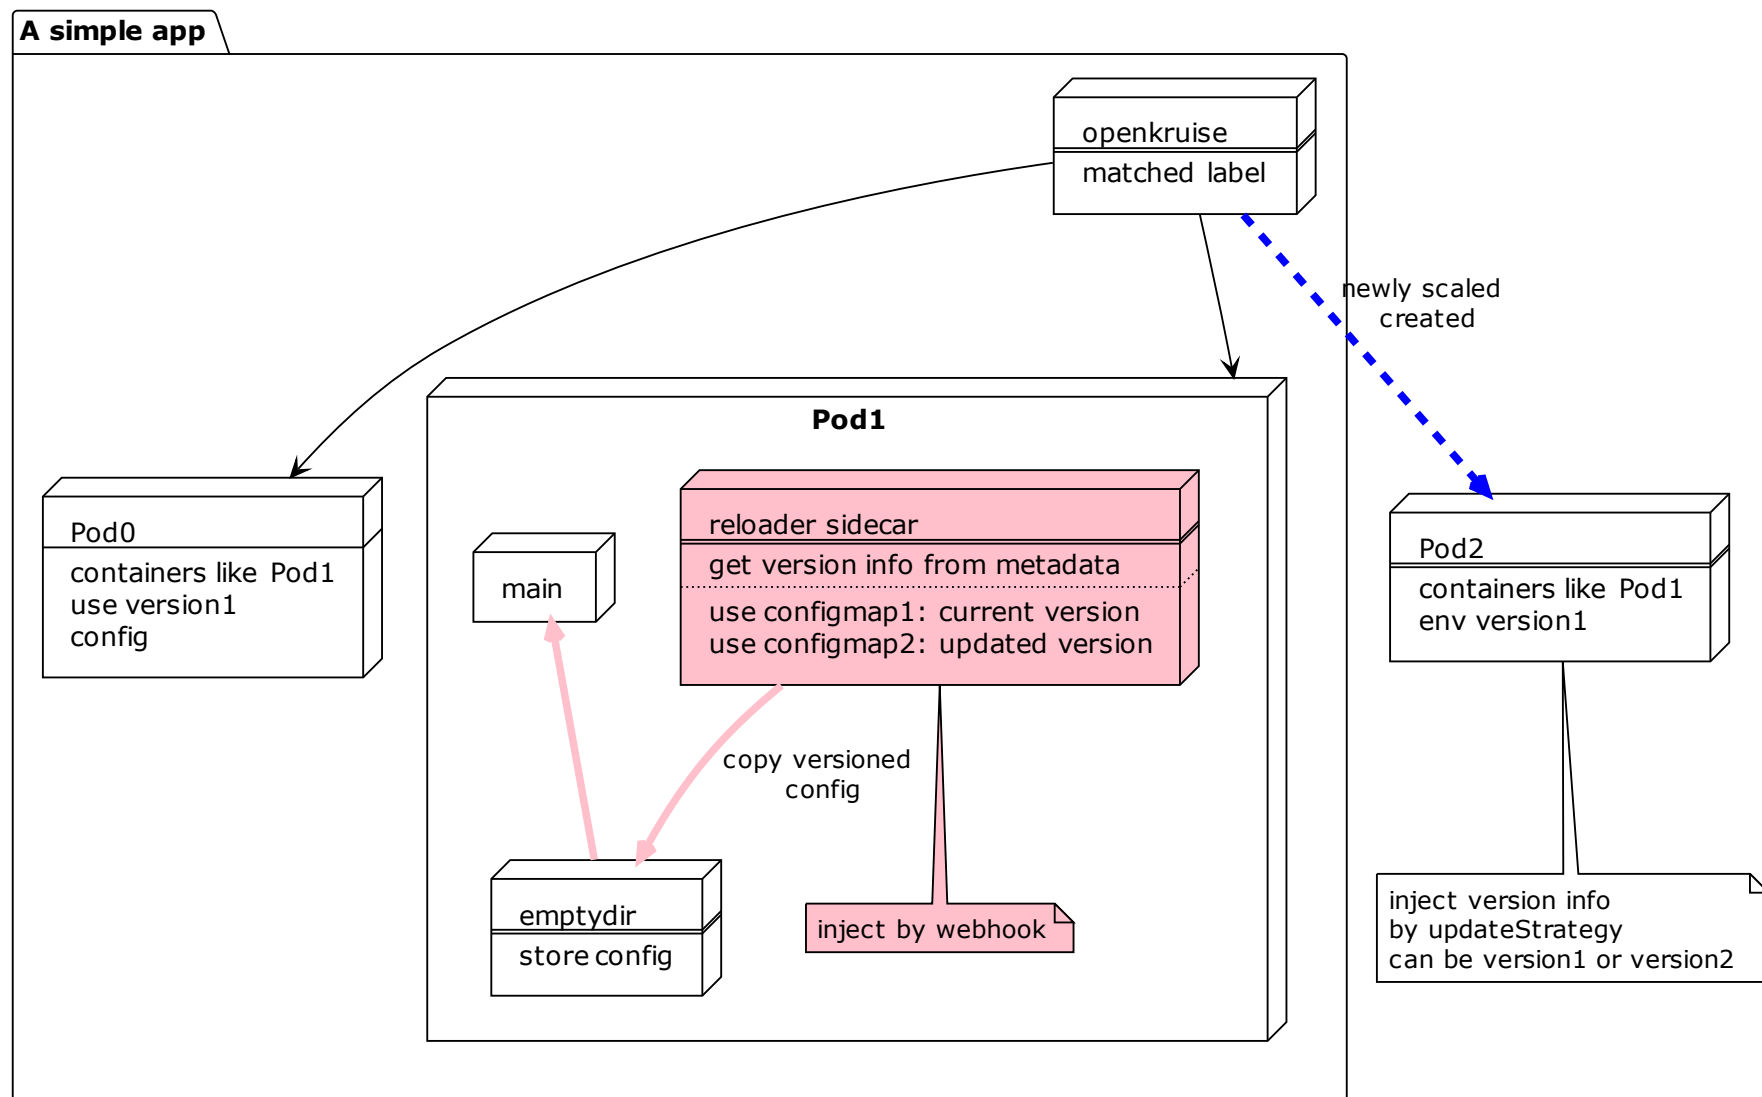
# New opportunity : progressive configmap reload



China 2024

## Solution:

1. Support canary
2. Support label selector



# Keep In Touch



China 2024

- Wechat Group: openkruise
- Dingding Group

<https://kubernetes.slack.com/archives/openkruise>


OpenKruise 社区交...

服务



1632人



 扫一扫群二维码，立刻加入该群。