

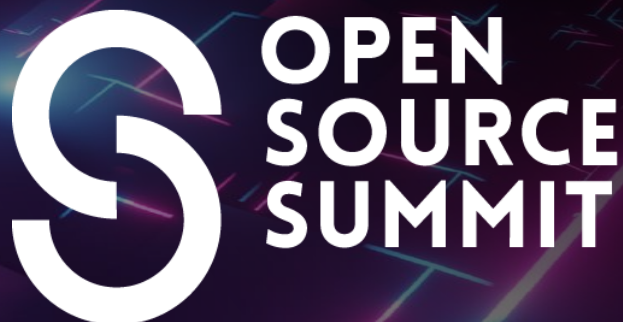


KubeCon



CloudNativeCon

THE LINUX FOUNDATION



AI_dev
Open Source GenAI & ML Summit

China 2024



KubeCon



CloudNativeCon



China 2024

Rollout Patterns: Smoothly Migrate and Roll Out Your Microservices

Tim Xiao, DaoCloud

About me



China 2024



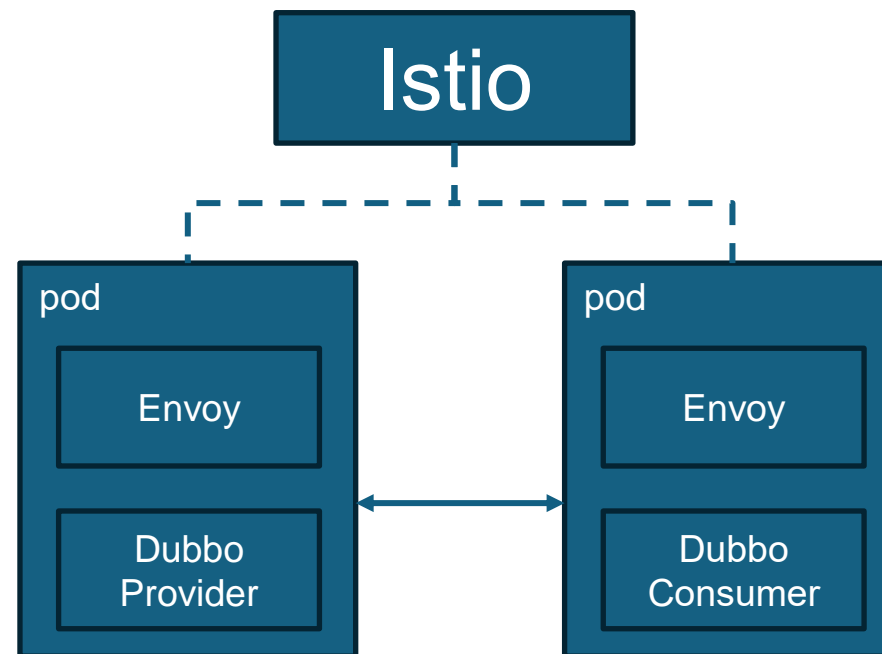
- Tim @muma378
- DaoCloud, Software Engineer, DevOps Platform Leader
- Argo Project Contributor
- Hiking 🏔️、Badminton 🏸 and Coffee ☕ Lover

Background



China 2024

- Services were built with Java and Dubbo2
- Microservices architecture
- Migrate to service mesh
- Want canary rollout



Rollout Patterns



China 2024

- PATTERN 1: One service at a time
- Usually happen for monolithic or compatible interfaces
- The delivery order could be very important
- Low requirement on Ops but high requirement on Dev

How you deliver your services?



KubeCon



CloudNativeCon



China 2024



Rollout Patterns

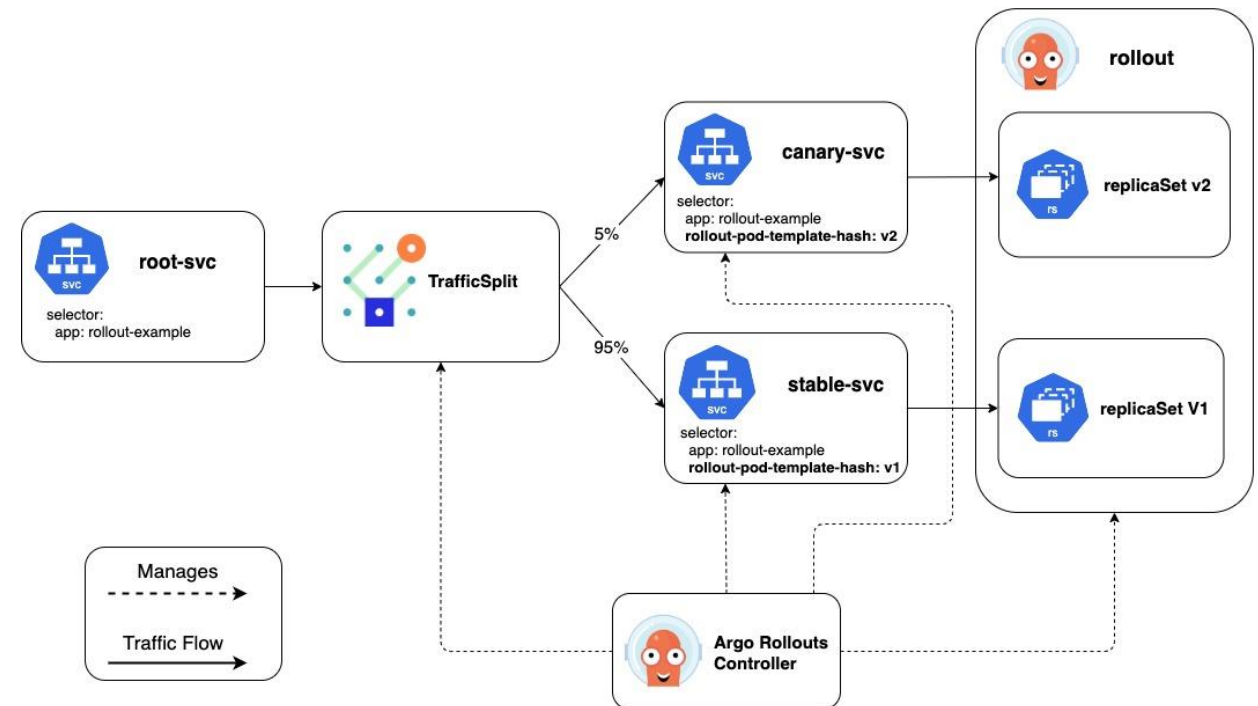


China 2024

- The ideal scenario for argo-rollouts

What is Argo-Rollouts

- (optionally) integrates with ingress controllers and service meshes, leveraging their traffic shaping abilities to gradually shift traffic to the new version during an update

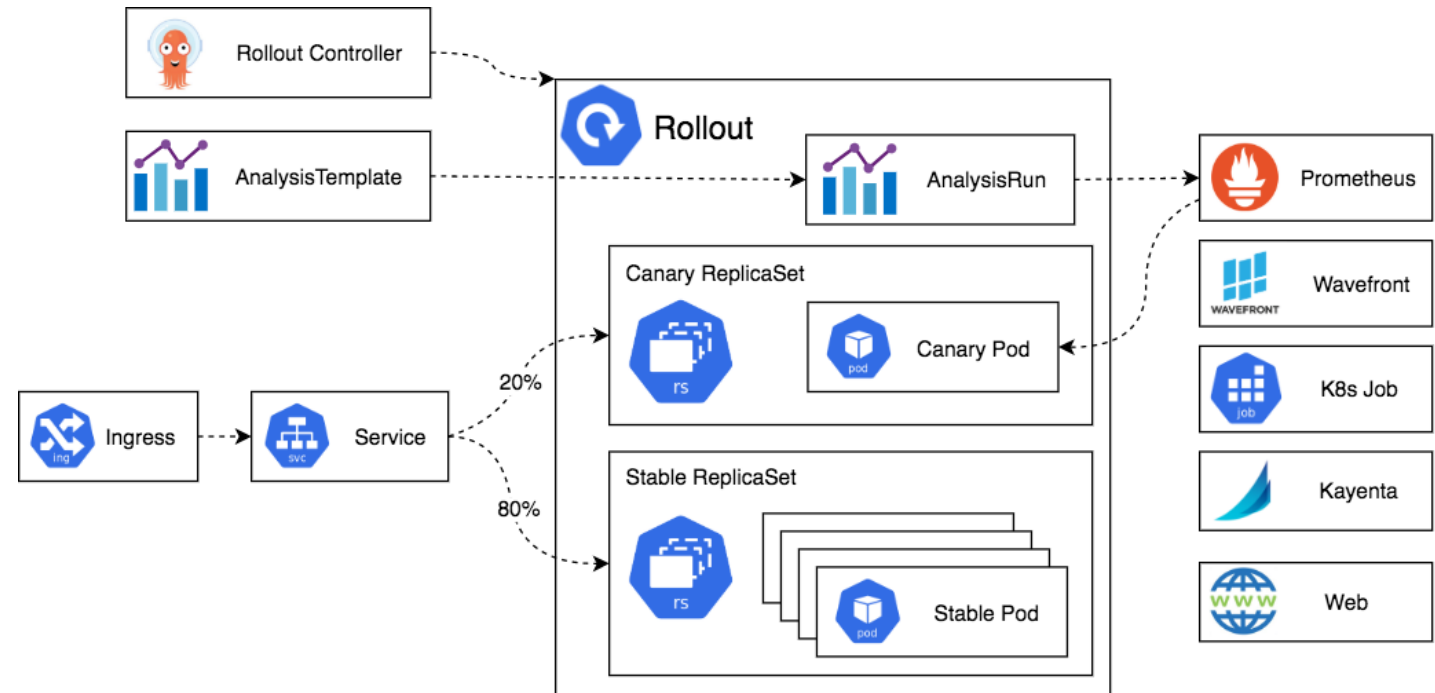


What is Argo-Rollouts





China 2024

- can query and interpret metrics from various providers to verify key KPIs and drive automated promotion or rollback during an update



Pattern 1



However...




China 2024

- When a change involves multiple components, they are preferred to be delivered in together
- Not only will the workloads be updated, but the configuration policies will also change

Rollout Patterns

- PATTERN 2: multiple services, each is backward-compatible
- Compared with pattern 1, applications are delivered as a whole
- Backward-compatible is for callers, not for callees
- Traffic become very complicate, hard to troubleshoot, hard to rollback



Pattern 2



China 2024

All you need is a delivery tool to control the order of deploying



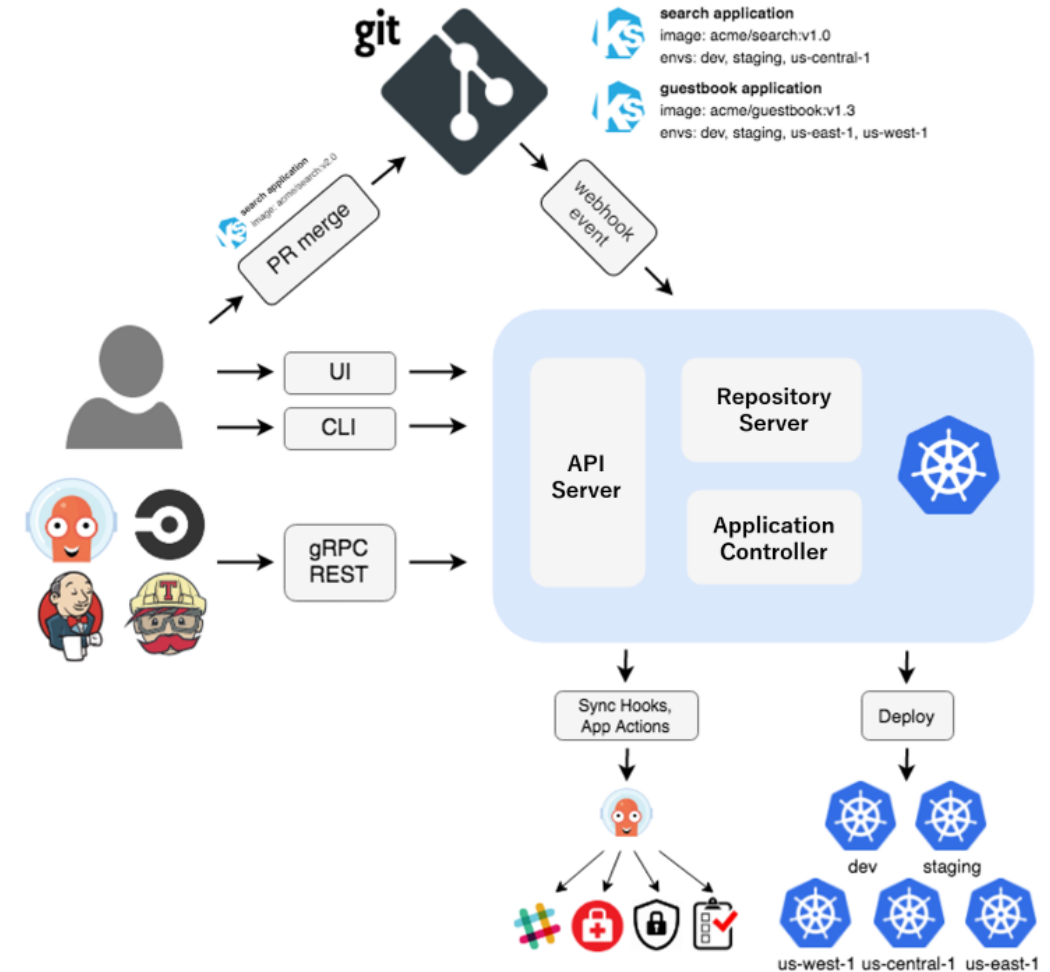
Argo CD

What is Argo CD



China 2024

- Argo CD follows the **GitOps** pattern of using Git repositories as the source of truth for defining the desired application state
- Argo CD automates the deployment of the desired application states in the specified target environments
- Argo CD reports & visualizes the differences, while providing facilities to automatically or manually sync the live state back to the desired target state.
- Support to define waves to apply resource



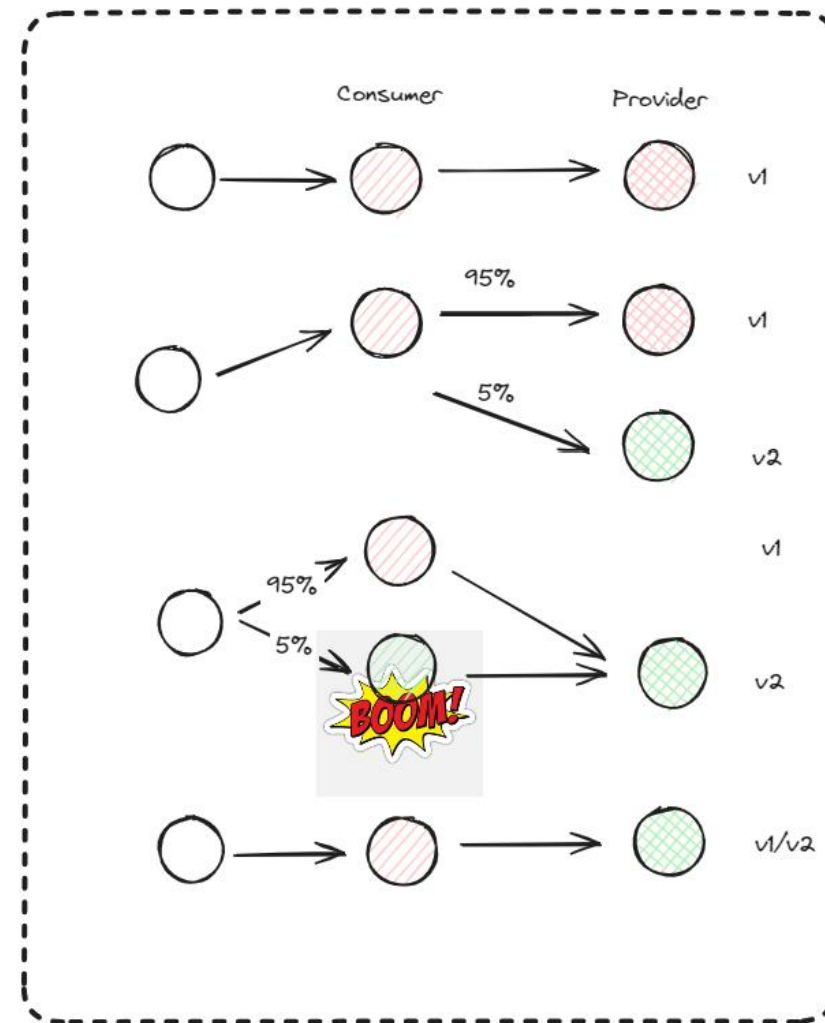

Pattern 2



China 2024

1. Identify the dependencies between services and design the order according to the direction
2. Progress one by one and monitor traffic shift
3. Rollback only the releasing service once outage occurs


Pattern 2



LIVE DEMO


Rollout Patterns

- PATTERN 3: multiple services with versions lock
- Compared with pattern 2, APIs are not compatible, requests are passed alone the lane, lane-cross is not permitted
- Rollout means shifting traffic in the entrance



First of all

- How to control the endpoint which consumer connects to?
- Headers (aka. A/B Testing)
- A non-invoke solution?
- Auto Inject headers via OpenTelemetry and WasmPlugin



Pattern 3



China 2024

1. Deploy all services but don't update traffic rules, no traffic will be directed to the new release at present
2. Shift traffic to the new release
3. Progressive promote till the new release take over all traffic
4. Scale the old replicas to 0
5. Reset the rules for the next rollout

How to update replicas but no update route?

setCanaryScale

```
24         port: 20880
25     strategy:
26       canary:
27         canaryService: dubbo3-provider-canary
28         stableService: dubbo3-provider
29         trafficRouting:
30           managedRoutes:
31             - name: test-v2
32           istio:
33             virtualService:
34               name: provider
35             routes:
36               - test-v1
37     steps:
38       - setCanaryScale:
39         weight: 100
40       - setHeaderRoute:
41         name: test-v2
42         match:
43           - headerName: lane
44             headerValue:
45               exact: test-v2
46       - pause: {}
```

More details



China 2024

How to add different headers to requests forwarded to different destination?

ask traffic management tools for help

```
1  apiVersion: networking.istio.io/v1beta1
2  kind: VirtualService
3  metadata:
4    name: consumer
5  spec:
6    gateways:
7      - dubbo
8    hosts:
9      - "*"
10   http:
11     - name: test-v1
12       route:
13         - destination:
14             host: dubbo3-consumer
15             port:
16               number: 9090
17             weight: 100
18             headers:
19               request:
20                 add:
21                   lane: test-v1
22         - destination:
23             host: dubbo3-consumer-canary
24             port:
25               number: 9090
26             weight: 0
27             headers:
28               request:
29                 add:
30                   lane: test-v2
```

add lane:test-v1

add lane:test-v2

How to update route to only accept requests with specified headers

setHeaderRoute

```
24         port: 20880
25     strategy:
26       canary:
27         canaryService: dubbo3-provider-canary
28         stableService: dubbo3-provider
29         trafficRouting:
30           managedRoutes:
31             - name: test-v2
32           istio:
33             virtualService:
34               name: provider
35             routes:
36               - test-v1
37     steps:
38       - setCanaryScale:
39         weight: 100
40       - setHeaderRoute:
41         name: test-v2
42         match:
43           - headerName: lane
44             headerValue:
45               exact: test-v2
46       - pause: {}
```

add new
header match
route

More details

```
1  apiVersion: networking.istio.io/v1beta1
2  kind: VirtualService
3  metadata:
4    name: provider
5  spec:
6    gateways:
7      - mesh
8    hosts:
9      - dubbo3-provider
10   http:
11     - name: test-v1
12       match:
13         - headers:
14             lane:
15               exact: test-v1
16         route:
17           - destination:
18               host: dubbo3-provider
19               port:
20                 number: 20880
21               weight: 100
22       - route:
23         - destination:
24             host: dubbo3-provider
25             port:
26               number: 20880
27
```




```
10  http:
11    - name: test-v1
12      match:
13        - headers:
14            lane:
15              exact: test-v1
16        route:
17          - destination:
18              host: dubbo3-provider
19              port:
20                number: 20880
21              weight: 100
22    - name: test-v2
23      match:
24        - headers:
25            lane:
26              exact: test-v2
27      route:
28        - destination:
29            host: dubbo3-provider-canary
30            port:
31              number: 20880
32            weight: 100
33    - route:
34      - destination:
35          host: dubbo3-provider
36          port:
37            number: 20880

```

More details

Once a rollout was finished, header route will be cleared, but requests are still with headers (because the entrance rollout is not finished yet), how to handle this?

add default route




```
1  apiVersion: networking.istio.io/v1beta1
2  kind: VirtualService
3  metadata:
4    name: provider
5  spec:
6    gateways:
7      - mesh
8    hosts:
9      - dubbo3-provider
10   http:
11     - name: test-v1
12       match:
13         - headers:
14             lane:
15               exact: test-v1
16         route:
17           - destination:
18               host: dubbo3-provider
19               port:
20                 number: 20880
21             weight: 100
22     - route:
23       - destination:
24           host: dubbo3-provider
25           port:
26             number: 20880
27
```

route if no match

More details

Why do we have to leave the entrance rollout finish last?

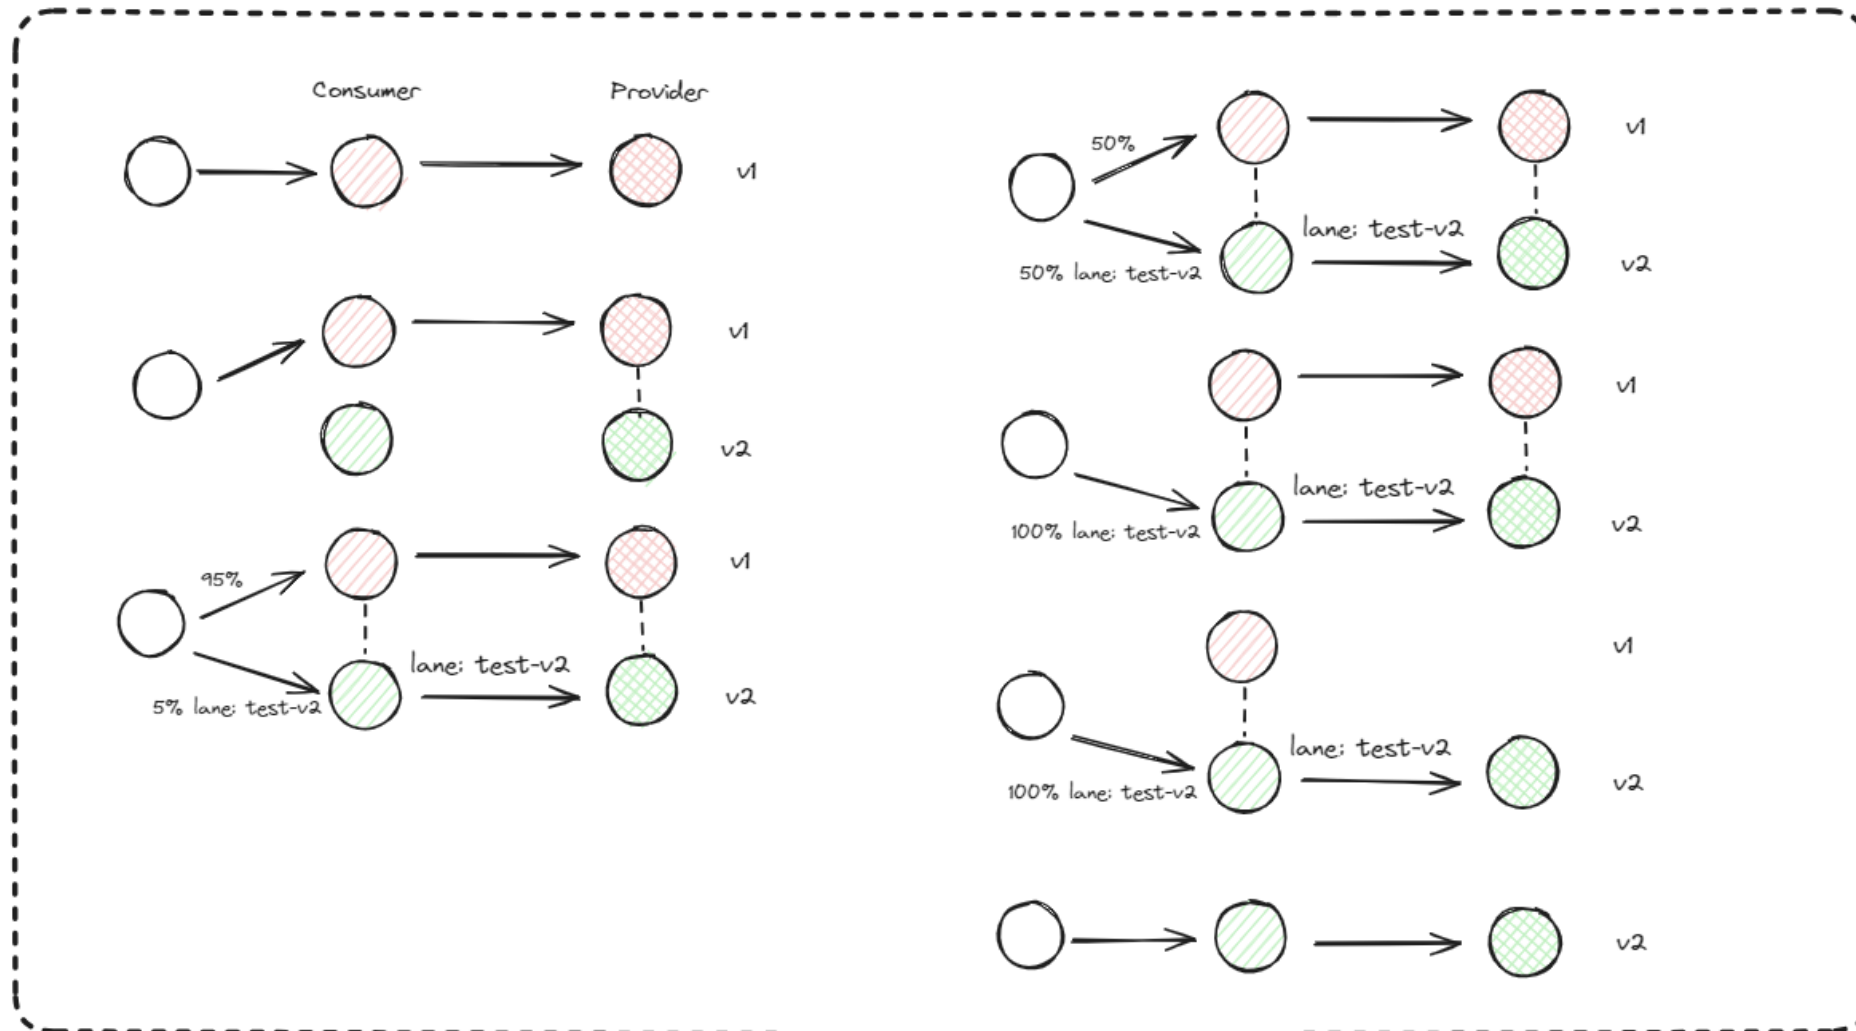
If we don't, the old stable service will be scaled down to 0 and the canary service will be the new stable. Which means all traffic will go the lane v1



```
29 strategy:
30   canary:
31     canaryService: dubbo3-consumer-canary
32     stableService: dubbo3-consumer
33     trafficRouting:
34       istio:
35         virtualService:
36           name: consumer
37           routes:
38             - test-v1
39     steps:
40       - setCanaryScale:
41         weight: 20
42       - pause: {}
43       - setWeight: 20
44       - pause:
45         duration: 10m
46       - setWeight: 50
47       - pause:
48         duration: 10m
49       - setWeight: 100
50       - pause: {}
```


wait here

Pattern 3



Pattern 3: rollback

- Once an outage happens, abort the entrance service at first, then the others



LIVE DEMO

Questions?



KubeCon



CloudNativeCon



China 2024



DaoCloud

Appreciate any feedback

wechat @muma378

github @muma378